

WebGPU

Kai Ninomiya (Google)

1. Introduction

This section is non-normative.

[Graphics Processing Units](#), or GPUs for short, have been essential in enabling rich rendering and computational applications in personal computing. WebGPU is an API that exposes the capabilities of GPU hardware for the Web. The API is designed from the ground up to efficiently map to (post-2014) native GPU APIs. WebGPU is not related to [WebGL](#) and does not explicitly target OpenGL ES.

WebGPU sees physical GPU hardware as [GPUAdapters](#). It provides a connection to an adapter via [GPUDevice](#), which manages resources, and the device's [GPUQueues](#), which execute commands. [GPUDevice](#) may have its own memory with high-speed access to the processing units. [GPUBuffer](#) and [GPUTexture](#) are the *physical resources* backed by GPU memory. [GPUCommandBuffer](#) and [GPURenderBundle](#) are containers for user-recorded commands. [GPUShaderModule](#) contains *shader* code. The other resources, such as [GPUSampler](#) or [GPUBindGroup](#), configure the way *physical resources* are used by the GPU.

GPUs execute commands encoded in [GPUCommandBuffers](#) by feeding data through a [pipeline](#), which is a mix of fixed-function and programmable stages. Programmable stages execute *shaders*, which are special programs designed to run on GPU hardware. Most of the state of a [pipeline](#) is defined by a [GPURenderPipeline](#) or a [GPUComputePipeline](#) object. The state not included in these [pipeline](#) objects is set during encoding with commands, such as [beginRenderPass\(\)](#) or [setBlendConstant\(\)](#).

2. Malicious use considerations

This section is non-normative. It describes the risks associated with exposing this API on the Web.

2.1. Security Considerations

The security requirements for WebGPU are the same as ever for the web, and are likewise non-negotiable. The general approach is strictly validating all the commands before they reach GPU, ensuring that a page can only work with its own data.

2.1.1. CPU-based undefined behavior

A WebGPU implementation translates the workloads issued by the user into API commands specific to the target platform. Native APIs specify the valid usage for the commands (for example, see [vkCreateDescriptorSetLayout](#)) and generally don't guarantee any outcome if the valid usage rules are not followed. This is called "undefined behavior", and it can be exploited by an attacker to access memory they don't own, or force the driver to execute arbitrary code.

In order to disallow insecure usage, the range of allowed WebGPU behaviors is defined for any input. An implementation has to validate all the input from the user and only reach the driver with the valid workloads. This document specifies all the error conditions and handling semantics. For example, specifying the same buffer with intersecting ranges in both "source" and "destination" of [copyBufferToBuffer\(\)](#) results in [GPUCommandEncoder](#) generating an error, and no other operation occurring.

See [§ 2.2 Errors & Debugging](#) for more information about error handling.

2.1.2. GPU-based undefined behavior

WebGPU [shaders](#) are executed by the compute units inside GPU hardware. In native APIs, some of the shader instructions may result in undefined behavior on the GPU. In order to address that, the shader instruction set and its defined behaviors are strictly defined by WebGPU. When a shader is provided to [createShaderModule\(\)](#), the WebGPU implementation has to validate it before doing any translation (to platform-specific shaders) or transformation passes.

2.1.3. Uninitialized data

Generally, allocating new memory may expose the leftover data of other applications running on the system. In order to address that, WebGPU conceptually initializes all the resources to zero, although in practice an implementation may skip this step if it sees the developer initializing the contents manually. This includes variables and shared workgroup memory inside shaders.

The precise mechanism of clearing the workgroup memory can differ between platforms. If the native API does not provide facilities to clear it, the WebGPU implementation transforms the compute shader to first do a clear across all invocations, synchronize them, and continue executing developer's code.

NOTE:

The initialization status of a resource used in a queue operation can only be known when the operation is enqueued (not when it is encoded into a command buffer, for example). Therefore, some implementations will require an unoptimized late-clear at enqueue time (e.g. clearing a texture, rather than changing [GPULoadOp "load"](#) to ["clear"](#)).

As a result, all implementations **should** issue a developer console warning about this potential performance penalty, even if there is no penalty in that implementation.

2.1.4. Out-of-bounds access in shaders

[Shaders](#) can access [physical resources](#) either directly (for example, as a ["uniform" GPUBufferBinding](#)), or via *texture units*, which are fixed-function hardware blocks that handle texture coordinate conversions. Validation in the WebGPU API can only guarantee that all the inputs to the shader are provided and they have the correct usage and types. The WebGPU API can not guarantee that the data is accessed within bounds if the [texture units](#) are not involved.

In order to prevent the shaders from accessing GPU memory an application doesn't own, the WebGPU implementation may enable a special mode (called "robust buffer access") in the driver that guarantees that the access is limited to buffer bounds.

Alternatively, an implementation may transform the shader code by inserting manual bounds checks. When this path is taken, the out-of-bound checks only apply to array indexing. They aren't needed for plain field access of shader structures due to the [minBindingSize](#) validation on the host side.

If the shader attempts to load data outside of [physical resource](#) bounds, the implementation is allowed to:

- return a value at a different location within the resource bounds

- return a value vector of "(0, 0, 0, X)" with any "X"

- partially discard the draw or dispatch call

If the shader attempts to write data outside of [physical resource](#) bounds, the implementation is allowed to:

- write the value to a different location within the resource bounds

- discard the write operation

- partially discard the draw or dispatch call

2.1.5. Invalid data

When uploading [floating-point](#) data from CPU to GPU, or generating it on the GPU, we may end up with a binary representation that doesn't correspond to a valid number, such as infinity or NaN (not-a-number). The GPU behavior in this case is subject to the accuracy of the GPU hardware implementation of the IEEE-754 standard. WebGPU guarantees that introducing invalid floating-point numbers would only affect the results of arithmetic computations and will not have other side effects.

2.1.6. Driver bugs

GPU drivers are subject to bugs like any other software. If a bug occurs, an attacker could possibly exploit the incorrect behavior of the driver to get access to unprivileged data. In order to reduce the risk, the WebGPU working group will coordinate with GPU vendors to integrate the WebGPU Conformance Test Suite (CTS) as part of their driver testing process, like it was done for WebGL. WebGPU implementations are expected to have workarounds for some of the discovered bugs, and disable WebGPU on drivers with known bugs that can't be worked around.

2.1.7. Timing attacks

2.1.7.1. Content-timeline timing

WebGPU does not expose new states to JavaScript (the [content timeline](#)) which are shared between [agents](#) in an [agent cluster](#). [Content timeline](#) states such as [\[\[mapping\]\]](#) only change during explicit [content timeline](#) tasks, like in plain JavaScript.

2.1.7.2. Device/queue-timeline timing

Writable storage buffers and other cross-invocation communication may be usable to construct high-precision timers on the [queue timeline](#).

The optional ["timestamp-query"](#) feature also provides high precision timing of GPU operations. To mitigate security and privacy concerns, the timing query values are aligned to a lower precision: see [current queue timestamp](#). Note in particular:

The [device timeline](#) typically runs in a process that is shared by multiple origins, so cross-origin isolation (provided by COOP/COEP) does not provide isolation of device/queue-timeline timers.

[Queue timeline](#) work is issued from the device timeline, and may execute on GPU hardware that does not provide the isolation expected of CPU processes (such as Meltdown mitigations).

GPU hardware is not typically susceptible to Spectre-style attacks, **but** WebGPU may be implemented in software, and software implementations may run in a shared process, preventing isolation-based mitigations.

2.1.8. Row hammer attacks

[Row hammer](#) is a class of attacks that exploit the leaking of states in DRAM cells. It could be used [on GPU](#). WebGPU does not have any specific mitigations in place, and relies on platform-level solutions, such as reduced memory refresh intervals.

2.1.9. Denial of service

WebGPU applications have access to GPU memory and compute units. A WebGPU implementation may limit the available GPU memory to an application, in order to keep other applications responsive. For GPU processing time, a WebGPU implementation may set up "watchdog" timer that makes sure an application doesn't cause GPU

unresponsiveness for more than a few seconds. These measures are similar to those used in WebGL.

2.1.10. Workload identification

WebGPU provides access to constrained global resources shared between different programs (and web pages) running on the same machine. An application can try to indirectly probe how constrained these global resources are, in order to reason about workloads performed by other open web pages, based on the patterns of usage of these shared resources. These issues are generally analogous to issues with Javascript, such as system memory and CPU execution throughput. WebGPU does not provide any additional mitigations for this.

2.1.11. Memory resources

WebGPU exposes fallible allocations from machine-global memory heaps, such as VRAM. This allows for probing the size of the system's remaining available memory (for a given heap type) by attempting to allocate and watching for allocation failures.

GPUs internally have one or more (typically only two) heaps of memory shared by all running applications. When a heap is depleted, WebGPU would fail to create a resource. This is observable, which may allow a malicious application to guess what heaps are used by other applications, and how much they allocate from them.

2.1.12. Computation resources

If one site uses WebGPU at the same time as another, it may observe the increase in time it takes to process some work. For example, if a site constantly submits compute workloads and tracks completion of work on the queue, it may observe that something else also started using the GPU.

A GPU has many parts that can be tested independently, such as the arithmetic units, texture sampling units, atomic units, etc. A malicious application may sense when some of these units are stressed, and attempt to guess the workload of another application by analyzing the stress patterns. This is analogous to the realities of CPU execution of Javascript.

2.1.13. Abuse of capabilities

Malicious sites could abuse the capabilities exposed by WebGPU to run computations that don't benefit the user or their experience and instead only benefit the site. Examples would be hidden crypto-mining, password cracking or rainbow tables computations.

It is not possible to guard against these types of uses of the API because the browser is not able to distinguish between valid workloads and abusive workloads. This is a general problem with all general-purpose computation capabilities on the Web: JavaScript, WebAssembly or WebGL. WebGPU only makes some workloads easier to implement, or slightly more efficient to run than using WebGL.

To mitigate this form of abuse, browsers can throttle operations on background tabs, could warn that a tab is using a lot of resource, and restrict which contexts are allowed to use WebGPU.

User agents can heuristically issue warnings to users about high power use, especially due to potentially malicious usage. If a user agent implements such a warning, it should include WebGPU usage in its heuristics, in addition to JavaScript, WebAssembly, WebGL, and so on.

2.2. Privacy Considerations

The privacy considerations for WebGPU are similar to those of WebGL. GPU APIs are complex and must expose various aspects of a device's capabilities out of necessity in order to enable developers to take advantage of those capabilities effectively. The general mitigation approach involves normalizing or binning potentially identifying information and enforcing uniform behavior where possible.

A user agent must not reveal more than 32 distinguishable configurations or buckets.

2.2.1. Machine-specific features and limits

WebGPU can expose a lot of detail on the underlying GPU architecture and the device geometry. This includes available physical adapters, many limits on the GPU and CPU resources that could be used (such as the maximum texture size), and any optional hardware-specific capabilities that are available.

User agents are not obligated to expose the real hardware limits, they are in full control of how much the machine specifics are exposed. One strategy to reduce fingerprinting is binning all the target platforms into a few number of bins. In general, the privacy impact of exposing the hardware limits matches the one of WebGL.

The [default](#) limits are also deliberately high enough to allow most applications to work without requesting higher limits. All the usage of the API is validated according to the requested limits, so the actual hardware capabilities are not exposed to the users by accident.

2.2.2. Machine-specific artifacts

There are some machine-specific rasterization/precision artifacts and performance differences that can be observed roughly in the same way as in WebGL. This applies to rasterization coverage and patterns, interpolation precision of the varyings between shader stages, compute unit scheduling, and more aspects of execution.

Generally, rasterization and precision fingerprints are identical across most or all of the devices of each vendor. Performance differences are relatively intractable, but also relatively low-signal (as with JS execution performance).

Privacy-critical applications and user agents should utilize software implementations to eliminate such artifacts.

2.2.3. Machine-specific performance

Another factor for differentiating users is measuring the performance of specific operations on the GPU. Even with low precision timing, repeated execution of an operation can show if the user's machine is fast at specific workloads. This is a fairly common vector (present in both WebGL and Javascript), but it's also low-signal and relatively intractable to truly normalize.

WebGPU compute pipelines expose access to GPU unobstructed by the fixed-function hardware. This poses an additional risk for unique device fingerprinting. User agents can take steps to dissociate logical GPU invocations with actual compute units to reduce this risk.

2.2.4. User Agent State

This specification doesn't define any additional user-agent state for an origin. However it is expected that user agents will have compilation caches for the result of expensive compilation like [GPUShaderModule](#), [GPURenderPipeline](#) and [GPUComputePipeline](#). These caches are important to improve the loading time of WebGPU applications after the first visit.

For the specification, these caches are indistinguishable from incredibly fast compilation, but for applications it would be easy to measure how long [createComputePipelineAsync\(\)](#) takes to resolve. This can leak information across origins (like "did the user access a site with this specific shader") so user agents should follow the best practices in [storage partitioning](#).

The system's GPU driver may also have its own cache of compiled shaders and pipelines. User agents may want to disable these when at all possible, or add per-partition data to shaders in ways that will make the GPU driver consider them different.

2.2.5. Driver bugs

In addition to the concerns outlined in [Security Considerations](#), driver bugs may introduce differences in behavior that can be observed as a method of differentiating users. The mitigations mentioned in Security Considerations apply here as well, including coordinating with GPU vendors and implementing workarounds for known issues in the user agent.

2.2.6. Adapter Identifiers

Past experience with WebGL has demonstrated that developers have a legitimate need to be able to identify the GPU their code is running on in order to create and maintain robust GPU-based content. For example, to identify adapters with known driver bugs in order to work around them or to avoid features that perform more poorly than expected on a given class of hardware.

But exposing adapter identifiers also naturally expands the amount of fingerprinting information available, so there's a desire to limit the precision with which we identify the adapter.

There are several mitigations that can be applied to strike a balance between enabling robust content and preserving privacy. First is that user agents can reduce the burden on developers by identifying and working around known driver issues, as they have since browsers began making use of GPUs.

When adapter identifiers are exposed by default they should be as broad as possible while still being useful. Possibly identifying, for example, the adapter's vendor and general architecture without identifying the specific adapter in use. Similarly, in some cases identifiers for an adapter that is considered a reasonable proxy for the actual adapter may be reported.

In cases where full and detailed information about the adapter is useful (for example: when filing bug reports) the user can be asked for consent to reveal additional information about their hardware to the page.

Finally, the user agent will always have the discretion to not report adapter identifiers at all if it considers it appropriate, such as in enhanced privacy modes.

3. Fundamentals

3.1. Conventions

3.1.1. Syntactic Shorthands

In this specification, the following syntactic shorthands are used:

The `.` ("dot") syntax, common in programming languages.

The phrasing "`Foo.Bar`" means "the `Bar` member of the value (or interface) `Foo`." If `Foo` is an [ordered map](#) and `Bar` does not [exist](#) in `Foo`, returns `undefined`.

The phrasing "`Foo.Bar` is [provided](#)" means "the `Bar` member [exists](#) in the [map](#) value `Foo`"

The `?.` ("optional chaining") syntax, adopted from JavaScript.

The phrasing "`Foo?.Bar`" means "if `Foo` is `null` or `undefined` or `Bar` does not [exist](#) in `Foo`, `undefined`; otherwise, `Foo.Bar`".

For example, where `buffer` is a [GPUBuffer](#), `buffer?.\[[device]].\[[adapter]]` means "if `buffer` is `null` or `undefined`, then `undefined`; otherwise, the `\[[adapter]]` internal slot of the `\[[device]]` internal slot of `buffer`."

The ?? ("nullish coalescing") syntax, adopted from JavaScript.

The phrasing "x ?? y" means "x, if x is not null or undefined, and y otherwise".

slot-backed attribute

A WebIDL attribute which is backed by an internal slot of the same name. It may or may not be mutable.

3.1.2. WebGPU Objects

A *WebGPU object* consists of a [WebGPU Interface](#) and an [internal object](#).

The *WebGPU interface* defines the public interface and state of the [WebGPU object](#). It can be used on the [content timeline](#) where it was created, where it is a JavaScript-exposed WebIDL interface.

Any interface which includes *GPUObjectBase* is a [WebGPU interface](#).

The *internal object* tracks the state of the [WebGPU object](#) on the [device timeline](#). All reads/writes to the mutable state of an [internal object](#) occur from steps executing on a single well-ordered [device timeline](#).

The following special property types can be defined on [WebGPU objects](#):

immutable property

A read-only slot set during initialization of the object. It can be accessed from any timeline.

Note: Since the slot is immutable, implementations may have a copy on multiple timelines, as needed. [Immutable properties](#) are defined in this way to avoid describing multiple copies in this spec.

If named `[[with brackets]]`, it is an internal slot.

If named `withoutBrackets`, it is a [readonly slot-backed attribute](#) of the [WebGPU interface](#).

content timeline property

A property which is only accessible from the [content timeline](#) where the object was created.

If named `[[with brackets]]`, it is an internal slot.

If named `withoutBrackets`, it is a [slot-backed attribute](#) of the [WebGPU interface](#).

device timeline property

A property which tracks state of the [internal object](#) and is only accessible from the [device timeline](#) where the object was created. [device timeline properties](#) may be mutable.

[Device timeline properties](#) are named `[[with brackets]]`, and are internal slots.

queue timeline property

A property which tracks state of the [internal object](#) and is only accessible from the [queue timeline](#) where the object was created. [queue timeline properties](#) may be mutable.

[Queue timeline properties](#) are named `[[with brackets]]`, and are internal slots.

```
interface mixin GPUObjectBase {  
  attribute USVString label;  
};
```

To create a new *WebGPU object* ([GPUObjectBase](#) *parent*, interface *T*, [GPUObjectDescriptorBase](#) *descriptor*) (where *T* extends [GPUObjectBase](#)), run the following [content timeline](#) steps:

Let *device* be *parent*.[\[\[device\]\]](#).

Let *object* be a new instance of *T*.

Set *object*.[\[\[device\]\]](#) to *device*.

Set *object*.[label](#) to *descriptor*.[label](#).

Return *object*.

[GPUObjectBase](#) has the following [immutable properties](#):

[\[\[device\]\]](#), of type [device](#), [readonly](#)

The [device](#) that owns the [internal object](#).

Operations on the contents of this object [assert](#) they are running on the [device timeline](#), and that the device is [valid](#).

[GPUObjectBase](#) has the following [content timeline properties](#):

[label](#), of type [USVString](#)

A developer-provided label which is used in an [implementation-defined](#) way. It can be used by the browser, OS, or other tools to help identify the underlying [internal](#)

[object](#) to the developer. Examples include displaying the label in [GPUError](#) messages, console warnings, browser developer tools, and platform debugging utilities.

NOTE:

Implementations **should** use labels to enhance error messages by using them to identify WebGPU objects.

However, this need not be the only way of identifying objects: implementations **should** also use other available information, especially when no label is available. For example:

- The label of the parent [GPUTexture](#) when printing a [GPUTextureView](#).
- The label of the parent [GPUCommandEncoder](#) when printing a [GPURenderPassEncoder](#) or [GPUComputePassEncoder](#).
- The label of the source [GPUCommandEncoder](#) when printing a [GPUCommandBuffer](#).
- The label of the source [GPURenderBundleEncoder](#) when printing a [GPURenderBundle](#).

NOTE:

The [label](#) is a property of the [GPUObjectBase](#). Two [GPUObjectBase](#) "wrapper" objects have completely separate label states, even if they refer to the same underlying object (for example returned by [getBindGroupLayout\(\)](#)). The [label](#) property will not change except by being set from JavaScript.

This means one underlying object could be associated with multiple labels. This specification does not define how the label is propagated to the [device timeline](#). How labels are used is completely [implementation-defined](#): error messages could show the most recently set label, all known labels, or no labels at all.

It is defined as a [USVString](#) because some user agents may supply it to the debug facilities of the underlying native APIs.

[GPUObjectBase](#) has the following [device timeline properties](#):

[\[\[valid\]\]](#), of type [boolean](#), initially `true`.

If `true`, indicates that the [internal object](#) is valid to use.

NOTE:

Ideally [WebGPU interfaces](#) should not prevent their parent objects, such as the [\[\[device\]\]](#) that owns them, from being garbage collected. This cannot be guaranteed, however, as holding a strong reference to a parent object may be required in some implementations.

As a result, developers should assume that a [WebGPU interface](#) may remain live until all child objects of that interface have also been garbage collected, causing some resources to remain allocated longer than anticipated.

Calling the `destroy` method on a [WebGPU interface](#) (such as [GPUDevice.destroy\(\)](#) or [GPUBuffer.destroy\(\)](#)) should be favored over relying on garbage collection if predictable release of allocated resources is needed.

3.1.3. Object Descriptors

An *object descriptor* holds the information needed to create an object, which is typically done via one of the `create*` methods of [GPUDevice](#).

dictionary

[GPUObjectDescriptorBase](#)

```
{  
  USVString label = "";  
};
```

[GPUObjectDescriptorBase](#) has the following members:

`label`, of type [USVString](#), defaulting to `""`

The initial value of [GPUObjectBase.label](#).

3.2. Asynchrony

3.2.1. Invalid Internal Objects & Contagious Invalidity

Object creation operations in WebGPU don't return promises, but nonetheless are internally asynchronous. Returned objects refer to [internal objects](#) which are manipulated on a [device timeline](#). Rather than fail with exceptions or rejections, most errors that occur on a [device timeline](#) are communicated through [GPUErrors](#) generated on the associated [device](#).

[Internal objects](#) are either [valid](#) or [invalid](#). An [invalid](#) object will never become [valid](#) at a later time, but some [valid](#) objects may be [invalidated](#).

Objects are [invalid](#) from creation if it wasn't possible to create them. This can happen, for example, if the [object descriptor](#) doesn't describe a valid object, or if there is not enough memory to allocate a resource. It can also happen if an object is created with or from another invalid object (for example calling [createView\(\)](#) on an invalid [GPUTexture](#)) (for example the [GPUTexture](#) of a [createView\(\)](#) call): this case is referred to as *contagious invalidity*.

[Internal objects](#) of *most* types cannot become [invalid](#) after they are created, but still may become unusable, e.g. if the owning device is [lost](#) or [destroyed](#), or the object has a special internal state, like buffer state `"destroyed"`.

[Internal objects](#) of some types can become [invalid](#) after they are created; specifically, [devices](#), [adapters](#), [GPUCommandBuffer](#)s, and command/pass/bundle encoders.

A given [GPUObjectBase](#) object is *valid* if `object.[[valid]]` is `true`.

A given [GPUObjectBase](#) object is *invalid* if `object.[[valid]]` is `false`.

A given [GPUObjectBase](#) object is *valid to use with a targetObject* if the all of the requirements in the following [device timeline](#) steps are met:

`object.[[valid]]` must be `true`.

`object.[[device]].[[valid]]` must be `true`.

`object.[[device]]` must equal `targetObject.[[device]]`.

To *invalidate* a [GPUObjectBase](#) object, run the following [device timeline](#) steps:

`object.[[valid]]` to `false`.

3.2.2. Promise Ordering

Several operations in WebGPU return promises.

[GPU.requestAdapter\(\)](#)

[GPUAdapter.requestDevice\(\)](#)

[GPUDevice.createComputePipelineAsync\(\)](#)

[GPUDevice.createRenderPipelineAsync\(\)](#)

[GPUShaderModule.getCompilationInfo\(\)](#)

[GPUQueue.onSubmittedWorkDone\(\)](#)

[GPUBuffer.mapAsync\(\)](#)

[GPUDevice.lost](#)

[GPUDevice.popErrorScope\(\)](#)

WebGPU does not make any guarantees about the order in which these promises settle (resolve or reject), except for the following:

For some [GPUQueue](#) *q*, if *p1* = *q.onSubmittedWorkDone()* is called before *p2* = *q.onSubmittedWorkDone()*, then *p1* must settle before *p2*.

For some [GPUQueue](#) *q* and [GPUBuffer](#) *b* on the same [GPUDevice](#), if *p1* = *b.mapAsync()* is called before *p2* = *q.onSubmittedWorkDone()*, then *p1* must settle before *p2*.

Applications must not rely on any other promise settlement ordering.

3.3. Coordinate Systems

Rendering operations use the following coordinate systems:

Normalized device coordinates (or NDC) have three dimensions, where:

$$-1.0 \leq x \leq 1.0$$

$$-1.0 \leq y \leq 1.0$$

$$0.0 \leq z \leq 1.0$$

The bottom-left corner is at (-1.0, -1.0, z).

Normalized device coordinates.

Note: Whether *z* = 0 or *z* = 1 is treated as the near plane is application specific. The above diagram presents *z* = 0 as the near plane but the observed behavior is determined by a combination of the projection matrices used by shaders, the [depthClearValue](#), and the [depthCompare](#) function.

Clip space coordinates have four dimensions: (x, y, z, w)

Clip space coordinates are used for the the [clip position](#) of a vertex (i.e. the [position](#) output of a vertex shader), and for the [clip volume](#).

[Normalized device coordinates](#) and clip space coordinates are related as follows: If point *p* = (*p.x*, *p.y*, *p.z*, *p.w*) is in the [clip volume](#), then its normalized device coordinates are (*p.x* ÷ *p.w*, *p.y* ÷ *p.w*, *p.z* ÷ *p.w*).

Framebuffer coordinates address the pixels in the [framebuffer](#)

They have two dimensions.

Each pixel extends 1 unit in x and y dimensions.

The top-left corner is at (0.0, 0.0).

x increases to the right.

y increases down.

See [§ 17 Render Passes](#) and [§ 23.2.5 Rasterization](#).

Framebuffer coordinates.

Viewport coordinates combine [framebuffer coordinates](#) in x and y dimensions, with depth in z.

Normally $0.0 \leq z \leq 1.0$, but this can be modified by setting `[[viewport]].minDepth` and `maxDepth` via [setViewport\(\)](#)

Fragment coordinates match [viewport coordinates](#).

Texture coordinates, sometimes called "UV coordinates" in 2D, are used to sample textures and have a number of components matching the [texture dimension](#).

$0 \leq u \leq 1.0$

$0 \leq v \leq 1.0$

$0 \leq w \leq 1.0$

(0.0, 0.0, 0.0) is in the first texel in texture memory address order.

(1.0, 1.0, 1.0) is in the last texel texture memory address order.

2D Texture coordinates.

Window coordinates, or *present coordinates*, match [framebuffer coordinates](#), and are used when interacting with an external display or conceptually similar interface.

Note: WebGPU's coordinate systems match DirectX's coordinate systems in a graphics pipeline.

3.4. Programming Model

3.4.1. Timelines

WebGPU's behavior is described in terms of "timelines". Each operation (defined as algorithms) occurs on a timeline. Timelines clearly define both the order of operations, and which state is available to which operations.

Note: This "timeline" model describes the constraints of the multi-process models of browser engines (typically with a "content process" and "GPU process"), as well as the GPU itself as a separate execution unit in many implementations. Implementing WebGPU does not require timelines to execute in parallel, so does not require multiple processes, or even multiple threads. (It does require concurrency for cases like [get a copy of the image contents of a context](#) which synchronously blocks on another timeline to complete.)

Content timeline

Associated with the execution of the Web script. It includes calling all methods described by this specification.

To issue steps to the content timeline from an operation on [GPUDevice](#) device, [queue a global task for GPUDevice](#) device with those steps.

Device timeline

Associated with the GPU device operations that are issued by the user agent. It includes creation of adapters, devices, and GPU resources and state objects, which are typically synchronous operations from the point of view of the user agent part that controls the GPU, but can live in a separate OS process.

Queue timeline

Associated with the execution of operations on the compute units of the GPU. It includes actual draw, copy, and compute jobs that run on the GPU.

Timeline-agnostic

Associated with any of the above timelines

Steps may be issued to any timeline if they only operate on [immutable properties](#) or arguments passed from the calling steps.

The following show the styling of steps and values associated with each timeline. This styling is non-normative; the specification text always describes the association.

Immutable value example term definition

Can be used on any timeline.

Content-timeline example term definition

Can only be used on the [content timeline](#).

Device-timeline example term definition

Can only be used on the [device timeline](#).

Queue-timeline example term definition

Can only be used on the [queue timeline](#).

Steps which are [timeline-agnostic](#) look like this.

[Immutable value example term](#) usage.

Steps executed on the [content timeline](#) look like this.

[Immutable value example term](#) usage. [Content-timeline example term](#) usage.

Steps executed on the [device timeline](#) look like this.

[Immutable value example term](#) usage. [Device-timeline example term](#) usage.

Steps executed on the [queue timeline](#) look like this.

[Immutable value example term](#) usage. [Queue-timeline example term](#) usage.

In this specification, asynchronous operations are used when the return value depends on work that happens on any timeline other than the [Content timeline](#). They are represented by promises and events in API.

[GPUComputePassEncoder.dispatchWorkgroups\(\)](#):

User encodes a `dispatchWorkgroups` command by calling a method of the [GPUComputePassEncoder](#) which happens on the [Content timeline](#).

User issues [GPUQueue.submit\(\)](#) that hands over the [GPUCommandBuffer](#) to the user agent, which processes it on the [Device timeline](#) by calling the OS driver to do a low-level submission.

The submit gets dispatched by the GPU invocation scheduler onto the actual compute units for execution, which happens on the [Queue timeline](#).

[GPUDevice.createBuffer\(\)](#):

User fills out a [GPUBufferDescriptor](#) and creates a [GPUBuffer](#) with it, which happens on the [Content timeline](#).

User agent creates a low-level buffer on the [Device timeline](#).

[GPUBuffer.mapAsync\(\)](#):

User requests to map a [GPUBuffer](#) on the [Content timeline](#) and gets a promise in return.

User agent checks if the buffer is currently used by the GPU and makes a reminder to itself to check back when this usage is over.

After the GPU operating on [Queue timeline](#) is done using the buffer, the user agent maps it to memory and [resolves](#) the promise.

3.4.2. Memory Model

This section is non-normative.

Once a [GPUDevice](#) has been obtained during an application initialization routine, we can describe the *WebGPU platform* as consisting of the following layers:

User agent implementing the specification.

Operating system with low-level native API drivers for this device.

Actual CPU and GPU hardware.

Each layer of the [WebGPU platform](#) may have different memory types that the user agent needs to consider when implementing the specification:

The script-owned memory, such as an [ArrayBuffer](#) created by the script, is generally not accessible by a GPU driver.

A user agent may have different processes responsible for running the content and communication to the GPU driver. In this case, it uses inter-process shared memory to transfer data.

Dedicated GPUs have their own memory with high bandwidth, while integrated GPUs typically share memory with the system.

Most [physical resources](#) are allocated in the memory of type that is efficient for computation or rendering by the GPU. When the user needs to provide new data to the GPU, the data may first need to cross the process boundary in order to reach the user agent part that communicates with the GPU driver. Then it may need to be made visible to the driver, which sometimes requires a copy into driver-allocated staging memory. Finally, it may need to be transferred to the dedicated GPU memory, potentially changing the internal layout into one that is most efficient for GPUs to operate on.

All of these transitions are done by the WebGPU implementation of the user agent.

Note: This example describes the worst case, while in practice the implementation might not need to cross the process boundary, or may be able to expose the driver-managed memory directly to the user behind an [ArrayBuffer](#), thus avoiding any data copies.

3.4.3. Resource Usages

A [physical resource](#) can be used with an *internal usage* by a [GPU command](#):

input

Buffer with input data for draw or dispatch calls. Preserves the contents. Allowed by buffer [INDEX](#), buffer [VERTEX](#), or buffer [INDIRECT](#).

constant

Resource bindings that are constant from the shader point of view. Preserves the contents. Allowed by buffer [UNIFORM](#) or texture [TEXTURE_BINDING](#).

storage

Read/write storage resource binding. Allowed by buffer [STORAGE](#) or texture [STORAGE_BINDING](#).

storage-read

Read-only storage resource bindings. Preserves the contents. Allowed by buffer [STORAGE](#) or texture [STORAGE_BINDING](#).

attachment

Texture used as a read/write output attachment or write-only resolve target in a render pass. Allowed by texture [RENDER_ATTACHMENT](#).

attachment-read

Texture used as a read-only attachment in a render pass. Preserves the contents. Allowed by texture [RENDER_ATTACHMENT](#).

We define *subresource* to be either a whole buffer, or a [texture subresource](#).

Some [internal usages](#) are compatible with others. A [subresource](#) can be in a state that combines multiple usages together. We consider a list *U* to be a *compatible usage list* if (and only if) it satisfies any of the following rules:

Each usage in *U* is [input](#), [constant](#), [storage-read](#), or [attachment-read](#).

Each usage in *U* is [storage](#).

Multiple such usages are allowed even though they are writable. This is the [usage scope storage exception](#).

Each usage in *U* is [attachment](#).

Multiple such usages are allowed even though they are writable. This is the [usage scope attachment exception](#).

Enforcing that the usages are only combined into a [compatible usage list](#) allows the API to limit when data races can occur in working with memory. That property makes applications written against WebGPU more likely to run without modification on different platforms.

EXAMPLE:

Binding the same buffer for [storage](#) as well as for [input](#) within the same [GPURenderPassEncoder](#) results in a non-[compatible usage list](#) for that buffer.

EXAMPLE:

These rules allow for *read-only depth-stencil*: a single depth/stencil texture can be used as two different read-only usages in a render pass simultaneously:

[attachment-read](#)

As a depth/stencil attachment with all aspects marked read-only (using [depthReadOnly](#) and/or [stencilReadOnly](#) as necessary).

[constant](#)

As a texture binding to a draw call.

EXAMPLE:

The *usage scope storage exception* allows two cases that would not be allowed otherwise:

A buffer or texture may be bound as [storage](#) to two different draw calls in a render pass.

Disjoint ranges of a single buffer may be bound to two different binding points as [storage](#).

Overlapping ranges must not be bound to a single dispatch/draw call; this is checked by "[Encoder bind groups alias a writable resource](#)".

EXAMPLE:

The *usage scope attachment exception* allows a texture subresource to be used as [attachment](#) more than once. This is necessary to allow disjoint slices of 3D textures to be bound as different attachments to a single render pass.

One slice must not be bound twice for two different attachments; this is checked by [beginRenderPass\(\)](#).

3.4.4. Synchronization

A *usage scope* is a [map](#) from [subresource](#) to [list<internal usage>>](#). Each usage scope covers a range of operations which may execute in a concurrent fashion with each other, and therefore may only use [subresources](#) in consistent [compatible usage lists](#) within the scope.

A [usage scope](#) *scope* passes *usage scope validation* if, for each [[subresource](#), *usageList*] in *scope*, *usageList* is a [compatible usage list](#).

To add a [subresource](#) *subresource* to [usage scope](#) *usageScope* with usage ([internal usage](#) or set of [internal usages](#)) *usage*:

If *usageScope*[*subresource*] does not [exist](#), set it to `[]`.

[Append](#) *usage* to *usageScope*[*subresource*].

To merge [usage scope](#) *A* into [usage scope](#) *B*:

For each [*subresource*, *usage*] in *A*:

[Add](#) *subresource* to *B* with usage *usage*.

[Usage scopes](#) are constructed and validated during encoding:

in [dispatchWorkgroups\(\)](#)

in [dispatchWorkgroupsIndirect\(\)](#)

at [GPURenderPassEncoder.end\(\)](#)

at [GPURenderBundleEncoder.finish\(\)](#)

The [usage scopes](#) are as follows:

In a compute pass, each dispatch command ([dispatchWorkgroups\(\)](#) or [dispatchWorkgroupsIndirect\(\)](#)) is one usage scope.

A subresource is used in the usage scope if it is potentially accessible by the dispatched invocations, including:

All [subresources](#) referenced by bind groups in slots used by the current [GPUComputePipeline](#)'s [\[\[layout\]\]](#)

Buffers used directly by dispatch calls (such as indirect buffers)

Note: State-setting compute pass commands, like [setBindGroup\(\)](#), do not contribute their bound resources directly to a usage scope: they only change the state that is checked in dispatch commands.

One render pass is one usage scope.

A subresource is used in the usage scope if it's referenced by any command, including state-setting commands (unlike in compute passes), including:

Buffers set by [setVertexBuffer\(\)](#)

Buffers set by [setIndexBuffer\(\)](#)

All [subresources](#) referenced by bind groups set by [setBindGroup\(\)](#)

Buffers used directly by draw calls (such as indirect buffers)

Note: Copy commands are standalone operations and don't use [usage scopes](#) for validation. They implement their own validation to prevent self-races.

EXAMPLE:

The following example resource usages *are* included in [usage scopes](#):

In a render pass, subresources used in any [setBindGroup\(\)](#) call, regardless of whether the currently bound pipeline's shader or layout actually depends on these bindings, or the bind group is shadowed by another 'set' call.

A buffer used in any [setVertexBuffer\(\)](#) call, regardless of whether any draw call depends on this buffer, or whether this buffer is shadowed by another 'set' call.

A buffer used in any [setIndexBuffer\(\)](#) call, regardless of whether any draw call depends on this buffer, or whether this buffer is shadowed by another 'set' call.

A texture subresource used as a color attachment, resolve attachment, or depth/stencil attachment in [GPURenderPassDescriptor](#) by [beginRenderPass\(\)](#), regardless of whether the shader actually depends on these attachments.

Resources used in bind group entries with visibility 0, or visible only to the compute stage but used in a render pass (or vice versa).

3.5. Core Internal Objects

3.5.1. Adapters

An *adapter* identifies an implementation of WebGPU on the system: both an instance of compute/rendering functionality on the platform underlying a browser, and an instance of a browser's implementation of WebGPU on top of that functionality.

[Adapters](#) are exposed via [GPUAdapter](#).

[Adapters](#) do not uniquely represent underlying implementations: calling [requestAdapter\(\)](#) multiple times returns a different [adapter](#) object each time.

Each [adapter](#) object can only be used to create one [device](#): upon a successful [requestDevice\(\)](#) call, the adapter's [\[\[state\]\]](#) changes to ["consumed"](#).

Additionally, [adapter](#) objects may [expire](#) at any time.

Note: This ensures applications use the latest system state for adapter selection when creating a device. It also encourages robustness to more scenarios by making them look similar: first initialization, reinitialization due to an unplugged adapter, reinitialization due to a test [GPUDevice.destroy\(\)](#) call, etc.

An [adapter](#) may be considered a *fallback adapter* if it has significant performance caveats in exchange for some combination of wider compatibility, more predictable behavior, or improved privacy. It is not required that a [fallback adapter](#) is available on every system.

[adapter](#) has the following [immutable properties](#):

[\[\[features\]\]](#), of type [ordered set<GPUFeatureName>](#), readonly

The [features](#) which can be used to create devices on this adapter.

[\[\[limits\]\]](#), of type [supported limits](#), readonly

The [best](#) limits which can be used to create devices on this adapter.

Each adapter limit must be the same or [better](#) than its default value in [supported limits](#).

[\[\[fallback\]\]](#), of type [boolean](#), readonly

If set to `true` indicates that the adapter is a [fallback adapter](#).

`[[xrCompatible]]`, of type boolean

If set to `true` indicates that the adapter was requested with compatibility with [WebXR sessions](#).

[adapter](#) has the following [device timeline properties](#):

`[[state]]`, initially ["valid"](#)

`"valid"`

The adapter can be used to create a device.

`"consumed"`

The adapter has already been used to create a device, and cannot be used again.

`"expired"`

The adapter has expired for some other reason.

To expire a [GPUAdapter](#) adapter, run the following [device timeline](#) steps:

Set `adapter.[[adapter]].[[state]]` to ["expired"](#).

3.5.2. Devices

A *device* is the logical instantiation of an [adapter](#), through which [internal objects](#) are created.

[Devices](#) are exposed via [GPUDevice](#).

A [device](#) is the exclusive owner of all [internal objects](#) created from it: when the [device](#) becomes [invalid](#) (is [lost](#) or [destroyed](#)), it and all objects created on it (directly, e.g. [createTexture\(\)](#), or indirectly, e.g. [createView\(\)](#)) become implicitly [unusable](#).

[device](#) has the following [immutable properties](#):

`[[adapter]]`, of type [adapter](#), readonly

The [adapter](#) from which this device was created.

`[[features]]`, of type [ordered set](#)<[GPUFeatureName](#)>, readonly

The [features](#) which can be used on this device, as computed [at creation](#). No additional features can be used, even if the underlying [adapter](#) can support them.

`[[limits]]`, of type [supported limits](#), readonly

The limits which can be used on this device, as computed [at creation](#). No [better](#) limits can be used, even if the underlying [adapter](#) can support them.

[device](#) has the following [content timeline properties](#):

`[[content device]]`, of type [GPUDevice](#), readonly

The [Content timeline GPUDevice](#) interface which this device is associated with.

To create a new device from [adapter](#) adapter with [GPUDeviceDescriptor](#) descriptor, run the following [device timeline](#) steps:

Let *features* be the [set](#) of values in `descriptor.requiredFeatures`.

If *features* contains ["texture-formats-tier2"](#):

Append ["texture-formats-tier1"](#) to *features*

If *features* contains ["texture-formats-tier1"](#):

Append ["rg11b10float-renderable"](#) to *features*

Append ["core-features-and-limits"](#) to *features*.

Let *limits* be a [supported limits](#) object with all values set to their defaults.

For each (key, value) pair in `descriptor.requiredLimits`:

If *value* is not `undefined` and *value* is [better](#) than `limits[key]`:

Set `limits[key]` to *value*.

Let *device* be a [device](#) object.

Set `device.[[adapter]]` to *adapter*.

Set `device.[[features]]` to *features*.

Set `device.[[limits]]` to *limits*.

Return *device*.

Any time the user agent needs to revoke access to a device, it calls [lose the device](#)(device, ["unknown"](#)) on the device's [device timeline](#), potentially ahead of other operations currently queued on that timeline.

If an operation fails with side effects that would observably change the state of objects on the device or potentially corrupt internal implementation/driver state, the device **should** be lost to prevent these changes from being observable.

Note: For all device losses not initiated by the application (via [destroy\(\)](#)), user agents should consider issuing developer-visible warnings *unconditionally*, even if the [lost](#) promise is handled. These scenarios should be rare, and the signal is vital to developers because most of the WebGPU API tries to behave like nothing is wrong to avoid interrupting the runtime flow of the application: no validation errors are raised, most promises resolve normally, etc.

To *lose the device*(device, reason) run the following [device timeline](#) steps:

[Invalidate](#) device.

Issue the following steps on the [content timeline](#) of device.[\[\[content_device\]\]](#):

Resolve device.[lost](#) with a new [GPUDeviceLostInfo](#) with [reason](#) set to reason and [message](#) set to an [implementation-defined](#) value.

Note: [message](#) should not disclose unnecessary user/system information and should never be parsed by applications.

Complete any outstanding steps that are waiting until device becomes lost.

Note: No errors are generated from a device which is lost. See [§ 2.2 Errors & Debugging](#).

To *listen for timeline event* event on [device](#) device, handled by steps on timeline timeline:

If or when the [device timeline](#) has been informed of the completion of event, or

If device is [lost](#) already, or when it [becomes lost](#):

Then issue steps on timeline.

3.6. Optional Capabilities

WebGPU [adapters](#) and [devices](#) have *capabilities*, which describe WebGPU functionality that differs between different implementations, typically due to hardware or system software constraints. A [capability](#) is either a [feature](#) or a [limit](#).

A user agent must not reveal more than 32 distinguishable configurations or buckets.

The capabilities of an [adapter](#) must conform to [§ 4.2.1 Adapter Capability Guarantees](#).

Only supported capabilities may be requested in [requestDevice\(\)](#); requesting unsupported capabilities results in failure.

The capabilities of a [device](#) are determined in "a new device" by starting with the adapter's defaults (no features and the default [supported limits](#)) and adding capabilities as requested in [requestDevice\(\)](#). These capabilities are enforced regardless of the capabilities of the [adapter](#).

For privacy considerations, see [§ 2.2.1 Machine-specific features and limits](#).

3.6.1. Features

A *feature* is a set of optional WebGPU functionality that is not supported on all implementations, typically due to hardware or system software constraints.

All [features](#) are optional, but [adapters](#) make some guarantees about their availability (see [§ 4.2.1 Adapter Capability Guarantees](#)).

A [device](#) supports the exact set of features determined at creation (see [§ 3.6 Optional Capabilities](#)). API calls perform validation according to these features (not the [adapter](#)'s features):

Using existing API surfaces in a new way **typically** results in a [validation error](#).

There are several types of *optional API surface*:

Using a new method or enum value always throws a [TypeError](#).

Using a new dictionary member with a (correctly-typed) non-default value **typically** results in a [validation error](#).

Using a new WGSL [enable](#) directive always results in a [createShaderModule\(\). validation error](#).

A [GPUFeatureName](#) feature is enabled for a [GPUObjectBase](#) object if and only if object.[\[\[device\]\].\[\[features\]\]](#) contains feature.

See the [Feature Index](#) for a description of the functionality each feature enables.

Note: Even where supported, enabling features is not necessarily desirable, as doing so may have a performance impact. Because of this, and to improve portability across devices and implementations, applications should generally only request features that they may actually require.

3.6.2. Limits

Each *limit* is a numeric limit on the usage of WebGPU on a device.

Note: Even where supported, setting "better" limits is not necessarily desirable, as doing so may have a performance impact. Because of this, and to improve portability across devices and implementations, applications should generally only request limits better than the defaults if they may actually require them.

Each limit has a *default* value.

[Adapters](#) are always guaranteed to support the defaults or [better](#) (see [§ 4.2.1 Adapter Capability Guarantees](#)).

A [device](#) supports the exact set of limits determined at creation (see [§ 3.6 Optional Capabilities](#)). API calls perform validation according to these limits (not the [adapter's](#) limits), no [better](#) or worse.

For any given limit, some values are *better* than others. A [better](#) limit value always relaxes validation, enabling strictly more programs to be valid. For each [limit class](#), "better" is defined.

Different limits have different *limit classes*:

maximum

The limit enforces a maximum on some value passed into the API.

Higher values are [better](#).

May only be set to values \geq the [default](#). Lower values are clamped to the [default](#).

alignment

The limit enforces a minimum alignment on some value passed into the API; that is, the value must be a multiple of the limit.

Lower values are [better](#).

May only be set to powers of 2 which are \leq the [default](#). Values which are not powers of 2 are invalid. Higher powers of 2 are clamped to the [default](#).

A *supported limits* object has a value for every limit defined by WebGPU:

Limit name	Type	Limit class	Default
<i>maxTextureDimension1D</i>	GPUSize32	maximum	8192
The maximum allowed value for the size.width of a texture created with dimension "1d" .			
<i>maxTextureDimension2D</i>	GPUSize32	maximum	8192
The maximum allowed value for the size.width and size.height of a texture created with dimension "2d" .			
<i>maxTextureDimension3D</i>	GPUSize32	maximum	2048
The maximum allowed value for the size.width , size.height and size.depthOrArrayLayers of a texture created with dimension "3d" .			
<i>maxTextureArrayLayers</i>	GPUSize32	maximum	256
The maximum allowed value for the size.depthOrArrayLayers of a texture created with dimension "2d" .			
<i>maxBindGroups</i>	GPUSize32	maximum	4
The maximum number of GPUBindGroupLayouts allowed in bindGroupLayouts when creating a GPUPipelineLayout .			
<i>maxBindGroupsPlusVertexBuffers</i>	GPUSize32	maximum	24
The maximum number of bind group and vertex buffer slots used simultaneously, counting any empty slots below the highest index. Validated in createRenderPipeline() and in draw calls .			
<i>maxBindingsPerBindGroup</i>	GPUSize32	maximum	1000
The number of binding indices available when creating a GPUBindGroupLayout . Note: This limit is normative, but arbitrary. With the default binding slot limits , it is impossible to use 1000 bindings in one bind group, but this allows GPUBindGroupLayoutEntry.binding values up to 999. This limit allows implementations to treat binding space as an array, within reasonable memory space, rather than a sparse map structure.			
<i>maxDynamicUniformBuffersPerPipelineLayout</i>	GPUSize32	maximum	8
The maximum number of GPUBindGroupLayoutEntry entries across a GPUPipelineLayout which are uniform buffers with dynamic offsets. See Exceeds the binding slot limits .			
<i>maxDynamicStorageBuffersPerPipelineLayout</i>	GPUSize32	maximum	4
The maximum number of GPUBindGroupLayoutEntry entries across a GPUPipelineLayout which are storage buffers with dynamic offsets. See Exceeds the binding slot limits .			
<i>maxSampledTexturesPerShaderStage</i>	GPUSize32	maximum	16
For each possible GPUShaderStage stage, the maximum number of GPUBindGroupLayoutEntry entries across a GPUPipelineLayout which are sampled textures. See Exceeds the binding slot limits .			
<i>maxSamplersPerShaderStage</i>	GPUSize32	maximum	16
For each possible GPUShaderStage stage, the maximum number of GPUBindGroupLayoutEntry entries across a GPUPipelineLayout which are samplers. See Exceeds the binding slot limits .			
<i>maxStorageBuffersPerShaderStage</i>	GPUSize32	maximum	8
For each possible GPUShaderStage stage, the maximum number of GPUBindGroupLayoutEntry entries across a GPUPipelineLayout which are storage buffers. See Exceeds the binding slot limits .			

Limit name	Type	Limit class	Default
<code>maxStorageTexturesPerShaderStage</code>	GPUSize32	maximum	4
For each possible GPUShaderStage stage, the maximum number of GPUBindGroupLayoutEntry entries across a GPUPipelineLayout which are storage textures. See Exceeds the binding slot limits .			
<code>maxUniformBuffersPerShaderStage</code>	GPUSize32	maximum	12
For each possible GPUShaderStage stage, the maximum number of GPUBindGroupLayoutEntry entries across a GPUPipelineLayout which are uniform buffers. See Exceeds the binding slot limits .			
<code>maxUniformBufferBindingSize</code>	GPUSize64	maximum	65536 bytes
The maximum GPUBufferBinding.size for bindings with a GPUBindGroupLayoutEntry entry for which <code>entry.buffer?.type</code> is <code>"uniform"</code> .			
<code>maxStorageBufferBindingSize</code>	GPUSize64	maximum	134217728 bytes (128 MiB)
The maximum GPUBufferBinding.size for bindings with a GPUBindGroupLayoutEntry entry for which <code>entry.buffer?.type</code> is <code>"storage"</code> or <code>"read-only-storage"</code> .			
<code>minUniformBufferOffsetAlignment</code>	GPUSize32	alignment	256 bytes
The required alignment for GPUBufferBinding.offset and the dynamic offsets provided in setBindGroup() , for bindings with a GPUBindGroupLayoutEntry entry for which <code>entry.buffer?.type</code> is <code>"uniform"</code> .			
<code>minStorageBufferOffsetAlignment</code>	GPUSize32	alignment	256 bytes
The required alignment for GPUBufferBinding.offset and the dynamic offsets provided in setBindGroup() , for bindings with a GPUBindGroupLayoutEntry entry for which <code>entry.buffer?.type</code> is <code>"storage"</code> or <code>"read-only-storage"</code> .			
<code>maxVertexBuffers</code>	GPUSize32	maximum	8
The maximum number of buffers when creating a GPURenderPipeline .			
<code>maxBufferSize</code>	GPUSize64	maximum	268435456 bytes (256 MiB)
The maximum size of size when creating a GPUBuffer .			
<code>maxVertexAttributes</code>	GPUSize32	maximum	16
The maximum number of attributes in total across buffers when creating a GPURenderPipeline .			
<code>maxVertexBufferArrayStride</code>	GPUSize32	maximum	2048 bytes
The maximum allowed arrayStride when creating a GPURenderPipeline .			
<code>maxInterStageShaderVariables</code>	GPUSize32	maximum	16
The maximum allowed number of input or output variables for inter-stage communication (like vertex outputs or fragment inputs).			
<code>maxColorAttachments</code>	GPUSize32	maximum	8
The maximum allowed number of color attachments in GPURenderPipelineDescriptor.fragment.targets , GPURenderPassDescriptor.colorAttachments , and GPURenderPassLayout.colorFormats .			
<code>maxColorAttachmentBytesPerSample</code>	GPUSize32	maximum	32
The maximum number of bytes necessary to hold one sample (pixel or subpixel) of render pipeline output data, across all color attachments.			
<code>maxComputeWorkgroupStorageSize</code>	GPUSize32	maximum	16384 bytes
The maximum number of bytes of workgroup storage used for a compute stage GPUShaderModule entry-point.			
<code>maxComputeInvocationsPerWorkgroup</code>	GPUSize32	maximum	256
The maximum value of the product of the <code>workgroup_size</code> dimensions for a compute stage GPUShaderModule entry-point.			
<code>maxComputeWorkgroupSizeX</code>	GPUSize32	maximum	256
The maximum value of the <code>workgroup_size</code> X dimension for a compute stage GPUShaderModule entry-point.			
<code>maxComputeWorkgroupSizeY</code>	GPUSize32	maximum	256
The maximum value of the <code>workgroup_size</code> Y dimensions for a compute stage GPUShaderModule entry-point.			
<code>maxComputeWorkgroupSizeZ</code>	GPUSize32	maximum	64
The maximum value of the <code>workgroup_size</code> Z dimensions for a compute stage GPUShaderModule entry-point.			
<code>maxComputeWorkgroupsPerDimension</code>	GPUSize32	maximum	65535
The maximum value for the arguments of dispatchWorkgroups(workgroupCountX, workgroupCountY, workgroupCountZ) .			

3.6.2.1. GPUSupportedLimits

[GPUSupportedLimits](#) exposes an adapter or device's [supported limits](#). See [GPUAdapter.limits](#) and [GPUDevice.limits](#).

[Exposed=(Window, Worker), [SecureContext](#)]

interface [GPUSupportedLimits](#) {

 readonly attribute [unsigned long](#)

`maxTextureDimension1D`

;

```
    readonly attribute unsigned long
maxTextureDimension2D
;
    readonly attribute unsigned long
maxTextureDimension3D
;
    readonly attribute unsigned long
maxTextureArrayLayers
;
    readonly attribute unsigned long
maxBindGroups
;
    readonly attribute unsigned long
maxBindGroupsPlusVertexBuffers
;
    readonly attribute unsigned long
maxBindingsPerBindGroup
;
    readonly attribute unsigned long
maxDynamicUniformBuffersPerPipelineLayout
;
    readonly attribute unsigned long
maxDynamicStorageBuffersPerPipelineLayout
;
    readonly attribute unsigned long
maxSampledTexturesPerShaderStage
;
    readonly attribute unsigned long
maxSamplersPerShaderStage
;
    readonly attribute unsigned long
maxStorageBuffersPerShaderStage
;
    readonly attribute unsigned long
maxStorageTexturesPerShaderStage
;
    readonly attribute unsigned long
maxUniformBuffersPerShaderStage
;
    readonly attribute unsigned long long
maxUniformBufferBindingSize
;
    readonly attribute unsigned long long
maxStorageBufferBindingSize
;
    readonly attribute unsigned long
minUniformBufferOffsetAlignment
;
    readonly attribute unsigned long
minStorageBufferOffsetAlignment
;
    readonly attribute unsigned long
maxVertexBuffers
```



```

;
    readonly attribute unsigned long long
maxBufferSize

;
    readonly attribute unsigned long
maxVertexAttributes

;
    readonly attribute unsigned long
maxVertexBufferArrayStride

;
    readonly attribute unsigned long
maxInterStageShaderVariables

;
    readonly attribute unsigned long
maxColorAttachments

;
    readonly attribute unsigned long
maxColorAttachmentBytesPerSample

;
    readonly attribute unsigned long
maxComputeWorkgroupStorageSize

;
    readonly attribute unsigned long
maxComputeInvocationsPerWorkgroup

;
    readonly attribute unsigned long
maxComputeWorkgroupSizeX

;
    readonly attribute unsigned long
maxComputeWorkgroupSizeY

;
    readonly attribute unsigned long
maxComputeWorkgroupSizeZ

;
    readonly attribute unsigned long
maxComputeWorkgroupsPerDimension

;
};

```

3.6.2.2. GPUSupportedFeatures

[GPUSupportedFeatures](#) is a [setlike](#) interface. Its [set entries](#) are the [GPUFeatureName](#) values of the [features](#) supported by an adapter or device. It must only contain strings from the [GPUFeatureName](#) enum.

[[Exposed](#)=(Window, Worker), [SecureContext](#)]

```

interface GPUSupportedFeatures {
    readonly setlike<DOMString>;
};

```

NOTE:

The type of the [GPUSupportedFeatures](#) [set entries](#) is [DOMString](#) to allow user agents to gracefully handle valid [GPUFeatureName](#)s which are added in later revisions of the spec but which the user agent has not been updated to recognize yet. If the [set entries](#) type was [GPUFeatureName](#) the following code would throw an [TypeError](#) rather than reporting `false`:

Check for support of an unrecognized feature:

```

if (adapter.features.has('unknown-feature')) {
    // Use unknown-feature

```

```

} else {
    console.warn('unknown-feature is not supported by this adapter.');
```

3.6.2.3. WGSLLanguageFeatures

[WGSLLanguageFeatures](#) is the [setlike](#) interface of `navigator.gpu.wgslLanguageFeatures`. Its [set entries](#) are the string names of the WGSL [language extensions](#) supported by the implementation (regardless of the adapter or device).

```

[Exposed=(Window, Worker), SecureContext]
interface WGSLLanguageFeatures {
    readonly setlike<DOMString>;
};
```

3.6.2.4. GPUAdapterInfo

[GPUAdapterInfo](#) exposes various identifying information about an adapter.

None of the members in [GPUAdapterInfo](#) are guaranteed to be populated with any particular value; if no value is provided, the attribute will return the empty string `""`. It is at the user agent's discretion which values to reveal, and it is likely that on some devices none of the values will be populated. As such, applications **must** be able to handle any possible [GPUAdapterInfo](#) values, including the absence of those values.

The [GPUAdapterInfo](#) for an adapter is exposed via [GPUAdapter.info](#) and [GPUDevice.adapterInfo](#). This info is immutable: for a given adapter, each [GPUAdapterInfo](#) attribute will return the same value every time it's accessed.

Note: Though the [GPUAdapterInfo](#) attributes are immutable *once accessed*, an implementation may delay the decision on what to expose for each attribute until the first time it is accessed.

Note: Other [GPUAdapter](#) instances, even if they represent the same physical adapter, may expose different values in [GPUAdapterInfo](#). However, they **should** expose the same values unless a specific event has increased the amount of identifying information the page is allowed to access. (No such events are defined by this specification.)

For privacy considerations, see [§ 2.2.6 Adapter Identifiers](#).

```

[Exposed=(Window, Worker), SecureContext]
interface GPUAdapterInfo {
    readonly attribute DOMString vendor;
    readonly attribute DOMString architecture;
    readonly attribute DOMString device;
    readonly attribute DOMString description;
    readonly attribute unsigned long subgroupMinSize;
    readonly attribute unsigned long subgroupMaxSize;
    readonly attribute boolean isFallbackAdapter;
};
```

[GPUAdapterInfo](#) has the following attributes:

vendor, of type [DOMString](#), readonly

The name of the vendor of the [adapter](#), if available. Empty string otherwise.

architecture, of type [DOMString](#), readonly

The name of the family or class of GPUs the [adapter](#) belongs to, if available. Empty string otherwise.

device, of type [DOMString](#), readonly

A vendor-specific identifier for the [adapter](#), if available. Empty string otherwise.

Note: This is a value that represents the type of adapter. For example, it may be a [PCI device ID](#). It does not uniquely identify a given piece of hardware like a serial number.

description, of type [DOMString](#), readonly

A human readable string describing the [adapter](#) as reported by the driver, if available. Empty string otherwise.

Note: Because no formatting is applied to [description](#) attempting to parse this value is not recommended. Applications which change their behavior based on the [GPUAdapterInfo](#), such as applying workarounds for known driver issues, should rely on the other fields when possible.

subgroupMinSize, of type [unsigned long](#), readonly

If the ["subgroups"](#) feature is supported, the minimum supported [subgroup size](#) for the [adapter](#).

subgroupMaxSize, of type [unsigned long](#), readonly

If the ["subgroups"](#) feature is supported, the maximum supported [subgroup size](#) for the [adapter](#).

`isFallbackAdapter`, of type [boolean](#), readonly

Whether the [adapter](#) is a [fallback adapter](#).

To create a new *adapter info* for a given [adapter](#), run the following [content timeline](#) steps:

Let `adapterInfo` be a new [GPUAdapterInfo](#).

If the vendor is known, set `adapterInfo.vendor` to the name of *adapter*'s vendor as a [normalized identifier string](#). To preserve privacy, the user agent may instead set `adapterInfo.vendor` to the empty string or a reasonable approximation of the vendor as a [normalized identifier string](#).

If the architecture is known, set `adapterInfo.architecture` to a [normalized identifier string](#) representing the family or class of adapters to which *adapter* belongs. To preserve privacy, the user agent may instead set `adapterInfo.architecture` to the empty string or a reasonable approximation of the architecture as a [normalized identifier string](#).

If the device is known, set `adapterInfo.device` to a [normalized identifier string](#) representing a vendor-specific identifier for *adapter*. To preserve privacy, the user agent may instead set `adapterInfo.device` to the empty string or a reasonable approximation of a vendor-specific identifier as a [normalized identifier string](#).

If a description is known, set `adapterInfo.description` to a description of the *adapter* as reported by the driver. To preserve privacy, the user agent may instead set `adapterInfo.description` to the empty string or a reasonable approximation of a description.

If ["subgroups"](#) is supported, set `subgroupMinSize` to the smallest supported subgroup size. Otherwise, set this value to 4.

Note: To preserve privacy, the user agent may choose to not support some features or provide values for the property which do not distinguish different devices, but are still usable (e.g. use the default value of 4 for all devices).

If ["subgroups"](#) is supported, set `subgroupMaxSize` to the largest supported subgroup size. Otherwise, set this value to 128.

Note: To preserve privacy, the user agent may choose to not support some features or provide values for the property which do not distinguish different devices, but are still usable (e.g. use the default value of 128 for all devices).

Set `adapterInfo.isFallbackAdapter` to `adapter.[[fallback]]`.

Return `adapterInfo`.

A *normalized identifier string* is one that follows the following pattern:

`[a-z0-9]+(-[a-z0-9]+)*`

⏏ ⏏ a-z 0-9 ⏏ ⏏
- ⏏ ⏏

Examples of valid normalized identifier strings include:

`gpu`

`3d`

`0x3b2f`

`next-gen`

`series-x20-ultra`

3.7. Feature Detection

This section is non-normative.

Fully implementing this specification requires implementation of everything it specifies, except where otherwise stated (like [§ 3.6 Optional Capabilities](#)).

However, since new "core" additions are added to this specification before being exposed by implementations, many features are designed to be feature-detectable by applications:

Interface support can be detected with `typeof InterfaceName !== 'undefined'`.

Method and attribute support can be detected with `'itemName' in InterfaceName.prototype`.

New dictionary members, if they need to be detectable, generally document a specific mechanism for feature detection. For example:

[unclippedDepth](#) support is part of a device feature, ["depth-clip-control"](#).

Canvas support for [toneMapping](#) is detected using [getConfigurations\(\)](#).

3.8. Extension Documents

"Extension Documents" are additional documents which describe new functionality which is non-normative and **not part of the WebGPU/WGSL specifications**. They describe functionality that builds upon these specifications, often including one or more new API [feature](#) flags and/or WGSL `enable` directives, or interactions with other draft web specifications.

WebGPU implementations **must not** expose extension functionality; doing so is a spec violation. New functionality does not become part of the WebGPU standard until it is integrated into the WebGPU specification (this document) and/or WGSL specification.

3.9. Origin Restrictions

WebGPU allows accessing image data stored in images, videos, and canvases. Restrictions are imposed on the use of cross-domain media, because shaders can be used to indirectly deduce the contents of textures which have been uploaded to the GPU.

WebGPU disallows uploading an image source if it [is not origin-clean](#).

This also implies that the [origin-clean](#) flag for a canvas rendered using WebGPU will never be set to `false`.

For more information on issuing CORS requests for image and video elements, consult:

[HTML § 2.5.4 CORS settings attributes](#)

[HTML § 4.8.3 The `img` element `img`](#)

[HTML § 4.8.11 Media elements](#) [HTMLMediaElement](#)

3.10. Task Sources

3.10.1. WebGPU Task Source

WebGPU defines a new [task source](#) called the *WebGPU task source*. It is used for the [uncapturederror](#) event and [GPUDevice.lost](#).

To *queue a global task* for [GPUDevice](#) device, with a series of steps *steps* on the [content timeline](#):

[Queue a global task](#) on the [WebGPU task source](#), with the global object that was used to create device, and the steps *steps*.

3.10.2. Automatic Expiry Task Source

WebGPU defines a new [task source](#) called the [automatic expiry task source](#). It is used for the automatic, timed expiry (destruction) of certain objects:

[GPUTextures](#) returned by [getCurrentTexture\(\)](#)

[GPUExternalTextures](#) created from [HTMLVideoElements](#)

To *queue an automatic expiry task* with [GPUDevice](#) device and a series of steps *steps* on the [content timeline](#):

[Queue a global task](#) on the [automatic expiry task source](#), with the global object that was used to create device, and the steps *steps*.

Tasks from the [automatic expiry task source](#) **should** be processed with high priority; in particular, once queued, they **should** run before user-defined (JavaScript) tasks.

NOTE:

This behavior is more predictable, and the strictness helps developers write more portable applications by eagerly detecting incorrect assumptions about implicit lifetimes that may be hard to detect. Developers are still strongly encouraged to test in multiple implementations.

Implementation note: It is valid to implement a high-priority expiry "task" by instead inserting additional steps at a fixed point inside the [event loop processing model](#) rather than running an actual task.

3.11. Color Spaces and Encoding

WebGPU does not provide color management. All values within WebGPU (such as texture elements) are raw numeric values, not color-managed color values.

WebGPU *does* interface with color-managed outputs (via [GPUCanvasConfiguration](#)) and inputs (via [copyExternalImageToTexture\(\)](#) and [importExternalTexture\(\)](#)). Thus, color conversion must be performed between the WebGPU numeric values and the external color values. Each such interface point locally defines an encoding (color space, transfer function, and alpha premultiplication) in which the WebGPU numeric values are to be interpreted.

WebGPU allows all of the color spaces in the [PredefinedColorSpace](#) enum. Note, each color space is defined over an extended range, as defined by the referenced CSS definitions, to represent color values outside of its space (in both chrominance and luminance).

An *out-of-gamut premultiplied RGBA value* is one where any of the R/G/B channel values exceeds the alpha channel value. For example, the premultiplied sRGB RGBA value [1.0, 0, 0, 0.5] represents the (unpremultiplied) color [2, 0, 0] with 50% alpha, written `rgb(srgb 2 0 0 / 50%)` in CSS. Just like any color value outside the sRGB color gamut, this is a well defined point in the extended color space (except when alpha is 0, in which case there is no color). However, when such values are output to a visible canvas, the result is undefined (see [GPUCanvasAlphaMode "premultiplied"](#)).

3.11.1. Color Space Conversions

A color is converted between spaces by translating its representation in one space to a representation in another according to the definitions above.

If the source value has fewer than 4 RGBA channels, the missing green/blue/alpha channels are set to 0, 0, 1, respectively, before converting for color space/encoding and alpha premultiplication. After conversion, if the destination needs fewer than 4 channels, the additional channels are ignored.

Note: Grayscale images generally represent RGB values (V, V, V), or RGBA values (V, V, V, A) in their color space.

Colors are not lossily clamped during conversion: converting from one color space to another will result in values outside the range [0, 1] if the source color values were outside the range of the destination color space's gamut. For an sRGB destination, for example, this can occur if the source is `rgba16float`, in a wider color space like Display-P3, or is premultiplied and contains [out-of-gamut values](#).

Similarly, if the source value has a high bit depth (e.g. PNG with 16 bits per component) or extended range (e.g. canvas with `float16` storage), these colors are preserved through color space conversion, with intermediate computations having at least the precision of the source.

3.11.2. Color Space Conversion Elision

If the source and destination of a color space/encoding conversion are the same, then conversion is not necessary. In general, if any given step of the conversion is an identity function (no-op), implementations **should** elide it, for performance.

For optimal performance, applications **should** set their color space and encoding options so that the number of necessary conversions is minimized throughout the process. For various image sources of [GPUCopyExternalImageSourceInfo](#):

[ImageBitmap](#):

Premultiplication is controlled via [premultiplyAlpha](#).

Color space is controlled via [colorSpaceConversion](#).

2d canvas:

[Always premultiplied](#).

Color space is controlled via the [colorSpace](#) context creation attribute.

WebGL canvas:

Premultiplication is controlled via the `premultipliedAlpha` option in [WebGLContextAttributes](#).

Color space is controlled via the [WebGLRenderingContextBase](#)'s `drawingBufferColorSpace` state.

Note: Check browser implementation support for these features before relying on them.

3.12. Numeric conversions from JavaScript to WGSL

Several parts of the WebGPU API ([pipeline-overridable constants](#) and render pass clear values) take numeric values from WebIDL ([double](#) or [float](#)) and convert them to WGSL values (`bool`, `i32`, `u32`, `f32`, `f16`).

To convert an IDL value *idlValue* of type [double](#) or [float](#) to WGSL type *T*, possibly throwing a [TypeError](#), run the following [device timeline](#) steps:

Note: This [TypeError](#) is generated in the [device timeline](#) and never surfaced to JavaScript.

[Assert](#) *idlValue* is a finite value, since it is not [unrestricted double](#) or [unrestricted float](#).

Let *v* be the ECMAScript Number resulting from [!](#) converting *idlValue* to [an ECMAScript value](#).

If *T* is `bool`

Return the WGSL `bool` value corresponding to the result of [!](#) converting *v* to [an IDL value](#) of type [boolean](#).

Note: This algorithm is called after the conversion from an ECMAScript value to an IDL [double](#) or [float](#) value. If the original ECMAScript value was a non-numeric, non-boolean value like `[]` or `{}`, then the WGSL `bool` result may be different than if the ECMAScript value had been converted to IDL [boolean](#) directly.

If *T* is `i32`

Return the WGSL `i32` value corresponding to the result of [?](#) converting *v* to [an IDL value](#) of type [\[EnforceRange\] long](#).

If *T* is `u32`

Return the WGSL `u32` value corresponding to the result of [?](#) converting *v* to [an IDL value](#) of type [\[EnforceRange\] unsigned long](#).

If *T* is `f32`

Return the WGSL `f32` value corresponding to the result of [?](#) converting *v* to [an IDL value](#) of type [float](#).

If *T* is `f16`

1. Let *wgslF32* be the WGSL `f32` value corresponding to the result of [?](#) converting *v* to [an IDL value](#) of type [float](#).
2. Return `f16(wgslF32)`, the result of [!](#) converting the WGSL `f32` value to `f16` as defined in [WGSL floating point conversion](#).

Note: As long as the value is in-range of `f32`, no error is thrown, even if the value is out-of-range of `f16`.

To convert a [GPUColor](#) color to a texel value of texture format *format*, possibly throwing a [TypeError](#), run the following [device timeline](#) steps:

Note: This [TypeError](#) is generated in the [device timeline](#) and never surfaced to JavaScript.

If the components of *format* ([assert](#) they all have the same type) are:

floating-point types or normalized types

Let *T* be f32.
signed integer types
Let *T* be i32.
unsigned integer types
Let *T* be u32.

Let *wgslColor* be a WGSL value of type `vec4<T>`, where the 4 components are the RGBA channels of *color*, each 2 converted to WGSL type *T*.
Convert *wgslColor* to *format* using the same conversion rules as the § 23.2.7 Output Merging step, and return the result.

Note: For non-integer types, the exact choice of value is implementation-defined. For normalized types, the value is clamped to the range of the type.
Note: In other words, the value written will be as if it was written by a WGSL shader that outputs the value represented as a `vec4` of f32, i32, or u32.

4. Initialization

4.1. navigator.gpu

A GPU object is available in the Window and WorkerGlobalScope contexts through the Navigator and WorkerNavigator interfaces respectively and is exposed via navigator.gpu:

```
interface mixin
NavigatorGPU
{
  [SameObject, SecureContext] readonly attribute GPU gpu;
};
Navigator includes NavigatorGPU;
WorkerNavigator includes NavigatorGPU;
NavigatorGPU has the following attributes:
```

gpu, of type GPU, readonly
A global singleton providing top-level entry points like requestAdapter().

4.2. GPU

GPU is the entry point to WebGPU.

```
[Exposed=(Window, Worker), SecureContext]
interface GPU {
  Promise<GPUAdapter?> requestAdapter(optional GPURequestAdapterOptions options = {});
  GPUTextureFormat getPreferredCanvasFormat();
  [SameObject] readonly attribute WGSLLanguageFeatures wgslLanguageFeatures;
};
```

GPU has the following methods:

requestAdapter(options)

Requests an adapter from the user agent. The user agent chooses whether to return an adapter, and, if so, chooses according to the provided options.

Called on: GPU this.

Arguments:

Arguments for the GPU.requestAdapter(options) method.

Parameter	Type	Nullable	Optional	Description
options	GPURequestAdapterOptions	✗	✓	Criteria used to select the adapter.

Returns: Promise<GPUAdapter?>

Content timeline steps:

- Let contentTimeline be the current Content timeline.
- Let promise be a new promise.
- Issue the initialization steps on the Device timeline of this.
- Return promise.

Device timeline initialization steps:

- All of the requirements in the following steps must be met.

1. *options.featureLevel* must be a [feature level string](#).

If they are met **and** the user agent chooses to return an adapter:

1. Set *adapter* to an [adapter](#) chosen according to the rules in [§ 4.2.2 Adapter Selection](#) and the criteria in *options*, adhering to [§ 4.2.1 Adapter Capability Guarantees](#). Initialize the properties of *adapter* according to their definitions:

1. Set *adapter.[]limits[]* and *adapter.[]features[]* according to the supported capabilities of the adapter. *adapter.[]features[]* must contain "[core-features-and-limits](#)".
2. If *adapter* meets the criteria of a [fallback adapter](#) set *adapter.[]fallback[]* to `true`. Otherwise, set it to `false`.
3. Set *adapter.[]xrCompatible[]* to *options.xrCompatible*.

Otherwise:

1. Let *adapter* be `null`.
2. Issue the subsequent steps on *contentTimeline*.

[Content timeline](#) steps:

1. If *adapter* is not `null`:
 1. [Resolve](#) *promise* with a new [GPUAdapter](#) encapsulating *adapter*.

Otherwise:

1. [Resolve](#) *promise* with `null`.

getPreferredCanvasFormat()

Returns an optimal [GPUTextureFormat](#) for displaying 8-bit depth, standard dynamic range content on this system. Must only return "[rgba8unorm](#)" or "[bgra8unorm](#)".

The returned value can be passed as the *format* to [configure\(\)](#) calls on a [GPUCanvasContext](#) to ensure the associated canvas is able to display its contents efficiently.

Note: Canvases which are not displayed to the screen may or may not benefit from using this format.

Called on: [GPU](#) this.

Returns: [GPUTextureFormat](#)

[Content timeline](#) steps:

1. Return either "[rgba8unorm](#)" or "[bgra8unorm](#)", depending on which format is optimal for displaying WebGPU canvases on this system.

[GPU](#) has the following attributes:

wgslLanguageFeatures, of type [WGSLLanguageFeatures](#), readonly

The names of supported WGSL [language extensions](#). Supported language extensions are automatically enabled.

[Adapters](#) may [expire](#) at any time. Upon any change in the system's state that could affect the result of any [requestAdapter\(\)](#) call, the user agent **should** [expire](#) all previously-returned [adapters](#). For example:

A physical adapter is added/removed (via plug/unplug, driver update, hang recovery, etc.)

The system's power configuration has changed (laptop unplugged, power settings changed, etc.)

Note: User agents may choose to [expire adapters](#) often, even when there has been no system state change (e.g. seconds or minutes after the adapter was created). This can help obfuscate real system state changes, and make developers more aware that calling [requestAdapter\(\)](#) again is always necessary before calling [requestDevice\(\)](#). If an application does encounter this situation, standard device-loss recovery handling should allow it to recover.

Requesting a [GPUAdapter](#) with no hints:

```
const gpuAdapter = await navigator.gpu.requestAdapter();
```

4.2.1. Adapter Capability Guarantees

Any [GPUAdapter](#) returned by [requestAdapter\(\)](#) must provide the following guarantees:

At least one of the following must be true:

["texture-compression-bc"](#) is supported.

Both ["texture-compression-etc2"](#) and ["texture-compression-astc"](#) are supported.

If ["texture-compression-bc-sliced-3d"](#) is supported, then ["texture-compression-bc"](#) must be supported.

If ["texture-compression-astc-sliced-3d"](#) is supported, then ["texture-compression-astc"](#) must be supported.

All supported limits must be either the [default](#) value or [better](#).

All [alignment-class](#) limits must be powers of 2.

[maxBindingsPerBindGroup](#) must be $\geq (\text{max bindings per shader stage} \times \text{max shader stages per pipeline})$, where:

max bindings per shader stage is $(\text{maxSampledTexturesPerShaderStage} + \text{maxSamplersPerShaderStage} + \text{maxStorageBuffersPerShaderStage} + \text{maxStorageTexturesPerShaderStage} + \text{maxUniformBuffersPerShaderStage})$.

max shader stages per pipeline is 2, because a [GPURenderPipeline](#) supports both a vertex and fragment shader.

Note: [maxBindingsPerBindGroup](#) does not reflect a fundamental limit; implementations should raise it to conform to this requirement, rather than lowering the other limits.

[maxBindGroups](#) must be $\leq \text{maxBindGroupsPlusVertexBuffers}$.

[maxVertexBuffers](#) must be $\leq \text{maxBindGroupsPlusVertexBuffers}$.

[minUniformBufferOffsetAlignment](#) and [minStorageBufferOffsetAlignment](#) must both be ≥ 32 bytes.

Note: 32 bytes would be the alignment of `vec4<f64>`. See [WebGPU Shading Language § 14.4.1 Alignment and Size](#).

[maxUniformBufferBindingSize](#) must be $\leq \text{maxBufferSize}$.

[maxStorageBufferBindingSize](#) must be $\leq \text{maxBufferSize}$.

[maxStorageBufferBindingSize](#) must be a multiple of 4 bytes.

[maxVertexBufferArrayStride](#) must be a multiple of 4 bytes.

[maxComputeWorkgroupSizeX](#) must be $\leq \text{maxComputeInvocationsPerWorkgroup}$.

[maxComputeWorkgroupSizeY](#) must be $\leq \text{maxComputeInvocationsPerWorkgroup}$.

[maxComputeWorkgroupSizeZ](#) must be $\leq \text{maxComputeInvocationsPerWorkgroup}$.

[maxComputeInvocationsPerWorkgroup](#) must be $\leq \text{maxComputeWorkgroupSizeX} \times \text{maxComputeWorkgroupSizeY} \times \text{maxComputeWorkgroupSizeZ}$.

4.2.2. Adapter Selection

[GPURequestAdapterOptions](#) provides hints to the user agent indicating what configuration is suitable for the application.

dictionary [GPURequestAdapterOptions](#) {

[DOMString](#) [featureLevel](#) = "core";

[GPUPowerPreference](#) [powerPreference](#);

[boolean](#) [forceFallbackAdapter](#) = false;

[boolean](#) [xrCompatible](#) = false;

};

enum

[GPUPowerPreference](#)

{

["low-power"](#),

["high-performance"](#),

};

[GPURequestAdapterOptions](#) has the following members:

featureLevel, of type [DOMString](#), defaulting to "core"

"Feature level" for the adapter request.

The allowed *feature level string* values are:

"core"

No effect.

"compatibility"

No effect.

Note: This value is reserved for future use as a way to opt into additional validation restrictions. Applications should not use this value at this time.

powerPreference, of type [GPUPowerPreference](#)

Optionally provides a hint indicating what class of [adapter](#) should be selected from the system's available adapters.

The value of this hint may influence which adapter is chosen, but it must not influence whether an adapter is returned or not.

Note: The primary utility of this hint is to influence which GPU is used in a multi-GPU system. For instance, some laptops have a low-power integrated GPU and a high-performance discrete GPU. This hint may also affect the power configuration of the selected GPU to match the requested power preference.

Note: Depending on the exact hardware configuration, such as battery status and attached displays or removable GPUs, the user agent may select different [adapters](#) given the same power preference. Typically, given the same hardware configuration and state and `powerPreference`, the user agent is likely to select the same adapter.

It must be one of the following values:

`undefined` (or not present)

Provides no hint to the user agent.

`"low-power"`

Indicates a request to prioritize power savings over performance.

Note: Generally, content should use this if it is unlikely to be constrained by drawing performance; for example, if it renders only one frame per second, draws only relatively simple geometry with simple shaders, or uses a small HTML canvas element. Developers are encouraged to use this value if their content allows, since it may significantly improve battery life on portable devices.

`"high-performance"`

Indicates a request to prioritize performance over power consumption.

Note: By choosing this value, developers should be aware that, for [devices](#) created on the resulting adapter, user agents are more likely to force device loss, in order to save power by switching to a lower-power adapter. Developers are encouraged to only specify this value if they believe it is absolutely necessary, since it may significantly decrease battery life on portable devices.

`forceFallbackAdapter`, of type [boolean](#), defaulting to `false`

When set to `true` indicates that only a [fallback adapter](#) may be returned. If the user agent does not support a [fallback adapter](#), will cause `requestAdapter()` to resolve to `null`.

Note: `requestAdapter()` may still return a [fallback adapter](#) if `forceFallbackAdapter` is set to `false` and either no other appropriate [adapter](#) is available or the user agent chooses to return a [fallback adapter](#). Developers that wish to prevent their applications from running on [fallback adapters](#) should check the [info.isFallbackAdapter](#) attribute prior to requesting a [GPUDevice](#).

`xrCompatible`, of type [boolean](#), defaulting to `false`

When set to `true` indicates that the best [adapter](#) for rendering to a [WebXR session](#) must be returned. If the user agent or system does not support [WebXR sessions](#) then adapter selection may ignore this value.

Note: If `xrCompatible` is not set to `true` when the adapter is requested, [GPUDevices](#) created from the adapter cannot be used to render for [WebXR sessions](#).

Requesting a `"high-performance" GPUAdapter`:

```
const gpuAdapter = await navigator.gpu.requestAdapter({
  powerPreference: 'high-performance'
});
```

4.3. GPUAdapter

A [GPUAdapter](#) encapsulates an [adapter](#), and describes its capabilities ([features](#) and [limits](#)).

To get a [GPUAdapter](#), use `requestAdapter()`.

[[Exposed](#)=(Window, Worker), [SecureContext](#)]

interface [GPUAdapter](#) {

 [[SameObject](#)] readonly attribute [GPUSupportedFeatures](#) [features](#);

 [[SameObject](#)] readonly attribute [GPUSupportedLimits](#) [limits](#);

 [[SameObject](#)] readonly attribute [GPUAdapterInfo](#) [info](#);

[Promise](#)<[GPUDevice](#)> [requestDevice](#)(optional [GPUDeviceDescriptor](#) [descriptor](#) = {});

[GPUAdapter](#) has the following [immutable properties](#)

`features`, of type [GPUSupportedFeatures](#), readonly

The set of values in `this.[[adapter]].[[features]]`.

`limits`, of type [GPUSupportedLimits](#), readonly

The limits in `this.[[adapter]].[[limits]]`.

`info`, of type [GPUAdapterInfo](#), readonly

Information about the physical adapter underlying this [GPUAdapter](#).

For a given [GPUAdapter](#), the [GPUAdapterInfo](#) values exposed are constant over time.

The same object is returned each time. To create that object for the first time:

Called on: [GPUAdapter](#) *this*.

Returns: [GPUAdapterInfo](#)

[Content timeline](#) steps:

1. Return a [new adapter info](#) for *this*.[\[\[adapter\]\]](#).

[\[\[adapter\]\]](#), of type [adapter](#), readonly

The [adapter](#) to which this [GPUAdapter](#) refers.

[GPUAdapter](#) has the following methods:

requestDevice(descriptor)

Requests a [device](#) from the [adapter](#).

This is a one-time action: if a device is returned successfully, the adapter becomes "[consumed](#)".

Called on: [GPUAdapter](#) *this*.

Arguments:

Arguments for the [GPUAdapter.requestDevice\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUDeviceDescriptor	✗	✓	Description of the GPUDevice to request.

Returns: [Promise](#)<[GPUDevice](#)>

[Content timeline](#) steps:

1. Let *contentTimeline* be the current [Content timeline](#).
2. Let *promise* be a [new promise](#).
3. Let *adapter* be *this*.[\[\[adapter\]\]](#).
4. Issue the *initialization steps* to the [Device timeline](#) of *this*.
5. Return *promise*.

[Device timeline](#) *initialization steps*:

1. If any of the following requirements are unmet:

- The set of values in *descriptor*.[requiredFeatures](#) must be a subset of those in *adapter*.[\[\[features\]\]](#).

Then issue the following steps on *contentTimeline* and return:

Note: This is the same error that is produced if a feature name isn't known by the browser at all (in its [GPUFeatureName](#) definition). This converges the behavior when the browser doesn't support a feature with the behavior when a particular adapter doesn't support a feature.

2. All of the requirements in the following steps *must* be met.

1. *adapter*.[\[\[state\]\]](#) must not be "[consumed](#)".
2. For each [*key*, *value*] in *descriptor*.[requiredLimits](#) for which *value* is not *undefined*:
 1. *key* *must* be the name of a member of [supported limits](#).
 2. *value* *must* be no [better](#) than *adapter*.[\[\[limits\]\]](#)[*key*].
3. If *key*'s [class](#) is [alignment](#), *value* *must* be a power of 2 less than 2³².

Note: User agents should consider issuing developer-visible warnings when *key* is not recognized, even when *value* is *undefined*.

If any are unmet, issue the following steps on *contentTimeline* and return:

3. If *adapter*.[\[\[state\]\]](#) is "[expired](#)" or the user agent otherwise cannot fulfill the request:
 1. Let *device* be a new [device](#).
 2. [Lose the device](#)(*device*, "[unknown](#)").
 3. [Assert](#) *adapter*.[\[\[state\]\]](#) is "[expired](#)".

Note: User agents should consider issuing developer-visible warnings in most or all cases when this occurs. Applications should perform reinitialization logic starting with [requestAdapter\(\)](#).

Otherwise:

1. Let *device* be a [new device](#) with the capabilities described by *descriptor*.

2. [Expire adapter](#).

4. Issue the subsequent steps on *contentTimeline*.

[Content timeline](#) steps:

1. Let *gpuDevice* be a new [GPUDevice](#) instance.
2. Set *gpuDevice*.[\[\[device\]\]](#) to *device*.
3. Set *device*.[\[\[content_device\]\]](#) to *gpuDevice*.
4. Set *gpuDevice*.[label](#) to *descriptor*.[label](#).
5. [Resolve](#) *promise* with *gpuDevice*.

Note: If the device is already lost because the adapter could not fulfill the request, *device*.[lost](#) has already resolved before *promise* resolves.

Requesting a [GPUDevice](#) with default features and limits:

```
const gpuAdapter = await navigator.gpu.requestAdapter();
const gpuDevice = await gpuAdapter.requestDevice();
```

4.3.1. GPUDeviceDescriptor

[GPUDeviceDescriptor](#) describes a device request.

dictionary [GPUDeviceDescriptor](#)

```
: GPUObjectDescriptorBase {
  sequence<GPUFeatureName> requiredFeatures = [];
  record<DOMString, (GPUSize64 or undefined)> requiredLimits = {};
  GPUQueueDescriptor defaultQueue = {};
};
```

[GPUDeviceDescriptor](#) has the following members:

requiredFeatures, of type sequence<[GPUFeatureName](#)>, defaulting to []

Specifies the [features](#) that are required by the device request. The request will fail if the adapter cannot provide these features.

Exactly the specified set of features, and no more or less, will be allowed in validation of API calls on the resulting device.

requiredLimits, of type record<DOMString, (GPUSize64 or undefined)>, defaulting to {}

Specifies the [limits](#) that are required by the device request. The request will fail if the adapter cannot provide these limits.

Each key with a non-undefined value must be the name of a member of [supported limits](#).

API calls on the resulting device perform validation according to the exact limits of the device (not the adapter; see [§3.6.2 Limits](#)).

defaultQueue, of type [GPUQueueDescriptor](#), defaulting to {}

The descriptor for the default [GPUQueue](#).

Requesting a [GPUDevice](#) with the ["texture-compression-astc"](#) feature if supported:

```
const gpuAdapter = await navigator.gpu.requestAdapter();
```

```
const requiredFeatures = [];
if (gpuAdapter.features.has('texture-compression-astc')) {
  requiredFeatures.push('texture-compression-astc')
}
```

```
const gpuDevice = await gpuAdapter.requestDevice({
  requiredFeatures
});
```

Requesting a [GPUDevice](#) with a higher [maxColorAttachmentBytesPerSample](#) limit:

```
const gpuAdapter = await navigator.gpu.requestAdapter();
```

```
if (gpuAdapter.limits.maxColorAttachmentBytesPerSample < 64) {
  // When the desired limit isn't supported, take action to either fall back to a code
  // path that does not require the higher limit or notify the user that their device
  // does not meet minimum requirements.
}
```

```
// Request higher limit of max color attachments bytes per sample.
const gpuDevice = await gpuAdapter.requestDevice({
  requiredLimits: { maxColorAttachmentBytesPerSample: 64 },
});
```

4.3.1.1. GPUFeatureName

Each [GPUFeatureName](#) identifies a set of functionality which, if available, allows additional usages of WebGPU that would have otherwise been invalid.

```
enum GPUFeatureName {
  "core-features-and-limits",
  "depth-clip-control",
  "depth32float-stencil8",
  "texture-compression-bc",
  "texture-compression-bc-sliced-3d",
  "texture-compression-etc2",
  "texture-compression-astc",
  "texture-compression-astc-sliced-3d",
  "timestamp-query",
  "indirect-first-instance",
  "shader-f16",
  "rg11b10float-renderable",
  "bgra8unorm-storage",
  "float32-filterable",
  "float32-blendable",
  "clip-distances",
  "dual-source-blending",
  "subgroups",
  "texture-formats-tier1",
  "texture-formats-tier2",
  "primitive-index",
};
```

4.4. GPUDevice

A [GPUDevice](#) encapsulates a [device](#) and exposes the functionality of that device.

[GPUDevice](#) is the top-level interface through which [WebGPU interfaces](#) are created.

To get a [GPUDevice](#), use [requestDevice\(\)](#).

[Exposed=(Window, Worker), [SecureContext](#)]

```
interface GPUDevice : EventTarget {
  [SameObject] readonly attribute GPUSupportedFeatures features;
  [SameObject] readonly attribute GPUSupportedLimits limits;
  [SameObject] readonly attribute GPUAdapterInfo adapterInfo;

  [SameObject] readonly attribute GPUQueue queue;

  undefined destroy();

  GPUBuffer createBuffer(GPUBufferDescriptor descriptor);
  GPUTexture createTexture(GPUTextureDescriptor descriptor);
  GPUSampler createSampler(optional GPUSamplerDescriptor descriptor = {});
  GPUExternalTexture importExternalTexture(GPUExternalTextureDescriptor descriptor);

  GPUBindGroupLayout createBindGroupLayout(GPUBindGroupLayoutDescriptor descriptor);
  GPUPipelineLayout createPipelineLayout(GPUPipelineLayoutDescriptor descriptor);
  GPUBindGroup createBindGroup(GPUBindGroupDescriptor descriptor);

  GPUShaderModule createShaderModule(GPUShaderModuleDescriptor descriptor);
  GPUComputePipeline createComputePipeline(GPUComputePipelineDescriptor descriptor);
  GPURenderPipeline createRenderPipeline(GPURenderPipelineDescriptor descriptor);
```

```
Promise<GPUComputePipeline> createComputePipelineAsync(GPUComputePipelineDescriptor descriptor);
Promise<GPURenderPipeline> createRenderPipelineAsync(GPURenderPipelineDescriptor descriptor);
```

```
GPUCommandEncoder createCommandEncoder(optional GPUCommandEncoderDescriptor descriptor = {});
GPURenderBundleEncoder createRenderBundleEncoder(GPURenderBundleEncoderDescriptor descriptor);
```

```
GPUQuerySet createQuerySet(GPUQuerySetDescriptor descriptor);
};
```

[GPUDevice](#) includes [GPUObjectBase](#);

[GPUDevice](#) has the following [immutable properties](#):

features, of type [GPUSupportedFeatures](#), readonly

A set containing the [GPUFeatureName](#) values of the features supported by the device (`[[device]].[[features]]`).

limits, of type [GPUSupportedLimits](#), readonly

The limits supported by the device (`[[device]].[[limits]]`).

queue, of type [GPUQueue](#), readonly

The primary [GPUQueue](#) for this device.

adapterInfo, of type [GPUAdapterInfo](#), readonly

Information about the physical adapter which created the [device](#) that this [GPUDevice](#) refers to.

For a given [GPUDevice](#), the [GPUAdapterInfo](#) values exposed are constant over time.

The same object is returned each time. To create that object for the first time:

Called on: [GPUDevice](#) *this*.

Returns: [GPUAdapterInfo](#)

[Content timeline](#) steps:

1. Return a [new adapter info](#) for *this*.`[[device]].[[adapter]]`.

The `[[device]]` for a [GPUDevice](#) is the [device](#) that the [GPUDevice](#) refers to.

[GPUDevice](#) has the following methods:

destroy()

Destroys the [device](#), preventing further operations on it. Outstanding asynchronous operations will fail.

Note: It is valid to destroy a device multiple times.

Called on: [GPUDevice](#) *this*.

[Content timeline](#) steps:

1. [unmap\(\)](#) all [GPUBuffer](#)s from this device.
2. Issue the subsequent steps on the [Device timeline](#) of *this*.
1. [Lose the device](#)(*this*.`[[device]]`, "destroyed").

Note: Since no further operations can be enqueued on this device, implementations can abort outstanding asynchronous operations immediately and free resource allocations, including mapped memory that was just unmapped.

A [GPUDevice](#)'s *allowed buffer usages* are:

Always allowed: [MAP_READ](#), [MAP_WRITE](#), [COPY_SRC](#), [COPY_DST](#), [INDEX](#), [VERTEX](#), [UNIFORM](#), [STORAGE](#), [INDIRECT](#), [QUERY_RESOLVE](#)

A [GPUDevice](#)'s *allowed texture usages* are:

Always allowed: [COPY_SRC](#), [COPY_DST](#), [TEXTURE_BINDING](#), [STORAGE_BINDING](#), [RENDER_ATTACHMENT](#)

4.5. Example

A more robust example of requesting a [GPUAdapter](#) and [GPUDevice](#) with error handling:

```
let gpuDevice = null;
```

```
async function initializeWebGPU() {
  // Check to ensure the user agent supports WebGPU.
  if (!('gpu' in navigator)) {
    console.error("User agent doesn't support WebGPU.");
  }
}
```

```

    return false;
}

// Request an adapter.
const gpuAdapter = await navigator.gpu.requestAdapter();

// requestAdapter may resolve with null if no suitable adapters are found.
if (!gpuAdapter) {
    console.error('No WebGPU adapters found.');
```

return false;

```

}

// Request a device.
// Note that the promise will reject if invalid options are passed to the optional
// dictionary. To avoid the promise rejecting always check any features and limits
// against the adapters features and limits prior to calling requestDevice().
gpuDevice = await gpuAdapter.requestDevice();

// requestDevice will never return null, but if a valid device request can't be
// fulfilled for some reason it may resolve to a device which has already been lost.
// Additionally, devices can be lost at any time after creation for a variety of reasons
// (ie: browser resource management, driver updates), so it's a good idea to always
// handle lost devices gracefully.
gpuDevice.lost.then((info) => {
    console.error(`WebGPU device was lost: ${info.message}`);

    gpuDevice = null;

    // Many causes for lost devices are transient, so applications should try getting a
    // new device once a previous one has been lost unless the loss was caused by the
    // application intentionally destroying the device. Note that any WebGPU resources
    // created with the previous device (buffers, textures, etc) will need to be
    // re-created with the new one.
    if (info.reason !== 'destroyed') {
        initializeWebGPU();
    }
});

onWebGPUInitialized();

return true;
}

function onWebGPUInitialized() {
    // Begin creating WebGPU resources here...
}
```

```
initializeWebGPU();
```

5. Buffers

5.1. GPUBuffer

A [GPUBuffer](#) represents a block of memory that can be used in GPU operations. Data is stored in linear layout, meaning that each byte of the allocation can be addressed by its offset from the start of the [GPUBuffer](#), subject to alignment restrictions depending on the operation. Some [GPUBuffers](#) can be mapped which makes the block of memory accessible via an [ArrayBuffer](#) called its mapping.

[GPUBuffer](#)s are created via [createBuffer\(\)](#). Buffers may be [mappedAtCreation](#).

[[Exposed](#)=(Window, Worker), [SecureContext](#)]

```

interface GPUBuffer {
    readonly attribute GPUSize64Out size;
    readonly attribute GPUFlagsConstant usage;
```

readonly attribute [GPUBufferMapState](#) `mapState`;

```
Promise<undefined> mapAsync(GPUMapModeFlags mode, optional GPUSize64 offset = 0, optional GPUSize64 size);
ArrayBuffer getMappedRange(optional GPUSize64 offset = 0, optional GPUSize64 size);
undefined unmap();
```

```
undefined destroy();
```

```
};
```

[GPUBuffer](#) includes [GPUObjectBase](#);

```
enum GPUBufferMapState {
```

```
    "unmapped",
```

```
    "pending",
```

```
    "mapped",
```

```
};
```

[GPUBuffer](#) has the following [immutable properties](#):

size, of type [GPUSize64Out](#), readonly

The length of the [GPUBuffer](#) allocation in bytes.

usage, of type [GPUFlagsConstant](#), readonly

The allowed usages for this [GPUBuffer](#).

[GPUBuffer](#) has the following [content timeline properties](#):

mapState, of type [GPUBufferMapState](#), readonly

The current *GPUBufferMapState* of the buffer:

"unmapped"

The buffer is not mapped for use by `this.getMappedRange()`.

"pending"

A mapping of the buffer has been requested, but is pending. It may succeed, or fail validation in `mapAsync()`.

"mapped"

The buffer is mapped and `this.getMappedRange()` may be used.

The [getter steps](#) are:

[Content timeline](#) steps:

1. If `this.[[mapping]]` is not null, return "mapped".
2. If `this.[[pending_map]]` is not null, return "pending".
3. Return "unmapped".

`[[pending_map]]`, of type [Promise](#)<void> or null, initially null

The [Promise](#) returned by the currently-pending `mapAsync()` call.

There is never more than one pending map, because `mapAsync()` will refuse immediately if a request is already in flight.

`[[mapping]]`, of type [active buffer mapping](#) or null, initially null

Set if and only if the buffer is currently mapped for use by `getMappedRange()`. Null otherwise (even if there is a `[[pending_map]]`).

An *active buffer mapping* is a structure with the following fields:

data, of type [Data Block](#)

The mapping for this [GPUBuffer](#). This data is accessed through [ArrayBuffer](#)s which are views onto this data, returned by `getMappedRange()` and stored in [views](#).

mode, of type [GPUMapModeFlags](#)

The [GPUMapModeFlags](#) of the map, as specified in the corresponding call to `mapAsync()` or `createBuffer()`.

range, of type tuple [[unsigned long long](#), [unsigned long long](#)]

The range of this [GPUBuffer](#) that is mapped.

views, of type [list](#)<[ArrayBuffer](#)>

The [ArrayBuffer](#)s returned via `getMappedRange()` to the application. They are tracked so they can be detached when `unmap()` is called.

To initialize an active buffer mapping with mode *mode* and range *range*, run the following [content timeline](#) steps:

1. Let *size* be `range[1] - range[0]`.
2. Let *data* be `? CreateByteDataBlock(size)`.

NOTE:

This may result in a [RangeError](#) being thrown. For consistency and predictability:

- For any size at which `new ArrayBuffer()` would succeed at a given moment, this allocation **should** succeed at that moment.
- For any size at which `new ArrayBuffer()` *deterministically* throws a [RangeError](#), this allocation **should** as well.

3. Return an [active buffer mapping](#) with:

- [data](#) set to *data*.
- [mode](#) set to *mode*.
- [range](#) set to *range*.
- [views](#) set to `[]`.

[GPUBuffer](#) has the following [device timeline properties](#):

`[[internal state]]`

The current internal state of the buffer:

"available"

The buffer can be used in queue operations (unless it is [invalid](#)).

"unavailable"

The buffer cannot be used in queue operations due to being mapped.

"destroyed"

The buffer cannot be used in any operations due to being [destroy\(\)](#)ed.

5.1.1. GPUBufferDescriptor

dictionary [GPUBufferDescriptor](#)

```
: GPUObjectDescriptorBase {
  required GPUSize64 size;
  required GPUBufferUsageFlags usage;
  boolean mappedAtCreation = false;
};
```

[GPUBufferDescriptor](#) has the following members:

size, of type [GPUSize64](#)

The size of the buffer in bytes.

usage, of type [GPUBufferUsageFlags](#)

The allowed usages for the buffer.

mappedAtCreation, of type [boolean](#), defaulting to `false`

If `true` creates the buffer in an already mapped state, allowing [getMappedRange\(\)](#) to be called immediately. It is valid to set [mappedAtCreation](#) to `true` even if [usage](#) does not contain [MAP_READ](#) or [MAP_WRITE](#). This can be used to set the buffer's initial data.

Guarantees that even if the buffer creation eventually fails, it will still appear as if the mapped range can be written/read to until it is unmapped.

5.1.2. Buffer Usages

typedef [[EnforceRange](#)] [unsigned long](#)

[GPUBufferUsageFlags](#)

;

[[Exposed](#)=(Window, Worker), [SecureContext](#)]

namespace

[GPUBufferUsage](#)

```
{
  const GPUFlagsConstant MAP_READ = 0x0001;
  const GPUFlagsConstant MAP_WRITE = 0x0002;
  const GPUFlagsConstant COPY_SRC = 0x0004;
  const GPUFlagsConstant COPY_DST = 0x0008;
```



```

const GPUFlagsConstant INDEX      = 0x0010;
const GPUFlagsConstant VERTEX     = 0x0020;
const GPUFlagsConstant UNIFORM    = 0x0040;
const GPUFlagsConstant STORAGE    = 0x0080;
const GPUFlagsConstant INDIRECT   = 0x0100;
const GPUFlagsConstant QUERY_RESOLVE = 0x0200;
};

```

The [GPUBufferUsage](#) flags determine how a [GPUBuffer](#) may be used after its creation:

MAP_READ

The buffer can be mapped for reading. (Example: calling [mapAsync\(\)](#) with [GPUMapMode.READ](#))

May only be combined with [COPY_DST](#).

MAP_WRITE

The buffer can be mapped for writing. (Example: calling [mapAsync\(\)](#) with [GPUMapMode.WRITE](#))

May only be combined with [COPY_SRC](#).

COPY_SRC

The buffer can be used as the source of a copy operation. (Examples: as the `source` argument of a [copyBufferToBuffer\(\)](#) or [copyBufferToTexture\(\)](#) call.)

COPY_DST

The buffer can be used as the destination of a copy or write operation. (Examples: as the `destination` argument of a [copyBufferToBuffer\(\)](#) or [copyTextureToBuffer\(\)](#) call, or as the target of a [writeBuffer\(\)](#) call.)

INDEX

The buffer can be used as an index buffer. (Example: passed to [setIndexBuffer\(\)](#).)

VERTEX

The buffer can be used as a vertex buffer. (Example: passed to [setVertexBuffer\(\)](#).)

UNIFORM

The buffer can be used as a uniform buffer. (Example: as a bind group entry for a [GPUBufferBindingLayout](#) with a `buffer.type` of `"uniform"`.)

STORAGE

The buffer can be used as a storage buffer. (Example: as a bind group entry for a [GPUBufferBindingLayout](#) with a `buffer.type` of `"storage"` or `"read-only-storage"`.)

INDIRECT

The buffer can be used as to store indirect command arguments. (Examples: as the `indirectBuffer` argument of a [drawIndirect\(\)](#) or [dispatchWorkgroupsIndirect\(\)](#) call.)

QUERY_RESOLVE

The buffer can be used to capture query results. (Example: as the `destination` argument of a [resolveQuerySet\(\)](#) call.)

5.1.3. Buffer Creation

createBuffer(descriptor)

Creates a [GPUBuffer](#).

Called on: [GPUDevice](#) *this*.

Arguments:

Arguments for the [GPUDevice.createBuffer\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUBufferDescriptor	✗	✗	Description of the GPUBuffer to create.

Returns: [GPUBuffer](#)

[Content timeline](#) steps:

- Let *b* be [! create a new WebGPU object](#)(*this*, [GPUBuffer](#), *descriptor*).
- Set *b.size* to *descriptor.size*.
- Set *b.usage* to *descriptor.usage*.
- If *descriptor.mappedAtCreation* is `true`:
 - If *descriptor.size* is not a multiple of 4, throw a [RangeError](#).
 - Set *b.[[mapping]]* to [? initialize an active buffer mapping](#) with mode [WRITE](#) and range `[0, descriptor.size]`.

5. Issue the *initialization steps* on the [Device timeline](#) of *this*.

6. Return *b*.

[Device timeline](#) *initialization steps*:

1. If any of the following requirements are unmet, [generate a validation error](#), [invalidate](#) *b* and return.

- *this* must not be [lost](#).
- *descriptor.usage* must not be 0.
- *descriptor.usage* must be a subset of the [allowed buffer usages](#) for *this*.
- If *descriptor.usage* contains [MAP_READ](#):
- *descriptor.usage* must contain no other flags except [COPY_DST](#).
- If *descriptor.usage* contains [MAP_WRITE](#):
- *descriptor.usage* must contain no other flags except [COPY_SRC](#).
- If *descriptor.size* must be \leq *this*.[\[\[device\]\].\[\[limits\]\].maxBufferSize](#).

Note: If buffer creation fails, and *descriptor.mappedAtCreation* is `false`, any calls to [mapAsync\(\)](#) will reject, so any resources allocated to enable mapping can and may be discarded or recycled.

1. If *descriptor.mappedAtCreation* is `true`:

1. Set *b*.[\[\[internal state\]\]](#) to `"unavailable"`.

Otherwise:

1. Set *b*.[\[\[internal state\]\]](#) to `"available"`.

2. Create a device allocation for *b* where each byte is zero.

If the allocation fails without side-effects, [generate an out-of-memory error](#), [invalidate](#) *b*, and return.

Creating a 128 byte uniform buffer that can be written into:

```
const buffer = gpuDevice.createBuffer({
  size: 128,
  usage: GPUBufferUsage.UNIFORM | GPUBufferUsage.COPY_DST
});
```

5.1.4. Buffer Destruction

An application that no longer requires a [GPUBuffer](#) can choose to lose access to it before garbage collection by calling [destroy\(\)](#). Destroying a buffer also unmaps it, freeing any memory allocated for the mapping.

Note: This allows the user agent to reclaim the GPU memory associated with the [GPUBuffer](#) once all previously submitted operations using it are complete.

[GPUBuffer](#) has the following methods:

destroy()

Destroys the [GPUBuffer](#).

Note: It is valid to destroy a buffer multiple times.

Called on: [GPUBuffer](#) *this*.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Call *this*.[unmap\(\)](#).

2. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. Set *this*.[\[\[internal state\]\]](#) to `"destroyed"`.

Note: Since no further operations can be enqueued using this buffer, implementations can free resource allocations, including mapped memory that was just unmapped.

5.2. Buffer Mapping

An application can request to map a [GPUBuffer](#) so that they can access its content via [ArrayBuffer](#)s that represent part of the [GPUBuffer](#)'s allocations. Mapping a

[GPUBuffer](#) is requested asynchronously with [mapAsync\(\)](#) so that the user agent can ensure the GPU finished using the [GPUBuffer](#) before the application can access its content. A mapped [GPUBuffer](#) cannot be used by the GPU and must be unmapped using [unmap\(\)](#) before work using it can be submitted to the [Queue timeline](#).

Once the [GPUBuffer](#) is mapped, the application can synchronously ask for access to ranges of its content with [getMappedRange\(\)](#). The returned [ArrayBuffer](#) can only be [detached](#) by [unmap\(\)](#) (directly, or via [GPUBuffer.destroy\(\)](#) or [GPUDevice.destroy\(\)](#)), and cannot be [transferred](#). A [TypeError](#) is thrown by any other operation that attempts to do so.

typedef [\[EnforceRange\]](#) unsigned long

[GPUMapModeFlags](#)

```
;  
\[Exposed=(Window, Worker), SecureContext]  
namespace  
GPUMapMode  
{  
  const GPUFlagsConstant READ = 0x0001;  
  const GPUFlagsConstant WRITE = 0x0002;  
};
```

The [GPUMapMode](#) flags determine how a [GPUBuffer](#) is mapped when calling [mapAsync\(\)](#):

READ

Only valid with buffers created with the [MAP_READ](#) usage.

Once the buffer is mapped, calls to [getMappedRange\(\)](#) will return an [ArrayBuffer](#) containing the buffer’s current values. Changes to the returned [ArrayBuffer](#) will be discarded after [unmap\(\)](#) is called.

WRITE

Only valid with buffers created with the [MAP_WRITE](#) usage.

Once the buffer is mapped, calls to [getMappedRange\(\)](#) will return an [ArrayBuffer](#) containing the buffer’s current values. Changes to the returned [ArrayBuffer](#) will be stored in the [GPUBuffer](#) after [unmap\(\)](#) is called.

Note: Since the [MAP_WRITE](#) buffer usage may only be combined with the [COPY_SRC](#) buffer usage, mapping for writing can never return values produced by the GPU, and the returned [ArrayBuffer](#) will only ever contain the default initialized data (zeros) or data written by the webpage during a previous mapping.

[GPUBuffer](#) has the following methods:

mapAsync(mode, offset, size)

Maps the given range of the [GPUBuffer](#) and resolves the returned [Promise](#) when the [GPUBuffer](#)’s content is ready to be accessed with [getMappedRange\(\)](#).

The resolution of the returned [Promise](#) **only** indicates that the buffer has been mapped. It does not guarantee the completion of any other operations visible to the [content timeline](#), and in particular does not imply that any other [Promise](#) returned from [onSubmittedWorkDone\(\)](#) or [mapAsync\(\)](#) on other [GPUBuffer](#)s have resolved.

The resolution of the [Promise](#) returned from [onSubmittedWorkDone\(\)](#) **does** imply the completion of [mapAsync\(\)](#) calls made prior to that call, on [GPUBuffer](#)s last used exclusively on that queue.

Called on: [GPUBuffer](#) *this*.

Arguments:

Arguments for the [GPUBuffer.mapAsync\(mode, offset, size\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>mode</i>	GPUMapModeFlags	✗	✗	Whether the buffer should be mapped for reading or writing.
<i>offset</i>	GPUSize64	✗	✓	Offset in bytes into the buffer to the start of the range to map.
<i>size</i>	GPUSize64	✗	✓	Size in bytes of the range to map.

Returns: [Promise](#)<undefined>

[Content timeline](#) steps:

1. Let *contentTimeline* be the current [Content timeline](#).
2. If *this.mapState* is not `"unmapped"`:
 1. Issue the *early-reject steps* on the [Device timeline](#) of *this.[[device]]*.
 2. Return [a promise rejected with OperationError](#).
 3. Let *p* be a new [Promise](#).
 4. Set *this.[[pending_map]]* to *p*.

5. Issue the *validation steps* on the [Device timeline](#) of this.[\[\[device\]\]](#).

6. Return *p*.

[Device timeline](#) *early-reject steps*:

1. [Generate a validation error](#).

2. Return.

[Device timeline](#) *validation steps*:

1. If *size* is **undefined**:

1. Let *rangeSize* be $\max(0, \text{this}.\text{size} - \text{offset})$.

Otherwise:

1. Let *rangeSize* be *size*.

2. If any of the following conditions are unsatisfied:

- *this* must be [valid](#).

1. Set *deviceLost* to **true**.

2. Issue the *map failure steps* on *contentTimeline*.

3. Return.

3. If any of the following conditions are unsatisfied:

- *this*.[\[\[internal state\]\]](#) is "[available](#)".
- *offset* is a multiple of 8.
- *rangeSize* is a multiple of 4.
- $\text{offset} + \text{rangeSize} \leq \text{this}.\text{size}$
- *mode* contains only bits defined in [GPUMapMode](#).
- *mode* contains exactly one of [READ](#) or [WRITE](#).
- If *mode* contains [READ](#) then *this*.*usage* must contain [MAP_READ](#).
- If *mode* contains [WRITE](#) then *this*.*usage* must contain [MAP_WRITE](#).

Then:

1. Set *deviceLost* to **false**.

2. Issue the *map failure steps* on *contentTimeline*.

3. [Generate a validation error](#).

4. Return.

4. Set *this*.[\[\[internal state\]\]](#) to "[unavailable](#)".

Note: Since the buffer is mapped, its contents cannot change between this step and [unmap\(\)](#).

5. When either of the following events occur (whichever comes first), or if either has already occurred:

- The [device timeline](#) becomes informed of the completion of an unspecified [queue timeline](#) point:
- after the completion of currently-enqueued operations that use *this*
- and no later than the completion of all currently-enqueued operations (regardless of whether they use *this*).
- *this*.[\[\[device\]\]](#) becomes [lost](#).

Then issue the subsequent steps on the [device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. Set *deviceLost* to **true** if *this*.[\[\[device\]\]](#) is [lost](#), and **false** otherwise.

Note: The device could have been lost between the previous block of steps and this one.

2. If *deviceLost*:

1. Issue the *map failure steps* on *contentTimeline*.

Otherwise:

1. Let *internalStateAtCompletion* be *this*.[\[\[internal state\]\]](#).

Note: If, and only if, at this point the buffer has become "available" again due to an `unmap()` call, then `[[pending_map]] != p` below, so mapping will not succeed in the steps below.

2. Let `dataForMappedRegion` be the contents of `this` starting at offset `offset`, for `rangeSize` bytes.
3. Issue the `map success steps` on the `contentTimeline`.

[Content timeline](#) `map success steps`:

1. If `this. [[pending_map]] != p`:

Note: The map has been cancelled by `unmap()`.

1. `Assert p` is rejected.
2. Return.
2. `Assert p` is pending.
3. `Assert internalStateAtCompletion` is "unavailable".
4. Let `mapping` be `initialize an active buffer mapping` with mode `mode` and range `[offset, offset + rangeSize]`.

If this allocation fails:

1. Set `this. [[pending_map]]` to `null`, and `reject p` with a `RangeError`.
2. Return.
5. Set the content of `mapping.data` to `dataForMappedRegion`.
6. Set `this. [[mapping]]` to `mapping`.
7. Set `this. [[pending_map]]` to `null`, and `resolve p`.

[Content timeline](#) `map failure steps`:

1. If `this. [[pending_map]] != p`:

Note: The map has been cancelled by `unmap()`.

1. `Assert p` is already rejected.
2. Return.
2. `Assert p` is still pending.
3. Set `this. [[pending_map]]` to `null`.
4. If `deviceLost`:

1. `Reject p` with an `AbortError`.

Note: This is the same error type produced by cancelling the map using `unmap()`.

Otherwise:

1. `Reject p` with an `OperationError`.

`getMappedRange(offset, size)`

Returns an `ArrayBuffer` with the contents of the `GPUBuffer` in the given mapped range.

Called on: `GPUBuffer` `this`.

Arguments:

Arguments for the `GPUBuffer.getMappedRange(offset, size)` method.

Parameter	Type	Nullable	Optional	Description
<code>offset</code>	GPUSize64	✗	✓	Offset in bytes into the buffer to return buffer contents from.
<code>size</code>	GPUSize64	✗	✓	Size in bytes of the ArrayBuffer to return.

Returns: [ArrayBuffer](#)

[Content timeline](#) `steps`:

1. If `size` is missing:
 1. Let `rangeSize` be `max(0, this.size - offset)`.
Otherwise, let `rangeSize` be `size`.
2. If any of the following conditions are unsatisfied, throw an `OperationError` and return.
 - `this. [[mapping]]` is not `null`.

- *offset* is a multiple of 8.
- *rangeSize* is a multiple of 4.
- *offset* \geq *this*.[\[\[mapping\]\].range\[0\]](#).
- *offset* + *rangeSize* \leq *this*.[\[\[mapping\]\].range\[1\]](#).
- (*offset*, *offset* + *rangeSize*) does not overlap another range in *this*.[\[\[mapping\]\].views](#).

Note: It is always valid to get mapped ranges of a [GPUBuffer](#) that is [mappedAtCreation](#), even if it is [invalid](#), because the [Content timeline](#) might not know it is invalid.

- Let *data* be *this*.[\[\[mapping\]\].data](#).
- Let *view* be [! create an ArrayBuffer](#) of size *rangeSize*, but with its pointer mutably referencing the content of *data* at offset (*offset* - [\[\[mapping\]\].range\[0\]](#)).

Note: A [RangeError](#) cannot be thrown here, because the *data* has already been allocated during [mapAsync\(\)](#) or [createBuffer\(\)](#).

- Set *view*.[\[\[ArrayBufferDetachKey\]\]](#) to "WebGPUBufferMapping".

Note: This causes a [TypeError](#) to be thrown if an attempt is made to [DetachArrayBuffer](#), except by [unmap\(\)](#).

- [Append](#) *view* to *this*.[\[\[mapping\]\].views](#).

- Return *view*.

Note: User agents should consider issuing a developer-visible warning if [getMappedRange\(\)](#) succeeds without having checked the status of the map, by waiting for [mapAsync\(\)](#) to succeed, querying a [mapState](#) of "mapped", or waiting for a later [onSubmittedWorkDone\(\)](#) call to succeed.

[unmap\(\)](#)

Unmaps the mapped range of the [GPUBuffer](#) and makes its contents available for use by the GPU again.

Called on: [GPUBuffer](#) *this*.

Returns: [undefined](#)

[Content timeline](#) steps:

- If *this*.[\[\[pending_map\]\]](#) is not null:
 - [Reject](#) *this*.[\[\[pending_map\]\]](#) with an [AbortError](#).
 - Set *this*.[\[\[pending_map\]\]](#) to null.
- If *this*.[\[\[mapping\]\]](#) is null:
 - Return.
- For each [ArrayBuffer](#) *ab* in *this*.[\[\[mapping\]\].views](#):
 - Perform [DetachArrayBuffer](#)(*ab*, "WebGPUBufferMapping").
- Let *bufferUpdate* be null.
- If *this*.[\[\[mapping\]\].mode](#) contains [WRITE](#):
 - Set *bufferUpdate* to { **data**: *this*.[\[\[mapping\]\].data](#), **offset**: *this*.[\[\[mapping\]\].range\[0\]](#) }.

Note: When a buffer is mapped without the [WRITE](#) mode, then unmapped, any local modifications done by the application to the mapped ranges [ArrayBuffer](#) are discarded and will not affect the content of later mappings.

- Set *this*.[\[\[mapping\]\]](#) to null.
- Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

- If any of the following conditions are unsatisfied, return.
 - *this* is [valid to use with](#) *this*.[\[\[device\]\]](#).
- [Assert](#) *this*.[\[\[internal_state\]\]](#) is "unavailable".
- If *bufferUpdate* is not null:
 - Issue the following steps on the [Queue timeline](#) of *this*.[\[\[device\]\].queue](#):

[Queue timeline](#) steps:

- Update the contents of *this* at offset *bufferUpdate.offset* with the data *bufferUpdate.data*.
- Set *this*.[\[\[internal_state\]\]](#) to "available".

6. Textures and Texture Views

6.1. GPUTexture

A *texture* is made up of [1d](#), [2d](#), or [3d](#) arrays of data which can contain multiple values per-element to represent things like colors. Textures can be read and written in many ways, depending on the [GPUTextureUsage](#) they are created with. For example, textures can be sampled, read, and written from render and compute pipeline shaders, and they can be written by render pass outputs. Internally, textures are often stored in GPU memory with a layout optimized for multidimensional access rather than linear access.

One [texture](#) consists of one or more *texture subresources*, each uniquely identified by a [mipmap level](#) and, for [2d](#) textures only, [array layer](#) and [aspect](#).

A [texture subresource](#) is a [subresource](#): each can be used in different [internal usages](#) within a single [usage scope](#).

Each subresource in a *mipmap level* is approximately half the size, in each spatial dimension, of the corresponding resource in the lesser level (see [logical miplevel-specific texture extent](#)). The subresource in level 0 has the dimensions of the texture itself. Smaller levels are typically used to store lower resolution versions of the same image. [GPUSampler](#) and WGSL provide facilities for selecting and interpolating between [levels of detail](#), explicitly or automatically.

A ["2d"](#) texture may be an array of *array layers*. Each subresource in a layer is the same size as the corresponding resources in other layers. For non-2d textures, all subresources have an array layer index of 0.

Each subresource has an *aspect*. Color textures have just one aspect: *color*. [Depth-or-stencil format](#) textures may have multiple aspects: a *depth* aspect, a *stencil* aspect, or both, and may be used in special ways, such as in [depthStencilAttachment](#) and in ["depth"](#) bindings.

A ["3d"](#) texture may have multiple *slices*, each being the two-dimensional image at a particular Z value in the texture. Slices are not separate subresources.

[Exposed=(Window, Worker), SecureContext]

```
interface GPUTexture {  
    GPUTextureView createView(optional GPUTextureViewDescriptor descriptor = {});
```

```
    undefined destroy();
```

```
    readonly attribute GPUIntegerCoordinateOut width;  
    readonly attribute GPUIntegerCoordinateOut height;  
    readonly attribute GPUIntegerCoordinateOut depthOrArrayLayers;  
    readonly attribute GPUIntegerCoordinateOut mipLevelCount;  
    readonly attribute GPUSize32Out sampleCount;  
    readonly attribute GPUTextureDimension dimension;  
    readonly attribute GPUTextureFormat format;  
    readonly attribute GPUFlagsConstant usage;
```

```
};
```

[GPUTexture](#) includes [GPUObjectBase](#);

[GPUTexture](#) has the following [immutable properties](#):

width, of type [GPUIntegerCoordinateOut](#), readonly

The width of this [GPUTexture](#).

height, of type [GPUIntegerCoordinateOut](#), readonly

The height of this [GPUTexture](#).

depthOrArrayLayers, of type [GPUIntegerCoordinateOut](#), readonly

The depth or layer count of this [GPUTexture](#).

mipLevelCount, of type [GPUIntegerCoordinateOut](#), readonly

The number of mip levels of this [GPUTexture](#).

sampleCount, of type [GPUSize32Out](#), readonly

The number of sample count of this [GPUTexture](#).

dimension, of type [GPUTextureDimension](#), readonly

The dimension of the set of texel for each of this [GPUTexture](#)'s subresources.

format, of type [GPUTextureFormat](#), readonly

The format of this [GPUTexture](#).

usage, of type [GPUFlagsConstant](#), readonly

The allowed usages for this [GPUTexture](#).

[[viewFormats]], of type [sequence](#)<[GPUTextureFormat](#)>

The set of [GPUTextureFormat](#)s that can be used as the [GPUTextureViewDescriptor.format](#) when creating views on this [GPUTexture](#).

[GPUTexture](#) has the following [device timeline properties](#):

`[[destroyed]]`, of type [boolean](#), initially `false`

If the texture is destroyed, it can no longer be used in any operation, and its underlying memory can be freed.

compute render extent(baseSize, mipLevel)

Arguments:

[GPUExtent3D](#) *baseSize*

[GPUSize32](#) *mipLevel*

Returns: [GPUExtent3DDict](#)

[Device timeline](#) steps:

Let *extent* be a new [GPUExtent3DDict](#) object.

Set *extent.width* to $\max(1, \text{baseSize.width} \gg \text{mipLevel})$.

Set *extent.height* to $\max(1, \text{baseSize.height} \gg \text{mipLevel})$.

Set *extent.depthOrArrayLayers* to 1.

Return *extent*.

The *logical miplevel-specific texture extent* of a [texture](#) is the size of the [texture](#) in texels at a specific miplevel. It is calculated by this procedure:

Logical miplevel-specific texture extent(descriptor, mipLevel)

Arguments:

[GPUTextureDescriptor](#) *descriptor*

[GPUSize32](#) *mipLevel*

Returns: [GPUExtent3DDict](#)

Let *extent* be a new [GPUExtent3DDict](#) object.

If *descriptor.dimension* is:

"1d"

- Set *extent.width* to $\max(1, \text{descriptor.size.width} \gg \text{mipLevel})$.
- Set *extent.height* to 1.
- Set *extent.depthOrArrayLayers* to 1.

"2d"

- Set *extent.width* to $\max(1, \text{descriptor.size.width} \gg \text{mipLevel})$.
- Set *extent.height* to $\max(1, \text{descriptor.size.height} \gg \text{mipLevel})$.
- Set *extent.depthOrArrayLayers* to *descriptor.size.depthOrArrayLayers*.

"3d"

- Set *extent.width* to $\max(1, \text{descriptor.size.width} \gg \text{mipLevel})$.
- Set *extent.height* to $\max(1, \text{descriptor.size.height} \gg \text{mipLevel})$.
- Set *extent.depthOrArrayLayers* to $\max(1, \text{descriptor.size.depthOrArrayLayers} \gg \text{mipLevel})$.

Return *extent*.

The *physical miplevel-specific texture extent* of a [texture](#) is the size of the [texture](#) in texels at a specific miplevel that includes the possible extra padding to form complete [texel blocks](#) in the [texture](#). It is calculated by this procedure:

Physical miplevel-specific texture extent(descriptor, mipLevel)

Arguments:

[GPUTextureDescriptor](#) *descriptor*

[GPUSize32](#) *mipLevel*

Returns: [GPUExtent3DDict](#)

Let *extent* be a new [GPUExtent3DDict](#) object.

Let *logicalExtent* be [logical miplevel-specific texture extent\(descriptor, mipLevel\)](#).

If *descriptor.dimension* is:

"1d"

- Set *extent.width* to *logicalExtent.width* rounded up to the nearest multiple of *descriptor*'s *texel block width*.
- Set *extent.height* to 1.
- Set *extent.depthOrArrayLayers* to 1.

"2d"

- Set *extent.width* to *logicalExtent.width* rounded up to the nearest multiple of *descriptor*'s *texel block width*.
- Set *extent.height* to *logicalExtent.height* rounded up to the nearest multiple of *descriptor*'s *texel block height*.
- Set *extent.depthOrArrayLayers* to *logicalExtent.depthOrArrayLayers*.

"3d"

- Set *extent.width* to *logicalExtent.width* rounded up to the nearest multiple of *descriptor*'s *texel block width*.
- Set *extent.height* to *logicalExtent.height* rounded up to the nearest multiple of *descriptor*'s *texel block height*.
- Set *extent.depthOrArrayLayers* to *logicalExtent.depthOrArrayLayers*.

Return *extent*.

6.1.1. GPUTextureDescriptor

dictionary [GPUTextureDescriptor](#)

```
: GPUObjectDescriptorBase {  
    required GPUExtent3D size;  
    GPUIntegerCoordinate mipLevelCount = 1;  
    GPUSize32 sampleCount = 1;  
    GPUTextureDimension dimension = "2d";  
    required GPUTextureFormat format;  
    required GPUTextureUsageFlags usage;  
    sequence<GPUTextureFormat> viewFormats = [];  
};
```

[GPUTextureDescriptor](#) has the following members:

size, of type [GPUExtent3D](#)

The width, height, and depth or layer count of the texture.

mipLevelCount, of type [GPUIntegerCoordinate](#), defaulting to 1

The number of mip levels the texture will contain.

sampleCount, of type [GPUSize32](#), defaulting to 1

The sample count of the texture. A *sampleCount* > 1 indicates a multisampled texture.

dimension, of type [GPUTextureDimension](#), defaulting to "2d"

Whether the texture is one-dimensional, an array of two-dimensional layers, or three-dimensional.

format, of type [GPUTextureFormat](#)

The format of the texture.

usage, of type [GPUTextureUsageFlags](#)

The allowed usages for the texture.

viewFormats, of type sequence<[GPUTextureFormat](#)>, defaulting to []

Specifies what view *format* values will be allowed when calling [createView\(\)](#) on this texture (in addition to the texture's actual *format*).

NOTE:

Adding a format to this list may have a significant performance impact, so it is best to avoid adding formats unnecessarily.

The actual performance impact is highly dependent on the target system; developers must test various systems to find out the impact on their particular application. For example, on some systems any texture with a *format* or *viewFormats* entry including "[rgba8unorm-srgb](#)" will perform less optimally than a "[rgba8unorm](#)" texture which does not. Similar caveats exist for other formats and pairs of formats on other systems.

Formats in this list must be [texture view format compatible](#) with the texture format.

Two [GPUTextureFormat](#)s *format* and *viewFormat* are *texture view format compatible* if:

- *format* equals *viewFormat*, or
- *format* and *viewFormat* differ only in whether they are *srgb* formats (have the *-srgb* suffix).

enum

```
GPUTextureDimension
```

```
{  
    "1d",  
    "2d",  
    "3d",  
};
```

"1d"

Specifies a texture that has one dimension, width. "1d" textures cannot have mipmaps, be multisampled, use compressed or depth/stencil formats, or be used as a render target.

"2d"

Specifies a texture that has a width and height, and may have layers.

"3d"

Specifies a texture that has a width, height, and depth. "3d" textures cannot be multisampled, and their format must support 3d textures (all [plain color formats](#) and some [packed/compressed formats](#)).

6.1.2. Texture Usages

typedef [[EnforceRange](#)] unsigned long

```
GPUTextureUsageFlags
```

```
;  
[Exposed=(Window, Worker), SecureContext]  
namespace
```

```
GPUTextureUsage
```

```
{  
    const GPUFlagsConstant COPY_SRC      = 0x01;  
    const GPUFlagsConstant COPY_DST      = 0x02;  
    const GPUFlagsConstant TEXTURE_BINDING = 0x04;  
    const GPUFlagsConstant STORAGE_BINDING = 0x08;  
    const GPUFlagsConstant RENDER_ATTACHMENT = 0x10;  
};
```

The [GPUTextureUsage](#) flags determine how a [GPUTexture](#) may be used after its creation:

COPY_SRC

The texture can be used as the source of a copy operation. (Examples: as the `SOURCE` argument of a [copyTextureToTexture\(\)](#) or [copyTextureToBuffer\(\)](#) call.)

COPY_DST

The texture can be used as the destination of a copy or write operation. (Examples: as the `destination` argument of a [copyTextureToTexture\(\)](#) or [copyBufferToTexture\(\)](#) call, or as the target of a [writeTexture\(\)](#) call.)

TEXTURE_BINDING

The texture can be bound for use as a sampled texture in a shader (Example: as a bind group entry for a [GPUTextureBindingLayout](#).)

STORAGE_BINDING

The texture can be bound for use as a storage texture in a shader (Example: as a bind group entry for a [GPUStorageTextureBindingLayout](#).)

RENDER_ATTACHMENT

The texture can be used as a color or depth/stencil attachment in a render pass. (Example: as a [GPURenderPassColorAttachment.view](#) or [GPURenderPassDepthStencilAttachment.view](#).)

maximum mipLevel count(dimension, size)

Arguments:

[GPUTextureDimension](#) *dimension*

[GPUTextureDimension](#) *size*

Calculate the max dimension value *m*:

If *dimension* is:

["1d"](#)

Return 1.

"2d"

Let $m = \max(\text{size.width}, \text{size.height})$.

"3d"

Let $m = \max(\max(\text{size.width}, \text{size.height}), \text{size.depthOrArrayLayers})$.

Return $\text{floor}(\log_2(m)) + 1$.

6.1.3. Texture Creation

createTexture(descriptor)

Creates a [GPUTexture](#).

Called on: [GPUDevice](#) this.

Arguments:

Arguments for the [GPUDevice.createTexture\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUTextureDescriptor	✗	✗	Description of the GPUTexture to create.

Returns: [GPUTexture](#)

[Content timeline](#) steps:

1. ? [validate GPUExtent3D shape\(descriptor.size\)](#).
2. ? [Validate texture format required features](#) of *descriptor.format* with *this.[[device]]*.
3. ? [Validate texture format required features](#) of each element of *descriptor.viewFormats* with *this.[[device]]*.
4. Let t be ! [create a new WebGPU object](#)(*this*, [GPUTexture](#), *descriptor*).
5. Set $t.\text{width}$ to *descriptor.size.width*.
6. Set $t.\text{height}$ to *descriptor.size.height*.
7. Set $t.\text{depthOrArrayLayers}$ to *descriptor.size.depthOrArrayLayers*.
8. Set $t.\text{mipLevelCount}$ to *descriptor.mipLevelCount*.
9. Set $t.\text{sampleCount}$ to *descriptor.sampleCount*.
10. Set $t.\text{dimension}$ to *descriptor.dimension*.
11. Set $t.\text{format}$ to *descriptor.format*.
12. Set $t.\text{usage}$ to *descriptor.usage*.
13. Issue the *initialization steps* on the [Device timeline](#) of *this*.
14. Return t .

[Device timeline](#) initialization steps:

1. If any of the following conditions are unsatisfied [generate a validation error](#), [invalidate](#) t and return.
 - [validating GPUTextureDescriptor](#)(*this*, *descriptor*) returns **true**.
2. Set $t.[[\text{viewFormats}]]$ to *descriptor.viewFormats*.
3. Create a device allocation for t where each block has an [equivalent texel representation](#) to a block with a bit representation of zero.

If the allocation fails without side-effects, [generate an out-of-memory error](#), [invalidate](#) t , and return.

validating GPUTextureDescriptor(this, descriptor):

Arguments:

[GPUDevice](#) *this*

[GPUTextureDescriptor](#) *descriptor*

[Device timeline](#) steps:

Let *limits* be *this.[[limits]]*.

Return **true** if all of the following requirements are met, and **false** otherwise:

this must not be [lost](#).

`descriptor.usage` must not be 0.

`descriptor.usage` must contain only bits present in this's [allowed texture usages](#).

`descriptor.size.width`, `descriptor.size.height`, and `descriptor.size.depthOrArrayLayers` must be > zero.

`descriptor.mipLevelCount` must be > zero.

`descriptor.sampleCount` must be either 1 or 4.

If `descriptor.dimension` is:

"1d"

- `descriptor.size.width` must be \leq `limits.maxTextureDimension1D`.
- `descriptor.size.height` must be 1.
- `descriptor.size.depthOrArrayLayers` must be 1.
- `descriptor.sampleCount` must be 1.
- `descriptor.format` must not be a [compressed format](#) or [depth-or-stencil format](#).

"2d"

- `descriptor.size.width` must be \leq `limits.maxTextureDimension2D`.
- `descriptor.size.height` must be \leq `limits.maxTextureDimension2D`.
- `descriptor.size.depthOrArrayLayers` must be \leq `limits.maxTextureArrayLayers`.

"3d"

- `descriptor.size.width` must be \leq `limits.maxTextureDimension3D`.
- `descriptor.size.height` must be \leq `limits.maxTextureDimension3D`.
- `descriptor.size.depthOrArrayLayers` must be \leq `limits.maxTextureDimension3D`.
- `descriptor.sampleCount` must be 1.
- `descriptor.format` must support "3d" textures according to [§ 26.1 Texture Format Capabilities](#).

`descriptor.size.width` must be multiple of [texel block width](#).

`descriptor.size.height` must be multiple of [texel block height](#).

If `descriptor.sampleCount` > 1:

`descriptor.mipLevelCount` must be 1.

`descriptor.size.depthOrArrayLayers` must be 1.

`descriptor.usage` must not include the [STORAGE_BINDING](#) bit.

`descriptor.usage` must include the [RENDER_ATTACHMENT](#) bit.

`descriptor.format` must support multisampling according to [§ 26.1 Texture Format Capabilities](#).

`descriptor.mipLevelCount` must be \leq [maximum mipLevel count](#)(`descriptor.dimension`, `descriptor.size`).

If `descriptor.usage` includes the [RENDER_ATTACHMENT](#) bit:

`descriptor.format` must be a [renderable format](#).

`descriptor.dimension` must be either "2d" or "3d".

If `descriptor.usage` includes the [STORAGE_BINDING](#) bit:

`descriptor.format` must be listed in [§ 26.1.1 Plain color formats](#) table with [STORAGE_BINDING](#) capability for at least one access mode.

For each `viewFormat` in `descriptor.viewFormats`, `descriptor.format` and `viewFormat` must be [texture view format compatible](#).

NOTE:

Implementations may consider issuing a developer-visible warning if `viewFormat` is not compatible with any of the given [usage](#) bits, as that `viewFormat` will be unusable.

Creating a 16x16, RGBA, 2D texture with one array layer and one mip level:

```
const texture = gpuDevice.createTexture({
    size: { width: 16, height: 16 },
    format: 'rgba8unorm',
    usage: GPUTextureUsage.TEXTURE_BINDING,
});
```

6.1.4. Texture Destruction

An application that no longer requires a [GPUTexture](#) can choose to lose access to it before garbage collection by calling [destroy\(\)](#).

Note: This allows the user agent to reclaim the GPU memory associated with the [GPUTexture](#) once all previously submitted operations using it are complete.

[GPUTexture](#) has the following methods:

destroy()

Destroys the [GPUTexture](#).

Called on: [GPUTexture](#) *this*.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [device timeline](#).

[Device timeline](#) steps:

1. Set *this*.[\[\[destroyed\]\]](#) to true.

6.2. GPUTextureView

A [GPUTextureView](#) is a view onto some subset of the [texture subresources](#) defined by a particular [GPUTexture](#).

[Exposed=(Window, Worker), [SecureContext](#)]

interface [GPUTextureView](#) {

};

[GPUTextureView](#) includes [GPUObjectBase](#);

[GPUTextureView](#) has the following [immutable properties](#):

[\[\[texture\]\]](#), readonly

The [GPUTexture](#) into which this is a view.

[\[\[descriptor\]\]](#), readonly

The [GPUTextureViewDescriptor](#) describing this texture view.

All optional fields of [GPUTextureViewDescriptor](#) are defined.

[\[\[renderExtent\]\]](#), readonly

For renderable views, this is the effective [GPUExtent3DDict](#) for rendering.

Note: this extent depends on the [baseMipLevel](#).

The set of *subresources* of a texture view *view*, with [\[\[descriptor\]\]](#) *desc*, is the subset of the subresources of *view*.[\[\[texture\]\]](#) for which each subresource *s* satisfies the following:

The [mipmap level](#) of *s* is $\geq \text{desc}.\text{baseMipLevel}$ and $< \text{desc}.\text{baseMipLevel} + \text{desc}.\text{mipLevelCount}$.

The [array layer](#) of *s* is $\geq \text{desc}.\text{baseArrayLayer}$ and $< \text{desc}.\text{baseArrayLayer} + \text{desc}.\text{arrayLayerCount}$.

The [aspect](#) of *s* is in the [set of aspects](#) of *desc*.[aspect](#).

Two [GPUTextureView](#) objects are *texture-view-aliasing* if and only if their sets of subresources intersect.

6.2.1. Texture View Creation

dictionary

[GPUTextureViewDescriptor](#)

: [GPUObjectDescriptorBase](#) {

[GPUTextureFormat](#) *format*;

[GPUTextureViewDimension](#) *dimension*;

[GPUTextureUsageFlags](#) *usage* = 0;

[GPUTextureAspect](#) *aspect* = "all";

[GPUIntegerCoordinate](#) *baseMipLevel* = 0;

[GPUIntegerCoordinate](#) *mipLevelCount*;

[GPUIntegerCoordinate](#) *baseArrayLayer* = 0;

[GPUIntegerCoordinate](#) *arrayLayerCount*;

};

[GPUTextureViewDescriptor](#) has the following members:

format, of type [GPUTextureFormat](#)

The format of the texture view. Must be either the [format](#) of the texture or one of the [viewFormats](#) specified during its creation.

dimension, of type [GPUTextureViewDimension](#)

The dimension to view the texture as.

usage, of type [GPUTextureUsageFlags](#), defaulting to 0

The allowed [usage\(s\)](#) for the texture view. Must be a subset of the [usage](#) flags of the texture. If 0, defaults to the full set of [usage](#) flags of the texture.

Note: If the view's [format](#) doesn't support all of the texture's [usages](#), the default will fail, and the view's [usage](#) must be specified explicitly.

aspect, of type [GPUTextureAspect](#), defaulting to "all"

Which [aspect\(s\)](#) of the texture are accessible to the texture view.

baseMipLevel, of type [GPUIntegerCoordinate](#), defaulting to 0

The first (most detailed) mipmap level accessible to the texture view.

mipLevelCount, of type [GPUIntegerCoordinate](#)

How many mipmap levels, starting with [baseMipLevel](#), are accessible to the texture view.

baseArrayLayer, of type [GPUIntegerCoordinate](#), defaulting to 0

The index of the first array layer accessible to the texture view.

arrayLayerCount, of type [GPUIntegerCoordinate](#)

How many array layers, starting with [baseArrayLayer](#), are accessible to the texture view.

enum

[GPUTextureViewDimension](#)

```
{  
  "1d",  
  "2d",  
  "2d-array",  
  "cube",  
  "cube-array",  
  "3d",  
};
```

"1d"

The texture is viewed as a 1-dimensional image.

Corresponding WGSL types:

- texture_1d
- texture_storage_1d

"2d"

The texture is viewed as a single 2-dimensional image.

Corresponding WGSL types:

- texture_2d
- texture_storage_2d
- texture_multisampled_2d
- texture_depth_2d
- texture_depth_multisampled_2d

"2d-array"

The texture view is viewed as an array of 2-dimensional images.

Corresponding WGSL types:

- texture_2d_array
- texture_storage_2d_array
- texture_depth_2d_array

"cube"

The texture is viewed as a cubemap.

The view has 6 array layers, each corresponding to a face of the cube in the order [+X, -X, +Y, -Y, +Z, -Z] and the following orientations:

Cubemap faces. The +U/+V axes indicate the individual faces' texture coordinates, and thus the [texel copy](#) memory layout of each face.

Note: When viewed from the inside, this results in a left-handed coordinate system where +X is right, +Y is up, and +Z is forward.

Sampling is done seamlessly across the faces of the cubemap.

Corresponding WGSL types:

- `texture_cube`
- `texture_depth_cube`

"cube-array"

The texture is viewed as a packed array of *n* cubemaps, each with 6 array layers behaving like one ["cube"](#) view, for 6*n* array layers in total.

Corresponding WGSL types:

- `texture_cube_array`
- `texture_depth_cube_array`

"3d"

The texture is viewed as a 3-dimensional image.

Corresponding WGSL types:

- `texture_3d`
- `texture_storage_3d`

Each `GPUTextureAspect` value corresponds to a set of [aspects](#). The *set of aspects* are defined for each value below.

```
enum GPUTextureAspect {  
    "all",  
    "stencil-only",  
    "depth-only",  
};
```

"all"

All available aspects of the texture format will be accessible to the texture view. For color formats the color aspect will be accessible. For [combined depth-stencil formats](#) both the depth and stencil aspects will be accessible. [Depth-or-stencil formats](#) with a single aspect will only make that aspect accessible.

The [set of aspects](#) is [[color](#), [depth](#), [stencil](#)].

"stencil-only"

Only the stencil aspect of a [depth-or-stencil format](#) format will be accessible to the texture view.

The [set of aspects](#) is [[stencil](#)].

"depth-only"

Only the depth aspect of a [depth-or-stencil format](#) format will be accessible to the texture view.

The [set of aspects](#) is [[depth](#)].

`createView(descriptor)`

Creates a [GPUTextureView](#).

NOTE:

By default [createView\(\)](#) will create a view with a dimension that can represent the entire texture. For example, calling [createView\(\)](#) without specifying a [dimension](#) on a ["2d"](#) texture with more than one layer will create a ["2d-array"](#) [GPUTextureView](#), even if an [arrayLayerCount](#) of 1 is specified.

For textures created from sources where the layer count is unknown at the time of development it is recommended that calls to [createView\(\)](#) are provided with an explicit [dimension](#) to ensure shader compatibility.

Called on: [GPUTexture](#) *this*.

Arguments:

Arguments for the [GPUTexture.createView\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUTextureViewDescriptor	✗	✓	Description of the GPUTextureView to create.

Returns: *view*, of type [GPUTextureView](#).

[Content timeline](#) steps:

1. [? Validate texture format required features](#) of `descriptor.format` with `this.[[device]]`.
2. Let `view` be [! create a new WebGPU object](#)(`this`, [GPUTextureView](#), `descriptor`).
3. Issue the *initialization steps* on the [Device timeline](#) of `this`.
4. Return `view`.

[Device timeline](#) *initialization steps*:

1. Set `descriptor` to the result of [resolving GPUTextureViewDescriptor defaults](#) for `this` with `descriptor`.
2. If any of the following conditions are unsatisfied [generate a validation error](#), [invalidate](#) `view` and return.
 - `this` is [valid to use with](#) `this.[[device]]`.
 - `descriptor.aspect` must be present in `this.format`.
 - If the `descriptor.aspect` is `"all"`:
 - `descriptor.format` must equal either `this.format` or one of the formats in `this.[[viewFormats]]`.Otherwise:
 - `descriptor.format` must equal the result of [resolving GPUTextureAspect](#)(`this.format`, `descriptor.aspect`).
 - `descriptor.usage` must be a subset of `this.usage`.
 - If `descriptor.usage` includes the `RENDER_ATTACHMENT` bit:
 - `descriptor.format` must be a [renderable format](#).
 - If `descriptor.usage` includes the `STORAGE_BINDING` bit:
 - `descriptor.format` must be listed in [§ 26.1.1 Plain color formats](#) table with `STORAGE_BINDING` capability for at least one access mode.
 - `descriptor.mipLevelCount` must be > 0 .
 - `descriptor.baseMipLevel` + `descriptor.mipLevelCount` must be \leq `this.mipLevelCount`.
 - `descriptor.arrayLayerCount` must be > 0 .
 - `descriptor.baseArrayLayer` + `descriptor.arrayLayerCount` must be \leq the [array layer count](#) of `this`.
 - If `this.sampleCount` > 1 , `descriptor.dimension` must be `"2d"`.
 - If `descriptor.dimension` is:
 - `"1d"`
 - `this.dimension` must be `"1d"`.
 - `descriptor.arrayLayerCount` must be 1.
 - `"2d"`
 - `this.dimension` must be `"2d"`.
 - `descriptor.arrayLayerCount` must be 1.
 - `"2d-array"`
 - `this.dimension` must be `"2d"`.
 - `"cube"`
 - `this.dimension` must be `"2d"`.
 - `descriptor.arrayLayerCount` must be 6.
 - `this.width` must equal `this.height`.
 - `"cube-array"`
 - `this.dimension` must be `"2d"`.
 - `descriptor.arrayLayerCount` must be a multiple of 6.
 - `this.width` must equal `this.height`.
 - `"3d"`
 - `this.dimension` must be `"3d"`.
 - `descriptor.arrayLayerCount` must be 1.
3. Let `view` be a new [GPUTextureView](#) object.
4. Set `view.[[texture]]` to `this`.

5. Set `view.\[\[descriptor\]\]` to `descriptor`.
6. If `descriptor.usage` contains `RENDER_ATTACHMENT`:
 1. Let `renderExtent` be `compute render extent([this.width, this.height, this.depthOrArrayLayers], descriptor.baseMipLevel)`.
 2. Set `view.\[\[renderExtent\]\]` to `renderExtent`.

When resolving `GPUTextureViewDescriptor` defaults for `GPUTextureView` texture with a `GPUTextureViewDescriptor` descriptor, run the following [device timeline](#) steps:

Let `resolved` be a copy of `descriptor`.

If `resolved.format` is not [provided](#):

Let `format` be the result of [resolving GPUTextureAspect](#)(`format`, `descriptor.aspect`).

If `format` is `null`:

Set `resolved.format` to `texture.format`.

Otherwise:

Set `resolved.format` to `format`.

If `resolved.mipLevelCount` is not [provided](#): set `resolved.mipLevelCount` to `texture.mipLevelCount – resolved.baseMipLevel`.

If `resolved.dimension` is not [provided](#) and `texture.dimension` is:

["1d"](#)

Set `resolved.dimension` to ["1d"](#).

["2d"](#)

If the `array layer count` of `texture` is 1:

- Set `resolved.dimension` to ["2d"](#).

Otherwise:

- Set `resolved.dimension` to ["2d-array"](#).

["3d"](#)

Set `resolved.dimension` to ["3d"](#).

If `resolved.arrayLayerCount` is not [provided](#) and `resolved.dimension` is:

["1d"](#), ["2d"](#), or ["3d"](#)

Set `resolved.arrayLayerCount` to 1.

["cube"](#)

Set `resolved.arrayLayerCount` to 6.

["2d-array"](#) or ["cube-array"](#)

Set `resolved.arrayLayerCount` to the `array layer count` of `texture` – `resolved.baseArrayLayer`.

If `resolved.usage` is 0: set `resolved.usage` to `texture.usage`.

Return `resolved`.

To determine the `array layer count` of `GPUTexture` texture, run the following steps:

If `texture.dimension` is:

["1d"](#) or ["3d"](#)

Return 1.

["2d"](#)

Return `texture.depthOrArrayLayers`.

6.3. Texture Formats

The name of the format specifies the order of components, bits per component, and data type for the component.

r, g, b, a = red, green, blue, alpha

unorm = unsigned normalized

snorm = signed normalized

uint = unsigned int

`sint` = signed int

`float` = floating point

If the format has the `-srgb` suffix, then sRGB conversions from gamma to linear and vice versa are applied during the reading and writing of color values in the shader. Compressed texture formats are provided by [features](#). Their naming should follow the convention here, with the texture name as a prefix. e.g. `etc2-rgba8unorm`.

The *texel block* is a single addressable element of the textures in pixel-based [GPUTextureFormats](#), and a single compressed block of the textures in block-based compressed [GPUTextureFormats](#).

The *texel block width* and *texel block height* specifies the dimension of one [texel block](#).

For pixel-based [GPUTextureFormats](#), the [texel block width](#) and [texel block height](#) are always 1.

For block-based compressed [GPUTextureFormats](#), the [texel block width](#) is the number of texels in each row of one [texel block](#), and the [texel block height](#) is the number of texel rows in one [texel block](#). See [§ 26.1 Texture Format Capabilities](#) for an exhaustive list of values for every texture format.

The *texel block copy footprint* of an [aspect](#) of a [GPUTextureFormat](#) is the number of bytes one texel block occupies during a [texel copy](#), if applicable.

Note: The *texel block memory cost* of a [GPUTextureFormat](#) is the number of bytes needed to store one [texel block](#). It is not fully defined for all formats. **This value is informative and non-normative.**

```
enum
GPUTextureFormat
{
    // 8-bit formats

    "r8unorm"

    ,

    "r8snorm"

    ,

    "r8uint"

    ,

    "r8sint"

    ,

    // 16-bit formats

    "r16unorm"

    ,

    "r16snorm"

    ,

    "r16uint"

    ,

    "r16sint"

    ,

    "r16float"

    ,

    "rg8unorm"

    ,

    "rg8snorm"
```

"rg8uint"

,

"rg8sint"

,

// 32-bit formats

"r32uint"

,

"r32sint"

,

"r32float"

,

"rg16unorm"

,

"rg16snorm"

,

"rg16uint"

,

"rg16sint"

,

"rg16float"

,

"rgba8unorm"

,

"rgba8unorm-srgb"

,

"rgba8snorm"

,

"rgba8uint"

,

"rgba8sint"

,

"bgra8unorm"

,

"bgra8unorm-srgb"

,

// Packed 32-bit formats

"*rgb9e5ufloat*"

,

"*rgb10a2uint*"

,

"*rgb10a2unorm*"

,

"*rg11b10ufloat*"

,

// 64-bit formats

"*rg32uint*"

,

"*rg32sint*"

,

"*rg32float*"

,

"*rgba16unorm*"

,

"*rgba16snorm*"

,

"*rgba16uint*"

,

"*rgba16sint*"

,

"*rgba16float*"

,

// 128-bit formats

"*rgba32uint*"

,

"*rgba32sint*"

,

"*rgba32float*"

,

// Depth/stencil formats

"*stencil8*"

,
"depth16unorm"

,
"depth24plus"

,
"depth24plus-stencil8"

,
"depth32float"

,
// "depth32float-stencil8" feature

"depth32float-stencil8"

,
// BC compressed formats usable if "texture-compression-bc" is both
// supported by the device/user agent and enabled in requestDevice.

"bc1-rgba-unorm"

,
"bc1-rgba-unorm-srgb"

,
"bc2-rgba-unorm"

,
"bc2-rgba-unorm-srgb"

,
"bc3-rgba-unorm"

,
"bc3-rgba-unorm-srgb"

,
"bc4-r-unorm"

,
"bc4-r-snorm"

,
"bc5-rg-unorm"

,
"bc5-rg-snorm"

,
"bc6h-rgb-ufloat"

"bc6h-rgb-float"

,

"bc7-rgba-unorm"

,

"bc7-rgba-unorm-srgb"

,

// ETC2 compressed formats usable if "texture-compression-etc2" is both
// supported by the device/user agent and enabled in requestDevice.

"etc2-rgb8unorm"

,

"etc2-rgb8unorm-srgb"

,

"etc2-rgb8a1unorm"

,

"etc2-rgb8a1unorm-srgb"

,

"etc2-rgba8unorm"

,

"etc2-rgba8unorm-srgb"

,

"eac-r11unorm"

,

"eac-r11snorm"

,

"eac-rg11unorm"

,

"eac-rg11snorm"

,

// ASTC compressed formats usable if "texture-compression-astc" is both
// supported by the device/user agent and enabled in requestDevice.

"astc-4x4-unorm"

,

"astc-4x4-unorm-srgb"

,

"astc-5x4-unorm"

,

"astc-5x4-unorm-srgb"

,

"astc-5x5-unorm"

,

"astc-5x5-unorm-srgb"

,

"astc-6x5-unorm"

,

"astc-6x5-unorm-srgb"

,

"astc-6x6-unorm"

,

"astc-6x6-unorm-srgb"

,

"astc-8x5-unorm"

,

"astc-8x5-unorm-srgb"

,

"astc-8x6-unorm"

,

"astc-8x6-unorm-srgb"

,

"astc-8x8-unorm"

,

"astc-8x8-unorm-srgb"

,

"astc-10x5-unorm"

,

"astc-10x5-unorm-srgb"

,

"astc-10x6-unorm"

,

"astc-10x6-unorm-srgb"

,

"astc-10x8-unorm"

```
,
"astc-10x8-unorm-srgb"
```

```
,
"astc-10x10-unorm"
```

```
,
"astc-10x10-unorm-srgb"
```

```
,
"astc-12x10-unorm"
```

```
,
"astc-12x10-unorm-srgb"
```

```
,
"astc-12x12-unorm"
```

```
,
"astc-12x12-unorm-srgb"
```

```
};
```

The depth component of the ["depth24plus"](#) and ["depth24plus-stencil8"](#) formats may be implemented as either a [24-bit depth](#) value or a ["depth32float"](#) value.

The [stencil8](#) format may be implemented as either a real "stencil8", or "depth24stencil8", where the depth aspect is hidden and inaccessible.

NOTE:

While the precision of depth32float channels is strictly higher than the precision of [24-bit depth](#) channels for all values in the representable range (0.0 to 1.0), note that the set of representable values is not an exact superset.

For [24-bit depth](#), 1 ULP has a constant value of $1 / (2^{24} - 1)$.

For depth32float, 1 ULP has a variable value no greater than $1 / (2^{24})$.

A format is *renderable* if it is either a *color renderable format*, or a [depth-or-stencil format](#). If a format is listed in [§ 26.1.1 Plain color formats](#) with [RENDER_ATTACHMENT](#) capability, it is a color renderable format. Any other format is not a color renderable format. All [depth-or-stencil formats](#) are renderable.

A [renderable format](#) is also *blendable* if it can be used with render pipeline blending. See [§ 26.1 Texture Format Capabilities](#).

A format is *filterable* if it supports the [GPUTextureSampleType "float"](#) (not just ["unfilterable-float"](#)); that is, it can be used with ["filtering"](#) [GPUSamplers](#). See [§ 26.1 Texture Format Capabilities](#).

resolving GPUTextureAspect(format, aspect)

Arguments:

[GPUTextureFormat](#) *format*

[GPUTextureAspect](#) *aspect*

Returns: [GPUTextureFormat](#) or null

If *aspect* is:

["all"](#)

Return *format*.

["depth-only"](#)

["stencil-only"](#)

If *format* is a depth-stencil-format: Return the [aspect-specific format](#) of *format* according to [§ 26.1.2 Depth-stencil formats](#) or null if the aspect is not present in *format*.

Return null.

Use of some texture formats require a feature to be enabled on the [GPUDevice](#). Because new formats can be added to the specification, those enum values might not be

known by the implementation. In order to normalize behavior across implementations, attempting to use a format that requires a feature will throw an exception if the associated feature is not enabled on the device. This makes the behavior the same as when the format is unknown to the implementation.

See [§ 26.1 Texture Format Capabilities](#) for information about which [GPUTextureFormat](#)s require features.

To *Validate texture format required features* of a [GPUTextureFormat](#) format with logical [device](#) device, run the following [content timeline](#) steps:

If *format* requires a feature and *device*.[\[\[features\]\]](#) does not [contain](#) the feature:

Throw a [TypeError](#).

6.4. GPUExternalTexture

A [GPUExternalTexture](#) is a sampleable 2D texture wrapping an external video frame. It is an immutable snapshot; its contents cannot change over time, either from inside WebGPU (it is only sampleable) or from outside WebGPU (e.g. due to video frame advancement).

[GPUExternalTexture](#)s can be bound into bind groups via the [externalTexture](#) bind group layout entry member. Note that member uses several binding slots, as defined there.

NOTE:

[GPUExternalTexture](#) *can* be implemented without creating a copy of the imported source, but this depends [implementation-defined](#) factors. Ownership of the underlying representation may either be exclusive or shared with other owners (such as a video decoder), but this is not visible to the application.

The underlying representation of an external texture is unobservable (except for precise sampling behavior), but typically may include:

Up to three 2D planes of data (e.g. RGBA, Y+UV, Y+U+V).

Metadata for converting coordinates before reading from those planes (crop and rotation).

Metadata for converting values into the specified output color space (matrices, gammas, 3D LUT).

The configuration used internally by an implementation may be inconsistent across time, systems, user agents, media sources, or even frames within a single video source. In order to account for many possible representations, the binding conservatively uses the following, for *each* external texture:

three sampled texture bindings (for up to 3 planes),

one sampled texture binding for a 3D LUT,

one sampler binding to sample the 3D LUT, and

one uniform buffer binding for metadata.

[\[Exposed=\(Window, Worker\), SecureContext\]](#)

```
interface GPUExternalTexture {
```

```
};
```

[GPUExternalTexture](#) includes [GPUObjectBase](#);

[GPUExternalTexture](#) has the following [immutable properties](#):

[\[\[descriptor\]\]](#), of type [GPUExternalTextureDescriptor](#), readonly

The descriptor with which the texture was created.

[GPUExternalTexture](#) has the following [immutable properties](#):

[\[\[expired\]\]](#), of type [boolean](#), initially **false**

Indicates whether the object has expired (can no longer be used).

Note: Unlike [\[\[destroyed\]\]](#) slots, which are similar, this can change from **true** back to **false**.

6.4.1. Importing External Textures

An external texture is created from an external video object using [importExternalTexture\(\)](#).

An external texture created from an [HTMLVideoElement](#) expires (is destroyed) automatically in a task after it is imported, instead of manually or upon garbage collection like other resources. When an external texture expires, its [\[\[expired\]\]](#) slot changes to **true**.

An external texture created from a [VideoFrame](#) expires (is destroyed) when, and only when, the source [VideoFrame](#) is [closed](#), either explicitly by [close\(\)](#), or by other means.

Note: As noted in [decode\(\)](#), authors **should** call [close\(\)](#) on output [VideoFrames](#) to avoid decoder stalls. If an imported [VideoFrame](#) is dropped without being closed, the imported [GPUExternalTexture](#) object will keep it alive until it is also dropped. The [VideoFrame](#) cannot be garbage collected until both objects are dropped. Garbage collection is unpredictable, so this may still stall the video decoder.

Once the [GPUExternalTexture](#) expires, [importExternalTexture\(\)](#) must be called again. However, the user agent may un-expire and return the same [GPUExternalTexture](#) again, instead of creating a new one. This will commonly happen unless the execution of the application is scheduled to match the video's frame

rate (e.g. using `requestVideoFrameCallback()`). If the same object is returned again, it will compare equal, and [GPUBindGroup](#), [GPURenderBundles](#), etc. referencing the previous object can still be used.

dictionary

GPUExternalTextureDescriptor

```
: GPUObjectDescriptorBase {
  required (HTMLVideoElement or VideoFrame) source;
  PredefinedColorSpace colorSpace = "srgb";
};
```

[GPUExternalTextureDescriptor](#) dictionaries have the following members:

- source**, of type ([HTMLVideoElement](#) or [VideoFrame](#))
The video source to import the external texture from. Source size is determined as described by the [external source dimensions](#) table.
- colorSpace**, of type [PredefinedColorSpace](#), defaulting to "srgb"
The color space the image contents of [source](#) will be converted into when reading.

importExternalTexture(descriptor)
Creates a [GPUExternalTexture](#) wrapping the provided image source.

Called on: [GPUDevice](#) *this*.

Arguments:

Arguments for the [GPUDevice.importExternalTexture\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUExternalTextureDescriptor	✗	✗	Provides the external image source object (and any creation options).

Returns: [GPUExternalTexture](#)

[Content timeline](#) steps:

- Let *source* be *descriptor*.[source](#).
- If the current image contents of *source* are the same as the most recent [importExternalTexture\(\)](#) call with the same *descriptor* (ignoring [label](#)), and the user agent chooses to reuse it:
 - Let *previousResult* be the [GPUExternalTexture](#) returned previously.
 - Set *previousResult*.[\[\[expired\]\]](#) to `false`, renewing ownership of the underlying resource.
 - Let *result* be *previousResult*.

Note: This allows the application to detect duplicate imports and avoid re-creating dependent objects (such as [GPUBindGroups](#)). Implementations still need to be able to handle a single frame being wrapped by multiple [GPUExternalTexture](#), since import metadata like [colorSpace](#) can change even for the same frame.

Otherwise:

- If *source* is `is not origin-clean`, throw a [SecurityError](#) and return.
- Let *usability* be [? check the usability of the image argument\(source\)](#).
- If *usability* is not `good`:
 - [Generate a validation error](#).
 - Return an `invalidated` [GPUExternalTexture](#).
- Let *data* be the result of converting the current image contents of *source* into the color space *descriptor*.[colorSpace](#) with unpremultiplied alpha.

This [may result](#) in values outside of the range [0, 1]. If clamping is desired, it may be performed after sampling.

Note: This is described like a copy, but may be implemented as a reference to read-only underlying data plus appropriate metadata to perform conversion later.

- Let *result* be a new [GPUExternalTexture](#) object wrapping *data*.
- If *source* is an [HTMLVideoElement](#), [queue an automatic expiry task](#) with device *this* and the following steps:
 - Set *result*.[\[\[expired\]\]](#) to `true`, releasing ownership of the underlying resource.

Note: An [HTMLVideoElement](#) should be imported in the same task that samples the texture (which should generally be scheduled using `requestVideoFrameCallback` or `requestAnimationFrame()` depending on the application). Otherwise, a texture could get destroyed by these steps before the application is finished using it.

- If *source* is a [VideoFrame](#), then when *source* is `closed`, run the following steps:

1. Set `result.\[\[expired\]\]` to `true`.
5. Set `result.label` to `descriptor.label`.
6. Return `result`.

Rendering using an video element external texture at the page animation frame rate:

```
const videoElement = document.createElement('video');
// ... set up videoElement, wait for it to be ready...

function frame() {
    requestAnimationFrame(frame);

    // Always re-import the video on every animation frame, because the
    // import is likely to have expired.
    // The browser may cache and reuse a past frame, and if it does it
    // may return the same GPUExternalTexture object again.
    // In this case, old bind groups are still valid.
    const externalTexture = gpuDevice.importExternalTexture({
        source: videoElement
    });

    // ... render using externalTexture...
}
```

Rendering using an video element external texture at the video's frame rate, if `requestVideoFrameCallback` is available:

```
const videoElement = document.createElement('video');
// ... set up videoElement...

function frame() {
    videoElement.requestVideoFrameCallback(frame);

    // Always re-import, because we know the video frame has advanced
    const externalTexture = gpuDevice.importExternalTexture({
        source: videoElement
    });

    // ... render using externalTexture...
}

videoElement.requestVideoFrameCallback(frame);
```

6.5. Sampling External Texture Bindings

The [externalTexture](#) binding point allows binding [GPUExternalTexture](#) objects (from dynamic image sources like videos). It also supports [GPUTexture](#) and [GPUTextureView](#).

Note: When a [GPUTexture](#) or a [GPUTextureView](#) is bound to an [externalTexture](#) binding, it is like a [GPUExternalTexture](#) with a single RGBA plane and no crop, rotation, or color conversion.

External textures are represented in WGSL with `texture_external` and may be read using `textureLoad` and `textureSampleBaseClampToEdge`.

The `sampler` provided to `textureSampleBaseClampToEdge` is used to sample the underlying textures.

When the [binding resource type](#) is a [GPUExternalTexture](#), the result is in the color space set by [colorSpace](#). It is implementation-dependent whether, for any given external texture, the sampler (and filtering) is applied before or after conversion from underlying values into the specified color space.

Note: If the internal representation is an RGBA plane, sampling behaves as on a regular 2D texture. If there are several underlying planes (e.g. Y+UV), the sampler is used to sample each underlying texture separately, prior to conversion from YUV to the specified color space.

7. Samplers

7.1. GPUSampler

A [GPUSampler](#) encodes transformations and filtering information that can be used in a shader to interpret texture resource data.

[GPUSampler](#)s are created via [createSampler\(\)](#).

[Exposed=(Window, Worker), SecureContext]

```
interface GPUSampler {  
};
```

GPUSampler includes GPUObjectBase;

GPUSampler has the following immutable properties:

[[descriptor]], of type GPUSamplerDescriptor, readonly

The GPUSamplerDescriptor with which the GPUSampler was created.

[[isComparison]], of type boolean, readonly

Whether the GPUSampler is used as a comparison sampler.

[[isFiltering]], of type boolean, readonly

Whether the GPUSampler weights multiple samples of a texture.

7.1.1. GPUSamplerDescriptor

A GPUSamplerDescriptor specifies the options to use to create a GPUSampler.

dictionary

GPUSamplerDescriptor

```
    : GPUObjectDescriptorBase {  
    GPUAddressMode addressModeU = "clamp-to-edge";  
    GPUAddressMode addressModeV = "clamp-to-edge";  
    GPUAddressMode addressModeW = "clamp-to-edge";  
    GPUFilterMode magFilter = "nearest";  
    GPUFilterMode minFilter = "nearest";  
    GPUTextureFilterMode mipmapFilter = "nearest";  
    float lodMinClamp = 0;  
    float lodMaxClamp = 32;  
    GPUCompareFunction compare;  
    [Clamp] unsigned short maxAnisotropy = 1;  
};
```

addressModeU, of type GPUAddressMode, defaulting to "clamp-to-edge"

addressModeV, of type GPUAddressMode, defaulting to "clamp-to-edge"

addressModeW, of type GPUAddressMode, defaulting to "clamp-to-edge"

Specifies the address modes for the texture width, height, and depth coordinates, respectively.

magFilter, of type GPUFilterMode, defaulting to "nearest"

Specifies the sampling behavior when the sampled area is smaller than or equal to one texel.

minFilter, of type GPUFilterMode, defaulting to "nearest"

Specifies the sampling behavior when the sampled area is larger than one texel.

mipmapFilter, of type GPUTextureFilterMode, defaulting to "nearest"

Specifies behavior for sampling between mipmap levels.

lodMinClamp, of type float, defaulting to 0

lodMaxClamp, of type float, defaulting to 32

Specifies the minimum and maximum levels of detail, respectively, used internally when sampling a texture.

compare, of type GPUCompareFunction

When provided the sampler will be a comparison sampler with the specified GPUCompareFunction.

Note: Comparison samplers may use filtering, but the sampling results will be implementation-dependent and may differ from the normal filtering rules.

maxAnisotropy, of type unsigned short, defaulting to 1

Specifies the maximum anisotropy value clamp used by the sampler. Anisotropic filtering is enabled when maxAnisotropy is > 1 and the implementation supports it.

Anisotropic filtering improves the image quality of textures sampled at oblique viewing angles. Higher maxAnisotropy values indicate the maximum ratio of anisotropy supported when filtering.

NOTE:

Most implementations support maxAnisotropy values in range between 1 and 16, inclusive. The used value of maxAnisotropy will be clamped to the maximum value that the platform supports.

The precise filtering behavior is implementation-dependent.

Level of detail (LOD) describes which mip level(s) are selected when sampling a texture. It may be specified explicitly through shader methods like [textureSampleLevel](#) or implicitly determined from the [texture coordinate](#) derivatives.

Note: See [Scale Factor Operation, LOD Operation and Image Level Selection](#) in the [Vulkan 1.3](#) spec for an example of how implicit LODs may be calculated.

[GPUAddressMode](#) describes the behavior of the sampler if the sampled texels extend beyond the bounds of the sampled texture.

enum

[GPUAddressMode](#)

```
{  
    "clamp-to-edge",  
    "repeat",  
    "mirror-repeat",  
};
```

"clamp-to-edge"

Texture coordinates are clamped between 0.0 and 1.0, inclusive.

"repeat"

Texture coordinates wrap to the other side of the texture.

"mirror-repeat"

Texture coordinates wrap to the other side of the texture, but the texture is flipped when the integer part of the coordinate is odd.

[GPUFilterMode](#) and [GPUMipmapFilterMode](#) describe the behavior of the sampler if the sampled area does not cover exactly one texel.

Note: See [Texel Filtering](#) in the [Vulkan 1.3](#) spec for an example of how samplers may determine which texels are sampled from for the various filtering modes.

enum

[GPUFilterMode](#)

```
{  
    "nearest",  
    "linear",  
};
```

enum

[GPUMipmapFilterMode](#)

```
{
```

"nearest"

,

"linear"

,

```
};
```

"nearest"

Return the value of the texel nearest to the texture coordinates.

"linear"

Select two texels in each dimension and return a linear interpolation between their values.

[GPUCompareFunction](#) specifies the behavior of a comparison sampler. If a comparison sampler is used in a shader, the **depth_ref** is compared to the fetched texel value, and the result of this comparison test is generated (**1.0f** for pass, or **0.0f** for fail).

After comparison, if texture filtering is enabled, the filtering step occurs, so that comparison results are mixed together resulting in values in the range **[0, 1]**. Filtering **should** behave as usual, however it **may** be computed with lower precision or not mix results at all.

enum

[GPUCompareFunction](#)

```
{  
    "never",  
    "less",  
    "equal",
```

["less-equal"](#),
["greater"](#),
["not-equal"](#),
["greater-equal"](#),
["always"](#),

};

"never"

Comparison tests never pass.

"less"

A provided value passes the comparison test if it is less than the sampled value.

"equal"

A provided value passes the comparison test if it is equal to the sampled value.

"less-equal"

A provided value passes the comparison test if it is less than or equal to the sampled value.

"greater"

A provided value passes the comparison test if it is greater than the sampled value.

"not-equal"

A provided value passes the comparison test if it is not equal to the sampled value.

"greater-equal"

A provided value passes the comparison test if it is greater than or equal to the sampled value.

"always"

Comparison tests always pass.

7.1.2. Sampler Creation

createSampler(descriptor)

Creates a [GPUSampler](#).

Called on: [GPUDevice](#) this.

Arguments:

Arguments for the [GPUDevice.createSampler\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUSamplerDescriptor	✗	✓	Description of the GPUSampler to create.

Returns: [GPUSampler](#)

[Content timeline](#) steps:

- Let *s* be [! create a new WebGPU object](#)(*this*, [GPUSampler](#), *descriptor*).
- Issue the *initialization steps* on the [Device timeline](#) of *this*.
- Return *s*.

[Device timeline](#) *initialization steps*:

- If any of the following conditions are unsatisfied [generate a validation error](#), [invalidate](#) *s* and return.
 - this* must not be [lost](#).
 - descriptor.lodMinClamp* ≥ 0.
 - descriptor.lodMaxClamp* ≥ *descriptor.lodMinClamp*.
 - descriptor.maxAnisotropy* ≥ 1.

Note: Most implementations support [maxAnisotropy](#) values in range between 1 and 16, inclusive. The provided [maxAnisotropy](#) value will be clamped to the maximum value that the platform supports.

- If *descriptor.maxAnisotropy* > 1:
 - descriptor.magFilter*, *descriptor.minFilter*, and *descriptor.mipmapFilter* must be ["linear"](#).
- Set *s.[[descriptor]]* to *descriptor*.
 - Set *s.[[isComparison]]* to **false** if the [compare](#) attribute of *s.[[descriptor]]* is **null** or **undefined**. Otherwise, set it to **true**.
 - Set *s.[[isFiltering]]* to **false** if none of [minFilter](#), [magFilter](#), or [mipmapFilter](#) has the value of ["linear"](#). Otherwise, set it to **true**.

Creating a [GPUSampler](#) that does trilinear filtering and repeats texture coordinates:

```
const sampler = gpuDevice.createSampler({
  addressModeU: 'repeat',
  addressModeV: 'repeat',
  magFilter: 'linear',
  minFilter: 'linear',
  mipmapFilter: 'linear',
});
```

8. Resource Binding

8.1. GPUBindGroupLayout

A [GPUBindGroupLayout](#) defines the interface between a set of resources bound in a [GPUBindGroup](#) and their accessibility in shader stages.

[[Exposed](#)=(Window, Worker), [SecureContext](#)]

```
interface GPUBindGroupLayout {
};
```

[GPUBindGroupLayout](#) includes [GPUObjectBase](#);

[GPUBindGroupLayout](#) has the following [immutable properties](#):

[[descriptor]], of type [GPUBindGroupLayoutDescriptor](#), readonly

8.1.1. Bind Group Layout Creation

A [GPUBindGroupLayout](#) is created via [GPUDevice.createBindGroupLayout\(\)](#).

dictionary

GPUBindGroupLayoutDescriptor

```
  : GPUObjectDescriptorBase {
    required sequence<GPUBindGroupLayoutEntry> entries;
  };
```

[GPUBindGroupLayoutDescriptor](#) dictionaries have the following members:

entries, of type sequence<[GPUBindGroupLayoutEntry](#)>

A list of entries describing the shader resource bindings for a bind group.

A [GPUBindGroupLayoutEntry](#) describes a single shader resource binding to be included in a [GPUBindGroupLayout](#).

dictionary

GPUBindGroupLayoutEntry

```
{
  required GPUIndex32 binding;
  required GPUShaderStageFlags visibility;

  GPUBufferBindingLayout buffer;
  GPUSamplerBindingLayout sampler;
  GPUTextureBindingLayout texture;
  GPUStorageTextureBindingLayout storageTexture;
  GPUExternalTextureBindingLayout externalTexture;
};
```

[GPUBindGroupLayoutEntry](#) dictionaries have the following members:

binding, of type [GPUIndex32](#)

A unique identifier for a resource binding within the [GPUBindGroupLayout](#), corresponding to a [GPUBindGroupEntry.binding](#) and a [@binding](#) attribute in the [GPUShaderModule](#).

visibility, of type [GPUShaderStageFlags](#)

A bitset of the members of [GPUShaderStage](#). Each set bit indicates that a [GPUBindGroupLayoutEntry](#)'s resource will be accessible from the associated shader stage.

buffer, of type [GPUBufferBindingLayout](#)

sampler, of type [GPUSamplerBindingLayout](#)

texture, of type [GPUTextureBindingLayout](#)
storageTexture, of type [GPUStorageTextureBindingLayout](#)
externalTexture, of type [GPUExternalTextureBindingLayout](#)

Exactly one of these members must be set, indicating the binding type. The contents of the member specify options specific to that type.

The corresponding resource in [createBindGroup\(\)](#) requires the corresponding [binding resource type](#) for this binding.

```
typedef [EnforceRange] unsigned long
GPUShaderStageFlags

;
[Exposed=(Window, Worker), SecureContext]
namespace
GPUShaderStage
{
    const GPUFlagsConstant VERTEX = 0x1;
    const GPUFlagsConstant FRAGMENT = 0x2;
    const GPUFlagsConstant COMPUTE = 0x4;
};
```

[GPUShaderStage](#) contains the following flags, which describe which shader stages a corresponding [GPUBindGroupEntry](#) for this [GPUBindGroupLayoutEntry](#) will be visible to:

- VERTEX**
The bind group entry will be accessible to vertex shaders.
- FRAGMENT**
The bind group entry will be accessible to fragment shaders.
- COMPUTE**
The bind group entry will be accessible to compute shaders.

The [binding member](#) of a [GPUBindGroupLayoutEntry](#) is determined by which member of the [GPUBindGroupLayoutEntry](#) is defined: [buffer](#), [sampler](#), [texture](#), [storageTexture](#), or [externalTexture](#). Only one may be defined for any given [GPUBindGroupLayoutEntry](#). Each member has an associated [GPUBindingResource](#) type and each [binding type](#) has an associated [internal usage](#), given by this table:

<i>Binding member</i>	<i>Resource type</i>	<i>Binding type</i>	<i>Binding usage</i>
buffer	GPUBufferBinding (or GPUBuffer as shorthand)	"uniform"	constant
		"storage"	storage
		"read-only-storage"	storage-read
sampler	GPUSampler	"filtering"	constant
		"non-filtering"	
		"comparison"	
texture	GPUTextureView (or GPUTexture as shorthand)	"float"	constant
		"unfilterable-float"	
		"depth"	
		"sint"	
		"uint"	
storageTexture	GPUTextureView (or GPUTexture as shorthand)	"write-only"	storage
		"read-write"	storage-read
		"read-only"	
externalTexture	GPUExternalTexture or GPUTextureView (or GPUTexture as shorthand)		constant

The [list](#) of [GPUBindGroupLayoutEntry](#) values *entries exceeds the binding slot limits* of [supported limits](#) *limits* if the number of slots used toward a limit exceeds the supported value in *limits*. Each entry may use multiple slots toward multiple limits.

[Device timeline](#) steps:

For each *entry* in *entries*, if:

- entry*.[buffer?.type](#) is "uniform" and *entry*.[buffer?.hasDynamicOffset](#) is true
Consider 1 [maxDynamicUniformBuffersPerPipelineLayout](#) slot to be used.
- entry*.[buffer?.type](#) is "storage" and *entry*.[buffer?.hasDynamicOffset](#) is true
Consider 1 [maxDynamicStorageBuffersPerPipelineLayout](#) slot to be used.

For each shader stage *stage* in « [VERTEX](#), [FRAGMENT](#), [COMPUTE](#) »:

For each *entry* in *entries* for which *entry.visibility* contains *stage*, if:

entry.buffer?.type is "[uniform](#)"

Consider 1 [maxUniformBuffersPerShaderStage](#) slot to be used.

entry.buffer?.type is "[storage](#)" or "[read-only-storage](#)"

Consider 1 [maxStorageBuffersPerShaderStage](#) slot to be used.

entry.sampler is [provided](#)

Consider 1 [maxSamplersPerShaderStage](#) slot to be used.

entry.texture is [provided](#)

Consider 1 [maxSampledTexturesPerShaderStage](#) slot to be used.

entry.storageTexture is [provided](#)

Consider 1 [maxStorageTexturesPerShaderStage](#) slot to be used.

entry.externalTexture is [provided](#)

Consider 4 [maxSampledTexturesPerShaderStage](#) slot, 1 [maxSamplersPerShaderStage](#) slot, and 1 [maxUniformBuffersPerShaderStage](#) slot to be used.

Note: See [GPUExternalTexture](#) for an explanation of this behavior.

enum

```
GPUBufferBindingType
```

```
{
```

```
"uniform"
```

```
,
```

```
"storage"
```

```
,
```

```
"read-only-storage"
```

```
,
```

```
};
```

dictionary

```
GPUBufferBindingLayout
```

```
{
```

```
  GPUBufferBindingType type = "uniform";
```

```
  boolean hasDynamicOffset = false;
```

```
  GPUSize64 minBindingSize = 0;
```

```
};
```

[GPUBufferBindingLayout](#) dictionaries have the following members:

type, of type [GPUBufferBindingType](#), defaulting to "[uniform](#)"

Indicates the type required for buffers bound to this bindings.

hasDynamicOffset, of type [boolean](#), defaulting to [false](#)

Indicates whether this binding requires a dynamic offset.

minBindingSize, of type [GPUSize64](#), defaulting to 0

Indicates the minimum [size](#) of a buffer binding used with this bind point.

Bindings are always validated against this size in [createBindGroup\(\)](#).

If this is *not* 0, pipeline creation additionally [validates](#) that this value ≥ the [minimum buffer binding size](#) of the variable.

If this is 0, it is ignored by pipeline creation, and instead draw/dispatch commands [validate](#) that each binding in the [GPUBindGroup](#) satisfies the [minimum buffer binding size](#) of the variable.

Note: Similar execution-time validation is theoretically possible for other binding-related fields specified for early validation, like [sampleType](#) and [format](#), which currently can only be validated in pipeline creation. However, such execution-time validation could be costly or unnecessarily complex, so it is available only for [minBindingSize](#) which is expected to have the most ergonomic impact.

enum

GPUSamplerBindingType

{

"*filtering*"

,

"*non-filtering*"

,

"*comparison*"

,

};

dictionary

GPUSamplerBindingLayout

{

[GPUSamplerBindingType](#) *type* = "filtering";

};

[GPUSamplerBindingLayout](#) dictionaries have the following members:

type, of type [GPUSamplerBindingType](#), defaulting to "**filtering**"

Indicates the required type of a sampler bound to this bindings.

enum

GPUTextureSampleType

{

"*float*"

,

"*unfilterable-float*"

,

"*depth*"

,

"*sint*"

,

"*uint*"

,

};

dictionary

GPUTextureBindingLayout

{

[GPUTextureSampleType](#) *sampleType* = "float";

[GPUTextureViewDimension](#) *viewDimension* = "2d";

[boolean](#) *multisampled* = false;

};

[GPUTextureBindingLayout](#) dictionaries have the following members:

sampleType, of type [GPUTextureSampleType](#), defaulting to "**float**"

Indicates the type required for texture views bound to this binding.

viewDimension, of type [GPUTextureViewDimension](#), defaulting to "2d"
Indicates the required [dimension](#) for texture views bound to this binding.

multisampled, of type [boolean](#), defaulting to false
Indicates whether or not texture views bound to this binding must be multisampled.

```
enum
GPUStorageTextureAccess
{
  "write-only"
,
  "read-only"
,
  "read-write"
,
};
```

```
dictionary
GPUStorageTextureBindingLayout
{
  GPUStorageTextureAccess access = "write-only";
  required GPUTextureFormat format;
  GPUTextureViewDimension viewDimension = "2d";
};
```

[GPUStorageTextureBindingLayout](#) dictionaries have the following members:

access, of type [GPUStorageTextureAccess](#), defaulting to "write-only"
The access mode for this binding, indicating readability and writability.

format, of type [GPUTextureFormat](#)
The required [format](#) of texture views bound to this binding.

viewDimension, of type [GPUTextureViewDimension](#), defaulting to "2d"
Indicates the required [dimension](#) for texture views bound to this binding.

```
dictionary
GPUExternalTextureBindingLayout
{
};
```

A [GPUBindGroupLayout](#) object has the following [device timeline properties](#):

[[entryMap]], of type [ordered map](#)<[GPUSize32](#), [GPUBindGroupLayoutEntry](#)>, readonly
The map of binding indices pointing to the [GPUBindGroupLayoutEntry](#)s, which this [GPUBindGroupLayout](#) describes.

[[dynamicOffsetCount]], of type [GPUSize32](#), readonly
The number of buffer bindings with dynamic offsets in this [GPUBindGroupLayout](#).

[[exclusivePipeline]], of type [GPUPipelineBase?](#), readonly
The pipeline that created this [GPUBindGroupLayout](#), if it was created as part of a [default pipeline layout](#). If not null, [GPUBindGroups](#) created with this [GPUBindGroupLayout](#) can only be used with the specified [GPUPipelineBase](#).

createBindGroupLayout(descriptor)
Creates a [GPUBindGroupLayout](#).

Called on: [GPUDevice](#) *this*.

Arguments:

Arguments for the GPUDevice.createBindGroupLayout(descriptor) method.				
Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUBindGroupLayoutDescriptor	✗	✗	Description of the GPUBindGroupLayout to create.

Returns: [GPUBindGroupLayout](#)

[Content timeline](#) steps:

1. For each [GPUBindGroupLayoutEntry](#) *entry* in *descriptor.entries*:
 1. If *entry.storageTexture* is [provided](#):
 1. ? [Validate texture format required features](#) for *entry.storageTexture.format* with *this.[[device]]*.
 2. Let *layout* be ! [create a new WebGPU object](#)(*this*, [GPUBindGroupLayout](#), *descriptor*).
 3. Issue the *initialization steps* on the [Device timeline](#) of *this*.
 4. Return *layout*.

[Device timeline](#) initialization steps:

1. If any of the following conditions are unsatisfied [generate a validation error](#), [invalidate](#) *layout* and return.
 - *this* must not be [lost](#).
 - Let *limits* be *this.[[device]].[[limits]]*.
 - The [binding](#) of each entry in *descriptor* is unique.
 - The [binding](#) of each entry in *descriptor* must be < *limits.maxBindingsPerBindGroup*.
 - *descriptor.entries* must not [exceed the binding slot limits](#) of *limits*.
 - For each [GPUBindGroupLayoutEntry](#) *entry* in *descriptor.entries*:
 - Exactly one of *entry.buffer*, *entry.sampler*, *entry.texture*, *entry.storageTexture*, and *entry.externalTexture* is [provided](#).
 - *entry.visibility* contains only bits defined in [GPUShaderStage](#).
 - If *entry.visibility* includes [VERTEX](#):
 - If *entry.buffer* is [provided](#), *entry.buffer.type* must be ["uniform"](#) or ["read-only-storage"](#).
 - If *entry.storageTexture* is [provided](#), *entry.storageTexture.access* must be ["read-only"](#).
 - If *entry.texture?.multisampled* is [true](#):
 - *entry.texture.viewDimension* is ["2d"](#).
 - *entry.texture.sampleType* is not ["float"](#).
 - If *entry.storageTexture* is [provided](#):
 - *entry.storageTexture.viewDimension* is not ["cube"](#) or ["cube-array"](#).
 - *entry.storageTexture.format* must be a format which can support storage usage for the given *entry.storageTexture.access* according to the [§ 26.1.1 Plain color formats](#) table.
 - 2. Set *layout.[[descriptor]]* to *descriptor*.
 - 3. Set *layout.[[dynamicOffsetCount]]* to the number of entries in *descriptor* where [buffer](#) is [provided](#) and [buffer.hasDynamicOffset](#) is [true](#).
 - 4. Set *layout.[[exclusivePipeline]]* to [null](#).
 - 5. For each [GPUBindGroupLayoutEntry](#) *entry* in *descriptor.entries*:
 1. Insert *entry* into *layout.[[entryMap]]* with the key of *entry.binding*.

8.1.2. Compatibility

Two [GPUBindGroupLayout](#) objects *a* and *b* are considered *group-equivalent* if and only if all of the following conditions are satisfied:

a.[[exclusivePipeline]] == *b.[[exclusivePipeline]]*.

for any binding number *binding*, one of the following conditions is satisfied:

it's missing from both *a.[[entryMap]]* and *b.[[entryMap]]*.

a.[[entryMap]][binding] == *b.[[entryMap]][binding]*

If bind groups layouts are [group-equivalent](#) they can be interchangeably used in all contents.

8.2. GPUBindGroup

A [GPUBindGroup](#) defines a set of resources to be bound together in a group and how the resources are used in shader stages.

[[Exposed](#)=([Window](#), [Worker](#)), [SecureContext](#)]

```
interface GPUBindGroup {  
};
```

[GPUBindGroup](#) includes [GPUObjectBase](#);

[GPUBindGroup](#) has the following [device timeline properties](#):

[\[\[layout\]\]](#), of type [GPUBindGroupLayout](#), readonly

The [GPUBindGroupLayout](#) associated with this [GPUBindGroup](#).

[\[\[entries\]\]](#), of type [sequence](#)<[GPUBindGroupEntry](#)>, readonly

The set of [GPUBindGroupEntry](#)s this [GPUBindGroup](#) describes.

[\[\[usedResources\]\]](#), of type [usage_scope](#), readonly

The set of buffer and texture [subresources](#) used by this bind group, associated with lists of the [internal usage](#) flags.

The *bound buffer ranges* of a [GPUBindGroup](#) *bindGroup*, given [list](#)<[GPUBufferDynamicOffset](#)> *dynamicOffsets*, are computed as follows:

Let *result* be a new [set](#)<(GPUBindGroupLayoutEntry, GPUBufferBinding)>.

Let *dynamicOffsetIndex* be 0.

For each [GPUBindGroupEntry](#) *bindGroupEntry* in *bindGroup*.[\[\[entries\]\]](#), sorted by *bindGroupEntry*.[binding](#):

Let *bindGroupLayoutEntry* be *bindGroup*.[\[\[layout\]\]](#).[\[\[entryMap\]\]](#)[*bindGroupEntry*.[binding](#)].

If *bindGroupLayoutEntry*.[buffer](#) is not [provided](#), **continue**.

Let *bound* be [get as buffer binding](#)(*bindGroupEntry*.[resource](#)).

If *bindGroupLayoutEntry*.[buffer.hasDynamicOffset](#):

Increment *bound*.[offset](#) by *dynamicOffsets*[*dynamicOffsetIndex*].

Increment *dynamicOffsetIndex* by 1.

[Append](#) (*bindGroupLayoutEntry*, *bound*) to *result*.

Return *result*.

8.2.1. Bind Group Creation

A [GPUBindGroup](#) is created via [GPUDevice.createBindGroup\(\)](#).

dictionary

[GPUBindGroupDescriptor](#)

```
    : GPUObjectDescriptorBase {  
    required GPUBindGroupLayout layout;  
    required sequence<GPUBindGroupEntry> entries;  
};
```

[GPUBindGroupDescriptor](#) dictionaries have the following members:

layout, of type [GPUBindGroupLayout](#)

The [GPUBindGroupLayout](#) the entries of this bind group will conform to.

entries, of type [sequence](#)<[GPUBindGroupEntry](#)>

A list of entries describing the resources to expose to the shader for each binding described by the [layout](#).

typedef ([GPUSampler](#) or

[GPUTexture](#) or

[GPUTextureView](#) or

[GPUBuffer](#) or

[GPUBufferBinding](#) or

[GPUExternalTexture](#))

[GPUBindingResource](#)

;

dictionary

[GPUBindGroupEntry](#)

```
{
```

```
required GPUIndex32 binding;
required GPUBindingResource resource;
};
```

A [GPUBindGroupEntry](#) describes a single resource to be bound in a [GPUBindGroup](#), and has the following members:

binding, of type [GPUIndex32](#)

A unique identifier for a resource binding within the [GPUBindGroup](#), corresponding to a [GPUBindGroupLayoutEntry.binding](#) and a [@binding](#) attribute in the [GPUShaderModule](#).

resource, of type [GPUBindingResource](#)

The resource to bind, which may be a [GPUSampler](#), [GPUTexture](#), [GPUTextureView](#), [GPUBuffer](#), [GPUBufferBinding](#), or [GPUExternalTexture](#).

[GPUBindGroupEntry](#) has the following [device timeline properties](#):

[[prevalidatedSize]], of type [boolean](#)

Whether or not this binding entry had its buffer size validated at time of creation.

dictionary

GPUBufferBinding

```
{
  required GPUBuffer buffer;
  GPUSize64 offset = 0;
  GPUSize64 size;
};
```

A [GPUBufferBinding](#) describes a buffer and optional range to bind as a resource, and has the following members:

buffer, of type [GPUBuffer](#)

The [GPUBuffer](#) to bind.

offset, of type [GPUSize64](#), defaulting to 0

The offset, in bytes, from the beginning of [buffer](#) to the beginning of the range exposed to the shader by the buffer binding.

size, of type [GPUSize64](#)

The size, in bytes, of the buffer binding. If not [provided](#), specifies the range starting at [offset](#) and ending at the end of [buffer](#).

createBindGroup(descriptor)

Creates a [GPUBindGroup](#).

Called on: [GPUDevice](#) *this*.

Arguments:

Arguments for the [GPUDevice.createBindGroup\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUBindGroupDescriptor	✗	✗	Description of the GPUBindGroup to create.

Returns: [GPUBindGroup](#)

[Content timeline](#) steps:

1. Let *bindGroup* be ! [create a new WebGPU object](#)(*this*, [GPUBindGroup](#), *descriptor*).
2. Issue the *initialization steps* on the [Device timeline](#) of *this*.
3. Return *bindGroup*.

[Device timeline](#) *initialization steps*:

1. Let *limits* be *this*.*[[device]].[[limits]]*.
2. If any of the following conditions are unsatisfied [generate a validation error](#), [invalidate](#) *bindGroup* and return.
 - *descriptor.layout* is [valid to use with](#) *this*.
 - The number of [entries](#) of *descriptor.layout* is exactly equal to the number of *descriptor.entries*.

For each [GPUBindGroupEntry](#) *bindingDescriptor* in *descriptor.entries*:

- Let *resource* be *bindingDescriptor.resource*.
- There is exactly one [GPUBindGroupLayoutEntry](#) *layoutBinding* in *descriptor.layout.entries* such that *layoutBinding.binding* equals to *bindingDescriptor.binding*.
- If the defined [binding member](#) for *layoutBinding* is:

sampler

- *resource* is a [GPUSampler](#).
- *resource* is [valid to use with this](#).
- If *layoutBinding.sampler.type* is:

"filtering"

resource.[[isComparison]] is false.

"non-filtering"

resource.[[isFiltering]] is false. *resource.[[isComparison]]* is false.

"comparison"

resource.[[isComparison]] is true.

texture

- *resource* is either a [GPUTexture](#) or a [GPUTextureView](#).
- *resource* is [valid to use with this](#).
- Let *textureView* be [get as texture view\(resource\)](#).
- Let *texture* be *textureView.[[texture]]*.
- *layoutBinding.texture.viewDimension* is equal to *textureView's dimension*.
- *layoutBinding.texture.sampleType* is [compatible](#) with *textureView's format*.
- *textureView.[[descriptor]].usage* includes [TEXTURE_BINDING](#).
- If *layoutBinding.texture.multisampled* is true, *texture's sampleCount* > 1, Otherwise *texture's sampleCount* is 1.

storageTexture

- *resource* is either a [GPUTexture](#) or a [GPUTextureView](#).
- *resource* is [valid to use with this](#).
- Let *storageTextureView* be [get as texture view\(resource\)](#).
- Let *texture* be *storageTextureView.[[texture]]*.
- *layoutBinding.storageTexture.viewDimension* is equal to *storageTextureView's dimension*.
- *layoutBinding.storageTexture.format* is equal to *storageTextureView.[[descriptor]].format*.
- *storageTextureView.[[descriptor]].usage* includes [STORAGE_BINDING](#).
- *storageTextureView.[[descriptor]].mipLevelCount* must be 1.

buffer

- *resource* is either a [GPUBuffer](#) or a [GPUBufferBinding](#).
- Let *bufferBinding* be [get as buffer binding\(resource\)](#).
- *bufferBinding.buffer* is [valid to use with this](#).
- The bound part designated by *bufferBinding.offset* and *bufferBinding.size* resides inside the buffer and has non-zero size.
- [effective buffer binding size\(bufferBinding\)](#) ≥ *layoutBinding.buffer.minBindingSize*.
- If *layoutBinding.buffer.type* is

"uniform"

- *bufferBinding.buffer.usage* includes [UNIFORM](#).
- [effective buffer binding size\(bufferBinding\)](#) ≤ *limits.maxUniformBufferBindingSize*.
- *bufferBinding.offset* is a multiple of *limits.minUniformBufferOffsetAlignment*.

"storage" or "read-only-storage"

- *bufferBinding.buffer.usage* includes [STORAGE](#).
- [effective buffer binding size\(bufferBinding\)](#) ≤ *limits.maxStorageBufferBindingSize*.
- [effective buffer binding size\(bufferBinding\)](#) is a multiple of 4.
- *bufferBinding.offset* is a multiple of *limits.minStorageBufferOffsetAlignment*.

externalTexture

- *resource* is either a [GPUExternalTexture](#), a [GPUTexture](#), or a [GPUTextureView](#).

- *resource* is [valid to use with this](#).

- If *resource* is a:

[GPUTexture](#) or [GPUTextureView](#)

- Let *view* be [get as texture view](#)(*resource*).
- *view*.[\[\[descriptor\]\].usage](#) must include [TEXTURE_BINDING](#).
- *view*.[\[\[descriptor\]\].dimension](#) must be ["2d"](#).
- *view*.[\[\[descriptor\]\].mipLevelCount](#) must be 1.
- *view*.[\[\[descriptor\]\].format](#) must be ["rgba8unorm"](#), ["bgra8unorm"](#), or ["rgba16float"](#).
- *view*.[\[\[texture\]\].sampleCount](#) must be 1.

3. Let *bindGroup*.[\[\[layout\]\]](#) = *descriptor*.[layout](#).

4. Let *bindGroup*.[\[\[entries\]\]](#) = *descriptor*.[entries](#).

5. Let *bindGroup*.[\[\[usedResources\]\]](#) = {}.

6. For each [GPUBindGroupEntry](#) *bindingDescriptor* in *descriptor*.[entries](#):

1. Let *internalUsage* be the [binding usage](#) for *layoutBinding*.

2. Each [subresource](#) seen by *resource* is added to [\[\[usedResources\]\]](#) as *internalUsage*.

3. Let *bindingDescriptor*.[\[\[prevalidatedSize\]\]](#) be `false` if the defined [binding member](#) for *layoutBinding* is [buffer](#) and *layoutBinding*.[buffer.minBindingSize](#) is 0, and `true` otherwise.

get as texture view(*resource*)

Arguments:

[GPUBindingResource](#) *resource*

Returns: [GPUTextureView](#)

[Assert](#) *resource* is either a [GPUTexture](#) or a [GPUTextureView](#).

If *resource* is a:

[GPUTexture](#)

1. Return *resource*.[createView\(\)](#).

[GPUTextureView](#)

1. Return *resource*.

get as buffer binding(*resource*)

Arguments:

[GPUBindingResource](#) *resource*

Returns: [GPUBufferBinding](#)

[Assert](#) *resource* is either a [GPUBuffer](#) or a [GPUBufferBinding](#).

If *resource* is a:

[GPUBuffer](#)

1. Let *bufferBinding* a new [GPUBufferBinding](#).

2. Set *bufferBinding*.[buffer](#) to *resource*.

3. Return *bufferBinding*.

[GPUBufferBinding](#)

1. Return *resource*.

effective buffer binding size(*binding*)

Arguments:

[GPUBufferBinding](#) *binding*

Returns: [GPUSize64](#)

If *binding*.[size](#) is not [provided](#):

Return `max(0, binding.buffer.size - binding.offset)`;

Return *binding.size*.

Two [GPUBufferBinding](#) objects *a* and *b* are considered *buffer-binding-aliasing* if and only if all of the following are true:

a.buffer == *b.buffer*

The range formed by *a.offset* and *a.size* intersects the range formed by *b.offset* and *b.size*, where if a *size* is [unspecified](#), the range goes to the end of the buffer.

Note: When doing this calculation, any dynamic offsets have already been applied to the ranges.

8.3. GPUPipelineLayout

A [GPUPipelineLayout](#) defines the mapping between resources of all [GPUBindGroup](#) objects set up during command encoding in [setBindGroup\(\)](#), and the shaders of the pipeline set by [GPURenderCommandsMixin.setPipeline](#) or [GPUComputePassEncoder.setPipeline](#).

The full binding address of a resource can be defined as a trio of:

shader stage mask, to which the resource is visible

bind group index

binding number

The components of this address can also be seen as the binding space of a pipeline. A [GPUBindGroup](#) (with the corresponding [GPUBindGroupLayout](#)) covers that space for a fixed bind group index. The contained bindings need to be a superset of the resources used by the shader at this bind group index.

[Exposed=(Window, Worker), SecureContext]

```
interface GPUPipelineLayout {  
};
```

[GPUPipelineLayout](#) includes [GPUObjectBase](#);

[GPUPipelineLayout](#) has the following [device timeline properties](#):

[[bindGroupLayouts]], of type [list<GPUBindGroupLayout>](#), readonly

The [GPUBindGroupLayout](#) objects provided at creation in [GPUPipelineLayoutDescriptor.bindGroupLayouts](#).

Note: using the same [GPUPipelineLayout](#) for many [GPURenderPipeline](#) or [GPUComputePipeline](#) pipelines guarantees that the user agent doesn't need to rebind any resources internally when there is a switch between these pipelines.

[GPUComputePipeline](#) object X was created with [GPUPipelineLayout.bindGroupLayouts](#) A, B, C. [GPUComputePipeline](#) object Y was created with [GPUPipelineLayout.bindGroupLayouts](#) A, D, C. Supposing the command encoding sequence has two dispatches:

[setBindGroup](#)(0, ...)

[setBindGroup](#)(1, ...)

[setBindGroup](#)(2, ...)

[setPipeline](#)(X)

[dispatchWorkgroups](#)()

[setBindGroup](#)(1, ...)

[setPipeline](#)(Y)

[dispatchWorkgroups](#)()

In this scenario, the user agent would have to re-bind the group slot 2 for the second dispatch, even though neither the [GPUBindGroupLayout](#) at index 2 of [GPUPipelineLayout.bindGroupLayouts](#), or the [GPUBindGroup](#) at slot 2, change.

Note: the expected usage of the [GPUPipelineLayout](#) is placing the most common and the least frequently changing bind groups at the "bottom" of the layout, meaning lower bind group slot numbers, like 0 or 1. The more frequently a bind group needs to change between draw calls, the higher its index should be. This general guideline allows the user agent to minimize state changes between draw calls, and consequently lower the CPU overhead.

8.3.1. Pipeline Layout Creation

A [GPUPipelineLayout](#) is created via [GPUDevice.createPipelineLayout\(\)](#).

dictionary

GPUPipelineLayoutDescriptor

```
: GPUObjectDescriptorBase {  
  required sequence<GPUBindGroupLayout?> bindGroupLayouts;  
};
```

[GPUPipelineLayoutDescriptor](#) dictionaries define all the [GPUBindGroupLayout](#)s used by a pipeline, and have the following members:

bindGroupLayouts, of type sequence<[GPUBindGroupLayout](#)?>

A list of optional [GPUBindGroupLayout](#)s the pipeline will use. Each element corresponds to a [@group](#) attribute in the [GPUShaderModule](#), with the Nth element corresponding with [@group\(N\)](#).

createPipelineLayout(descriptor)

Creates a [GPUPipelineLayout](#).

Called on: [GPUDevice](#) *this*.

Arguments:

Arguments for the [GPUDevice.createPipelineLayout\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUPipelineLayoutDescriptor	✗	✗	Description of the GPUPipelineLayout to create.

Returns: [GPUPipelineLayout](#)

[Content timeline](#) steps:

- Let *pl* be [! create a new WebGPU object](#)(*this*, [GPUPipelineLayout](#), *descriptor*).
- Issue the *initialization steps* on the [Device timeline](#) of *this*.
- Return *pl*.

[Device timeline](#) *initialization steps*:

- Let *limits* be *this*.[\[\[device\]\].\[\[limits\]\]](#).
- Let *bindGroupLayouts* be a [list](#) of null [GPUBindGroupLayout](#)s with [size](#) equal to *limits*.[maxBindGroups](#).
- For each *bindGroupLayout* at index *i* in *descriptor*.[bindGroupLayouts](#):
 - If *bindGroupLayout* is not null and *bindGroupLayout*.[\[\[descriptor\]\].entries](#) is not [empty](#):
 - Set *bindGroupLayouts*[*i*] to *bindGroupLayout*.
 - Let *allEntries* be the result of concatenating *bgl*.[\[\[descriptor\]\].entries](#) for all non-null *bgl* in *bindGroupLayouts*.
 - If any of the following conditions are unsatisfied [generate a validation error](#), [invalidate](#) *pl* and return.
 - Every non-null [GPUBindGroupLayout](#) in *bindGroupLayouts* must be [valid to use with](#) *this* and have a [\[\[exclusivePipeline\]\]](#) of null.
 - The [size](#) of *descriptor*.[bindGroupLayouts](#) must be ≤ *limits*.[maxBindGroups](#).
 - allEntries* must not [exceed the binding slot limits](#) of *limits*.
- Set the *pl*.[\[\[bindGroupLayouts\]\]](#) to *bindGroupLayouts*.

Note: two [GPUPipelineLayout](#) objects are considered equivalent for any usage if their internal [\[\[bindGroupLayouts\]\]](#) sequences contain [GPUBindGroupLayout](#) objects that are [group-equivalent](#).

8.4. Example

Create a [GPUBindGroupLayout](#) that describes a binding with a uniform buffer, a texture, and a sampler. Then create a [GPUBindGroup](#) and a [GPUPipelineLayout](#) using the [GPUBindGroupLayout](#).

```
const bindGroupLayout = gpuDevice.createBindGroupLayout({
  entries: [{
    binding: 0,
    visibility: GPUShaderStage.VERTEX | GPUShaderStage.FRAGMENT,
    buffer: {}
  }, {
    binding: 1,
    visibility: GPUShaderStage.FRAGMENT,
    texture: {}
  }, {
    binding: 2,
    visibility: GPUShaderStage.FRAGMENT,
    sampler: {}
  }]
});
```

```
const bindGroup = gpuDevice.createBindGroup({
  layout: bindGroupLayout,
  entries: [{
    binding: 0,
    resource: { buffer: buffer },
  }, {
    binding: 1,
    resource: texture
  }, {
    binding: 2,
    resource: sampler
  }]
});

const pipelineLayout = gpuDevice.createPipelineLayout({
  bindGroupLayouts: [bindGroupLayout]
});
```

9. Shader Modules

9.1. GPUShaderModule

[[Exposed](#)=(Window, Worker), [SecureContext](#)]

```
interface GPUShaderModule {
  Promise<GPUCompilationInfo> getCompilationInfo();
};
```

[GPUShaderModule](#) includes [GPUObjectBase](#);

[GPUShaderModule](#) is a reference to an internal shader module object.

9.1.1. Shader Module Creation

dictionary

GPUShaderModuleDescriptor

```
    : GPUObjectDescriptorBase {
  required USVString code;
  sequence<GPUShaderModuleCompilationHint> compilationHints = [];
};
```

code, of type [USVString](#)

The [WGLSL](#) source code for the shader module.

compilationHints, of type sequence<[GPUShaderModuleCompilationHint](#)>, defaulting to []

A list of [GPUShaderModuleCompilationHints](#).

Any hint provided by an application **should** contain information about one entry point of a pipeline that will eventually be created from the entry point.

Implementations **should** use any information present in the [GPUShaderModuleCompilationHint](#) to perform as much compilation as is possible within [createShaderModule\(\)](#).

Aside from type-checking, these hints are not validated in any way.

NOTE:

Supplying information in [compilationHints](#) does not have any observable effect, other than performance. It may be detrimental to performance to provide hints for pipelines that never end up being created.

Because a single shader module can hold multiple entry points, and multiple pipelines can be created from a single shader module, it can be more performant for an implementation to do as much compilation as possible once in [createShaderModule\(\)](#) rather than multiple times in the multiple calls to [createComputePipeline\(\)](#) or [createRenderPipeline\(\)](#).

Hints are only applied to the entry points they explicitly name. Unlike [GPUProgrammableStage.entryPoint](#), there is no default, even if only one entry point is present in the module.

Note: Hints are not validated in an observable way, but user agents **may** surface identifiable errors (like unknown entry point names or incompatible pipeline layouts) to developers, for example in the browser developer console.

createShaderModule(descriptor)

Creates a [GPUShaderModule](#).

Called on: [GPUDevice](#) this.

Arguments:

Arguments for the [GPUDevice.createShaderModule\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
descriptor	GPUShaderModuleDescriptor	✗	✗	Description of the GPUShaderModule to create.

Returns: [GPUShaderModule](#)

[Content timeline](#) steps:

- Let *sm* be [! create a new WebGPU object](#)(*this*, [GPUShaderModule](#), *descriptor*).
- Issue the *initialization steps* on the [Device timeline](#) of *this*.
- Return *sm*.

[Device timeline](#) initialization steps:

- Let *error* be any error that results from [shader module creation](#) with the WGSL source *descriptor.code*, or null if no errors occurred.
- If any of the following requirements are unmet, [generate a validation error](#), [invalidate](#) *sm*, and return.
 - this* must not be [lost](#).
 - error* must not be a [shader-creation program error](#).
 - For each **enable** extension in *descriptor.code*, the corresponding [GPUFeatureName](#) must be enabled (see the [Feature Index](#)).

Note: [Uncategorized errors](#) cannot arise from shader module creation. Implementations which detect such errors during shader module creation must behave as if the shader module is valid, and defer surfacing the error until pipeline creation.

NOTE:

User agents **should not** include detailed compiler error messages or shader text in the [message](#) text of validation errors arising here; these details are accessible via [getCompilationInfo\(\)](#). User agents **should** surface human-readable, formatted error details *to developers* for easier debugging (for example as a warning in the browser developer console, expandable to show full shader source).

As shader compilation errors should be rare in production applications, user agents could choose to surface them *to developers* regardless of error handling ([GPU error scopes](#) or [uncapturederror](#) event handlers), e.g. as an expandable warning. If not, they should provide and document another way for developers to access human-readable error details, for example by adding a checkbox to show errors unconditionally, or by showing human-readable details when logging a [GPUCompilationInfo](#) object to the console.

Create a [GPUShaderModule](#) from WGSL code:

```
// A simple vertex and fragment shader pair that will fill the viewport with red.
const shaderSource = `
var<private> pos : array<vec2<f32>, 3> = array<vec2<f32>, 3>(
    vec2(-1.0, -1.0), vec2(-1.0, 3.0), vec2(3.0, -1.0));

@vertex
fn vertexMain(@builtin(vertex_index) vertexIndex : u32) -> @builtin(position) vec4<f32> {
    return vec4(pos[vertexIndex], 1.0, 1.0);
}

@fragment
fn fragmentMain() -> @location(0) vec4<f32> {
    return vec4(1.0, 0.0, 0.0, 1.0);
}
`;

const shaderModule = gpuDevice.createShaderModule({
    code: shaderSource,
});
```

9.1.1.1. Shader Module Compilation Hints

Shader module compilation hints are optional, additional information indicating how a given [GPUShaderModule](#) entry point is intended to be used in the future. For some implementations this information may aid in compiling the shader module earlier, potentially increasing performance.

dictionary

GPUShaderModuleCompilationHint

```
{
    required USVString
    entryPoint
```

```
;
    (GPUPipelineLayout or GPUAutoLayoutMode) layout;
};
```

layout, of type ([GPUPipelineLayout](#) or [GPUAutoLayoutMode](#))

A [GPUPipelineLayout](#) that the [GPUShaderModule](#) may be used with in a future [createComputePipeline\(\)](#) or [createRenderPipeline\(\)](#) call.

If set to ["auto"](#) the layout will be the [default pipeline layout](#) for the entry point associated with this hint will be used.

NOTE:

If possible, authors should be supplying the same information to [createShaderModule\(\)](#) and [createComputePipeline\(\)](#) / [createRenderPipeline\(\)](#).

If an application is unable to provide hint information at the time of calling [createShaderModule\(\)](#), it should usually not delay calling [createShaderModule\(\)](#), but instead just omit the unknown information from the [compilationHints](#) sequence or the individual members of [GPUShaderModuleCompilationHint](#). Omitting this information may cause compilation to be deferred to [createComputePipeline\(\)](#) / [createRenderPipeline\(\)](#).

If an author is not confident that the hint information passed to [createShaderModule\(\)](#) will match the information later passed to [createComputePipeline\(\)](#) / [createRenderPipeline\(\)](#) with that same module, they should avoid passing that information to [createShaderModule\(\)](#), as passing mismatched information to [createShaderModule\(\)](#) may cause unnecessary compilations to occur.

9.1.2. Shader Module Compilation Information

enum

GPUCompilationMessageType

```
{
    "error"
    ,
    "warning"
    ,
    "info"
    ,
};
```

[[Exposed](#)=(Window, Worker), [Serializable](#), [SecureContext](#)]

interface

GPUCompilationMessage

```
{
    readonly attribute DOMString message;
    readonly attribute GPUCompilationMessageType type;
    readonly attribute unsigned long long lineNum;
    readonly attribute unsigned long long linePos;
    readonly attribute unsigned long long offset;
    readonly attribute unsigned long long length;
};
```

[[Exposed](#)=(Window, Worker), [Serializable](#), [SecureContext](#)]

interface

GPUCompilationInfo

```
{
    readonly attribute FrozenArray<GPUCompilationMessage>
    messages
```

```
;
};
```

A [GPUCompilationMessage](#) is an informational, warning, or error message generated by the [GPUShaderModule](#) compiler. The messages are intended to be human readable to help developers diagnose issues with their shader [code](#). Each message may correspond to a single point or range of the shader source, or may be unassociated with any specific part of the code.

[GPUCompilationMessage](#) has the following attributes:

message, of type [DOMString](#), readonly

The human-readable, [localizable text](#) for this compilation message.

Note: The [message](#) should follow the [best practices for language and direction information](#). This includes making use of any future standards which may emerge regarding the reporting of string language and direction metadata.

Editorial note: At the time of this writing, no language/direction recommendation is available that provides compatibility and consistency with legacy APIs, but when there is, adopt it formally.

type, of type [GPUCompilationMessageType](#), readonly

The severity level of the message.

If the [type](#) is ["error"](#), it corresponds to a [shader-creation error](#).

lineNum, of type [unsigned long long](#), readonly

The line number in the shader [code](#) the [message](#) corresponds to. Value is one-based, such that a `lineNum` of 1 indicates the first line of the shader [code](#). Lines are delimited by [line breaks](#).

If the [message](#) corresponds to a substring this points to the line on which the substring begins. Must be 0 if the [message](#) does not correspond to any specific point in the shader [code](#).

linePos, of type [unsigned long long](#), readonly

The offset, in UTF-16 code units, from the beginning of line [lineNum](#) of the shader [code](#) to the point or beginning of the substring that the [message](#) corresponds to. Value is one-based, such that a `linePos` of 1 indicates the first code unit of the line.

If [message](#) corresponds to a substring this points to the first UTF-16 code unit of the substring. Must be 0 if the [message](#) does not correspond to any specific point in the shader [code](#).

offset, of type [unsigned long long](#), readonly

The offset from the beginning of the shader [code](#) in UTF-16 code units to the point or beginning of the substring that [message](#) corresponds to. Must reference the same position as [lineNum](#) and [linePos](#). Must be 0 if the [message](#) does not correspond to any specific point in the shader [code](#).

length, of type [unsigned long long](#), readonly

The number of UTF-16 code units in the substring that [message](#) corresponds to. If the message does not correspond with a substring then [length](#) must be 0.

Note: [GPUCompilationMessage.lineNum](#) and [GPUCompilationMessage.linePos](#) are one-based since the most common use for them is expected to be printing human readable messages that can be correlated with the line and column numbers shown in many text editors.

Note: [GPUCompilationMessage.offset](#) and [GPUCompilationMessage.length](#) are appropriate to pass to `substr()` in order to retrieve the substring of the shader [code](#) the [message](#) corresponds to.

getCompilationInfo()

Returns any messages generated during the [GPUShaderModule](#)'s compilation.

The locations, order, and contents of messages are [implementation-defined](#). In particular, messages aren't necessarily ordered by [lineNum](#).

Called on: [GPUShaderModule](#) this

Returns: [Promise](#)<[GPUCompilationInfo](#)>

[Content timeline](#) steps:

1. Let *contentTimeline* be the current [Content timeline](#).
2. Let *promise* be [a new promise](#).
3. Issue the *synchronization steps* on the [Device timeline](#) of *this*.
4. Return *promise*.

[Device timeline](#) synchronization steps:

1. Let *event* occur upon the (successful or unsuccessful) completion of [shader module creation](#) for *this*.
2. [Listen for timeline event](#) *event* on *this*.[\[\[device\]\]](#), handled by the subsequent steps on *contentTimeline*.

[Content timeline](#) steps:

1. Let *info* be a new [GPUCompilationInfo](#).
2. Let *messages* be a list of any errors, warnings, or informational messages generated during [shader module creation](#) for *this*, or the empty list `[]` if the device was lost.
3. For each *message* in *messages*:
 1. Let *m* be a new [GPUCompilationMessage](#).
 2. Set *m.message* to be the text of *message*.
 3. If *message* is a [shader-creation error](#):
 - Set *m.type* to `"error"`
 - If *message* is a warning:
 - Set *m.type* to `"warning"`
 - Otherwise:
 - Set *m.type* to `"info"`
4. If *message* is associated with a specific substring or position within the shader [code](#):
 1. Set *m.lineNum* to the one-based number of the first line that the message refers to.
 2. Set *m.linePos* to the one-based number of the first UTF-16 code units on *m.lineNum* that the message refers to, or 1 if the *message* refers to the entire line.
 3. Set *m.offset* to the number of UTF-16 code units from the beginning of the shader to beginning of the substring or position that *message* refers to.
 4. Set *m.length* the length of the substring in UTF-16 code units that *message* refers to, or 0 if *message* refers to a position
- Otherwise:
 1. Set *m.lineNum* to 0.
 2. Set *m.linePos* to 0.
 3. Set *m.offset* to 0.
 4. Set *m.length* to 0.
5. [Append](#) *m* to *info.messages*.
4. [Resolve](#) *promise* with *info*.

10. Pipelines

A *pipeline*, be it [GPUComputePipeline](#) or [GPURenderPipeline](#), represents the complete function done by a combination of the GPU hardware, the driver, and the user agent, that process the input data in the shape of bindings and vertex buffers, and produces some output, like the colors in the output render targets.

Structurally, the [pipeline](#) consists of a sequence of programmable stages (shaders) and fixed-function states, such as the blending modes.

Note: Internally, depending on the target platform, the driver may convert some of the fixed-function states into shader code, and link it together with the shaders provided by the user. This linking is one of the reason the object is created as a whole.

This combination state is created as a single object (a [GPUComputePipeline](#) or [GPURenderPipeline](#)) and switched using one command ([GPUComputePassEncoder.setPipeline\(\)](#) or [GPURenderCommandsMixin.setPipeline\(\)](#) respectively).

There are two ways to create pipelines:

immediate pipeline creation

[createComputePipeline\(\)](#) and [createRenderPipeline\(\)](#) return a pipeline object which can be used immediately in a pass encoder.

When this fails, the pipeline object will be invalid and the call will generate either a [validation error](#) or an [internal error](#).

Note: A handle object is returned immediately, but actual pipeline creation is not synchronous. If pipeline creation takes a long time, this can incur a stall in the [device timeline](#) at some point between the creation call and execution of the [submit\(\)](#) in which it is first used. The point is unspecified, but most likely to be one of: at creation, at the first usage of the pipeline in [setPipeline\(\)](#), at the corresponding [finish\(\)](#) of that [GPUCommandEncoder](#) or [GPURenderBundleEncoder](#), or at [submit\(\)](#) of that [GPUCommandBuffer](#).

async pipeline creation

[createComputePipelineAsync\(\)](#) and [createRenderPipelineAsync\(\)](#) return a `Promise` which resolves to a pipeline object when creation of the pipeline has completed.

When this fails, the `Promise` rejects with a [GPUPipelineError](#).

[GPUPipelineError](#) describes a pipeline creation failure.

[Exposed=(Window, Worker), SecureContext, Serializable]

interface [GPUPipelineError](#) : [DOMException](#) {

[constructor](#)(optional [DOMString](#) *message* = "", [GPUPipelineErrorInit](#) options);

readonly attribute [GPUPipelineErrorReason](#) *reason*;

};

dictionary

GPUPipelineErrorInit

{

required GPUPipelineErrorReason

reason

;

};

enum GPUPipelineErrorReason {

"validation",

"internal",

};

GPUPipelineError constructor:

constructor()

Arguments:

Arguments for the GPUPipelineError.constructor() method.

Parameter	Type	Nullable	Optional	Description
message	DOMString	✗	✓	Error message of the base DOMException.
options	GPUPipelineErrorInit	✗	✗	Options specific to GPUPipelineError.

Content timeline steps:

1. Set this.name to "GPUPipelineError".
2. Set this.message to message.
3. Set this.reason to options.reason.

GPUPipelineError has the following attributes:

reason, of type GPUPipelineErrorReason, readonly

A read-only slot-backed attribute exposing the type of error encountered in pipeline creation as a GPUPipelineErrorReason:

- "validation": A validation error.
- "internal": An internal error.

GPUPipelineError objects are serializable objects.

10.1. Base pipelines

enum

GPUAutoLayoutMode

{

"auto"

,

};

dictionary

GPUPipelineDescriptorBase

: GPUObjectDescriptorBase {

required (GPUPipelineLayout or GPUAutoLayoutMode) layout;

};

layout, of type (GPUPipelineLayout or GPUAutoLayoutMode)

The GPUPipelineLayout for this pipeline, or "auto" to generate the pipeline layout automatically.

Note: If "auto" is used the pipeline cannot share GPUBindGroups with any other pipelines.

interface mixin

GPUPipelineBase

```
{
  [NewObject] GPUBindGroupLayout getBindGroupLayout(unsigned long index);
};
```

GPUPipelineBase has the following device timeline properties:

[[layout]], of type GPUPipelineLayout

The definition of the layout of resources which can be used with this.

GPUPipelineBase has the following methods:

getBindGroupLayout(index)

Gets a GPUBindGroupLayout that is compatible with the GPUPipelineBase's GPUBindGroupLayout at index.

Called on: GPUPipelineBase this

Arguments:

Arguments for the GPUPipelineBase.getBindGroupLayout(index) method.

Parameter	Type	Nullable	Optional	Description
index	unsigned long	✗	✗	Index into the pipeline layout's [[bindGroupLayouts]] sequence.

Returns: GPUBindGroupLayout

Content timeline steps:

1. Let layout be a new GPUBindGroupLayout object.
2. Issue the initialization steps on the Device timeline of this.
3. Return layout.

Device timeline initialization steps:

1. Let limits be this. [[device]]. [[limits]].
2. If any of the following conditions are unsatisfied generate a validation error, invalidate layout and return.
 - this must be valid.
 - index < limits.maxBindGroups.
3. Initialize layout so it is a copy of this. [[layout]]. [[bindGroupLayouts]]. [index].

Note: GPUBindGroupLayout is only ever used by-value, not by-reference, so this is equivalent to returning the same internal object with a new WebGPU interface. A new GPUBindGroupLayout WebGPU interface is returned each time to avoid a round-trip between the Content timeline and the Device timeline.

10.1.1. Default pipeline layout

A GPUPipelineBase object that was created with a layout set to "auto" has a default layout created and used instead.

Note: Default layouts are provided as a convenience for simple pipelines, but use of explicit layouts is recommended in most cases. Bind groups created from default layouts cannot be used with other pipelines, and the structure of the default layout may change when altering shaders, causing unexpected bind group creation errors.

To create a default pipeline layout for GPUPipelineBase pipeline, run the following device timeline steps:

Let groupCount be 0.

Let groupDescs be a sequence of device. [[limits]]. maxBindGroups new GPUBindGroupLayoutDescriptor objects.

For each groupDesc in groupDescs:

Set groupDesc.entries to an empty sequence.

For each GPUProgrammableStage stageDesc in the descriptor used to create pipeline:

Let shaderStage be the GPUShaderStageFlags for the shader stage at which stageDesc is used in pipeline.

Let entryPoint be get the entry point(shaderStage, stageDesc). Assert entryPoint is not null.

For each resource resource statically used by entryPoint:

Let group be resource's "group" decoration.

Let binding be resource's "binding" decoration.

Let entry be a new GPUBindGroupLayoutEntry.

Set entry.[binding](#) to *binding*.

Set entry.[visibility](#) to *shaderStage*.

If *resource* is for a sampler binding:

Let *samplerLayout* be a new [GPUSamplerBindingLayout](#).

Set entry.[sampler](#) to *samplerLayout*.

If *resource* is for a comparison sampler binding:

Let *samplerLayout* be a new [GPUSamplerBindingLayout](#).

Set *samplerLayout.type* to ["comparison"](#).

Set entry.[sampler](#) to *samplerLayout*.

If *resource* is for a buffer binding:

Let *bufferLayout* be a new [GPUBufferBindingLayout](#).

Set *bufferLayout.minBindingSize* to *resource*'s [minimum buffer binding size](#).

If *resource* is for a read-only storage buffer:

Set *bufferLayout.type* to ["read-only-storage"](#).

If *resource* is for a storage buffer:

Set *bufferLayout.type* to ["storage"](#).

Set entry.[buffer](#) to *bufferLayout*.

If *resource* is for a sampled texture binding:

Let *textureLayout* be a new [GPUTextureBindingLayout](#).

If *resource* is a depth texture binding:

Set *textureLayout.sampleType* to ["depth"](#)

Otherwise, if the sampled type of *resource* is:

f32 and there exists a [static use](#) of *resource* by *stageDesc* in a texture builtin function call that also uses a sampler

Set *textureLayout.sampleType* to ["float"](#)

f32 otherwise

Set *textureLayout.sampleType* to ["unfilterable-float"](#)

i32

Set *textureLayout.sampleType* to ["sint"](#)

u32

Set *textureLayout.sampleType* to ["uint"](#)

Set *textureLayout.viewDimension* to *resource*'s dimension.

If *resource* is for a multisampled texture:

Set *textureLayout.multisampled* to `true`.

Set entry.[texture](#) to *textureLayout*.

If *resource* is for a storage texture binding:

Let *storageTextureLayout* be a new [GPUStorageTextureBindingLayout](#).

Set *storageTextureLayout.format* to *resource*'s format.

Set *storageTextureLayout.viewDimension* to *resource*'s dimension.

If the access mode is:

read

Set *textureLayout.access* to ["read-only"](#).

write

Set *textureLayout.access* to ["write-only"](#).

read_write

Set *textureLayout.access* to ["read-write"](#).

Set entry.[storageTexture](#) to *storageTextureLayout*.

Set *groupCount* to $\max(\text{groupCount}, \text{group} + 1)$.

If *groupDescs*[*group*] has an entry *previousEntry* with [binding](#) equal to *binding*:

If entry has different [visibility](#) than *previousEntry*:

Add the bits set in entry.[visibility](#) into *previousEntry.visibility*.

If resource is for a buffer binding and entry has greater [buffer.minBindingSize](#) than *previousEntry*:

Set *previousEntry.buffer.minBindingSize* to entry.[buffer.minBindingSize](#).

If resource is a sampled texture binding and entry has different [texture.sampleType](#) than *previousEntry* and both entry and *previousEntry* have [texture.sampleType](#) of either "[float](#)" or "[unfilterable-float](#)":

Set *previousEntry.texture.sampleType* to "[float](#)".

If any other property is unequal between entry and *previousEntry*:

Return `null` (which will cause the creation of the pipeline to fail).

If resource is a storage texture binding, entry.*storageTexture.access* is "[read-write](#)", *previousEntry.storageTexture.access* is "[write-only](#)", and *previousEntry.storageTexture.format* is compatible with [STORAGE_BINDING](#) and "[read-write](#)" according to the [§ 26.1.1 Plain color formats](#) table:

Set *previousEntry.storageTexture.access* to "[read-write](#)".

Otherwise:

Append entry to *groupDescs*[*group*].

Let *groupLayouts* be a new [list](#).

For each *i* from 0 to *groupCount* - 1, inclusive:

Let *groupDesc* be *groupDescs*[*i*].

Let *bindGroupLayout* be the result of calling *device.createBindGroupLayout()*(*groupDesc*).

Set *bindGroupLayout*.[\[\[exclusivePipeline\]\]](#) to *pipeline*.

Append *bindGroupLayout* to *groupLayouts*.

Let *desc* be a new [GPUPipelineLayoutDescriptor](#).

Set *desc.bindGroupLayouts* to *groupLayouts*.

Return *device.createPipelineLayout()*(*desc*).

10.1.2. GPUProgrammableStage

A [GPUProgrammableStage](#) describes the entry point in the user-provided [GPUShaderModule](#) that controls one of the programmable stages of a [pipeline](#). Entry point names follow the rules defined in [WGSL identifier comparison](#).

```
dictionary GPUProgrammableStage {  
    required GPUShaderModule module;  
    USVString entryPoint;  
    record<USVString, GPUPipelineConstantValue> constants = {};  
};
```

typedef [double GPUPipelineConstantValue](#); // May represent WGSL's bool, f32, i32, u32, and f16 if enabled.

[GPUProgrammableStage](#) has the following members:

module, of type [GPUShaderModule](#)

The [GPUShaderModule](#) containing the code that this programmable stage will execute.

entryPoint, of type [USVString](#)

The name of the function in *module* that this stage will use to perform its work.

NOTE: Since the [entryPoint](#) dictionary member is not required, methods which consume a [GPUProgrammableStage](#) must use the "[get the entry point](#)" algorithm to determine which entry point it refers to.

constants, of type record<[USVString](#), [GPUPipelineConstantValue](#)>, defaulting to {}

Specifies the values of [pipeline-overridable](#) constants in the shader module *module*.

Each such [pipeline-overridable](#) constant is uniquely identified by a single [pipeline-overridable constant identifier string](#), representing the [pipeline constant ID](#) of the constant if its declaration specifies one, and otherwise the constant's identifier name.

The key of each key-value pair must equal the [identifier string](#) of one such constant, with the comparison performed according to the rules for [WGSL identifier comparison](#). When the pipeline is executed, that constant will have the specified value.

Values are specified as *GPUPipelineConstantValue*, which is a [double](#). They are converted [to WGSL type](#) of the pipeline-overridable constant (bool/i32/u32/f32/f16). If conversion fails, a validation error is generated.

Pipeline-overridable constants defined in WGSL:

```
@id(0)    override has_point_light: bool = true; // Algorithmic control.
@id(1200) override specular_param: f32 = 2.3;    // Numeric control.
@id(1300) override gain: f32;                    // Must be overridden.
          override width: f32 = 0.0;              // Specified at the API level
          // using the name "width".
          override depth: f32;                    // Specified at the API level
          // using the name "depth".
          // Must be overridden.
          override height = 2 * depth;            // The default value
          // (if not set at the API level),
          // depends on another
          // overridable constant.
```

Corresponding JavaScript code, providing only the overrides which are required (have no defaults):

```
{
  // ...
  constants: {
    1300: 2.0, // "gain"
    depth: -1, // "depth"
  }
}
```

Corresponding JavaScript code, overriding all constants:

```
{
  // ...
  constants: {
    0: false, // "has_point_light"
    1200: 3.0, // "specular_param"
    1300: 2.0, // "gain"
    width: 20, // "width"
    depth: -1, // "depth"
    height: 15, // "height"
  }
}
```

To get the entry point([GPUShaderStage](#) *stage*, [GPUProgrammableStage](#) *descriptor*), run the following [device timeline](#) steps:

If *descriptor*.[entryPoint](#) is [provided](#):

If *descriptor*.[module](#) contains an entry point whose name equals *descriptor*.[entryPoint](#), and whose shader stage equals *stage*, return that entry point.

Otherwise, return `null`.

Otherwise:

If there is exactly one entry point in *descriptor*.[module](#) whose shader stage equals *stage*, return that entry point.

Otherwise, return `null`.

validating GPUProgrammableStage(stage, descriptor, layout, device)

Arguments:

[GPUShaderStage](#) *stage*

[GPUProgrammableStage](#) *descriptor*

[GPUPipelineLayout](#) *layout*

[GPUDevice](#) *device*

All of the requirements in the following steps *must* be met. If any are unmet, return `false`; otherwise, return `true`.

descriptor.[module](#) *must* be [valid to use with](#) *device*.

Let *entryPoint* be [get the entry point](#)(*stage*, *descriptor*).

entryPoint must not be `null`.

For each *binding* that is [statically used](#) by *entryPoint*:

[validating shader binding](#)(*binding*, *layout*) must return `true`.

For each texture builtin function call in any of the [functions in the shader stage](#) rooted at *entryPoint*, if it uses a *textureBinding* of [sampled texture](#) or [depth texture](#) type together with a *samplerBinding* of `sampler` type (excluding `sampler_comparison`):

Let *texture* be the [GPUBindGroupLayoutEntry](#) corresponding to *textureBinding*.

Let *sampler* be the [GPUBindGroupLayoutEntry](#) corresponding to *samplerBinding*.

If *sampler.type* is `"filtering"`, then *texture.sampleType* must be `"float"`.

Note: `"comparison"` samplers can also only be used with `"depth"` textures, because they are the only texture type that can be bound to WGSL `texture_depth_*` bindings.

For each *key* → *value* in *descriptor.constants*:

key must equal the [pipeline-overridable constant identifier string](#) of some [pipeline-overridable](#) constant defined in the shader module *descriptor.module* by the rules defined in [WGSL identifier comparison](#). The pipeline-overridable constant is *not* required to be [statically used](#) by *entryPoint*. Let the type of that constant be *T*.

Converting the IDL value *value* to WGSL type *T* must not throw a `TypeError`.

For each [pipeline-overridable constant identifier string](#) *key* which is [statically used](#) by *entryPoint*:

If the pipeline-overridable constant identified by *key* [does not have a default value](#), *descriptor.constants* must [contain](#) *key*.

[Pipeline-creation program errors](#) must not result from the rules of the [\[WGSL\]](#) specification.

[validating shader binding](#)(*variable*, *layout*)

Arguments:

shader binding declaration *variable*, a module-scope variable declaration reflected from a shader module

[GPUPipelineLayout](#) *layout*

Let *bindGroup* be the bind group index, and *bindIndex* be the binding index, of the shader binding declaration *variable*.

Return `true` if all of the following conditions are satisfied:

layout.[[bindGroupLayouts]][bindGroup] contains a [GPUBindGroupLayoutEntry](#) *entry* whose *entry.binding* == *bindIndex*.

If the defined [binding member](#) for *entry* is:

[buffer](#)

If *entry.buffer.type* is:

["uniform"](#)

variable is declared with address space `uniform`.

["storage"](#)

variable is declared with address space `storage` and access mode `read_write`.

["read-only-storage"](#)

variable is declared with address space `storage` and access mode `read`.

If *entry.buffer.minBindingSize* is not 0, then it must be at least the [minimum buffer binding size](#) for the associated buffer binding variable in the shader.

[sampler](#)

If *entry.sampler.type* is:

["filtering"](#) or ["non-filtering"](#)

variable has type `sampler`.

["comparison"](#)

variable has type `sampler_comparison`.

[texture](#)

If, and only if, *entry.texture.multisampled* is `true`, *variable* has type `texture_multisampled_2d<T>` or `texture_depth_multisampled_2d<T>`.

If *entry.texture.sampleType* is:

["float"](#), ["unfilterable-float"](#), ["sint"](#) or ["uint"](#)

variable has one of the types:

- texture_1d<T>
- texture_2d<T>
- texture_2d_array<T>
- texture_cube<T>
- texture_cube_array<T>
- texture_3d<T>
- texture_multisampled_2d<T>

If entry [texture.sampleType](#) is:

["float"](#) or ["unfilterable-float"](#)

The sampled type T is f32.

["sint"](#)

The sampled type T is i32.

["uint"](#)

The sampled type T is u32.

["depth"](#)

variable has one of the types:

- texture_2d<T>
- texture_2d_array<T>
- texture_cube<T>
- texture_cube_array<T>
- texture_multisampled_2d<T>
- texture_depth_2d
- texture_depth_2d_array
- texture_depth_cube
- texture_depth_cube_array
- texture_depth_multisampled_2d

where the sampled type T is f32.

If entry [texture.viewDimension](#) is:

["1d"](#)

variable has type texture_1d<T>.

["2d"](#)

variable has type texture_2d<T> or texture_multisampled_2d<T>.

["2d-array"](#)

variable has type texture_2d_array<T>.

["cube"](#)

variable has type texture_cube<T>.

["cube-array"](#)

variable has type texture_cube_array<T>.

["3d"](#)

variable has type texture_3d<T>.

[storageTexture](#)

If entry [storageTexture.viewDimension](#) is:

["1d"](#)

variable has type texture_storage_1d<T, A>.

["2d"](#)

variable has type texture_storage_2d<T, A>.

["2d-array"](#)

variable has type `texture_storage_2d_array<T, A>`.

["3d"](#)

variable has type `texture_storage_3d<T, A>`.

If entry.[storageTexture.access](#) is:

["write-only"](#)

The access mode A is `write`.

["read-only"](#)

The access mode A is `read`.

["read-write"](#)

The access mode A is `read_write` or `write`.

The texel format T equals entry.[storageTexture.format](#).

The *minimum buffer binding size* for a buffer binding variable *var* is computed as follows:

Let *T* be the [store type](#) of *var*.

If *T* is a [runtime-sized](#) array, or contains a runtime-sized array, replace that `array<E>` with `array<E, 1>`.

Note: This ensures there's always enough memory for one element, which allows array indices to be clamped to the length of the array resulting in an in-memory access.

Return [SizeOf\(T\)](#).

Note: Enforcing this lower bound ensures reads and writes via the buffer variable only access memory locations within the bound region of the buffer.

A resource binding, [pipeline-overrideable](#) constant, shader stage input, or shader stage output is considered to be *statically used* by an entry point if it is present in the [interface of the shader stage](#) for that entry point.

10.2. GPUComputePipeline

A [GPUComputePipeline](#) is a kind of [pipeline](#) that controls the compute shader stage, and can be used in [GPUComputePassEncoder](#).

Compute inputs and outputs are all contained in the bindings, according to the given [GPUPipelineLayout](#). The outputs correspond to [buffer](#) bindings with a type of ["storage"](#) and [storageTexture](#) bindings with a type of ["write-only"](#) or ["read-write"](#).

Stages of a compute [pipeline](#):

Compute shader

[[Exposed](#)=(Window, Worker), [SecureContext](#)]

interface [GPUComputePipeline](#) {

};

[GPUComputePipeline](#) includes [GPUObjectBase](#);

[GPUComputePipeline](#) includes [GPUPipelineBase](#);

10.2.1. Compute Pipeline Creation

A [GPUComputePipelineDescriptor](#) describes a compute [pipeline](#). See [§ 23.1 Computing](#) for additional details.

dictionary

[GPUComputePipelineDescriptor](#)

: [GPUPipelineDescriptorBase](#) {

required [GPUProgrammableStage](#) `compute`;

};

[GPUComputePipelineDescriptor](#) has the following members:

compute, of type [GPUProgrammableStage](#)

Describes the compute shader entry point of the [pipeline](#).

createComputePipeline(descriptor)

Creates a [GPUComputePipeline](#) using [immediate pipeline creation](#).

Called on: [GPUDevice](#) *this*.

Arguments:

Arguments for the [GPUDevice.createComputePipeline\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
-----------	------	----------	----------	-------------

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUComputePipelineDescriptor	✗	✗	Description of the GPUComputePipeline to create.

Returns: [GPUComputePipeline](#)

[Content timeline](#) steps:

1. Let *pipeline* be ! [create a new WebGPU object](#)(*this*, [GPUComputePipeline](#), *descriptor*).
2. Issue the *initialization steps* on the [Device timeline](#) of *this*.
3. Return *pipeline*.

[Device timeline](#) initialization steps:

1. Let *layout* be a new [default pipeline layout](#) for *pipeline* if *descriptor.layout* is "[auto](#)", and *descriptor.layout* otherwise.
2. All of the requirements in the following steps *must* be met. If any are unmet, [generate a validation error](#), [invalidate pipeline](#) and return.
 1. *layout* *must* be [valid to use with this](#).
 2. [validating GPUProgrammableStage](#)([COMPUTE](#), *descriptor.compute*, *layout*, *this*) *must* succeed.
 3. Let *entryPoint* be [get the entry point](#)([COMPUTE](#), *descriptor.compute*).
3. Let *entryPoint* be [get the entry point](#)([COMPUTE](#), *descriptor.compute*).
4. Let *workgroupStorageUsed* be the sum of [roundUp](#)(16, [SizeOf](#)(*T*)) over each type *T* of all variables with address space "[workgroup](#)" [statically used](#) by *entryPoint*.
workgroupStorageUsed *must* be ≤ *device.limits.maxComputeWorkgroupStorageSize*.
5. *entryPoint* *must* use ≤ *device.limits.maxComputeInvocationsPerWorkgroup* per workgroup.
6. Each component of *entryPoint*'s *workgroup_size* attribute *must* be ≤ the corresponding component in [*device.limits.maxComputeWorkgroupSizeX*, *device.limits.maxComputeWorkgroupSizeY*, *device.limits.maxComputeWorkgroupSizeZ*].
3. If any [pipeline-creation uncategorized errors](#) result from the implementation of pipeline creation, [generate an internal error](#), [invalidate pipeline](#) and return.
Note: Even if the implementation detected [uncategorized errors](#) in shader module creation, the error is surfaced here.
4. Set *pipeline*.[\[\[layout\]\]](#) to *layout*.

createComputePipelineAsync(descriptor)

Creates a [GPUComputePipeline](#) using [async pipeline creation](#). The returned [Promise](#) resolves when the created pipeline is ready to be used without additional delay.

If pipeline creation fails, the returned [Promise](#) rejects with an [GPUPipelineError](#). (A [GPUError](#) is not dispatched to the device.)

Note: Use of this method is preferred whenever possible, as it prevents blocking the [queue timeline](#) work on pipeline compilation.

Called on: [GPUDevice](#) *this*.

Arguments:

Arguments for the [GPUDevice.createComputePipelineAsync\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUComputePipelineDescriptor	✗	✗	Description of the GPUComputePipeline to create.

Returns: [Promise](#)<[GPUComputePipeline](#)>

[Content timeline](#) steps:

1. Let *contentTimeline* be the current [Content timeline](#).
2. Let *promise* be [a new promise](#).
3. Issue the *initialization steps* on the [Device timeline](#) of *this*.
4. Return *promise*.

[Device timeline](#) initialization steps:

1. Let *pipeline* be a new [GPUComputePipeline](#) created as if *this.createComputePipeline()* was called with *descriptor*, except capturing any errors as *error*, rather than dispatching them to the device.
2. Let *event* occur upon the (successful or unsuccessful) completion of [pipeline creation](#) for *pipeline*.
3. [Listen for timeline event](#) *event* on *this*.[\[\[device\]\]](#), handled by the subsequent steps on the [device timeline](#) of *this*.

[Device timeline](#) steps:

1. If *pipeline* is [valid](#) or *this* is [lost](#):

1. Issue the following steps on *contentTimeline*:

2. Return.

Note: No errors are generated from a device which is lost. See [§ 22 Errors & Debugging](#).

2. If *pipeline* is [invalid](#) and *error* is an [internal error](#), issue the following steps on *contentTimeline*, and return.

3. If *pipeline* is [invalid](#) and *error* is a [validation error](#), issue the following steps on *contentTimeline*, and return.

Creating a simple [GPUComputePipeline](#):

```
const computePipeline = gpuDevice.createComputePipeline({
  layout: pipelineLayout,
  compute: {
    module: computeShaderModule,
    entryPoint: 'computeMain',
  }
});
```

10.3. GPURenderPipeline

A [GPURenderPipeline](#) is a kind of [pipeline](#) that controls the vertex and fragment shader stages, and can be used in [GPURenderPassEncoder](#) as well as [GPURenderBundleEncoder](#).

Render [pipeline](#) inputs are:

bindings, according to the given [GPUPipelineLayout](#)

vertex and index buffers, described by [GPUVertexState](#)

the color attachments, described by [GPUColorTargetState](#)

optionally, the depth-stencil attachment, described by [GPUDepthStencilState](#)

Render [pipeline](#) outputs are:

[buffer](#) bindings with a [type](#) of ["storage"](#)

[storageTexture](#) bindings with a [access](#) of ["write-only"](#) or ["read-write"](#)

the color attachments, described by [GPUColorTargetState](#)

optionally, depth-stencil attachment, described by [GPUDepthStencilState](#)

A render [pipeline](#) is comprised of the following *render stages*:

Vertex fetch, controlled by [GPUVertexState.buffers](#)

Vertex shader, controlled by [GPUVertexState](#)

Primitive assembly, controlled by [GPUPrimitiveState](#)

Rasterization, controlled by [GPUPrimitiveState](#), [GPUDepthStencilState](#), and [GPUMultisampleState](#)

Fragment shader, controlled by [GPUFragmentState](#)

Stencil test and operation, controlled by [GPUDepthStencilState](#)

Depth test and write, controlled by [GPUDepthStencilState](#)

Output merging, controlled by [GPUFragmentState.targets](#)

[[Exposed](#)=(Window, Worker), [SecureContext](#)]

```
interface GPURenderPipeline {
};
```

[GPURenderPipeline](#) includes [GPUObjectBase](#);

[GPURenderPipeline](#) includes [GPUPipelineBase](#);

[GPURenderPipeline](#) has the following [device timeline properties](#):

[[\[descriptor\]](#)], of type [GPURenderPipelineDescriptor](#), readonly

The [GPURenderPipelineDescriptor](#) describing this pipeline.

All optional fields of [GPURenderPipelineDescriptor](#) are defined.

[[\[writesDepth\]](#)], of type [boolean](#), readonly

True if the pipeline writes to the depth component of the depth/stencil attachment

`[[writesStencil]]`, of type [boolean](#), readonly
True if the pipeline writes to the stencil component of the depth/stencil attachment

10.3.1. Render Pipeline Creation

A [GPURenderPipelineDescriptor](#) describes a render [pipeline](#) by configuring each of the [render stages](#). See [§ 23.2 Rendering](#) for additional details.

dictionary

`GPURenderPipelineDescriptor`

```
: GPUPipelineDescriptorBase {
required GPUVertexState vertex;
GPUPrimitiveState primitive = {};
GPUDepthStencilState depthStencil;
GPUMultisampleState multisample = {};
GPUFragmentState fragment;
};
```

[GPURenderPipelineDescriptor](#) has the following members:

- vertex*, of type [GPUVertexState](#)
Describes the vertex shader entry point of the [pipeline](#) and its input buffer layouts.
- primitive*, of type [GPUPrimitiveState](#), defaulting to {}
Describes the primitive-related properties of the [pipeline](#).
- depthStencil*, of type [GPUDepthStencilState](#)
Describes the optional depth-stencil properties, including the testing, operations, and bias.
- multisample*, of type [GPUMultisampleState](#), defaulting to {}
Describes the multi-sampling properties of the [pipeline](#).
- fragment*, of type [GPUFragmentState](#)
Describes the fragment shader entry point of the [pipeline](#) and its output colors. If not [provided](#), the [§ 23.2.8 No Color Output](#) mode is enabled.
- createRenderPipeline(descriptor)*
Creates a [GPURenderPipeline](#) using [immediate pipeline creation](#).
Called on: [GPUDevice](#) *this*.

Arguments:

Arguments for the GPUDevice.createRenderPipeline(descriptor) method.				
Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPURenderPipelineDescriptor	✗	✗	Description of the GPURenderPipeline to create.

Returns: [GPURenderPipeline](#)

[Content timeline](#) steps:

- If *descriptor*.[fragment](#) is [provided](#):
 - For each non-null *colorState* of *descriptor*.[fragment](#).[targets](#):
 - [? Validate texture format required features](#) of *colorState*.[format](#) with *this*.[\[\[device\]\]](#).
- If *descriptor*.[depthStencil](#) is [provided](#):
 - [? Validate texture format required features](#) of *descriptor*.[depthStencil](#).[format](#) with *this*.[\[\[device\]\]](#).
- Let *pipeline* be [! create a new WebGPU object](#)(*this*, [GPURenderPipeline](#), *descriptor*).
- Issue the *initialization steps* on the [Device timeline](#) of *this*.
- Return *pipeline*.

[Device timeline](#) *initialization steps*:

- Let *layout* be a new [default pipeline layout](#) for *pipeline* if *descriptor*.[layout](#) is ["auto"](#), and *descriptor*.[layout](#) otherwise.
- All of the requirements in the following steps *must* be met. If any are unmet, [generate a validation error](#), [invalidate pipeline](#), and return.
 - layout* *must* be [valid to use with this](#).
 - [validating GPURenderPipelineDescriptor](#)(*descriptor*, *layout*, *this*) *must* succeed.

- Let `vertexBufferCount` be the index of the last non-null entry in `descriptor.vertex.buffers`, plus 1; or 0 if there are none.
- layout. `[[bindGroupLayouts]].size + vertexBufferCount` must be \leq this. `[[device]].[[limits]].maxBindGroupsPlusVertexBuffers`.
- If any [pipeline-creation uncategorized errors](#) result from the implementation of pipeline creation, [generate an internal error](#), [invalidate pipeline](#) and return.

Note: Even if the implementation detected [uncategorized errors](#) in shader module creation, the error is surfaced here.

- Set `pipeline. [[descriptor]]` to `descriptor`.
- Set `pipeline. [[writesDepth]]` to false.
- Set `pipeline. [[writesStencil]]` to false.
- Let `depthStencil` be `descriptor.depthStencil`.
- If `depthStencil` is not null:
 - If `depthStencil.depthWriteEnabled` is [provided](#):
 - Set `pipeline. [[writesDepth]]` to `depthStencil.depthWriteEnabled`.
 - If `depthStencil.stencilWriteMask` is not 0:
 - Let `stencilFront` be `depthStencil.stencilFront`.
 - Let `stencilBack` be `depthStencil.stencilBack`.
 - Let `cullMode` be `descriptor.primitive.cullMode`.
 - If `cullMode` is not ["front"](#), and any of `stencilFront.passOp`, `stencilFront.depthFailOp`, or `stencilFront.failOp` is not ["keep"](#):
 - Set `pipeline. [[writesStencil]]` to true.
 - If `cullMode` is not ["back"](#), and any of `stencilBack.passOp`, `stencilBack.depthFailOp`, or `stencilBack.failOp` is not ["keep"](#):
 - Set `pipeline. [[writesStencil]]` to true.
 - Set `pipeline. [[layout]]` to `layout`.

`createRenderPipelineAsync(descriptor)`

Creates a [GPURenderPipeline](#) using [async pipeline creation](#). The returned [Promise](#) resolves when the created pipeline is ready to be used without additional delay.

If pipeline creation fails, the returned [Promise](#) rejects with an [GPUPipelineError](#). (A [GPUError](#) is not dispatched to the device.)

Note: Use of this method is preferred whenever possible, as it prevents blocking the [queue timeline](#) work on pipeline compilation.

Called on: [GPUDevice](#) *this*.

Arguments:

Arguments for the [GPUDevice.createRenderPipelineAsync\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<code>descriptor</code>	GPURenderPipelineDescriptor	✗	✗	Description of the GPURenderPipeline to create.

Returns: [Promise](#)<[GPURenderPipeline](#)>

[Content timeline](#) steps:

- Let `contentTimeline` be the current [Content timeline](#).
- Let `promise` be [a new promise](#).
- Issue the *initialization steps* on the [Device timeline](#) of *this*.
- Return `promise`.

[Device timeline](#) *initialization steps*:

- Let `pipeline` be a new [GPURenderPipeline](#) created as if this. [createRenderPipeline\(\)](#) was called with `descriptor`, except capturing any errors as *error*, rather than dispatching them to the device.
- Let *event* occur upon the (successful or unsuccessful) completion of [pipeline creation](#) for `pipeline`.
- [Listen for timeline event](#) *event* on this. `[[device]]`, handled by the subsequent steps on the [device timeline](#) of *this*.

[Device timeline](#) steps:

- If `pipeline` is [valid](#) or *this* is [lost](#):
 - Issue the following steps on `contentTimeline`:
 - Return.

Note: No errors are generated from a device which is lost. See [§ 22 Errors & Debugging](#).

2. If *pipeline* is [invalid](#) and *error* is an [internal error](#), issue the following steps on *contentTimeline*, and return.
3. If *pipeline* is [invalid](#) and *error* is a [validation error](#), issue the following steps on *contentTimeline*, and return.

validating GPURenderPipelineDescriptor(descriptor, layout, device)

Arguments:

[GPURenderPipelineDescriptor](#) *descriptor*

[GPUPipelineLayout](#) *layout*

[GPUDevice](#) *device*

[Device timeline](#) steps:

Return `true` if all of the following conditions are satisfied:

validating GPUVertexState(device, descriptor.[vertex](#), layout) succeeds.

If *descriptor.[fragment](#)* is [provided](#):

validating GPUFragmentState(device, descriptor.[fragment](#), layout) succeeds.

If the [sample_mask](#) builtin is a [shader stage output](#) of *descriptor.[fragment](#)*:

descriptor.[multisample.alphaToCoverageEnabled](#) is `false`.

If the [frag_depth](#) builtin is a [shader stage output](#) of *descriptor.[fragment](#)*:

descriptor.[depthStencil](#) must be [provided](#), and *descriptor.[depthStencil.format](#)* must have a [depth](#) aspect.

validating GPUPrimitiveState(descriptor.[primitive](#), device) succeeds.

If *descriptor.[depthStencil](#)* is [provided](#):

validating GPUDepthStencilState(descriptor.[depthStencil](#), descriptor.[primitive.topology](#)) succeeds.

validating GPUMultisampleState(descriptor.[multisample](#)) succeeds.

If *descriptor.[multisample.alphaToCoverageEnabled](#)* is `true`:

descriptor.[fragment](#) must be [provided](#).

descriptor.[fragment.targets](#)[0] must [exist](#) and be non-null.

descriptor.[fragment.targets](#)[0].[format](#) must be a [GPUTextureFormat](#) which is [blendable](#) and has an alpha channel.

There must exist at least one attachment, either:

A non-null value in *descriptor.[fragment.targets](#)*, or

A *descriptor.[depthStencil](#)*.

validating inter-stage interfaces(device, descriptor) returns `true`.

validating inter-stage interfaces(device, descriptor)

Arguments:

[GPUDevice](#) *device*

[GPURenderPipelineDescriptor](#) *descriptor*

Returns: [boolean](#)

[Device timeline](#) steps:

Let *maxVertexShaderOutputVariables* be *device.limits.[maxInterStageShaderVariables](#)*.

Let *maxVertexShaderOutputLocation* be *device.limits.[maxInterStageShaderVariables](#)* - 1.

If *descriptor.[primitive.topology](#)* is ["point-list"](#):

Decrement *maxVertexShaderOutputVariables* by 1.

If [clip_distances](#) is declared in the output of *descriptor.[vertex](#)*:

Let *clipDistancesSize* be the array size of [clip_distances](#).

Decrement *maxVertexShaderOutputVariables* by `ceil(clipDistancesSize / 4)`.

Decrement *maxVertexShaderOutputLocation* by `ceil(clipDistancesSize / 4)`.

Return **false** if any of the following requirements are unmet:

There must be no more than *maxVertexShaderOutputVariables* user-defined outputs for *descriptor.vertex*.

The [location](#) of each user-defined output of *descriptor.vertex* must be \leq *maxVertexShaderOutputLocation*.

If *descriptor.fragment* is provided:

Let *maxFragmentShaderInputVariables* be *device.limits.maxInterStageShaderVariables*.

If any of the `front_facing`, `sample_index`, or `sample_mask` [builtins](#) are an input of *descriptor.fragment*:

Decrement *maxFragmentShaderInputVariables* by 1.

Return **false** if any of the following requirements are unmet:

For each user-defined input of *descriptor.fragment* there must be a user-defined output of *descriptor.vertex* that [location](#), type, and [interpolation](#) of the input.

Note: Vertex-only pipelines **can** have user-defined outputs in the vertex stage; their values will be discarded.

There must be no more than *maxFragmentShaderInputVariables* user-defined inputs for *descriptor.fragment*.

[Assert](#) that the [location](#) of each user-defined input of *descriptor.fragment* is less than *device.limits.maxInterStageShaderVariables*. (This follows from the above rules.)

Return **true**.

Creating a simple [GPURenderPipeline](#):

```
const renderPipeline = gpuDevice.createRenderPipeline({
  layout: pipelineLayout,
  vertex: {
    module: shaderModule,
    entryPoint: 'vertexMain'
  },
  fragment: {
    module: shaderModule,
    entryPoint: 'fragmentMain',
    targets: [{
      format: 'bgra8unorm',
    }],
  }
});
```

10.3.2. Primitive State

dictionary

GPUPrimitiveState

```
{
  GPUPrimitiveTopology topology = "triangle-list";
  GPUIndexFormat stripIndexFormat;
  GPUFrontFace frontFace = "ccw";
  GPUCullMode cullMode = "none";

  // Requires "depth-clip-control" feature.
  boolean unclippedDepth = false;
};
```

[GPUPrimitiveState](#) has the following members, which describe how a [GPURenderPipeline](#) constructs and rasterizes primitives from its vertex inputs:

topology, of type [GPUPrimitiveTopology](#), defaulting to "triangle-list"

The type of primitive to be constructed from the vertex inputs.

stripIndexFormat, of type [GPUIndexFormat](#)

For pipelines with strip topologies ("[line-strip](#)" or "[triangle-strip](#)"), this determines the index buffer format and primitive restart value ("[uint16](#)"/0xFFFF or "[uint32](#)"/0xFFFFFFFF). It is not allowed on pipelines with non-strip topologies.

Note: Some implementations require knowledge of the primitive restart value to compile pipeline state objects.

To use a strip-topology pipeline with an indexed draw call ([drawIndexed\(\)](#) or [drawIndexedIndirect\(\)](#)), this must be set, and it must match the index buffer format used with the draw call (set in [setIndexBuffer\(\)](#)).

See [§ 23.2.3 Primitive Assembly](#) for additional details.

frontFace, of type [GPUFrontFace](#), defaulting to "ccw"

Defines which polygons are considered [front-facing](#).

cullMode, of type [GPUCullMode](#), defaulting to "none"

Defines which polygon orientation will be culled, if any.

unclippedDepth, of type [boolean](#), defaulting to false

If true, indicates that [depth clipping](#) is disabled.

Requires the "[depth-clip-control](#)" feature to be enabled.

validating GPUPrimitiveState(descriptor, device) **Arguments:**

[GPUPrimitiveState](#) *descriptor*

[GPUDevice](#) *device*

[Device timeline](#) steps:

Return **true** if all of the following conditions are satisfied:

If *descriptor.topology* is not "[line-strip](#)" or "[triangle-strip](#)":

descriptor.stripIndexFormat must not be [provided](#).

If *descriptor.unclippedDepth* is **true**:

["depth-clip-control"](#) must be [enabled for](#) *device*.

enum

[GPUPrimitiveTopology](#)

```
{  
  "point-list",  
  "line-list",  
  "line-strip",  
  "triangle-list",  
  "triangle-strip",  
};
```

[GPUPrimitiveTopology](#) defines the primitive type draw calls made with a [GPURenderPipeline](#) will use. See [§ 23.2.5 Rasterization](#) for additional details:

"point-list"

Each vertex defines a point primitive.

"line-list"

Each consecutive pair of two vertices defines a line primitive.

"line-strip"

Each vertex after the first defines a line primitive between it and the previous vertex.

"triangle-list"

Each consecutive triplet of three vertices defines a triangle primitive.

"triangle-strip"

Each vertex after the first two defines a triangle primitive between it and the previous two vertices.

enum

[GPUFrontFace](#)

```
{  
  "ccw",  
  "cw",  
};
```

[GPUFrontFace](#) defines which polygons are considered [front-facing](#) by a [GPURenderPipeline](#). See [§ 23.2.5.4 Polygon Rasterization](#) for additional details:

"ccw"

Polygons with vertices whose framebuffer coordinates are given in counter-clockwise order are considered [front-facing](#).

"cw"

Polygons with vertices whose framebuffer coordinates are given in clockwise order are considered [front-facing](#).

enum

GPUCullMode

```
{
    "none",
    "front",
    "back",
};
```

[GPUPrimitiveTopology](#) defines which polygons will be culled by draw calls made with a [GPURenderPipeline](#). See [§ 23.2.5.4 Polygon Rasterization](#) for additional details:

"none"
No polygons are discarded.

"front"
[Front-facing](#) polygons are discarded.

"back"
[Back-facing](#) polygons are discarded.

Note: [GPUFrontFace](#) and [GPUCullMode](#) have no effect on ["point-list"](#), ["line-list"](#), or ["line-strip"](#) topologies.

10.3.3. Multisample State

dictionary

GPUMultisampleState

```
{
    GPUSize32 count = 1;
    GPUSampleMask mask = 0xFFFFFFFF;
    boolean alphaToCoverageEnabled = false;
};
```

[GPUMultisampleState](#) has the following members, which describe how a [GPURenderPipeline](#) interacts with a render pass's multisampled attachments.

count, of type [GPUSize32](#), defaulting to 1
Number of samples per pixel. This [GPURenderPipeline](#) will be compatible only with attachment textures ([colorAttachments](#) and [depthStencilAttachment](#)) with matching [sampleCounts](#).

mask, of type [GPUSampleMask](#), defaulting to 0xFFFFFFFF
Mask determining which samples are written to.

alphaToCoverageEnabled, of type [boolean](#), defaulting to `false`
When `true` indicates that a fragment's alpha channel should be used to generate a sample coverage mask.

validating GPUMultisampleState(descriptor) **Arguments:**

[GPUMultisampleState](#) *descriptor*

[Device timeline](#) steps:

Return `true` if all of the following conditions are satisfied:

descriptor.count must be either 1 or 4.

If *descriptor.alphaToCoverageEnabled* is `true`:

descriptor.count > 1.

10.3.4. Fragment State

dictionary

GPUFragmentState

```
    : GPUProgrammableStage {
        required sequence<GPUColorTargetState?> targets;
    };
```

targets, of type `sequence<GPUColorTargetState?>`
A list of [GPUColorTargetState](#) defining the formats and behaviors of the color targets this pipeline writes to.

validating GPUFragmentState(device, descriptor, layout)

Arguments:

[GPUDevice](#) device

[GPUFragmentState](#) descriptor

[GPUPipelineLayout](#) layout

[Device timeline](#) steps:

Return `true` if all of the following requirements are met:

[validating GPUProgrammableStage](#)([FRAGMENT](#), descriptor, layout, device) succeeds.

descriptor.targets.size must be \leq device.[[limits]].maxColorAttachments.

Let entryPoint be [get the entry point](#)([FRAGMENT](#), descriptor).

Let usesDualSourceBlending be `false`.

For each index of the [indices](#) of descriptor.targets containing a non-null value colorState:

colorState.format must be listed in [§ 26.1.1 Plain color formats](#) with [RENDER_ATTACHMENT](#) capability.

colorState.writeMask must be < 16 .

If colorState.blend is [provided](#):

The colorState.format must be [blendable](#).

colorState.blend.color must be a [valid GPUBlendComponent](#).

colorState.blend.alpha must be a [valid GPUBlendComponent](#).

If colorState.blend.color.srcFactor or colorState.blend.color.dstFactor or colorState.blend.alpha.srcFactor or colorState.blend.alpha.dstFactor uses the second input of the corresponding blending unit (is any of "[src1](#)", "[one-minus-src1](#)", "[src1-alpha](#)", "[one-minus-src1-alpha](#)"), then:

Set usesDualSourceBlending to `true`.

For each [shader stage output](#) value output with [location](#) attribute equal to index in entryPoint:

For each component in colorState.format, there must be a corresponding component in output. (That is, RGBA requires vec4, RGB requires vec3 or vec4, RG requires vec2 or vec3 or vec4.)

If the [GPUTextureSampleTypes](#) for colorState.format (defined in [§ 26.1 Texture Format Capabilities](#)) are:

["float"](#) and/or ["unfilterable-float"](#)

output must have a floating-point scalar type.

["sint"](#)

output must have a signed integer scalar type.

["uint"](#)

output must have an unsigned integer scalar type.

If colorState.blend is [provided](#) and colorState.blend.color.srcFactor or .dstFactor uses the source alpha (is any of "[src-alpha](#)", "[one-minus-src-alpha](#)", "[src-alpha-saturated](#)", "[src1-alpha](#)" or "[one-minus-src1-alpha](#)"), then:

output must have an alpha channel (that is, it must be a vec4).

If colorState.writeMask is not 0:

entryPoint must have a [shader stage output](#) with [location](#) equal to index and [blend_src](#) omitted or equal to 0.

If usesDualSourceBlending is `true`:

descriptor.targets.size must be 1.

All the [shader stage outputs](#) with [location](#) in entryPoint must be in one struct and [use dual source blending](#).

[Validating GPUFragmentState's color attachment bytes per sample](#)(device, descriptor.targets) succeeds.

[Validating GPUFragmentState's color attachment bytes per sample](#)(device, targets)

Arguments:

[GPUDevice](#) device

[sequence](#)<[GPUColorTargetState](#)?> targets

[Device timeline](#) steps:

Let *formats* be an empty [list](#)<[GPUTextureFormat](#)?>

For each *target* in *targets*:

If *target* is `undefined`, continue.

[Append](#) *target.format* to *formats*.

[Calculating color attachment bytes per sample](#)(*formats*) must be \leq *device*.[\[\[limits\]\].maxColorAttachmentBytesPerSample](#).

Note: The fragment shader may output more values than what the pipeline uses. If that is the case the values are ignored.

[GPUBlendComponent](#) *component* is a valid [GPUBlendComponent](#) with logical [device](#) *device* if it meets the following requirements:

If *component.operation* is `"min"` or `"max"`:

component.srcFactor and *component.dstFactor* must both be `"one"`.

If *component.srcFactor* or *component.dstFactor* requires a feature according to the [GPUBlendFactor](#) table and *device*.[\[\[features\]\]](#) does not [contain](#) the feature:

Throw a [TypeError](#).

10.3.5. Color Target State

dictionary
[GPUColorTargetState](#)

```
{
  required GPUTextureFormat format;

  GPUBlendState blend;
  GPUColorWriteFlags writeMask = 0xF; // GPUColorWrite.ALL
};
```

format, of type [GPUTextureFormat](#)
The [GPUTextureFormat](#) of this color target. The pipeline will only be compatible with [GPURenderPassEncoder](#)s which use a [GPUTextureView](#) of this format in the corresponding color attachment.

blend, of type [GPUBlendState](#)
The blending behavior for this color target. If left undefined, disables blending for this color target.

writeMask, of type [GPUColorWriteFlags](#), defaulting to 0xF
Bitmask controlling which channels are written to when drawing to this color target.

dictionary
[GPUBlendState](#)

```
{
  required GPUBlendComponent color;
  required GPUBlendComponent alpha;
};
```

color, of type [GPUBlendComponent](#)
Defines the blending behavior of the corresponding render target for color channels.

alpha, of type [GPUBlendComponent](#)
Defines the blending behavior of the corresponding render target for the alpha channel.

typedef [\[EnforceRange\]](#) [unsigned long](#)
[GPUColorWriteFlags](#)

```
;
```

[\[Exposed=\(Window, Worker\), SecureContext\]](#)

namespace
[GPUColorWrite](#)

```
{
  const GPUFlagsConstant
```

[RED](#)

```
= 0x1;
const GPUFlagsConstant
GREEN

= 0x2;
const GPUFlagsConstant
BLUE

= 0x4;
const GPUFlagsConstant
ALPHA

= 0x8;
const GPUFlagsConstant
ALL

= 0xF;
};
```

10.3.5.1. Blend State

```
dictionary
GPUBlendComponent

{
    GPUBlendOperation operation = "add";
    GPUBlendFactor srcFactor = "one";
    GPUBlendFactor dstFactor = "zero";
};
```

[GPUBlendComponent](#) has the following members, which describe how the color or alpha components of a fragment are blended:

- operation*, of type [GPUBlendOperation](#), defaulting to "add"
Defines the [GPUBlendOperation](#) used to calculate the values written to the target attachment components.
- srcFactor*, of type [GPUBlendFactor](#), defaulting to "one"
Defines the [GPUBlendFactor](#) operation to be performed on values from the fragment shader.
- dstFactor*, of type [GPUBlendFactor](#), defaulting to "zero"
Defines the [GPUBlendFactor](#) operation to be performed on values from the target attachment.

The following tables use this notation to describe color components for a given fragment location:

RGBA _{src}	Color output by the fragment shader for the color attachment. If the shader doesn't return an alpha channel, src-alpha blend factors cannot be used.
RGBA _{src1}	Color output by the fragment shader for the color attachment with " @blend_src " attribute equal to 1. If the shader doesn't return an alpha channel, src1-alpha blend factors cannot be used.
RGBA _{dst}	Color currently in the color attachment. Missing green/blue/alpha channels default to 0, 0, 1, respectively.
RGBA _{const}	The current [[blendConstant]] .
RGBA _{srcFactor}	The source blend factor components, as defined by srcFactor .
RGBA _{dstFactor}	The destination blend factor components, as defined by dstFactor .

```
enum
GPUBlendFactor

{
    "zero",
    "one",
    "src",
    "one-minus-src",
    "src-alpha",
    "one-minus-src-alpha",
    "dst",
    "one-minus-dst",
    "dst-alpha",
    "one-minus-dst-alpha",
    "src-alpha-saturated",
    "constant",
}
```

```

"one-minus-constant",
"src1",
"one-minus-src1",
"src1-alpha",
"one-minus-src1-alpha",
};

```

[GPUBlendFactor](#) defines how either a source or destination blend factors is calculated:

GPUBlendFactor	Blend factor RGBA components	Feature
"zero"	(0, 0, 0, 0)	
"one"	(1, 1, 1, 1)	
"src"	(R _{src} , G _{src} , B _{src} , A _{src})	
"one-minus-src"	(1 - R _{src} , 1 - G _{src} , 1 - B _{src} , 1 - A _{src})	
"src-alpha"	(A _{src} , A _{src} , A _{src} , A _{src})	
"one-minus-src-alpha"	(1 - A _{src} , 1 - A _{src} , 1 - A _{src} , 1 - A _{src})	
"dst"	(R _{dst} , G _{dst} , B _{dst} , A _{dst})	
"one-minus-dst"	(1 - R _{dst} , 1 - G _{dst} , 1 - B _{dst} , 1 - A _{dst})	
"dst-alpha"	(A _{dst} , A _{dst} , A _{dst} , A _{dst})	
"one-minus-dst-alpha"	(1 - A _{dst} , 1 - A _{dst} , 1 - A _{dst} , 1 - A _{dst})	
"src-alpha-saturated"	(min(A _{src} , 1 - A _{dst}), min(A _{src} , 1 - A _{dst}), min(A _{src} , 1 - A _{dst}), 1)	
"constant"	(R _{const} , G _{const} , B _{const} , A _{const})	
"one-minus-constant"	(1 - R _{const} , 1 - G _{const} , 1 - B _{const} , 1 - A _{const})	
"src1"	(R _{src1} , G _{src1} , B _{src1} , A _{src1})	dual-source-blending
"one-minus-src1"	(1 - R _{src1} , 1 - G _{src1} , 1 - B _{src1} , 1 - A _{src1})	
"src1-alpha"	(A _{src1} , A _{src1} , A _{src1} , A _{src1})	
"one-minus-src1-alpha"	(1 - A _{src1} , 1 - A _{src1} , 1 - A _{src1} , 1 - A _{src1})	

enum

GPUBlendOperation

```

{
"add",
"subtract",
"reverse-subtract",
"min",
"max",
};

```

[GPUBlendOperation](#) defines the algorithm used to combine source and destination blend factors:

GPUBlendOperation	RGBA Components
"add"	$RGBA_{src} \times RGBA_{srcFactor} + RGBA_{dst} \times RGBA_{dstFactor}$
"subtract"	$RGBA_{src} \times RGBA_{srcFactor} - RGBA_{dst} \times RGBA_{dstFactor}$
"reverse-subtract"	$RGBA_{dst} \times RGBA_{dstFactor} - RGBA_{src} \times RGBA_{srcFactor}$
"min"	$\min(RGBA_{src}, RGBA_{dst})$
"max"	$\max(RGBA_{src}, RGBA_{dst})$

10.3.6. Depth/Stencil State

dictionary

GPUDepthStencilState

```

{
    required GPUTextureFormat format;

    boolean depthWriteEnabled;
    GPUCompareFunction depthCompare;

    GPUStencilFaceState stencilFront = {};
    GPUStencilFaceState stencilBack = {};
}

```

```

GPUStencilValue stencilReadMask = 0xFFFFFFFF;
GPUStencilValue stencilWriteMask = 0xFFFFFFFF;

GPUDepthBias depthBias = 0;
float depthBiasSlopeScale = 0;
float depthBiasClamp = 0;
};

```

[GPUDepthStencilState](#) has the following members, which describe how a [GPURenderPipeline](#) will affect a render pass's [depthStencilAttachment](#):

format, of type [GPUTextureFormat](#)

The *format* of [depthStencilAttachment](#) this [GPURenderPipeline](#) will be compatible with.

depthWriteEnabled, of type [boolean](#)

Indicates if this [GPURenderPipeline](#) can modify [depthStencilAttachment](#) depth values.

depthCompare, of type [GPUCompareFunction](#)

The comparison operation used to test fragment depths against [depthStencilAttachment](#) depth values.

stencilFront, of type [GPUStencilFaceState](#), defaulting to `{}`

Defines how stencil comparisons and operations are performed for front-facing primitives.

stencilBack, of type [GPUStencilFaceState](#), defaulting to `{}`

Defines how stencil comparisons and operations are performed for back-facing primitives.

stencilReadMask, of type [GPUStencilValue](#), defaulting to `0xFFFFFFFF`

Bitmask controlling which [depthStencilAttachment](#) stencil value bits are read when performing stencil comparison tests.

stencilWriteMask, of type [GPUStencilValue](#), defaulting to `0xFFFFFFFF`

Bitmask controlling which [depthStencilAttachment](#) stencil value bits are written to when performing stencil operations.

depthBias, of type [GPUDepthBias](#), defaulting to `0`

Constant depth bias added to each triangle fragment. See [biased fragment depth](#) for details.

depthBiasSlopeScale, of type [float](#), defaulting to `0`

Depth bias that scales with the triangle fragment's slope. See [biased fragment depth](#) for details.

depthBiasClamp, of type [float](#), defaulting to `0`

The maximum depth bias of a triangle fragment. See [biased fragment depth](#) for details.

Note: [depthBias](#), [depthBiasSlopeScale](#), and [depthBiasClamp](#) have no effect on ["point-list"](#), ["line-list"](#), and ["line-strip"](#) primitives, and must be `0`.

The *biased fragment depth* for a fragment being written to [depthStencilAttachment](#) *attachment* when drawing using [GPUDepthStencilState](#) *state* is calculated by running the following [queue timeline](#) steps:

Let *format* be *attachment.view.format*.

Let *r* be the minimum positive representable value > 0 in the *format* converted to a 32-bit float.

Let *maxDepthSlope* be the maximum of the horizontal and vertical slopes of the fragment's depth value.

If *format* is a **unorm** format:

Let *bias* be $(\text{float})\text{state}.\text{depthBias} * r + \text{state}.\text{depthBiasSlopeScale} * \text{maxDepthSlope}$.

Otherwise, if *format* is a **float** format:

Let *bias* be $(\text{float})\text{state}.\text{depthBias} * 2^{(\exp(\text{max depth in primitive}) - r)} + \text{state}.\text{depthBiasSlopeScale} * \text{maxDepthSlope}$.

If *state.depthBiasClamp* > 0 :

Set *bias* to $\min(\text{state}.\text{depthBiasClamp}, \text{bias})$.

Otherwise, if *state.depthBiasClamp* < 0 :

Set *bias* to $\max(\text{state}.\text{depthBiasClamp}, \text{bias})$.

If *state.depthBias* $\neq 0$ or *state.depthBiasSlopeScale* $\neq 0$:

Set the fragment depth value to `fragment depth value + bias`

validating GPUDepthStencilState(descriptor, topology)

Arguments:

[GPUDepthStencilState](#) *descriptor*

[GPUPrimitiveTopology](#) *topology*

[Device timeline](#) steps:

Return `true` if, and only if, all of the following conditions are satisfied:

descriptor.[format](#) is a [depth-or-stencil format](#).

If *descriptor*.[depthWriteEnabled](#) is `true` or *descriptor*.[depthCompare](#) is [provided](#) and not ["always"](#):

descriptor.[format](#) must have a depth component.

If *descriptor*.[stencilFront](#) or *descriptor*.[stencilBack](#) are not the default values:

descriptor.[format](#) must have a stencil component.

If *descriptor*.[format](#) has a depth component:

descriptor.[depthWriteEnabled](#) must be [provided](#).

descriptor.[depthCompare](#) must be [provided](#) if:

descriptor.[depthWriteEnabled](#) is `true`, or

descriptor.[stencilFront.depthFailOp](#) is not ["keep"](#), or

descriptor.[stencilBack.depthFailOp](#) is not ["keep"](#).

If *topology* is ["point-list"](#), ["line-list"](#), or ["line-strip"](#):

descriptor.[depthBias](#) must be 0.

descriptor.[depthBiasSlopeScale](#) must be 0.

descriptor.[depthBiasClamp](#) must be 0.

dictionary

[GPUStencilFaceState](#)

```
{
    GPUCompareFunction compare = "always";
    GPUStencilOperation failOp = "keep";
    GPUStencilOperation depthFailOp = "keep";
    GPUStencilOperation passOp = "keep";
};
```

[GPUStencilFaceState](#) has the following members, which describe how stencil comparisons and operations are performed:

compare, of type [GPUCompareFunction](#), defaulting to ["always"](#)

The [GPUCompareFunction](#) used when testing the [\[\[stencilReference\]\]](#) value against the fragment's [depthStencilAttachment](#) stencil values.

failOp, of type [GPUStencilOperation](#), defaulting to ["keep"](#)

The [GPUStencilOperation](#) performed if the fragment stencil comparison test described by [compare](#) fails.

depthFailOp, of type [GPUStencilOperation](#), defaulting to ["keep"](#)

The [GPUStencilOperation](#) performed if the fragment depth comparison described by [depthCompare](#) fails.

passOp, of type [GPUStencilOperation](#), defaulting to ["keep"](#)

The [GPUStencilOperation](#) performed if the fragment stencil comparison test described by [compare](#) passes.

enum

[GPUStencilOperation](#)

```
{
    "keep",
    "zero",
    "replace",
    "invert",
    "increment-clamp",
    "decrement-clamp",
    "increment-wrap",
    "decrement-wrap",
};
```

[GPUStencilOperation](#) defines the following operations:

["keep"](#)

Keep the current stencil value.

"zero"

Set the stencil value to 0.

"replace"

Set the stencil value to [\[\[stencilReference\]\]](#).

"invert"

Bitwise-invert the current stencil value.

"increment-clamp"

Increments the current stencil value, clamping to the maximum representable value of the [depthStencilAttachment](#)'s stencil aspect.

"decrement-clamp"

Decrement the current stencil value, clamping to 0.

"increment-wrap"

Increments the current stencil value, wrapping to zero if the value exceeds the maximum representable value of the [depthStencilAttachment](#)'s stencil aspect.

"decrement-wrap"

Decrement the current stencil value, wrapping to the maximum representable value of the [depthStencilAttachment](#)'s stencil aspect if the value goes below 0.

10.3.7. Vertex State

enum

[GPUIndexFormat](#)

```
{  
    "uint16",  
    "uint32",  
};
```

The index format determines both the data type of index values in a buffer and, when used with strip primitive topologies ("[line-strip](#)" or "[triangle-strip](#)") also specifies the primitive restart value. The *primitive restart value* indicates which index value indicates that a new primitive should be started rather than continuing to construct the triangle strip with the prior indexed vertices.

[GPUPrimitiveState](#) that specify a strip primitive topology must specify a [stripIndexFormat](#) if they are used for indexed draws so that the [primitive restart value](#) that will be used is known at pipeline creation time. [GPUPrimitiveState](#) that specify a list primitive topology will use the index format passed to [setIndexBuffer\(\)](#) when doing indexed rendering.

Index format	Byte size	Primitive restart value
"uint16"	2	0xFFFF
"uint32"	4	0xFFFFFFFF

10.3.7.1. Vertex Formats

The [GPUVertexFormat](#) of a vertex attribute indicates how data from a vertex buffer will be interpreted and exposed to the shader. The name of the format specifies the order of components, bits per component, and [vertex data type](#) for the component.

Each *vertex data type* can map to any [WGSL scalar type](#) of the same base type, regardless of the bits per component:

Vertex format prefix	Vertex data type	Compatible WGSL types
uint	unsigned int	u32
sint	signed int	i32
unorm	unsigned normalized	f16, f32
snorm	signed normalized	
float	floating point	

The multi-component formats specify the number of components after "x". Mismatches in the number of components between the vertex format and shader type are allowed, with components being either dropped or filled with default values to compensate.

A vertex attribute with a format of "[unorm8x2](#)" and byte values [0x7F, 0xFF] can be accessed in the shader with the following types:

Shader type	Shader value
f16	0.5h
f32	0.5f
vec2<f16>	vec2(0.5h, 1.0h)

Shader type	Shader value
vec2<f32>	vec2(0.5f, 1.0f)
vec3<f16>	vec2(0.5h, 1.0h, 0.0h)
vec3<f32>	vec2(0.5f, 1.0f, 0.0f)
vec4<f16>	vec2(0.5h, 1.0h, 0.0h, 1.0h)
vec4<f32>	vec2(0.5f, 1.0f, 0.0f, 1.0f)

See [§ 23.2.2 Vertex Processing](#) for additional information about how vertex formats are exposed in the shader.

```
enum
GPUVertexFormat
```

```
{
    "uint8",
    "uint8x2",
    "uint8x4",
    "sint8",
    "sint8x2",
    "sint8x4",
    "unorm8",
    "unorm8x2",
    "unorm8x4",
    "snorm8",
    "snorm8x2",
    "snorm8x4",
    "uint16",
    "uint16x2",
    "uint16x4",
    "sint16",
    "sint16x2",
    "sint16x4",
    "unorm16",
    "unorm16x2",
    "unorm16x4",
    "snorm16",
    "snorm16x2",
    "snorm16x4",
    "float16",
    "float16x2",
    "float16x4",
    "float32",
    "float32x2",
    "float32x3",
    "float32x4",
    "uint32",
    "uint32x2",
    "uint32x3",
    "uint32x4",
    "sint32",
    "sint32x2",
    "sint32x3",
    "sint32x4",
    "unorm10-10-10-2",
    "unorm8x4-bgra",
};
```

Vertex format	Data type	Components	byteSize	Example WGSL type
"uint8"	unsigned int	1	1	u32
"uint8x2"	unsigned int	2	2	vec2<u32>
"uint8x4"	unsigned int	4	4	vec4<u32>
"sint8"	signed int	1	1	i32

Vertex format	Data type	Components	byteSize	Example WGSL type
" <i>sint8x2</i> "	signed int	2	2	vec2<i32>
" <i>sint8x4</i> "	signed int	4	4	vec4<i32>
" <i>unorm8</i> "	unsigned normalized	1	1	f32
" <i>unorm8x2</i> "	unsigned normalized	2	2	vec2<f32>
" <i>unorm8x4</i> "	unsigned normalized	4	4	vec4<f32>
" <i>snorm8</i> "	signed normalized	1	1	f32
" <i>snorm8x2</i> "	signed normalized	2	2	vec2<f32>
" <i>snorm8x4</i> "	signed normalized	4	4	vec4<f32>
" <i>uint16</i> "	unsigned int	1	2	u32
" <i>uint16x2</i> "	unsigned int	2	4	vec2<u32>
" <i>uint16x4</i> "	unsigned int	4	8	vec4<u32>
" <i>sint16</i> "	signed int	1	2	i32
" <i>sint16x2</i> "	signed int	2	4	vec2<i32>
" <i>sint16x4</i> "	signed int	4	8	vec4<i32>
" <i>unorm16</i> "	unsigned normalized	1	2	f32
" <i>unorm16x2</i> "	unsigned normalized	2	4	vec2<f32>
" <i>unorm16x4</i> "	unsigned normalized	4	8	vec4<f32>
" <i>snorm16</i> "	signed normalized	1	2	f32
" <i>snorm16x2</i> "	signed normalized	2	4	vec2<f32>
" <i>snorm16x4</i> "	signed normalized	4	8	vec4<f32>
" <i>float16</i> "	float	1	2	f32
" <i>float16x2</i> "	float	2	4	vec2<f16>
" <i>float16x4</i> "	float	4	8	vec4<f16>
" <i>float32</i> "	float	1	4	f32
" <i>float32x2</i> "	float	2	8	vec2<f32>
" <i>float32x3</i> "	float	3	12	vec3<f32>
" <i>float32x4</i> "	float	4	16	vec4<f32>
" <i>uint32</i> "	unsigned int	1	4	u32
" <i>uint32x2</i> "	unsigned int	2	8	vec2<u32>
" <i>uint32x3</i> "	unsigned int	3	12	vec3<u32>
" <i>uint32x4</i> "	unsigned int	4	16	vec4<u32>
" <i>sint32</i> "	signed int	1	4	i32
" <i>sint32x2</i> "	signed int	2	8	vec2<i32>
" <i>sint32x3</i> "	signed int	3	12	vec3<i32>
" <i>sint32x4</i> "	signed int	4	16	vec4<i32>
" <i>unorm10-10-10-2</i> "	unsigned normalized	4	4	vec4<f32>
" <i>unorm8x4-bgra</i> "	unsigned normalized	4	4	vec4<f32>

enum

GPUVertexStepMode

```

{
  "vertex",
  "instance",
};

```

The step mode configures how an address for vertex buffer data is computed, based on the current vertex or instance index:

"vertex"

The address is advanced by [arrayStride](#) for each vertex, and reset between instances.

"instance"

The address is advanced by [arrayStride](#) for each instance.

dictionary

GPUVertexState

```
: GPUProgrammableStage {  
    sequence<GPUVertexBufferLayout?> buffers = [];  
};
```

buffers, of type `sequence<GPUVertexBufferLayout?>`, defaulting to `[]`

A list of [GPUVertexBufferLayout](#)s, each defining the layout of vertex attribute data in a vertex buffer used by this pipeline.

A *vertex buffer* is, conceptually, a view into buffer memory as an *array of structures*. [arrayStride](#) is the stride, in bytes, between *elements* of that array. Each element of a vertex buffer is like a *structure* with a memory layout defined by its [attributes](#), which describe the *members* of the structure.

Each [GPUVertexAttribute](#) describes its [format](#) and its [offset](#), in bytes, within the structure.

Each attribute appears as a separate input in a vertex shader, each bound by a numeric *location*, which is specified by [shaderLocation](#). Every location must be unique within the [GPUVertexState](#).

dictionary

GPUVertexBufferLayout

```
{  
    required GPUSize64 arrayStride;  
    GPUVertexStepMode stepMode = "vertex";  
    required sequence<GPUVertexAttribute> attributes;  
};
```

arrayStride, of type [GPUSize64](#)

The stride, in bytes, between elements of this array.

stepMode, of type [GPUVertexStepMode](#), defaulting to `"vertex"`

Whether each element of this array represents per-vertex data or per-instance data

attributes, of type `sequence<GPUVertexAttribute>`

An array defining the layout of the vertex attributes within each element.

dictionary

GPUVertexAttribute

```
{  
    required GPUVertexFormat format;  
    required GPUSize64 offset;  
  
    required GPUIndex32 shaderLocation;  
};
```

format, of type [GPUVertexFormat](#)

The [GPUVertexFormat](#) of the attribute.

offset, of type [GPUSize64](#)

The offset, in bytes, from the beginning of the element to the data for the attribute.

shaderLocation, of type [GPUIndex32](#)

The numeric location associated with this attribute, which will correspond with a `"@location"` [attribute](#) declared in the [vertex.module](#).

validating `GPUVertexBufferLayout(device, descriptor)`

Arguments:

[GPUDevice](#) *device*

[GPUVertexBufferLayout](#) *descriptor*

[Device timeline](#) steps:

Return `true`, if and only if, all of the following conditions are satisfied:

`descriptor.arrayStride ≤ device.[[device]].[[limits]].maxVertexBufferArrayStride`.

`descriptor.arrayStride` is a multiple of 4.

For each attribute *attrib* in the list `descriptor.attributes`:

If *descriptor.arrayStride* is zero:

attrib.offset + *byteSize*(*attrib.format*) ≤ *device*.*[[device]].[[limits]].maxVertexBufferArrayStride*.

Otherwise:

attrib.offset + *byteSize*(*attrib.format*) ≤ *descriptor.arrayStride*.

attrib.offset is a multiple of the minimum of 4 and *byteSize*(*attrib.format*).

attrib.shaderLocation is < *device*.*[[device]].[[limits]].maxVertexAttributes*.

validating GPUVertexState(*device*, *descriptor*, *layout*)

Arguments:

[GPUDevice](#) *device*

[GPUVertexState](#) *descriptor*

[GPUPipelineLayout](#) *layout*

[Device timeline](#) steps:

Let *entryPoint* be [get the entry point](#)([VERTEX](#), *descriptor*).

[Assert](#) *entryPoint* is not null.

All of the requirements in the following steps *must* be met.

[validating GPUProgrammableStage](#)([VERTEX](#), *descriptor*, *layout*, *device*) *must* succeed.

descriptor.buffers.size *must* be ≤ *device*.*[[device]].[[limits]].maxVertexBuffers*.

Each *vertexBuffer* layout descriptor in the list *descriptor.buffers* *must* pass [validating GPUVertexBufferLayout](#)(*device*, *vertexBuffer*).

The sum of *vertexBuffer.attributes.size*, over every *vertexBuffer* in *descriptor.buffers*, *must* be ≤ *device*.*[[device]].[[limits]].maxVertexAttributes*.

For every vertex attribute declaration (at location *location* with type *T*) that is [statically used](#) by *entryPoint*, there *must* be exactly one pair (*i*, *j*) for which *descriptor.buffers*[*i*].*attributes*[*j*].*shaderLocation* == *location*.

Let *attrib* be that [GPUVertexAttribute](#).

T *must* be compatible with *attrib.format*'s [vertex data type](#):

"unorm", "snorm", or "float"

T *must* be f32 or vecN<f32>.

"uint"

T *must* be u32 or vecN<u32>.

"sint"

T *must* be i32 or vecN<i32>.

11. Copies

11.1. Buffer Copies

Buffer copy operations operate on raw bytes.

WebGPU provides "buffered" [GPUCommandEncoder](#) commands:

[copyBufferToBuffer\(\)](#)

[clearBuffer\(\)](#)

and "immediate" [GPUQueue](#) operations:

[writeBuffer\(\)](#), for [ArrayBuffer](#)-to-[GPUBuffer](#) writes

11.2. Texel Copies

Texel copy operations operate on texture/"image" data, rather than bytes.

WebGPU provides "buffered" [GPUCommandEncoder](#) commands:

[copyTextureToTexture\(\)](#)

[copyBufferToTexture\(\)](#)

[copyTextureToBuffer\(\)](#)

and "immediate" [GPUQueue](#) operations:

[writeTexture\(\)](#), for [ArrayBuffer](#)-to-[GPUTexture](#) writes

[copyExternalImageToTexture\(\)](#), for copies from Web Platform image sources to textures

In a texel copy, the bytes written to the destination texel blocks will have an *equivalent texel representation* to the source value.

Texel copies only guarantee that valid, finite, non-subnormal numeric values in the source have the same numeric value in the destination. Specifically, the texel block may be decoded and re-encoded in a way that preserves only those values. Where multiple byte representations are possible, the choice of representation is implementation-defined.

Any floating-point zero value may be represented as either -0.0 or +0.0.

Any floating-point subnormal value may be either preserved or replaced by -0.0 or +0.0.

Any floating-point NaN or Infinity value may be replaced by an [indeterminate value](#).

Packed formats and `snorm` formats may change bit-representation as long as the represented values follow the rules above, for example:

`snorm` formats may represent -1.0 as either -127 or -128.

Formats like `"rgb9e5ufloat"` have multiple bit-representations of some values.

Note: For formats supporting [RENDER_ATTACHMENT](#) or [STORAGE_BINDING](#), this can be thought of as similar to, and may be implemented as, writing the texture using a WGLSL shader. In general, any [WGLSL floating point behaviors](#) may be observed.

The following definitions are used by these methods:

11.2.1. GPUTextureCopyBufferLayout

"[GPUTextureCopyBufferLayout](#)" describes the "layout" of texels in a "buffer" of bytes ([GPUBuffer](#) or [AllowSharedBufferSource](#)) in a "texel copy" operation.

dictionary [GPUTextureCopyBufferLayout](#) {

[GPUSize64](#) `offset` = 0;

[GPUSize32](#) `bytesPerRow`;

[GPUSize32](#) `rowsPerImage`;

};

A *texel image* is comprised of one or more rows of [texel blocks](#), referred to here as *texel block rows*. Each [texel block row](#) of a [texel image](#) must contain the same number of [texel blocks](#), and all [texel blocks](#) in a [texel image](#) are of the same [GPUTextureFormat](#).

A [GPUTextureCopyBufferLayout](#) is a layout of [texel images](#) within some linear memory. It's used when copying data between a [texture](#) and a [GPUBuffer](#), or when scheduling a write into a [texture](#) from the [GPUQueue](#).

For [2d](#) textures, data is copied between one or multiple contiguous [texel images](#) and [array layers](#).

For [3d](#) textures, data is copied between one or multiple contiguous [texel images](#) and depth [slices](#).

Operations that copy between byte arrays and textures always operate on whole [texel block](#). It's not possible to update only a part of a [texel block](#).

[Texel blocks](#) are tightly packed within each [texel block row](#) in the linear memory layout of a [texel copy](#), with each subsequent [texel block](#) immediately following the previous [texel block](#), with no padding. This includes [copies](#) to/from specific aspects of [depth-or-stencil format](#) textures: stencil values are tightly packed in an array of bytes; depth values are tightly packed in an array of the appropriate type ("depth16unorm" or "depth32float").

offset, of type [GPUSize64](#), defaulting to 0

The offset, in bytes, from the beginning of the texel data source (such as a [GPUTextureCopyBufferInfo.buffer](#)) to the start of the texel data within that source.

bytesPerRow, of type [GPUSize32](#)

The stride, in bytes, between the beginning of each [texel block row](#) and the subsequent [texel block row](#).

Required if there are multiple [texel block rows](#) (i.e. the copy height or depth is more than one block).

rowsPerImage, of type [GPUSize32](#)

Number of [texel block rows](#) per single [texel image](#) of the [texture](#). `rowsPerImage` × `bytesPerRow` is the stride, in bytes, between the beginning of each [texel image](#) of data and the subsequent [texel image](#).

Required if there are multiple [texel images](#) (i.e. the copy depth is more than one).

11.2.2. GPUTextureCopyBufferInfo

"[GPUTextureCopyBufferInfo](#)" describes the "info" ([GPUBuffer](#) and [GPUTextureCopyBufferLayout](#)) about a "buffer" source or destination of a "texel copy" operation. Together with the `copySize`, it describes the footprint of a region of texels in a [GPUBuffer](#).

dictionary [GPUTextureCopyBufferInfo](#)

 : [GPUTextureCopyBufferLayout](#) {

```
required GPUBuffer buffer;  
};
```

buffer, of type [GPUBuffer](#)

A buffer which either contains texel data to be copied or will store the texel data being copied, depending on the method it is being passed to.

validating [GPUTexelCopyBufferInfo](#)

Arguments:

[GPUTexelCopyBufferInfo](#) *imageCopyBuffer*

Returns: [boolean](#)

[Device timeline](#) steps:

Return `true` if and only if all of the following conditions are satisfied:

imageCopyBuffer.buffer must be a [valid GPUBuffer](#).

imageCopyBuffer.bytesPerRow must be a multiple of 256.

11.2.3. GPUTexelCopyTextureInfo

"[GPUTexelCopyTextureInfo](#)" describes the "info" ([GPUTexture](#), etc.) about a "texture" source or destination of a "texel copy" operation. Together with the `copySize`, it describes a sub-region of a texture (spanning one or more contiguous [texture subresources](#) at the same mip-map level).

dictionary [GPUTexelCopyTextureInfo](#) {

required [GPUTexture](#) *texture*;

[GPUIntegerCoordinate](#) *mipLevel* = 0;

[GPUOrigin3D](#) *origin* = {};

[GPUTextureAspect](#) *aspect* = "all";

};

texture, of type [GPUTexture](#)

Texture to copy to/from.

mipLevel, of type [GPUIntegerCoordinate](#), defaulting to 0

Mip-map level of the [texture](#) to copy to/from.

origin, of type [GPUOrigin3D](#), defaulting to {}

Defines the origin of the copy - the minimum corner of the texture sub-region to copy to/from. Together with `copySize`, defines the full copy sub-region.

aspect, of type [GPUTextureAspect](#), defaulting to "all"

Defines which aspects of the [texture](#) to copy to/from.

The texture copy sub-region for depth slice or array layer index of [GPUTexelCopyTextureInfo](#) *copyTexture* is determined by running the following steps:

Let *texture* be *copyTexture.texture*.

If *texture.dimension* is:

[1d](#)

1. [Assert](#) *index* is 0

2. Let *depthSliceOrLayer* be *texture*

[2d](#)

Let *depthSliceOrLayer* be array layer index of *texture*

[3d](#)

Let *depthSliceOrLayer* be depth slice index of *texture*

Let *textureMip* be mip level *copyTexture.mipLevel* of *depthSliceOrLayer*.

Return aspect *copyTexture.aspect* of *textureMip*.

The texel block byte offset of data described by [GPUTexelCopyBufferLayout](#) *bufferLayout* corresponding to [texel block](#) *x*, *y* of depth slice or array layer *z* of a [GPUTexture](#) *texture* is determined by running the following steps:

Let *blockBytes* be the [texel block copy footprint](#) of *texture.format*.

Let *imageOffset* be $(z \times \text{bufferLayout.rowsPerImage} \times \text{bufferLayout.bytesPerRow}) + \text{bufferLayout.offset}$.

Let *rowOffset* be $(y \times \text{bufferLayout.bytesPerRow}) + \text{imageOffset}$.

Let *blockOffset* be $(x \times \text{blockBytes}) + \text{rowOffset}$.

Return *blockOffset*.

validating GPUTextureCopyTextureInfo(*texelCopyTextureInfo*, *copySize*)

Arguments:

[GPUTextureCopyTextureInfo](#) *texelCopyTextureInfo*

[GPUExtent3D](#) *copySize*

Returns: [boolean](#)

[Device timeline](#) steps:

Let *blockWidth* be the [texel block width](#) of *texelCopyTextureInfo.texture.format*.

Let *blockHeight* be the [texel block height](#) of *texelCopyTextureInfo.texture.format*.

Return **true** if and only if all of the following conditions apply:

[validating texture copy range](#)(*texelCopyTextureInfo*, *copySize*) returns **true**.

texelCopyTextureInfo.texture must be a [valid GPUTexture](#).

texelCopyTextureInfo.mipLevel must be < *texelCopyTextureInfo.texture.mipLevelCount*.

texelCopyTextureInfo.origin.x must be a multiple of *blockWidth*.

texelCopyTextureInfo.origin.y must be a multiple of *blockHeight*.

The [GPUTextureCopyTextureInfo physical subresource size](#) of *texelCopyTextureInfo* is equal to *copySize* if either of the following conditions is true:

texelCopyTextureInfo.texture.format is a depth-stencil format.

texelCopyTextureInfo.texture.sampleCount > 1.

validating texture buffer copy(*texelCopyTextureInfo*, *bufferLayout*, *dataLength*, *copySize*, *textureUsage*, *aligned*)

Arguments:

[GPUTextureCopyTextureInfo](#) *texelCopyTextureInfo*

[GPUTextureCopyBufferLayout](#) *bufferLayout*

[GPUSize64Out](#) *dataLength*

[GPUExtent3D](#) *copySize*

[GPUTextureUsage](#) *textureUsage*

[boolean](#) *aligned*

Returns: [boolean](#)

[Device timeline](#) steps:

Let *texture* be *texelCopyTextureInfo.texture*

Let *aspectSpecificFormat* = *texture.format*.

Let *offsetAlignment* = [texel block copy footprint](#) of *texture.format*.

Return **true** if and only if all of the following conditions apply:

[validating GPUTextureCopyTextureInfo](#)(*texelCopyTextureInfo*, *copySize*) returns **true**.

texture.sampleCount is 1.

texture.usage contains *textureUsage*.

If *texture.format* is a [depth-or-stencil format](#) format:

texelCopyTextureInfo.aspect must refer to a single aspect of *texture.format*.

If *textureUsage* is:

[COPY_SRC](#)

That aspect must be a valid [texel copy](#) source according to [§ 26.1.2 Depth-stencil formats](#).

[COPY_DST](#)

That aspect must be a valid [texel copy](#) destination according to [§ 26.1.2 Depth-stencil formats](#).

Set *aspectSpecificFormat* to the [aspect-specific format](#) according to [§ 26.1.2 Depth-stencil formats](#).

Set *offsetAlignment* to 4.

If *aligned* is `true`:

bufferLayout.offset is a multiple of *offsetAlignment*.

[validating linear texture data](#)(*bufferLayout*, *dataLength*, *aspectSpecificFormat*, *copySize*) succeeds.

11.2.4. GPUCopyExternalImageDestInfo

WebGPU textures hold raw numeric data, and are not tagged with semantic metadata describing colors. However, [copyExternalImageToTexture\(\)](#) copies from sources that describe colors.

"GPUCopyExternalImageDestInfo" describes the "info" about the "destination" of a "[copyExternalImageToTexture\(\)](#)" operation. It is a [GPUTexelCopyTextureInfo](#) which is additionally tagged with color space/encoding and alpha-premultiplication metadata, so that semantic color data may be preserved during copies. This metadata affects only the semantics of the copy operation operation, not the state or semantics of the destination texture object.

dictionary [GPUCopyExternalImageDestInfo](#)

```
: GPUTexelCopyTextureInfo {  
  PredefinedColorSpace colorSpace = "srgb";  
  boolean premultipliedAlpha = false;  
};
```

colorSpace, of type [PredefinedColorSpace](#), defaulting to "srgb"

Describes the color space and encoding used to encode data into the destination texture.

This [may result](#) in values outside of the range [0, 1] being written to the target texture, if its format can represent them. Otherwise, the results are clamped to the target texture format's range.

Note: If [colorSpace](#) matches the source image, conversion might not be necessary. See [§ 3.11.2 Color Space Conversion Elision](#).

premultipliedAlpha, of type [boolean](#), defaulting to `false`

Describes whether the data written into the texture should have its RGB channels premultiplied by the alpha channel, or not.

If this option is set to `true` and the [source](#) is also premultiplied, the source RGB values must be preserved even if they exceed their corresponding alpha values.

Note: If [premultipliedAlpha](#) matches the source image, conversion might not be necessary. See [§ 3.11.2 Color Space Conversion Elision](#).

11.2.5. GPUCopyExternalImageSourceInfo

"GPUCopyExternalImageSourceInfo" describes the "info" about the "source" of a "[copyExternalImageToTexture\(\)](#)" operation.

typedef ([ImageBitmap](#) or

[ImageData](#) or
[HTMLImageElement](#) or
[HTMLVideoElement](#) or
[VideoFrame](#) or
[HTMLCanvasElement](#) or
[OffscreenCanvas](#))

[GPUCopyExternalImageSource](#)

;

dictionary [GPUCopyExternalImageSourceInfo](#) {

required [GPUCopyExternalImageSource](#) source;

[GPUOrigin2D](#) origin = {};

[boolean](#) flipY = false;

};

[GPUCopyExternalImageSourceInfo](#) has the following members:

source, of type [GPUCopyExternalImageSource](#)

The source of the [texel copy](#). The copy source data is captured at the moment that [copyExternalImageToTexture\(\)](#) is issued. Source size is determined as described by the [external source dimensions](#) table.

origin, of type [GPUOrigin2D](#), defaulting to {}

Defines the origin of the copy - the minimum (top-left) corner of the source sub-region to copy from. Together with *copySize*, defines the full copy sub-region.

flipY, of type [boolean](#), defaulting to `false`

Describes whether the source image is vertically flipped, or not.

If this option is set to `true`, the copy is flipped vertically: the bottom row of the source region is copied into the first row of the destination region, and so on. The [origin](#) option is still relative to the top-left corner of the source image, increasing downward.

When external sources are used when creating or copying to textures, the *external source dimensions* are defined by the source type, given by this table:

External Source type	Dimensions
ImageBitmap	ImageBitmap.width , ImageBitmap.height
HTMLImageElement	HTMLImageElement.naturalWidth , HTMLImageElement.naturalHeight
HTMLVideoElement	intrinsic width of the frame , intrinsic height of the frame
VideoFrame	VideoFrame.displayWidth , VideoFrame.displayHeight
ImageData	ImageData.width , ImageData.height
HTMLCanvasElement or OffscreenCanvas with CanvasRenderingContext2D or GPUCanvasContext	HTMLCanvasElement.width , HTMLCanvasElement.height
HTMLCanvasElement or OffscreenCanvas with WebGLRenderingContextBase	WebGLRenderingContextBase.drawingBufferWidth , WebGLRenderingContextBase.drawingBufferHeight
HTMLCanvasElement or OffscreenCanvas with ImageBitmapRenderingContext	ImageBitmapRenderingContext 's internal output bitmap ImageBitmap.width , ImageBitmap.height

11.2.6. Subroutines

GPUTexelCopyTextureInfo physical subresource size

Arguments:

[GPUTexelCopyTextureInfo](#) *texelCopyTextureInfo*

Returns: [GPUExtent3D](#)

The [GPUTexelCopyTextureInfo](#) *physical subresource size* of *texelCopyTextureInfo* is calculated as follows:

Its [width](#), [height](#) and [depthOrArrayLayers](#) are the width, height, and depth, respectively, of the [physical miplevel-specific texture extent](#) of *texelCopyTextureInfo*. [texture subresource](#) at [mipmap level](#) *texelCopyTextureInfo.mipLevel*.

validating linear texture data(*layout*, *byteSize*, *format*, *copyExtent*)

Arguments:

[GPUTexelCopyBufferLayout](#) *layout*

Layout of the linear texture data.

[GPUSize64](#) *byteSize*

Total size of the linear data, in bytes.

[GPUTextureFormat](#) *format*

Format of the texture.

[GPUExtent3D](#) *copyExtent*

Extent of the texture to copy.

[Device timeline](#) steps:

Let:

widthInBlocks be *copyExtent.width* ÷ the [texel block width](#) of *format*. [Assert](#) this is an integer.

heightInBlocks be *copyExtent.height* ÷ the [texel block height](#) of *format*. [Assert](#) this is an integer.

bytesInLastRow be *widthInBlocks* × the [texel block copy footprint](#) of *format*.

Fail if the following input validation requirements are not met:

If *heightInBlocks* > 1, *layout.bytesPerRow* must be specified.

If *copyExtent.depthOrArrayLayers* > 1, *layout.bytesPerRow* and *layout.rowsPerImage* must be specified.

If specified, *layout.bytesPerRow* must be ≥ *bytesInLastRow*.

If specified, *layout.rowsPerImage* must be ≥ *heightInBlocks*.

Let:

bytesPerRow be *layout.bytesPerRow* ?? 0.

rowsPerImage be *layout.rowsPerImage* ?? 0.

Note: These default values have no effect, as they're always multiplied by 0.

Let *requiredBytesInCopy* be 0.

If *copyExtent.depthOrArrayLayers* > 0:

Increment *requiredBytesInCopy* by $\text{bytesPerRow} \times \text{rowsPerImage} \times (\text{copyExtent.depthOrArrayLayers} - 1)$.

If *heightInBlocks* > 0:

Increment *requiredBytesInCopy* by $\text{bytesPerRow} \times (\text{heightInBlocks} - 1) + \text{bytesInLastRow}$.

Fail if the following condition is not satisfied:

The layout fits inside the linear data: $\text{layout.offset} + \text{requiredBytesInCopy} \leq \text{byteSize}$.

validating texture copy range

Arguments:

[GPUTexelCopyTextureInfo](#) *texelCopyTextureInfo*

The texture subresource being copied into and copy origin.

[GPUExtent3D](#) *copySize*

The size of the texture.

[Device timeline](#) steps:

Let *blockWidth* be the [texel block width](#) of *texelCopyTextureInfo.texture.format*.

Let *blockHeight* be the [texel block height](#) of *texelCopyTextureInfo.texture.format*.

Let *subresourceSize* be the [GPUTexelCopyTextureInfo.physical subresource size](#) of *texelCopyTextureInfo*.

Return whether all the conditions below are satisfied:

$(\text{texelCopyTextureInfo.origin.x} + \text{copySize.width}) \leq \text{subresourceSize.width}$

$(\text{texelCopyTextureInfo.origin.y} + \text{copySize.height}) \leq \text{subresourceSize.height}$

$(\text{texelCopyTextureInfo.origin.z} + \text{copySize.depthOrArrayLayers}) \leq \text{subresourceSize.depthOrArrayLayers}$

copySize.width must be a multiple of *blockWidth*.

copySize.height must be a multiple of *blockHeight*.

Note: The texture copy range is validated against the *physical* (rounded-up) size for [compressed formats](#), allowing copies to access texture blocks which are not fully inside the texture.

Two [GPUTextureFormats](#) *format1* and *format2* are *copy-compatible* if:

format1 equals *format2*, or

format1 and *format2* differ only in whether they are **srgb** formats (have the **-srgb** suffix).

The set of subresources for texture copy(*texelCopyTextureInfo*, *copySize*) is the subset of subresources of *texture* = *texelCopyTextureInfo.texture* for which each subresource *s* satisfies the following:

The [mipmap level](#) of *s* equals *texelCopyTextureInfo.mipLevel*.

The [aspect](#) of *s* is in the [set of aspects](#) of *texelCopyTextureInfo.aspect*.

If *texture.dimension* is "2d":

The [array layer](#) of *s* is $\geq \text{texelCopyTextureInfo.origin.z}$ and $< \text{texelCopyTextureInfo.origin.z} + \text{copySize.depthOrArrayLayers}$.

12. Command Buffers

Command buffers are pre-recorded lists of [GPU commands](#) (blocks of [queue timeline](#) steps) that can be submitted to a [GPUQueue](#) for execution. Each *GPU command* represents a task to be performed on the [queue timeline](#), such as setting state, drawing, copying resources, etc.

A [GPUCommandBuffer](#) can only be submitted once, at which point it becomes [invalidated](#). To reuse rendering commands across multiple submissions, use [GPURenderBundle](#).

12.1. GPUCommandBuffer

[Exposed=(Window, Worker), [SecureContext](#)]

```
interface GPUCommandBuffer {  
};
```

[GPUCommandBuffer](#) includes [GPUObjectBase](#);

[GPUCommandBuffer](#) has the following [device timeline properties](#):

[[command_list]], of type [list](#)<[GPU command](#)>, readonly

A [list](#) of [GPU commands](#) to be executed on the [Queue timeline](#) when this command buffer is submitted.

`[[renderState]]`, of type [RenderState](#), initially `null`

The current state used by any render pass commands being executed.

12.1.1. Command Buffer Creation

dictionary

`GPUCommandBufferDescriptor`

```
: GPUObjectDescriptorBase {  
};
```

13. Command Encoding

13.1. GPUCommandsMixin

[GPUCommandsMixin](#) defines state common to all interfaces which encode commands. It has no methods.

```
interface mixin GPUCommandsMixin {  
};
```

[GPUCommandsMixin](#) has the following [device timeline properties](#):

`[[state]]`, of type [encoder state](#), initially `"open"`

The current state of the encoder.

`[[commands]]`, of type [list](#)<[GPU command](#)>, initially `[]`

A [list](#) of [GPU commands](#) to be executed on the [Queue timeline](#) when a [GPUCommandBuffer](#) containing these commands is submitted.

The *encoder state* may be one of the following:

`"open"`

The encoder is available to encode new commands.

`"locked"`

The encoder cannot be used, because it is locked by a child encoder: it is a [GPUCommandEncoder](#), and a [GPURenderPassEncoder](#) or [GPUComputePassEncoder](#) is active. The encoder becomes `"open"` again when the pass is ended.

Any command issued in this state [invalidates](#) the encoder.

`"ended"`

The encoder has been ended and new commands can no longer be encoded.

Any command issued in this state will [generate a validation error](#).

To Validate the *encoder state* of [GPUCommandsMixin](#) encoder run the following [device timeline](#) steps:

If encoder.`[[state]]` is:

`"open"`

Return `true`.

`"locked"`

[Invalidate](#) encoder and return `false`.

`"ended"`

[Generate a validation error](#), and return `false`.

To *Enqueue a command* on [GPUCommandsMixin](#) encoder which issues the steps of a [GPU Command command](#), run the following [device timeline](#) steps:

[Append](#) command to encoder.`[[commands]]`.

When *command* is executed as part of a [GPUCommandBuffer](#):

Issue the steps of *command*.

13.2. GPUCommandEncoder

[[Exposed](#)=(Window, Worker), [SecureContext](#)]

```
interface GPUCommandEncoder {
```

```
  GPURenderPassEncoder beginRenderPass(GPURenderPassDescriptor descriptor);
```

```
  GPUComputePassEncoder beginComputePass(optional GPUComputePassDescriptor descriptor = {});
```

```

    undefined copyBufferToBuffer(
        GPUBuffer
source
    ,
        GPUBuffer
destination
    ,
        optional GPUSize64
size
    );

undefined copyBufferToBuffer(
    GPUBuffer source,
    GPUSize64 sourceOffset,
    GPUBuffer destination,
    GPUSize64 destinationOffset,
    optional GPUSize64 size);

undefined copyBufferToTexture(
    GPUTexelCopyBufferInfo source,
    GPUTexelCopyTextureInfo destination,
    GPUExtent3D copySize);

undefined copyTextureToBuffer(
    GPUTexelCopyTextureInfo source,
    GPUTexelCopyBufferInfo destination,
    GPUExtent3D copySize);

undefined copyTextureToTexture(
    GPUTexelCopyTextureInfo source,
    GPUTexelCopyTextureInfo destination,
    GPUExtent3D copySize);

undefined clearBuffer(
    GPUBuffer buffer,
    optional GPUSize64 offset = 0,
    optional GPUSize64 size);

undefined resolveQuerySet(
    GPUQuerySet querySet,
    GPUSize32 firstQuery,
    GPUSize32 queryCount,
    GPUBuffer destination,
    GPUSize64 destinationOffset);

GPUCommandBuffer finish(optional GPUCommandBufferDescriptor descriptor = {});
};

GPUCommandEncoder includes GPUObjectBase;
GPUCommandEncoder includes GPUCommandsMixin;
GPUCommandEncoder includes GPUDebugCommandsMixin;

```

13.2.1. Command Encoder Creation

```

dictionary
GPUCommandEncoderDescriptor

: GPUObjectDescriptorBase {
};

createCommandEncoder(descriptor)
    Creates a GPUCommandEncoder.

```

Called on: [GPUDevice](#) this.

Arguments:

Arguments for the [GPUDevice.createCommandEncoder\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUCommandEncoderDescriptor	✗	✓	Description of the GPUCommandEncoder to create.

Returns: [GPUCommandEncoder](#)

[Content timeline](#) steps:

1. Let *e* be [! create a new WebGPU object](#)(*this*, [GPUCommandEncoder](#), *descriptor*).
2. Issue the *initialization steps* on the [Device timeline](#) of *this*.
3. Return *e*.

[Device timeline](#) *initialization steps*:

1. If any of the following conditions are unsatisfied [generate a validation error](#), [invalidate](#) *e* and return.
 - *this* must not be [lost](#).

Creating a [GPUCommandEncoder](#), encoding a command to clear a buffer, finishing the encoder to get a [GPUCommandBuffer](#), then submitting it to the [GPUQueue](#).

```
const commandEncoder = gpuDevice.createCommandEncoder();
commandEncoder.clearBuffer(buffer);
const commandBuffer = commandEncoder.finish();
gpuDevice.queue.submit([commandBuffer]);
```

13.3. Pass Encoding

beginRenderPass(descriptor)

Begins encoding a render pass described by *descriptor*.

Called on: [GPUCommandEncoder](#) *this*.

Arguments:

Arguments for the [GPUCommandEncoder.beginRenderPass\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPURenderPassDescriptor	✗	✗	Description of the GPURenderPassEncoder to create.

Returns: [GPURenderPassEncoder](#)

[Content timeline](#) steps:

1. For each non-null *colorAttachment* in *descriptor.colorAttachments*:
 1. If *colorAttachment.clearValue* is [provided](#):
 1. [? validate GPUColor shape](#)(*colorAttachment.clearValue*).
 2. Let *pass* be a new [GPURenderPassEncoder](#) object.
 3. Issue the *initialization steps* on the [Device timeline](#) of *this*.
 4. Return *pass*.

[Device timeline](#) *initialization steps*:

1. [Validate the encoder state](#) of *this*. If it returns false, [invalidate](#) *pass* and return.
2. Set *this.[[state]]* to "locked".
3. Let *attachmentRegions* be a [list](#) of [[texture subresource](#), [depthSlice](#)?] pairs, initially empty. Each pair describes the region of the texture to be rendered to, which includes a single depth slice for "3d" textures only.
4. For each non-null *colorAttachment* in *descriptor.colorAttachments*:
 1. Add [*colorAttachment.view*, *colorAttachment.depthSlice* ?? null] to *attachmentRegions*.
 2. If *colorAttachment.resolveTarget* is not null:
 1. Add [*colorAttachment.resolveTarget*, undefined] to *attachmentRegions*.
5. If any of the following requirements are unmet, [invalidate](#) *pass* and return.
 - *descriptor* must meet the [Valid Usage](#) rules given device *this.[[device]]*.
 - The set of texture regions in *attachmentRegions* must be pairwise disjoint. That is, no two texture regions may overlap.

6. Add each [texture subresource](#) in `attachmentRegions` to `pass.[[usage_scope]]` with usage [attachment](#).
7. Let `depthStencilAttachment` be `descriptor.depthStencilAttachment`.
8. If `depthStencilAttachment` is not `null`:
 1. Let `depthStencilView` be `depthStencilAttachment.view`.
 2. Add the [depth subresource](#) of `depthStencilView`, if any, to `pass.[[usage_scope]]` with usage [attachment-read](#) if `depthStencilAttachment.depthReadOnly` is true, or [attachment](#) otherwise.
 3. Add the [stencil subresource](#) of `depthStencilView`, if any, to `pass.[[usage_scope]]` with usage [attachment-read](#) if `depthStencilAttachment.stencilReadOnly` is true, or [attachment](#) otherwise.
 4. Set `pass.[[depthReadOnly]]` to `depthStencilAttachment.depthReadOnly`.
 5. Set `pass.[[stencilReadOnly]]` to `depthStencilAttachment.stencilReadOnly`.
 9. Set `pass.[[layout]]` to [derive render targets layout from pass\(descriptor\)](#).
10. If `descriptor.timestampWrites` is [provided](#):
 1. Let `timestampWrites` be `descriptor.timestampWrites`.
 2. If `timestampWrites.beginningOfPassWriteIndex` is [provided](#), [append](#) a [GPU command](#) to this `[[commands]]` with the following steps:
 1. Before the pass commands begin executing, write the [current queue timestamp](#) into index `timestampWrites.beginningOfPassWriteIndex` of `timestampWrites.querySet`.
 3. If `timestampWrites.endOfPassWriteIndex` is [provided](#), set `pass.[[endTimestampWrite]]` to a [GPU command](#) with the following steps:
 1. After the pass commands finish executing, write the [current queue timestamp](#) into index `timestampWrites.endOfPassWriteIndex` of `timestampWrites.querySet`.
11. Set `pass.[[drawCount]]` to 0.
12. Set `pass.[[maxDrawCount]]` to `descriptor.maxDrawCount`.
13. Set `pass.[[maxDrawCount]]` to `descriptor.maxDrawCount`.
14. [Enqueue a command](#) on this which issues the subsequent steps on the [Queue timeline](#) when executed.

[Queue timeline](#) steps:

 1. Let the `[[renderState]]` of the currently executing [GPUCommandBuffer](#) be a new [RenderState](#).
 2. Set `[[renderState]].[[colorAttachments]]` to `descriptor.colorAttachments`.
 3. Set `[[renderState]].[[depthStencilAttachment]]` to `descriptor.depthStencilAttachment`.
 4. For each non-`null` `colorAttachment` in `descriptor.colorAttachments`:
 1. Let `colorView` be `colorAttachment.view`.
 2. If `colorView.[[descriptor]].dimension` is:

["3d"](#)

Let `colorSubregion` be `colorAttachment.depthSlice` of `colorView`.

Otherwise

Let `colorSubregion` be `colorView`.
 3. If `colorAttachment.loadOp` is:

["load"](#)

Ensure the contents of `colorSubregion` are loaded into the [framebuffer memory](#) associated with `colorSubregion`.

["clear"](#)

Set every [texel](#) of the [framebuffer memory](#) associated with `colorSubregion` to `colorAttachment.clearValue`.
5. If `depthStencilAttachment` is not `null`:
 1. If `depthStencilAttachment.depthLoadOp` is:

Not [provided](#)

[Assert](#) that `depthStencilAttachment.depthReadOnly` is `true` and ensure the contents of the [depth subresource](#) of `depthStencilView` are loaded into the [framebuffer memory](#) associated with `depthStencilView`.

["load"](#)

Ensure the contents of the [depth subresource](#) of `depthStencilView` are loaded into the [framebuffer memory](#) associated with `depthStencilView`.

["clear"](#)

Set every [texel](#) of the [framebuffer memory](#) associated with the [depth subresource](#) of [depthStencilView](#) to [depthStencilAttachment.depthClearValue](#).

2. If [depthStencilAttachment.stencilLoadOp](#) is:

Not [provided](#)

[Assert](#) that [depthStencilAttachment.stencilReadOnly](#) is `true` and ensure the contents of the [stencil subresource](#) of [depthStencilView](#) are loaded into the [framebuffer memory](#) associated with [depthStencilView](#).

["load"](#)

Ensure the contents of the [stencil subresource](#) of [depthStencilView](#) are loaded into the [framebuffer memory](#) associated with [depthStencilView](#).

["clear"](#)

Set every [texel](#) of the [framebuffer memory](#) associated with the [stencil subresource](#) of [depthStencilView](#) to [depthStencilAttachment.stencilClearValue](#).

Note: [Read-only depth-stencil](#) attachments are implicitly treated as though the ["load"](#) operation was used. Validation that requires the load op to not be provided for read-only attachments is done in [GPURenderPassDepthStencilAttachment Valid Usage](#).

beginComputePass(descriptor)

Begins encoding a compute pass described by *descriptor*.

Called on: [GPUCommandEncoder](#) *this*.

Arguments:

Arguments for the [GPUCommandEncoder.beginComputePass\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUComputePassDescriptor	✗	✓	

Returns: [GPUComputePassEncoder](#)

[Content timeline](#) steps:

1. Let *pass* be a new [GPUComputePassEncoder](#) object.
2. Issue the *initialization steps* on the [Device timeline](#) of *this*.
3. Return *pass*.

[Device timeline](#) *initialization steps*:

1. [Validate the encoder state](#) of *this*. If it returns false, [invalidate](#) *pass* and return.
2. Set *this*.[\[\[state\]\]](#) to "locked".
3. If any of the following requirements are unmet, [invalidate](#) *pass* and return.
 - If *descriptor.timestampWrites* is [provided](#):
 - [Validate timestampWrites\(this. \[\[device\]\], descriptor.timestampWrites\)](#) must return true.
4. If *descriptor.timestampWrites* is [provided](#):
 1. Let *timestampWrites* be *descriptor.timestampWrites*.
 2. If *timestampWrites.beginningOfPassWriteIndex* is [provided](#), [append](#) a [GPU command](#) to *this. [[commands]]* with the following steps:
 1. Before the pass commands begin executing, write the [current queue timestamp](#) into index *timestampWrites.beginningOfPassWriteIndex* of *timestampWrites.querySet*.
 3. If *timestampWrites.endOfPassWriteIndex* is [provided](#), set *pass. [[endTimestampWrite]]* to a [GPU command](#) with the following steps:
 1. After the pass commands finish executing, write the [current queue timestamp](#) into index *timestampWrites.endOfPassWriteIndex* of *timestampWrites.querySet*.

13.4. Buffer Copy Commands

[copyBufferToBuffer\(\)](#) has two overloads:

copyBufferToBuffer(source, destination, size)

Shorthand, equivalent to [copyBufferToBuffer\(source, 0, destination, 0, size\)](#).

copyBufferToBuffer(source, sourceOffset, destination, destinationOffset, size)

Encode a command into the [GPUCommandEncoder](#) that copies data from a sub-region of a [GPUBuffer](#) to a sub-region of another [GPUBuffer](#).

Called on: [GPUCommandEncoder](#) *this*.

Arguments:

Arguments for the [GPUCommandEncoder.copyBufferToBuffer\(source, sourceOffset, destination, destinationOffset, size\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>source</i>	GPUBuffer	✗	✗	The GPUBuffer to copy from.
<i>sourceOffset</i>	GPUSize64	✗	✗	Offset in bytes into <i>source</i> to begin copying from.
<i>destination</i>	GPUBuffer	✗	✗	The GPUBuffer to copy to.
<i>destinationOffset</i>	GPUSize64	✗	✗	Offset in bytes into <i>destination</i> to place the copied data.
<i>size</i>	GPUSize64	✗	✓	Bytes to copy.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of this. [\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of this. If it returns false, return.
2. If *size* is [undefined](#), set it to *source.size* – *sourceOffset*.
3. If any of the following conditions are unsatisfied, [invalidate](#) this and return.
 - *source* is [valid to use with](#) this.
 - *destination* is [valid to use with](#) this.
 - *source.usage* contains [COPY_SRC](#).
 - *destination.usage* contains [COPY_DST](#).
 - *size* is a multiple of 4.
 - *sourceOffset* is a multiple of 4.
 - *destinationOffset* is a multiple of 4.
 - *source.size* ≥ (*sourceOffset* + *size*).
 - *destination.size* ≥ (*destinationOffset* + *size*).
 - *source* and *destination* are not the same [GPUBuffer](#).
4. [Enqueue a command](#) on this which issues the subsequent steps on the [Queue timeline](#) when executed.

[Queue timeline](#) steps:

1. Copy *size* bytes of *source*, beginning at *sourceOffset*, into *destination*, beginning at *destinationOffset*.

clearBuffer(buffer, offset, size)

Encode a command into the [GPUCommandEncoder](#) that fills a sub-region of a [GPUBuffer](#) with zeros.

Called on: [GPUCommandEncoder](#) this.

Arguments:

Arguments for the [GPUCommandEncoder.clearBuffer\(buffer, offset, size\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>buffer</i>	GPUBuffer	✗	✗	The GPUBuffer to clear.
<i>offset</i>	GPUSize64	✗	✓	Offset in bytes into <i>buffer</i> where the sub-region to clear begins.
<i>size</i>	GPUSize64	✗	✓	Size in bytes of the sub-region to clear. Defaults to the size of the buffer minus <i>offset</i> .

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of this. [\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of this. If it returns false, return.
2. If *size* is missing, set *size* to $\max(0, \text{buffer.size} - \text{offset})$.
3. If any of the following conditions are unsatisfied, [invalidate](#) this and return.
 - *buffer* is [valid to use with](#) this.
 - *buffer.usage* contains [COPY_DST](#).

- *size* is a multiple of 4.
 - *offset* is a multiple of 4.
 - *buffer.size* \geq (*offset* + *size*).
4. [Enqueue a command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) when executed.

[Queue timeline](#) steps:

1. Set *size* bytes of *buffer* to 0 starting at *offset*.

13.5. Texel Copy Commands

copyBufferToTexture(source, destination, copySize)

Encode a command into the [GPUCommandEncoder](#) that copies data from a sub-region of a [GPUBuffer](#) to a sub-region of one or multiple continuous [texture subresources](#).

Called on: [GPUCommandEncoder](#) *this*.

Arguments:

Arguments for the [GPUCommandEncoder.copyBufferToTexture\(source, destination, copySize\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>source</i>	GPUTexelCopyBufferInfo	✗	✗	Combined with <i>copySize</i> , defines the region of the source buffer.
<i>destination</i>	GPUTexelCopyTextureInfo	✗	✗	Combined with <i>copySize</i> , defines the region of the destination texture subresource .
<i>copySize</i>	GPUExtent3D	✗	✗	

Returns: [undefined](#)

[Content timeline](#) steps:

1. ? [validate GPUOrigin3D shape\(destination.origin\)](#).
2. ? [validate GPUExtent3D shape\(copySize\)](#).
3. Issue the subsequent steps on the [Device timeline](#) of *this.device*:

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. Let *aligned* be `true`.
3. Let *dataLength* be *source.buffer.size*.
4. If any of the following conditions are unsatisfied, [invalidate this](#) and return.
 - [validating GPUTexelCopyBufferInfo\(source\)](#) returns `true`.
 - *source.buffer.usage* contains `COPY_SRC`.
 - [validating texture buffer copy\(destination, source, dataLength, copySize, COPY_DST, aligned\)](#) returns `true`.
5. [Enqueue a command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) when executed.

[Queue timeline](#) steps:

1. Let *blockWidth* be the [texel block width](#) of *destination.texture*.
2. Let *blockHeight* be the [texel block height](#) of *destination.texture*.
3. Let *dstOrigin* be *destination.origin*.
4. Let *dstBlockOriginX* be (*dstOrigin.x* \div *blockWidth*).
5. Let *dstBlockOriginY* be (*dstOrigin.y* \div *blockHeight*).
6. Let *blockColumns* be (*copySize.width* \div *blockWidth*).
7. Let *blockRows* be (*copySize.height* \div *blockHeight*).
8. [Assert](#) that *dstBlockOriginX*, *dstBlockOriginY*, *blockColumns*, and *blockRows* are integers.
9. For each *z* in the range [0, *copySize.depthOrArrayLayers* - 1]:
 1. Let *dstSubregion* be [texture copy sub-region](#) (*z* + *dstOrigin.z*) of *destination*.
 2. For each *y* in the range [0, *blockRows* - 1]:
 1. For each *x* in the range [0, *blockColumns* - 1]:
 1. Let *blockOffset* be the [texel block byte offset](#) of *source* for (*x*, *y*, *z*) of *destination.texture*.

- Set [texel block](#) ($dstBlockOriginX + x$, $dstBlockOriginY + y$) of $dstSubregion$ to be an [equivalent texel representation](#) to the [texel block](#) described by $source.buffer$ at offset $blockOffset$.

copyTextureToBuffer(*source*, *destination*, *copySize*)

Encode a command into the [GPUCommandEncoder](#) that copies data from a sub-region of one or multiple continuous [texture subresources](#) to a sub-region of a [GPUBuffer](#).

Called on: [GPUCommandEncoder](#) *this*.

Arguments:

Arguments for the [GPUCommandEncoder.copyTextureToBuffer\(source, destination, copySize\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>source</i>	GPUTextureCopyTextureInfo	✗	✗	Combined with <i>copySize</i> , defines the region of the source texture subresources .
<i>destination</i>	GPUTextureCopyBufferInfo	✗	✗	Combined with <i>copySize</i> , defines the region of the destination buffer.
<i>copySize</i>	GPUExtent3D	✗	✗	

Returns: [undefined](#)

[Content timeline](#) steps:

- [? validate GPUOrigin3D shape](#)($source.origin$).
- [? validate GPUExtent3D shape](#)(*copySize*).
- Issue the subsequent steps on the [Device timeline](#) of $this.[[device]]$:

[Device timeline](#) steps:

- [Validate the encoder state](#) of *this*. If it returns false, return.
- Let *aligned* be `true`.
- Let *dataLength* be $destination.buffer.size$.
- If any of the following conditions are unsatisfied, [invalidate](#) *this* and return.
 - [validating GPUTextureCopyBufferInfo](#)(*destination*) returns `true`.
 - $destination.buffer.usage$ contains `COPY_DST`.
 - [validating texture buffer copy](#)(*source*, *destination*, *dataLength*, *copySize*, `COPY_SRC`, *aligned*) returns `true`.
- [Enqueue a command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) when executed.

[Queue timeline](#) steps:

- Let *blockWidth* be the [texel block width](#) of $source.texture$.
- Let *blockHeight* be the [texel block height](#) of $source.texture$.
- Let *srcOrigin* be $source.origin$.
- Let *srcBlockOriginX* be $(srcOrigin.x \div blockWidth)$.
- Let *srcBlockOriginY* be $(srcOrigin.y \div blockHeight)$.
- Let *blockColumns* be $(copySize.width \div blockWidth)$.
- Let *blockRows* be $(copySize.height \div blockHeight)$.
- [Assert](#) that *srcBlockOriginX*, *srcBlockOriginY*, *blockColumns*, and *blockRows* are integers.
- For each *z* in the range $[0, copySize.depthOrArrayLayers - 1]$:
 - Let *srcSubregion* be [texture copy sub-region](#) ($z + srcOrigin.z$) of *source*.
 - For each *y* in the range $[0, blockRows - 1]$:
 - For each *x* in the range $[0, blockColumns - 1]$:
 - Let *blockOffset* be the [texel block byte offset](#) of *destination* for (*x*, *y*, *z*) of $source.texture$.
 - Set $destination.buffer$ at offset *blockOffset* to be an [equivalent texel representation](#) to [texel block](#) ($srcBlockOriginX + x$, $srcBlockOriginY + y$) of *srcSubregion*.

copyTextureToTexture(*source*, *destination*, *copySize*)

Encode a command into the [GPUCommandEncoder](#) that copies data from a sub-region of one or multiple contiguous [texture subresources](#) to another sub-region of one or multiple continuous [texture subresources](#).

Called on: [GPUCommandEncoder](#) *this*.

Arguments:

Arguments for the [GPUCommandEncoder.copyTextureToTexture\(source, destination, copySize\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>source</i>	GPUTextureCopyTextureInfo	✗	✗	Combined with <i>copySize</i> , defines the region of the source texture subresources .
<i>destination</i>	GPUTextureCopyTextureInfo	✗	✗	Combined with <i>copySize</i> , defines the region of the destination texture subresources .
<i>copySize</i>	GPUExtent3D	✗	✗	

Returns: [undefined](#)

[Content timeline](#) steps:

1. [? validate GPUOrigin3D shape](#)(*source*.[origin](#)).
2. [? validate GPUOrigin3D shape](#)(*destination*.[origin](#)).
3. [? validate GPUExtent3D shape](#)(*copySize*).
4. Issue the subsequent steps on the [Device timeline](#) of this.[\[\[device\]\]](#):

[Device timeline](#) steps:

1. [Validate the encoder state](#) of this. If it returns false, return.
2. If any of the following conditions are unsatisfied, [invalidate](#) this and return.
 - Let *srcTexture* be *source*.[texture](#).
 - Let *dstTexture* be *destination*.[texture](#).
 - [validating GPUTextureCopyTextureInfo](#)(*source*, *copySize*) returns **true**.
 - *srcTexture*.[usage](#) contains [COPY_SRC](#).
 - [validating GPUTextureCopyTextureInfo](#)(*destination*, *copySize*) returns **true**.
 - *dstTexture*.[usage](#) contains [COPY_DST](#).
 - *srcTexture*.[sampleCount](#) is equal to *dstTexture*.[sampleCount](#).
 - *srcTexture*.[format](#) and *dstTexture*.[format](#) must be [copy-compatible](#).
 - If *srcTexture*.[format](#) is a depth-stencil format:
 - *source*.[aspect](#) and *destination*.[aspect](#) must both refer to all aspects of *srcTexture*.[format](#) and *dstTexture*.[format](#), respectively.
 - The [set of subresources for texture copy](#)(*source*, *copySize*) and the [set of subresources for texture copy](#)(*destination*, *copySize*) are disjoint.
3. [Enqueue a command](#) on this which issues the subsequent steps on the [Queue timeline](#) when executed.

[Queue timeline](#) steps:

1. Let *blockWidth* be the [texel block width](#) of *source*.[texture](#).
2. Let *blockHeight* be the [texel block height](#) of *source*.[texture](#).
3. Let *srcOrigin* be *source*.[origin](#).
4. Let *srcBlockOriginX* be (*srcOrigin*.[x](#) ÷ *blockWidth*).
5. Let *srcBlockOriginY* be (*srcOrigin*.[y](#) ÷ *blockHeight*).
6. Let *dstOrigin* be *destination*.[origin](#).
7. Let *dstBlockOriginX* be (*dstOrigin*.[x](#) ÷ *blockWidth*).
8. Let *dstBlockOriginY* be (*dstOrigin*.[y](#) ÷ *blockHeight*).
9. Let *blockColumns* be (*copySize*.[width](#) ÷ *blockWidth*).
10. Let *blockRows* be (*copySize*.[height](#) ÷ *blockHeight*).
11. [Assert](#) that *srcBlockOriginX*, *srcBlockOriginY*, *dstBlockOriginX*, *dstBlockOriginY*, *blockColumns*, and *blockRows* are integers.
12. For each *z* in the range [0, *copySize*.[depthOrArrayLayers](#) − 1]:
 1. Let *srcSubregion* be [texture copy sub-region](#) (*z* + *srcOrigin*.[z](#)) of *source*.
 2. Let *dstSubregion* be [texture copy sub-region](#) (*z* + *dstOrigin*.[z](#)) of *destination*.
 3. For each *y* in the range [0, *blockRows* − 1]:
 1. For each *x* in the range [0, *blockColumns* − 1]:
 1. Set [texel block](#) (*dstBlockOriginX* + *x*, *dstBlockOriginY* + *y*) of *dstSubregion* to be an [equivalent texel representation](#) to [texel block](#) (*srcBlockOriginX* + *x*, *srcBlockOriginY* + *y*) of *srcSubregion*.

13.6. Queries

resolveQuerySet(querySet, firstQuery, queryCount, destination, destinationOffset)

Resolves query results from a [GPUQuerySet](#) out into a range of a [GPUBuffer](#).

Called on: [GPUCommandEncoder](#) this.

Arguments:

Arguments for the [GPUCommandEncoder.resolveQuerySet\(querySet, firstQuery, queryCount, destination, destinationOffset\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>querySet</i>	GPUQuerySet	✗	✗	
<i>firstQuery</i>	GPUSize32	✗	✗	
<i>queryCount</i>	GPUSize32	✗	✗	
<i>destination</i>	GPUBuffer	✗	✗	
<i>destinationOffset</i>	GPUSize64	✗	✗	

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. If any of the following conditions are unsatisfied, [invalidate](#) *this* and return.
 - *querySet* is [valid to use with](#) *this*.
 - *destination* is [valid to use with](#) *this*.
 - *destination.usage* contains [QUERY_RESOLVE](#).
 - *firstQuery* < the number of queries in *querySet*.
 - $(firstQuery + queryCount) \leq$ the number of queries in *querySet*.
 - *destinationOffset* is a multiple of 256.
 - $destinationOffset + 8 \times queryCount \leq destination.size$.
3. [Enqueue a command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) when executed.

[Queue timeline](#) steps:

1. Let *queryIndex* be *firstQuery*.
2. Let *offset* be *destinationOffset*.
3. While *queryIndex* < *firstQuery* + *queryCount*:
 1. Set 8 bytes of *destination*, beginning at *offset*, to be the value of *querySet* at *queryIndex*.
 2. Set *queryIndex* to be *queryIndex* + 1.
 3. Set *offset* to be *offset* + 8.

13.7. Finalization

A [GPUCommandBuffer](#) containing the commands recorded by the [GPUCommandEncoder](#) can be created by calling [finish\(\)](#). Once [finish\(\)](#) has been called the command encoder can no longer be used.

finish(descriptor)

Completes recording of the commands sequence and returns a corresponding [GPUCommandBuffer](#).

Called on: [GPUCommandEncoder](#) *this*.

Arguments:

Arguments for the [GPUCommandEncoder.finish\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUCommandBufferDescriptor	✗	✓	

Returns: [GPUCommandBuffer](#)

[Content timeline](#) steps:

1. Let *commandBuffer* be a new [GPUCommandBuffer](#).
2. Issue the *finish steps* on the [Device timeline](#) of this.[\[\[device\]\]](#).
3. Return *commandBuffer*.

[Device timeline](#) *finish steps*:

1. Let *validationSucceeded* be `true` if all of the following requirements are met, and `false` otherwise.
 - *this* must be [valid](#).
 - *this*.[\[\[state\]\]](#) must be "[open](#)".
 - *this*.[\[\[debug_group_stack\]\]](#) must be [empty](#).
2. Set *this*.[\[\[state\]\]](#) to "[ended](#)".
3. If *validationSucceeded* is `false`, then:
 1. [Generate a validation error](#).
2. Return an [invalidated GPUCommandBuffer](#).
4. Set *commandBuffer*.[\[\[command_list\]\]](#) to *this*.[\[\[commands\]\]](#).

14. Programmable Passes

interface mixin

[GPUBindingCommandsMixin](#)

```
{
  undefined setBindGroup(GPUIndex32 index, GPUBindGroup? bindGroup,
    optional sequence<GPUBufferDynamicOffset> dynamicOffsets = []);

  undefined setBindGroup(GPUIndex32 index, GPUBindGroup? bindGroup,
    [AllowShared] Uint32Array dynamicOffsetsData,
    GPUSize64 dynamicOffsetsDataStart,
    GPUSize32 dynamicOffsetsDataLength);
};
```

[GPUBindingCommandsMixin](#) assumes the presence of [GPUObjectBase](#) and [GPUCommandsMixin](#) members on the same object. It must only be included by interfaces which also include those mixins.

[GPUBindingCommandsMixin](#) has the following [device timeline properties](#):

[\[\[bind_groups\]\]](#), of type [ordered map](#)<GPUIndex32, GPUBindGroup>, initially empty
The current [GPUBindGroup](#) for each index.

[\[\[dynamic_offsets\]\]](#), of type [ordered map](#)<GPUIndex32, list<GPUBufferDynamicOffset>>, initially empty
The current dynamic offsets for each [\[\[bind_groups\]\]](#) entry.

14.1. Bind Groups

setBindGroup() has two overloads:

setBindGroup(index, bindGroup, dynamicOffsets)

Sets the current [GPUBindGroup](#) for the given index.

Called on: [GPUBindingCommandsMixin](#) this.

Arguments:

index, of type [GPUIndex32](#), non-nullable, required
The index to set the bind group at.

bindGroup, of type [GPUBindGroup](#), nullable, required
Bind group to use for subsequent render or compute commands.

dynamicOffsets, of type [sequence](#)<GPUBufferDynamicOffset>, non-nullable, defaulting to []
Array containing buffer offsets in bytes for each entry in *bindGroup* marked as [buffer.hasDynamicOffset](#), ordered by [GPUBindGroupLayoutEntry.binding](#). See [note](#) for additional details.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of this.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of this. If it returns false, return.
2. Let *dynamicOffsetCount* be 0 if *bindGroup* is null, or *bindGroup*.[\[\[layout\]\]](#).[\[\[dynamicOffsetCount\]\]](#) if not.
3. If any of the following requirements are unmet, [invalidate](#) this and return.

- *index* must be < this.[\[\[device\]\]](#).[\[\[limits\]\]](#).[maxBindGroups](#).
- *dynamicOffsets*.[size](#) must equal *dynamicOffsetCount*.

4. If *bindGroup* is null:

1. [Remove](#) this.[\[\[bind_groups\]\]](#)[*index*].
2. [Remove](#) this.[\[\[dynamic_offsets\]\]](#)[*index*].

Otherwise:

1. If any of the following requirements are unmet, [invalidate](#) this and return.

- *bindGroup* must be [valid to use with](#) this.
- [For each dynamic binding](#) (*bufferBinding*, *bufferLayout*, *dynamicOffsetIndex*) in *bindGroup*:
- *bufferBinding*.[offset](#) + *dynamicOffsets*[*dynamicOffsetIndex*] + *bufferLayout*.[minBindingSize](#) must be ≤ *bufferBinding*.[buffer.size](#).
- If *bufferLayout*.[type](#) is ["uniform"](#):
- *dynamicOffset* must be a multiple of [minUniformBufferOffsetAlignment](#).
- If *bufferLayout*.[type](#) is ["storage"](#) or ["read-only-storage"](#):
- *dynamicOffset* must be a multiple of [minStorageBufferOffsetAlignment](#).

2. Set this.[\[\[bind_groups\]\]](#)[*index*] to be *bindGroup*.

3. Set this.[\[\[dynamic_offsets\]\]](#)[*index*] to be a copy of *dynamicOffsets*.

4. If this is a [GPURenderCommandsMixin](#):

1. For each *bindGroup* in this.[\[\[bind_groups\]\]](#), [merge](#) *bindGroup*.[\[\[usedResources\]\]](#) into this.[\[\[usage_scope\]\]](#).

setBindGroup(*index*, *bindGroup*, *dynamicOffsetsData*, *dynamicOffsetsDataStart*, *dynamicOffsetsDataLength*)

Sets the current [GPUBindGroup](#) for the given index, specifying dynamic offsets as a subset of a [Uint32Array](#).

Called on: [GPUBindingCommandsMixin](#) this.

Arguments:

Arguments for the [GPUBindingCommandsMixin.setBindGroup](#)(*index*, *bindGroup*, *dynamicOffsetsData*, *dynamicOffsetsDataStart*, *dynamicOffsetsDataLength*) method.

Parameter	Type	Nullable	Optional	Description
<i>index</i>	GPUIndex32	✗	✗	The index to set the bind group at.
<i>bindGroup</i>	GPUBindGroup?	✓	✗	Bind group to use for subsequent render or compute commands.
<i>dynamicOffsetsData</i>	Uint32Array	✗	✗	Array containing buffer offsets in bytes for each entry in <i>bindGroup</i> marked as buffer.hasDynamicOffset , ordered by GPUBindGroupLayoutEntry.binding . See note for additional details.
<i>dynamicOffsetsDataStart</i>	GPUSize64	✗	✗	Offset in elements into <i>dynamicOffsetsData</i> where the buffer offset data begins.
<i>dynamicOffsetsDataLength</i>	GPUSize32	✗	✗	Number of buffer offsets to read from <i>dynamicOffsetsData</i> .

Returns: [undefined](#)

[Content timeline](#) steps:

1. If any of the following requirements are unmet, throw a [RangeError](#) and return.

- *dynamicOffsetsDataStart* must be ≥ 0.
- *dynamicOffsetsDataStart* + *dynamicOffsetsDataLength* must be ≤ *dynamicOffsetsData*.[length](#).

2. Let *dynamicOffsets* be a [list](#) containing the range, starting at index *dynamicOffsetsDataStart*, of *dynamicOffsetsDataLength* elements of [a copy of](#) *dynamicOffsetsData*.

3. Call this.[setBindGroup](#)(*index*, *bindGroup*, *dynamicOffsets*).

NOTE:

Dynamic offset are applied in [GPUBindGroupLayoutEntry.binding](#) order.

This means that if `dynamic bindings` is the list of each [GPUBindGroupLayoutEntry](#) in the [GPUBindGroupLayout](#) with `buffer?.hasDynamicOffset` set to `true`, sorted by [GPUBindGroupLayoutEntry.binding](#), then `dynamic offset[i]`, as supplied to [setBindGroup\(\)](#), will correspond to `dynamic bindings[i]`.

For a [GPUBindGroupLayout](#) created with the following call:

```
// Note the bindings are listed out-of-order in this array, but it
// doesn't matter because they will be sorted by binding index.
let layout = gpuDevice.createBindGroupLayout({
  entries: [{
    binding: 1,
    buffer: {}},
  ], {
    binding: 2,
    buffer: { dynamicOffset: true }},
  ], {
    binding: 0,
    buffer: { dynamicOffset: true }},
  ]});
```

Used by a [GPUBindGroup](#) created with the following call:

```
// Like above, the array order doesn't matter here.
// It doesn't even need to match the order used in the layout.
let bindGroup = gpuDevice.createBindGroup({
  layout: layout,
  entries: [{
    binding: 1,
    resource: { buffer: bufferA, offset: 256 }},
  ], {
    binding: 2,
    resource: { buffer: bufferB, offset: 512 }},
  ], {
    binding: 0,
    resource: { buffer: bufferC }},
  ]});
```

And bound with the following call:

```
pass.setBindGroup(0, bindGroup, [1024, 2048]);
```

The following buffer offsets will be applied:

Binding	Buffer	Offset
0	bufferC	1024 (Dynamic)
1	bufferA	256 (Static)
2	bufferB	2560 (Static + Dynamic)

To Iterate over each dynamic binding offset in a given [GPUBindGroup](#) *bindGroup* with a given list of *steps* to be executed for each dynamic offset, run the following [device timeline](#) steps:

Let *dynamicOffsetIndex* be 0.

Let *layout* be *bindGroup*.[\[\[layout\]\]](#).

For each [GPUBindGroupEntry](#) *entry* in *bindGroup*.[\[\[entries\]\]](#) ordered in increasing values of *entry*.[binding](#):

Let *bindingDescriptor* be the [GPUBindGroupLayoutEntry](#) at *layout*.[\[\[entryMap\]\]](#)[*entry*.[binding](#)]:

If *bindingDescriptor*.[buffer?.hasDynamicOffset](#) is `true`:

Let *bufferBinding* be [get as buffer binding](#)(*entry*.[resource](#)).

Let *bufferLayout* be *bindingDescriptor*.[buffer](#).

Call *steps* with *bufferBinding*, *bufferLayout*, and *dynamicOffsetIndex*.

Let *dynamicOffsetIndex* be *dynamicOffsetIndex* + 1

Validate encoder bind groups(encoder, pipeline)

Arguments:

[GPUBindingCommandsMixin](#) *encoder*

Encoder whose bind groups are being validated.

[GPUPipelineBase](#) *pipeline*

Pipeline to validate *encoders* bind groups are compatible with.

[Device timeline](#) steps:

If any of the following conditions are unsatisfied, return **false**:

pipeline must not be **null**.

All bind groups used by the pipeline must be set and compatible with the pipeline layout: For each pair of ([GPUIndex32](#) *index*, [GPUBindGroupLayout](#) *bindGroupLayout*) in *pipeline*.[\[\[layout\]\]](#).[\[\[bindGroupLayouts\]\]](#):

If *bindGroupLayout* is **null**, [continue](#).

Let *bindGroup* be *encoder*.[\[\[bind_groups\]\]](#)[*index*].

Let *dynamicOffsets* be *encoder*.[\[\[dynamic_offsets\]\]](#)[*index*].

bindGroup must not be **null**.

bindGroup.[\[\[layout\]\]](#) must be [group-equivalent](#) with *bindGroupLayout*.

Let *dynamicOffsetIndex* be 0.

For each [GPUBindGroupEntry](#) *bindGroupEntry* in *bindGroup*.[\[\[entries\]\]](#), sorted by *bindGroupEntry*.[binding](#):

Let *bindGroupLayoutEntry* be *bindGroup*.[\[\[layout\]\]](#).[\[\[entryMap\]\]](#)[*bindGroupEntry*.[binding](#)].

If *bindGroupLayoutEntry*.[buffer](#) is not [provided](#), [continue](#).

Let *bound* be [get as buffer binding](#)(*bindGroupEntry*.[resource](#)).

If *bindGroupLayoutEntry*.[buffer.hasDynamicOffset](#):

Increment *bound*.[offset](#) by *dynamicOffsets*[*dynamicOffsetIndex*].

Increment *dynamicOffsetIndex* by 1.

If *bindGroupEntry*.[\[\[prevalidatedSize\]\]](#) is **false**:

[effective buffer binding size](#)(*bound*) must be ≥ [minimum buffer binding size](#) of the binding variable in *pipeline*'s shader that corresponds to *bindGroupEntry*.

[Encoder bind groups alias a writable resource](#)(*encoder*, *pipeline*) must be **false**.

Otherwise return **true**.

[Encoder bind groups alias a writable resource](#)(*encoder*, *pipeline*) if any writable buffer binding range overlaps with any other binding range of the same buffer, or any writable texture binding overlaps in [texture subresources](#) with any other texture binding (which may use the same or a different [GPUTextureView](#) object).

Note: This algorithm limits the use of the [usage scope storage exception](#).

Arguments:

[GPUBindingCommandsMixin](#) *encoder*

Encoder whose bind groups are being validated.

[GPUPipelineBase](#) *pipeline*

Pipeline to validate *encoders* bind groups are compatible with.

[Device timeline](#) steps:

For each *stage* in [[VERTEX](#), [FRAGMENT](#), [COMPUTE](#)]:

Let *bufferBindings* be a [list](#) of ([GPUBufferBinding](#), [boolean](#)) pairs, where the latter indicates whether the resource was used as writable.

Let *textureViews* be a [list](#) of ([GPUTextureView](#), [boolean](#)) pairs, where the latter indicates whether the resource was used as writable.

For each pair of ([GPUIndex32](#) *bindGroupIndex*, [GPUBindGroupLayout](#) *bindGroupLayout*) in *pipeline*.[\[\[layout\]\]](#).[\[\[bindGroupLayouts\]\]](#):

Let *bindGroup* be *encoder*.[\[\[bind_groups\]\]](#)[*bindGroupIndex*].

Let *bindGroupLayoutEntries* be *bindGroupLayout*.[\[\[descriptor\]\]](#).[entries](#).

Let *bufferRanges* be the [bound buffer ranges](#) of *bindGroup*, given dynamic offsets *encoder*.[\[\[dynamic_offsets\]\]](#)*[bindGroupIndex]*

For each ([GPUBindGroupLayoutEntry](#) *bindGroupLayoutEntry*, [GPUBufferBinding](#) *resource*) in *bufferRanges*, in which *bindGroupLayoutEntry.visibility* contains *stage*:

Let *resourceWritable* be (*bindGroupLayoutEntry.buffer.type* == "[storage](#)").

For each pair ([GPUBufferBinding](#) *pastResource*, [boolean](#) *pastResourceWritable*) in *bufferBindings*:

If (*resourceWritable* or *pastResourceWritable*) is true, and *pastResource* and *resource* are [buffer-binding-aliasing](#), return **true**.

[Append](#) (*resource*, *resourceWritable*) to *bufferBindings*.

For each [GPUBindGroupLayoutEntry](#) *bindGroupLayoutEntry* in *bindGroupLayoutEntries*, and corresponding [GPUTextureView](#) *resource* in *bindGroup*, in which *bindGroupLayoutEntry.visibility* contains *stage*:

If *bindGroupLayoutEntry.storageTexture* is not [provided](#), **continue**.

Let *resourceWritable* be whether *bindGroupLayoutEntry.storageTexture.access* is a writable access mode.

For each pair ([GPUTextureView](#) *pastResource*, [boolean](#) *pastResourceWritable*) in *textureViews*,

If (*resourceWritable* or *pastResourceWritable*) is true, and *pastResource* and *resource* is [texture-view-aliasing](#), return **true**.

[Append](#) (*resource*, *resourceWritable*) to *textureViews*.

Return **false**.

Note: Implementations are strongly encouraged to optimize this algorithm.

15. Debug Markers

GPUDebugCommandsMixin provides methods to apply debug labels to groups of commands or insert a single label into the command sequence.

Debug groups can be nested to create a hierarchy of labeled commands, and must be well-balanced.

Like [object labels](#), these labels have no required behavior, but may be shown in error messages and browser developer tools, and may be passed to native API backends.

```
interface mixin GPUDebugCommandsMixin {  
  undefined pushDebugGroup(USVString groupLabel);  
  undefined popDebugGroup();  
  undefined insertDebugMarker(USVString markerLabel);  
};
```

[GPUDebugCommandsMixin](#) assumes the presence of [GPUObjectBase](#) and [GPUCommandsMixin](#) members on the same object. It must only be included by interfaces which also include those mixins.

[GPUDebugCommandsMixin](#) has the following [device timeline properties](#):

[\[\[debug_group_stack\]\]](#), of type [stack<USVString>](#)
A stack of active debug group labels.

[GPUDebugCommandsMixin](#) has the following methods:

pushDebugGroup(groupLabel)
Begins a labeled debug group containing subsequent commands.

Called on: [GPUDebugCommandsMixin](#) *this*.

Arguments:

Arguments for the [GPUDebugCommandsMixin.pushDebugGroup\(groupLabel\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>groupLabel</i>	USVString	✗	✗	The label for the command group.

Returns: [undefined](#)

[Content timeline](#) steps:

- Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).
- [Device timeline](#) steps:
- [Validate the encoder state](#) of *this*. If it returns false, return.
 - [Push](#) *groupLabel* onto *this*.[\[\[debug_group_stack\]\]](#).

popDebugGroup()

Ends the labeled debug group most recently started by [pushDebugGroup\(\)](#).

Called on: [GPUDebugCommandsMixin](#) *this*.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).
- [Device timeline](#) steps:
1. [Validate the encoder state](#) of *this*. If it returns false, return.
 2. If any of the following requirements are unmet, [invalidate](#) *this* and return.
 - *this*.[\[\[debug_group_stack\]\]](#) must not [be empty](#).
 3. [Pop](#) an entry off of *this*.[\[\[debug_group_stack\]\]](#).

insertDebugMarker(markerLabel)

Marks a point in a stream of commands with a label.

Called on: [GPUDebugCommandsMixin](#) *this*.

Arguments:

Arguments for the [GPUDebugCommandsMixin.insertDebugMarker\(markerLabel\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>markerLabel</i>	USVString	✗	✗	The label to insert.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.

16. Compute Passes

16.1. GPUComputePassEncoder

[[Exposed](#)=(Window, Worker), [SecureContext](#)]

interface [GPUComputePassEncoder](#) {

[undefined](#) [setPipeline](#)(GPUComputePipeline pipeline);

[undefined](#) [dispatchWorkgroups](#)(GPUSize32 workgroupCountX, optional GPUSize32 workgroupCountY = 1, optional GPUSize32 workgroupCountZ = 1);

[undefined](#) [dispatchWorkgroupsIndirect](#)(GPUBuffer indirectBuffer, GPUSize64 indirectOffset);

[undefined](#) [end](#)();

};

[GPUComputePassEncoder](#) includes [GPUObjectBase](#);

[GPUComputePassEncoder](#) includes [GPUCommandsMixin](#);

[GPUComputePassEncoder](#) includes [GPUDebugCommandsMixin](#);

[GPUComputePassEncoder](#) includes [GPUBindingCommandsMixin](#);

[GPUComputePassEncoder](#) has the following [device timeline properties](#):

[\[\[command_encoder\]\]](#), of type [GPUCommandEncoder](#), readonly

 The [GPUCommandEncoder](#) that created this compute pass encoder.

[\[\[endTimeStampWrite\]\]](#), of type [GPU command?](#), readonly, defaulting to null

[GPU command](#), if any, writing a timestamp when the pass ends.

[\[\[pipeline\]\]](#), of type [GPUComputePipeline](#), initially null

 The current [GPUComputePipeline](#).

16.1.1. Compute Pass Encoder Creation

dictionary

GPUComputePassTimestampWrites

```
{
  required GPUQuerySet querySet;
  GPUSize32 beginningOfPassWriteIndex;
  GPUSize32 endOfPassWriteIndex;
};
```

querySet, of type [GPUQuerySet](#)

The [GPUQuerySet](#), of type ["timestamp"](#), that the query results will be written to.

beginningOfPassWriteIndex, of type [GPUSize32](#)

If defined, indicates the query index in [querySet](#) into which the timestamp at the beginning of the compute pass will be written.

endOfPassWriteIndex, of type [GPUSize32](#)

If defined, indicates the query index in [querySet](#) into which the timestamp at the end of the compute pass will be written.

Note: Timestamp query values are written in nanoseconds, but how the value is determined is [implementation-defined](#). See [§ 20.4 Timestamp Query](#) for details.

dictionary

GPUComputePassDescriptor

```
: GPUObjectDescriptorBase {
  GPUComputePassTimestampWrites timestampWrites;
};
```

timestampWrites, of type [GPUComputePassTimestampWrites](#)

Defines which timestamp values will be written for this pass, and where to write them to.

16.1.2. Dispatch

setPipeline(pipeline)

Sets the current [GPUComputePipeline](#).

Called on: [GPUComputePassEncoder](#) this.

Arguments:

Arguments for the [GPUComputePassEncoder.setPipeline\(pipeline\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>pipeline</i>	GPUComputePipeline	✗	✗	The compute pipeline to use for subsequent dispatch commands.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. If any of the following conditions are unsatisfied, [invalidate](#) *this* and return.
 - *pipeline* is [valid to use with](#) *this*.
3. Set *this*.[\[\[pipeline\]\]](#) to be *pipeline*.

dispatchWorkgroups(workgroupCountX, workgroupCountY, workgroupCountZ)

Dispatch work to be performed with the current [GPUComputePipeline](#). See [§ 23.1 Computing](#) for the detailed specification.

Called on: [GPUComputePassEncoder](#) this.

Arguments:

Arguments for the [GPUComputePassEncoder.dispatchWorkgroups\(workgroupCountX, workgroupCountY, workgroupCountZ\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>workgroupCountX</i>	GPUSize32	✗	✗	X dimension of the grid of workgroups to dispatch.
<i>workgroupCountY</i>	GPUSize32	✗	✓	Y dimension of the grid of workgroups to dispatch.
<i>workgroupCountZ</i>	GPUSize32	✗	✓	Z dimension of the grid of workgroups to dispatch.

NOTE:

The *x*, *y*, and *z* values passed to [dispatchWorkgroups\(\)](#) and [dispatchWorkgroupsIndirect\(\)](#) are the number of *workgroups* to dispatch for each dimension, *not* the number of shader invocations to perform across each dimension. This matches the behavior of modern native GPU APIs, but differs from the behavior of OpenCL.

This means that if a [GPUShaderModule](#) defines an entry point with `@workgroup_size(4, 4)`, and work is dispatched to it with the call `computePass.dispatchWorkgroups(8, 8)`; the entry point will be invoked 1024 times total: Dispatching a 4x4 workgroup 8 times along both the X and Y axes. (4*4*8*8=1024)

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. Let *usageScope* be an empty [usage scope](#).
3. For each *bindGroup* in *this*.[\[\[bind_groups\]\]](#), merge *bindGroup*.[\[\[usedResources\]\]](#) into *this*.[\[\[usage_scope\]\]](#).
4. If any of the following conditions are unsatisfied, [invalidate this](#) and return.
 - *usageScope* must satisfy [usage scope validation](#).
 - [Validate encoder bind groups](#)(*this*, *this*.[\[\[pipeline\]\]](#)) is `true`.
 - all of *workgroupCountX*, *workgroupCountY* and *workgroupCountZ* are \leq *this*.device.limits.[maxComputeWorkgroupsPerDimension](#).
5. Let *bindingState* be a snapshot of *this*'s current state.
6. [Enqueue a command](#) on *this* which issues the subsequent steps on the [Queue timeline](#).

[Queue timeline](#) steps:

1. Execute a grid of workgroups with dimensions [*workgroupCountX*, *workgroupCountY*, *workgroupCountZ*] with *bindingState*.[\[\[pipeline\]\]](#) using *bindingState*.[\[\[bind_groups\]\]](#).

dispatchWorkgroupsIndirect(indirectBuffer, indirectOffset)

Dispatch work to be performed with the current [GPUComputePipeline](#) using parameters read from a [GPUBuffer](#). See [§ 23.1 Computing](#) for the detailed specification.

The *indirect dispatch parameters* encoded in the buffer must be a tightly packed block of **three 32-bit unsigned integer values (12 bytes total)**, given in the same order as the arguments for [dispatchWorkgroups\(\)](#). For example:

```
let dispatchIndirectParameters = new Uint32Array(3);
dispatchIndirectParameters[0] = workgroupCountX;
dispatchIndirectParameters[1] = workgroupCountY;
dispatchIndirectParameters[2] = workgroupCountZ;
```

Called on: [GPUComputePassEncoder](#) *this*.

Arguments:

Arguments for the [GPUComputePassEncoder.dispatchWorkgroupsIndirect\(indirectBuffer, indirectOffset\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>indirectBuffer</i>	GPUBuffer	✗	✗	Buffer containing the indirect dispatch parameters .
<i>indirectOffset</i>	GPUSize64	✗	✗	Offset in bytes into <i>indirectBuffer</i> where the dispatch data begins.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. Let *usageScope* be an empty [usage scope](#).
3. For each *bindGroup* in *this*.[\[\[bind_groups\]\]](#), merge *bindGroup*.[\[\[usedResources\]\]](#) into *this*.[\[\[usage_scope\]\]](#).
4. [Add](#) *indirectBuffer* to *usageScope* with usage [input](#).
5. If any of the following conditions are unsatisfied, [invalidate this](#) and return.
 - *usageScope* must satisfy [usage scope validation](#).

- [Validate encoder bind groups](#)(*this*, *this*.[\[\[pipeline\]\]](#)) is `true`.
- *indirectBuffer* is [valid to use with this](#).
- *indirectBuffer*.[usage](#) contains [INDIRECT](#).
- $\text{indirectOffset} + \text{sizeof}(\text{indirect dispatch parameters}) \leq \text{indirectBuffer}.\text{size}$.
- *indirectOffset* is a multiple of 4.

6. Let *bindingState* be a snapshot of *this*'s current state.

7. [Enqueue a command](#) on *this* which issues the subsequent steps on the [Queue timeline](#).

[Queue timeline](#) steps:

1. Let *workgroupCountX* be an unsigned 32-bit integer read from *indirectBuffer* at *indirectOffset* bytes.
2. Let *workgroupCountY* be an unsigned 32-bit integer read from *indirectBuffer* at (*indirectOffset* + 4) bytes.
3. Let *workgroupCountZ* be an unsigned 32-bit integer read from *indirectBuffer* at (*indirectOffset* + 8) bytes.
4. If *workgroupCountX*, *workgroupCountY*, or *workgroupCountZ* is greater than *this*.*device*.limits.[maxComputeWorkgroupsPerDimension](#), return.
5. Execute a grid of workgroups with dimensions [*workgroupCountX*, *workgroupCountY*, *workgroupCountZ*] with *bindingState*.[\[\[pipeline\]\]](#) using *bindingState*.[\[\[bind_groups\]\]](#).

16.1.3. Finalization

The compute pass encoder can be ended by calling [end\(\)](#) once the user has finished recording commands for the pass. Once [end\(\)](#) has been called the compute pass encoder can no longer be used.

[end\(\)](#)

Completes recording of the compute pass commands sequence.

Called on: [GPUComputePassEncoder](#) *this*.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. Let *parentEncoder* be *this*.[\[\[command_encoder\]\]](#).
2. If any of the following requirements are unmet, [generate a validation error](#) and return.
 - *this*.[\[\[state\]\]](#) must be `"open"`.
 - *parentEncoder*.[\[\[state\]\]](#) must be `"locked"`.
3. Set *this*.[\[\[state\]\]](#) to `"ended"`.
4. Set *parentEncoder*.[\[\[state\]\]](#) to `"open"`.
5. If any of the following requirements are unmet, [invalidate](#) *parentEncoder* and return.
 - *this* must be [valid](#).
 - *this*.[\[\[debug_group_stack\]\]](#) must be empty.
6. [Extend](#) *parentEncoder*.[\[\[commands\]\]](#) with *this*.[\[\[commands\]\]](#).
7. If *this*.[\[\[endTimeStampWrite\]\]](#) is not `null`:
 1. [Extend](#) *parentEncoder*.[\[\[commands\]\]](#) with *this*.[\[\[endTimeStampWrite\]\]](#).

17. Render Passes

17.1. GPURenderPassEncoder

[Exposed=(Window, Worker), [SecureContext](#)]

```
interface GPURenderPassEncoder {
  undefined setViewport(float x, float y,
    float width, float height,
    float minDepth, float maxDepth);
```

```
  undefined setScissorRect(GPUIntegerCoordinate x, GPUIntegerCoordinate y,
```

[GPUIntegerCoordinate width, GPUIntegerCoordinate height](#));

[undefined setBlendConstant\(GPUColor color\)](#);

[undefined setStencilReference\(GPUStencilValue reference\)](#);

[undefined beginOcclusionQuery\(GPUSize32 queryIndex\)](#);

[undefined endOcclusionQuery\(\)](#);

[undefined executeBundles\(sequence<GPURenderBundle> bundles\)](#);

[undefined end\(\)](#);

};

[GPURenderPassEncoder](#) includes [GPUObjectBase](#);

[GPURenderPassEncoder](#) includes [GPUCommandsMixin](#);

[GPURenderPassEncoder](#) includes [GPUDebugCommandsMixin](#);

[GPURenderPassEncoder](#) includes [GPUBindingCommandsMixin](#);

[GPURenderPassEncoder](#) includes [GPURenderCommandsMixin](#);

[GPURenderPassEncoder](#) has the following [device timeline properties](#):

[\[\[command_encoder\]\]](#), of type [GPUCommandEncoder](#), readonly

The [GPUCommandEncoder](#) that created this render pass encoder.

[\[\[attachment_size\]\]](#), readonly

Set to the following extents:

- **width, height** = the dimensions of the pass's render attachments

[\[\[occlusion_query_set\]\]](#), of type [GPUQuerySet](#), readonly

The [GPUQuerySet](#) to store occlusion query results for the pass, which is initialized with [GPURenderPassDescriptor.occlusionQuerySet](#) at pass creation time.

[\[\[endTimestampWrite\]\]](#), of type [GPU command?](#), readonly, defaulting to null

[GPU command](#), if any, writing a timestamp when the pass ends.

[\[\[maxDrawCount\]\]](#) of type [GPUSize64](#), readonly

The maximum number of draws allowed in this pass.

[\[\[occlusion_query_active\]\]](#), of type [boolean](#)

Whether the pass's [\[\[occlusion_query_set\]\]](#) is being written.

When executing encoded render pass commands as part of a [GPUCommandBuffer](#), an internal *RenderState* object is used to track the current state required for rendering.

[RenderState](#) has the following [queue timeline properties](#):

[\[\[occlusionQueryIndex\]\]](#), of type [GPUSize32](#)

The index into [\[\[occlusion_query_set\]\]](#) at which to store the occlusion query results.

[\[\[viewport\]\]](#)

Current viewport rectangle and depth range. Initially set to the following values:

- **x, y** = 0.0, 0.0
- **width, height** = the dimensions of the pass's render targets
- **minDepth, maxDepth** = 0.0, 1.0

[\[\[scissorRect\]\]](#)

Current scissor rectangle. Initially set to the following values:

- **x, y** = 0, 0
- **width, height** = the dimensions of the pass's render targets

[\[\[blendConstant\]\]](#), of type [GPUColor](#)

Current blend constant value, initially [0, 0, 0, 0].

[\[\[stencilReference\]\]](#), of type [GPUStencilValue](#)

Current stencil reference value, initially 0.

[\[\[colorAttachments\]\]](#), of type [sequence<GPURenderPassColorAttachment?>](#)

The color attachments and state for this render pass.

[\[\[depthStencilAttachment\]\]](#), of type [GPURenderPassDepthStencilAttachment?](#)

The depth/stencil attachment and state for this render pass.

Render passes also have *framebuffer memory*, which contains the [texel](#) data associated with each attachment that is written into by draw commands and read from for blending and depth/stencil testing.

Note: Depending on the GPU hardware, [framebuffer memory](#) may be the memory allocated by the attachment textures or may be a separate area of memory that the texture data is copied to and from, such as with tile-based architectures.

17.1.1.1. Render Pass Encoder Creation

dictionary

GPURenderPassTimestampWrites

```
{
    required GPUQuerySet querySet;
    GPUSize32 beginningOfPassWriteIndex;
    GPUSize32 endOfPassWriteIndex;
};
```

querySet, of type [GPUQuerySet](#)

The [GPUQuerySet](#), of type ["timestamp"](#), that the query results will be written to.

beginningOfPassWriteIndex, of type [GPUSize32](#)

If defined, indicates the query index in [querySet](#) into which the timestamp at the beginning of the render pass will be written.

endOfPassWriteIndex, of type [GPUSize32](#)

If defined, indicates the query index in [querySet](#) into which the timestamp at the end of the render pass will be written.

Note: Timestamp query values are written in nanoseconds, but how the value is determined is [implementation-defined](#). See [§ 20.4 Timestamp Query](#) for details.

dictionary

GPURenderPassDescriptor

```
: GPUObjectDescriptorBase {
    required sequence<GPURenderPassColorAttachment?> colorAttachments;
    GPURenderPassDepthStencilAttachment depthStencilAttachment;
    GPUQuerySet occlusionQuerySet;
    GPURenderPassTimestampWrites timestampWrites;
    GPUSize64 maxDrawCount = 50000000;
};
```

colorAttachments, of type [sequence<GPURenderPassColorAttachment?>](#)

The set of [GPURenderPassColorAttachment](#) values in this sequence defines which color attachments will be output to when executing this render pass.

Due to [usage compatibility](#), no color attachment may alias another attachment or any resource used inside the render pass.

depthStencilAttachment, of type [GPURenderPassDepthStencilAttachment](#)

The [GPURenderPassDepthStencilAttachment](#) value that defines the depth/stencil attachment that will be output to and tested against when executing this render pass.

Due to [usage compatibility](#), no writable depth/stencil attachment may alias another attachment or any resource used inside the render pass.

occlusionQuerySet, of type [GPUQuerySet](#)

The [GPUQuerySet](#) value defines where the occlusion query results will be stored for this pass.

timestampWrites, of type [GPURenderPassTimestampWrites](#)

Defines which timestamp values will be written for this pass, and where to write them to.

maxDrawCount, of type [GPUSize64](#), defaulting to 50000000

The maximum number of draw calls that will be done in the render pass. Used by some implementations to size work injected before the render pass. Keeping the default value is a good default, unless it is known that more draw calls will be done.

Valid Usage

Given a [GPUDevice](#) *device* and [GPURenderPassDescriptor](#) *this*, the following validation rules apply:

this.colorAttachments.size must be \leq *device*.[\[\[limits\]\].maxColorAttachments](#).

For each non-`null` *colorAttachment* in *this.colorAttachments*:

colorAttachment.view must be [valid to use with](#) *device*.

If `colorAttachment.resolveTarget` is [provided](#):

`colorAttachment.resolveTarget` must be [valid to use with device](#).

`colorAttachment` must meet the [GPURenderPassColorAttachment Valid Usage](#) rules.

If `this.depthStencilAttachment` is [provided](#):

`this.depthStencilAttachment.view` must be [valid to use with device](#).

`this.depthStencilAttachment` must meet the [GPURenderPassDepthStencilAttachment Valid Usage](#) rules.

There must exist at least one attachment, either:

A non-`null` value in `this.colorAttachments`, or

A `this.depthStencilAttachment`.

[Validating GPURenderPassDescriptor's color attachment bytes per sample](#)(`device`, `this.colorAttachments`) succeeds.

All [views](#) in non-`null` members of `this.colorAttachments`, and `this.depthStencilAttachment.view` if present, must have equal [sampleCounts](#).

For each [view](#) in non-`null` members of `this.colorAttachments` and `this.depthStencilAttachment.view`, if present, the [\[\[renderExtent\]\]](#) must match.

If `this.occlusionQuerySet` is [provided](#):

`this.occlusionQuerySet` must be [valid to use with device](#).

`this.occlusionQuerySet.type` must be [occlusion](#).

If `this.timestampWrites` is [provided](#):

[Validate timestampWrites](#)(`device`, `this.timestampWrites`) must return true.

[Validating GPURenderPassDescriptor's color attachment bytes per sample](#)(`device`, `colorAttachments`)

Arguments:

[GPUDevice](#) `device`

[sequence](#)<[GPURenderPassColorAttachment](#)?> `colorAttachments`

[Device timeline](#) steps:

Let `formats` be an empty [list](#)<[GPUTextureFormat](#)?>

For each `colorAttachment` in `colorAttachments`:

If `colorAttachment` is `undefined`, continue.

[Append](#) `colorAttachment.view[[descriptor]].format` to `formats`.

[Calculating color attachment bytes per sample](#)(`formats`) must be \leq `device.[[limits]].maxColorAttachmentBytesPerSample`.

17.1.1.1. Color Attachments

dictionary

[GPURenderPassColorAttachment](#)

```
{
  required (GPUTexture or GPUTextureView) view;
  GPUIntegerCoordinate depthSlice;
  (GPUTexture or GPUTextureView) resolveTarget;

  GPUColor clearValue;
  required GPULoadOp loadOp;
  required GPUStoreOp storeOp;
};
```

`view`, of type ([GPUTexture](#) or [GPUTextureView](#))

Describes the texture [subresource](#) that will be output to for this color attachment. The [subresource](#) is determined by calling [get as texture view](#)(`view`).

`depthSlice`, of type [GPUIntegerCoordinate](#)

Indicates the depth slice index of ["3d" view](#) that will be output to for this color attachment.

`resolveTarget`, of type ([GPUTexture](#) or [GPUTextureView](#))

Describes the texture [subresource](#) that will receive the resolved output for this color attachment if `view` is multisampled. The [subresource](#) is determined by calling [get as texture view](#)(`resolveTarget`).

`clearValue`, of type [GPUColor](#)

Indicates the value to clear [view](#) prior to executing the render pass. If not [provided](#), defaults to {r: 0, g: 0, b: 0, a: 0}. Ignored if [loadOp](#) is not ["clear"](#).

The components of [clearValue](#) are all double values. They are converted [to a texel value of texture format](#) matching the render attachment. If conversion fails, a validation error is generated.

`loadOp`, of type [GPULoadOp](#)

Indicates the load operation to perform on [view](#) prior to executing the render pass.

Note: It is recommended to prefer clearing; see ["clear"](#) for details.

`storeOp`, of type [GPUStoreOp](#)

The store operation to perform on [view](#) after executing the render pass.

GPURenderPassColorAttachment Valid Usage

Given a [GPURenderPassColorAttachment](#) *this*:

Let *renderViewDescriptor* be *this.view.[[descriptor]]*.

Let *renderTexture* be *this.view.[[texture]]*.

All of the requirements in the following steps *must* be met.

renderViewDescriptor.format *must* be a [color renderable format](#).

this.view *must* be a [renderable texture view](#).

If *renderViewDescriptor.dimension* is ["3d"](#):

this.depthSlice *must be provided* and *must* be < the [depthOrArrayLayers](#) of the [logical miplevel-specific texture extent](#) of the *renderTexture subresource* at [mipmap level](#) *renderViewDescriptor.baseMipLevel*.

Otherwise:

this.depthSlice *must not be provided*.

If *this.loadOp* is ["clear"](#):

Converting the IDL value *this.clearValue to a texel value of texture format* *renderViewDescriptor.format* *must not throw* a [TypeError](#).

Note: An error is not thrown if the value is out-of-range for the format but in-range for the corresponding WGSL primitive type (f32, i32, or u32).

If *this.resolveTarget* is [provided](#):

Let *resolveViewDescriptor* be *this.resolveTarget.[[descriptor]]*.

Let *resolveTexture* be *this.resolveTarget.[[texture]]*.

renderTexture.sampleCount *must be* > 1.

resolveTexture.sampleCount *must be* 1.

this.resolveTarget *must be* a non-3d [renderable texture view](#).

this.resolveTarget.[[renderExtent]] and *this.view.[[renderExtent]]* *must match*.

resolveViewDescriptor.format *must equal* *renderViewDescriptor.format*.

resolveTexture.format *must equal* *renderTexture.format*.

resolveViewDescriptor.format *must support resolve according to* [§ 26.1.1 Plain color formats](#).

A [GPUTextureView](#) *view* is a [renderable texture view](#) if the all of the requirements in the following [device timeline](#) steps are met:

Let *descriptor* be *view.[[descriptor]]*.

descriptor.usage *must contain* [RENDER_ATTACHMENT](#).

descriptor.dimension *must be* ["2d"](#) or ["2d-array"](#) or ["3d"](#).

descriptor.mipLevelCount *must be* 1.

descriptor.arrayLayerCount *must be* 1.

descriptor.aspect *must refer to all* [aspects](#) of *view.[[texture]]*.

Calculating color attachment bytes per sample(formats)

Arguments:

[sequence<GPUTextureFormat?>](#) *formats*

Returns: [GPUSize32](#)

Let *total* be 0.

For each non-null *format* in *formats*

[Assert](#): *format* is a [color renderable format](#).

Let *renderTargetPixelByteCost* be the [render target pixel byte cost](#) of *format*.

Let *renderTargetComponentAlignment* be the [render target component alignment](#) of *format*.

Round *total* up to the smallest multiple of *renderTargetComponentAlignment* greater than or equal to *total*.

Add *renderTargetPixelByteCost* to *total*.

Return *total*.

17.1.1.2. Depth/Stencil Attachments

dictionary

GPURenderPassDepthStencilAttachment

```
{
  required (GPUTexture or GPUTextureView) view;
```

```
  float depthClearValue;
```

```
  GPULoadOp depthLoadOp;
```

```
  GPUStoreOp depthStoreOp;
```

```
  boolean depthReadOnly = false;
```

```
  GPUStencilValue stencilClearValue = 0;
```

```
  GPULoadOp stencilLoadOp;
```

```
  GPUStoreOp stencilStoreOp;
```

```
  boolean stencilReadOnly = false;
```

```
};
```

view, of type ([GPUTexture](#) or [GPUTextureView](#))

Describes the texture [subresource](#) that will be output to and read from for this depth/stencil attachment. The [subresource](#) is determined by calling [get as texture view](#)([view](#)).

depthClearValue, of type [float](#)

Indicates the value to clear [view](#)'s depth component to prior to executing the render pass. Ignored if [depthLoadOp](#) is not "[clear](#)". Must be between 0.0 and 1.0, inclusive.

depthLoadOp, of type [GPULoadOp](#)

Indicates the load operation to perform on [view](#)'s depth component prior to executing the render pass.

Note: It is recommended to prefer clearing; see "[clear](#)" for details.

depthStoreOp, of type [GPUStoreOp](#)

The store operation to perform on [view](#)'s depth component after executing the render pass.

depthReadOnly, of type [boolean](#), defaulting to `false`

Indicates that the depth component of [view](#) is read only.

stencilClearValue, of type [GPUStencilValue](#), defaulting to 0

Indicates the value to clear [view](#)'s stencil component to prior to executing the render pass. Ignored if [stencilLoadOp](#) is not "[clear](#)".

The value will be converted to the type of the stencil aspect of [view](#) by taking the same number of LSBs as the number of bits in the stencil aspect of one [texel](#) of [view](#).

stencilLoadOp, of type [GPULoadOp](#)

Indicates the load operation to perform on [view](#)'s stencil component prior to executing the render pass.

Note: It is recommended to prefer clearing; see "[clear](#)" for details.

stencilStoreOp, of type [GPUStoreOp](#)

The store operation to perform on [view](#)'s stencil component after executing the render pass.

stencilReadOnly, of type [boolean](#), defaulting to `false`

Indicates that the stencil component of [view](#) is read only.

GPURenderPassDepthStencilAttachment Valid Usage

Given a [GPURenderPassDepthStencilAttachment](#) *this*, the following validation rules apply:

this.view must have a [depth-or-stencil format](#).

this.view must be a [renderable texture view](#).

Let *format* be *this.view.[[descriptor]].format*.

If *this.depthLoadOp* is ["clear"](#), *this.depthClearValue* must [be provided](#) and must be between 0.0 and 1.0, inclusive.

If *format* has a depth aspect and *this.depthReadOnly* is [false](#):

this.depthLoadOp must [be provided](#).

this.depthStoreOp must [be provided](#).

Otherwise:

this.depthLoadOp must not [be provided](#).

this.depthStoreOp must not [be provided](#).

If *format* has a stencil aspect and *this.stencilReadOnly* is [false](#):

this.stencilLoadOp must [be provided](#).

this.stencilStoreOp must [be provided](#).

Otherwise:

this.stencilLoadOp must not [be provided](#).

this.stencilStoreOp must not [be provided](#).

17.1.1.3. Load & Store Operations

enum

[GPULoadOp](#)

```
{  
  "load",  
  "clear",  
};
```

["load"](#)

Loads the existing value for this attachment into the render pass.

["clear"](#)

Loads a clear value for this attachment into the render pass.

Note: On some GPU hardware (primarily mobile), ["clear"](#) is significantly cheaper because it avoids loading data from main memory into tile-local memory. On other GPU hardware, there isn't a significant difference. As a result, it is recommended to use ["clear"](#) rather than ["load"](#) in cases where the initial value doesn't matter (e.g. the render target will be cleared using a skybox).

enum

[GPUStoreOp](#)

```
{  
  "store",  
  "discard",  
};
```

["store"](#)

Stores the resulting value of the render pass for this attachment.

["discard"](#)

Discards the resulting value of the render pass for this attachment.

Note: Discarded attachments behave as if they are cleared to zero, but implementations are not required to perform a clear at the end of the render pass.

Implementations which do not explicitly clear discarded attachments at the end of a pass must lazily clear them prior to the reading the attachment contents, which occurs via sampling, copies, attaching to a later render pass with ["load"](#), displaying or reading back the canvas ([get a copy of the image contents of a context](#)), etc.

17.1.1.4. Render Pass Layout

[GPURenderPassLayout](#) declares the layout of the render targets of a [GPURenderBundle](#). It is also used internally to describe [GPURenderPassEncoder layouts](#)

and [GPURenderPipeline layouts](#). It determines compatibility between render passes, render bundles, and render pipelines.

dictionary

[GPURenderPassLayout](#)

```
: GPUObjectDescriptorBase {  
    required sequence<GPUTextureFormat?> colorFormats;  
    GPUTextureFormat depthStencilFormat;  
    GPUSize32 sampleCount = 1;  
};
```

[colorFormats](#), of type [sequence](#)<GPUTextureFormat?>

A list of the [GPUTextureFormat](#)s of the color attachments for this pass or bundle.

[depthStencilFormat](#), of type [GPUTextureFormat](#)

The [GPUTextureFormat](#) of the depth/stencil attachment for this pass or bundle.

[sampleCount](#), of type [GPUSize32](#), defaulting to 1

Number of samples per pixel in the attachments for this pass or bundle.

Two [GPURenderPassLayout](#) values are *equal* if:

Their [depthStencilFormat](#) and [sampleCount](#) are equal, and

Their [colorFormats](#) are equal ignoring any trailing `null`s.

derive render targets layout from pass

Arguments:

[GPURenderPassDescriptor](#) *descriptor*

Returns: [GPURenderPassLayout](#)

[Device timeline](#) steps:

Let *layout* be a new [GPURenderPassLayout](#) object.

For each *colorAttachment* in *descriptor*.[colorAttachments](#):

If *colorAttachment* is not `null`:

Set *layout*.[sampleCount](#) to *colorAttachment*.[view](#).[\[\[texture\]\]](#).[sampleCount](#).

Append *colorAttachment*.[view](#).[\[\[descriptor\]\]](#).[format](#) to *layout*.[colorFormats](#).

Otherwise:

Append `null` to *layout*.[colorFormats](#).

Let *depthStencilAttachment* be *descriptor*.[depthStencilAttachment](#).

If *depthStencilAttachment* is not `null`:

Let *view* be *depthStencilAttachment*.[view](#).

Set *layout*.[sampleCount](#) to *view*.[\[\[texture\]\]](#).[sampleCount](#).

Set *layout*.[depthStencilFormat](#) to *view*.[\[\[descriptor\]\]](#).[format](#).

Return *layout*.

derive render targets layout from pipeline

Arguments:

[GPURenderPipelineDescriptor](#) *descriptor*

Returns: [GPURenderPassLayout](#)

[Device timeline](#) steps:

Let *layout* be a new [GPURenderPassLayout](#) object.

Set *layout*.[sampleCount](#) to *descriptor*.[multisample.count](#).

If *descriptor*.[depthStencil](#) is *provided*:

Set *layout*.[depthStencilFormat](#) to *descriptor*.[depthStencil.format](#).

If *descriptor*.[fragment](#) is *provided*:

For each *colorTarget* in *descriptor.fragment.targets*:

Append *colorTarget.format* to *layout.colorFormats* if *colorTarget* is not `null`, or append `null` otherwise.

Return *layout*.

17.1.2. Finalization

The render pass encoder can be ended by calling `end()` once the user has finished recording commands for the pass. Once `end()` has been called the render pass encoder can no longer be used.

`end()`

Completes recording of the render pass commands sequence.

Called on: `GPURenderPassEncoder` *this*.

Returns: `undefined`

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this.device*.

[Device timeline](#) steps:

1. Let *parentEncoder* be *this.command_encoder*.
2. If any of the following requirements are unmet, [generate a validation error](#) and return.
 - *this.state* must be "open".
 - *parentEncoder.state* must be "locked".
3. Set *this.state* to "ended".
4. Set *parentEncoder.state* to "open".
5. If any of the following requirements are unmet, [invalidate](#) *parentEncoder* and return.
 - *this* must be [valid](#).
 - *this.usage_scope* must satisfy [usage scope validation](#).
 - *this.debug_group_stack* must be empty.
 - *this.occlusion_query_active* must be false.
 - *this.drawCount* must be \leq *this.maxDrawCount*.
6. [Extend](#) *parentEncoder.commands* with *this.commands*.
7. If *this.endTimestampWrite* is not `null`:
 1. [Extend](#) *parentEncoder.commands* with *this.endTimestampWrite*.
 8. [Enqueue a render command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) with *renderState* when executed.

[Queue timeline](#) steps:

1. For each non-`null` *colorAttachment* in *renderState.colorAttachments*:
 1. Let *colorView* be *colorAttachment.view*.
 2. If *colorView.descriptor.dimension* is:
 - ["3d"](#)
Let *colorSubregion* be *colorAttachment.depthSlice* of *colorView*.
 - Otherwise
Let *colorSubregion* be *colorView*.
 3. If *colorAttachment.resolveTarget* is not `null`:
 1. Resolve the multiple samples of every [texel](#) of *colorSubregion* to a single sample and copy to *colorAttachment.resolveTarget*.
 4. If *colorAttachment.storeOp* is:
 - ["store"](#)
Ensure the contents of the [framebuffer memory](#) associated with *colorSubregion* are stored in *colorSubregion*.
 - ["discard"](#)
Set every [texel](#) of *colorSubregion* to zero.

2. Let `depthStencilAttachment` be `renderState.[[depthStencilAttachment]]`.

3. If `depthStencilAttachment` is not `null`:

1. If `depthStencilAttachment.depthStoreOp` is:

Not `provided`

`Assert` that `depthStencilAttachment.depthReadOnly` is `true` and leave the `depth subresource` of `depthStencilView` unchanged.

`"store"`

Ensure the contents of the `framebuffer memory` associated with the `depth subresource` of `depthStencilView` are stored in `depthStencilView`.

`"discard"`

Set every `texel` in the `depth subresource` of `depthStencilView` to zero.

2. If `depthStencilAttachment.stencilStoreOp` is:

Not `provided`

`Assert` that `depthStencilAttachment.stencilReadOnly` is `true` and leave the `stencil subresource` of `depthStencilView` unchanged.

`"store"`

Ensure the contents of the `framebuffer memory` associated with the `stencil subresource` of `depthStencilView` are stored in `depthStencilView`.

`"discard"`

Set every `texel` in the `stencil subresource` of `depthStencilView` to zero.

4. Let `renderState` be `null`.

Note: Discarded attachments behave as if they are cleared to zero, but implementations are not required to perform a clear at the end of the render pass. See the note on `"discard"` for additional details.

Note: `Read-only depth-stencil` attachments can be thought of as implicitly using the `"store"` operation, but since their content is unchanged during the render pass implementations don't need to update the attachment. Validation that requires the store op to not be provided for read-only attachments is done in `GPURenderPassDepthStencilAttachment Valid Usage`.

17.2. GPURenderCommandsMixin

`GPURenderCommandsMixin` defines rendering commands common to `GPURenderPassEncoder` and `GPURenderBundleEncoder`.

interface mixin `GPURenderCommandsMixin` {

`undefined setPipeline(GPURenderPipeline pipeline);`

`undefined setIndexBuffer(GPUBuffer buffer, GPUIndexFormat indexFormat, optional GPUSize64 offset = 0, optional GPUSize64 size);`

`undefined setVertexBuffer(GPUIndex32 slot, GPUBuffer? buffer, optional GPUSize64 offset = 0, optional GPUSize64 size);`

`undefined draw(GPUSize32 vertexCount, optional GPUSize32 instanceCount = 1, optional GPUSize32 firstVertex = 0, optional GPUSize32 firstInstance = 0);`
`undefined drawIndexed(GPUSize32 indexCount, optional GPUSize32 instanceCount = 1, optional GPUSize32 firstIndex = 0, optional GPUSignedOffset32 baseVertex = 0, optional GPUSize32 firstInstance = 0);`

`undefined drawIndirect(GPUBuffer indirectBuffer, GPUSize64 indirectOffset);`

`undefined drawIndexedIndirect(GPUBuffer indirectBuffer, GPUSize64 indirectOffset);`

};

`GPURenderCommandsMixin` assumes the presence of `GPUObjectBase`, `GPUCommandsMixin`, and `GPUBindingCommandsMixin` members on the same object. It must only be included by interfaces which also include those mixins.

`GPURenderCommandsMixin` has the following `device timeline properties`:

`[[layout]]`, of type `GPURenderPassLayout`, readonly

The layout of the render pass.

`[[depthReadOnly]]`, of type `boolean`, readonly

If `true`, indicates that the depth component is not modified.

`[[stencilReadOnly]]`, of type `boolean`, readonly

If `true`, indicates that the stencil component is not modified.

`[[usage scope]]`, of type `usage scope`, initially empty

The `usage scope` for this render pass or bundle.

[[pipeline]], of type [GPURenderPipeline](#), initially null
The current [GPURenderPipeline](#).

[[index_buffer]], of type [GPUBuffer](#), initially null
The current buffer to read index data from.

[[index_format]], of type [GPUIndexFormat](#)
The format of the index data in [\[\[index_buffer\]\]](#).

[[index_buffer_offset]], of type [GPUSize64](#)
The offset in bytes of the section of [\[\[index_buffer\]\]](#) currently set.

[[index_buffer_size]], of type [GPUSize64](#)
The size in bytes of the section of [\[\[index_buffer\]\]](#) currently set, initially 0.

[[vertex_buffers]], of type [ordered_map](#)<slot, [GPUBuffer](#)>, initially empty
The current [GPUBuffer](#)s to read vertex data from for each slot.

[[vertex_buffer_sizes]], of type [ordered_map](#)<slot, [GPUSize64](#)>, initially empty
The size in bytes of the section of [GPUBuffer](#) currently set for each slot.

[[drawCount]], of type [GPUSize64](#)
The number of draw commands recorded in this encoder.

To *Enqueue a render command* on [GPURenderCommandsMixin](#) encoder which issues the steps of a [GPU Command](#) command with [RenderState](#) renderState, run the following [device timeline](#) steps:

[Append](#) command to encoder.[\[\[commands\]\]](#).

When *command* is executed as part of a [GPUCommandBuffer](#) *commandBuffer*:

Issue the steps of *command* with *commandBuffer*.[\[\[renderState\]\]](#) as *renderState*.

17.2.1. Drawing

setPipeline(pipeline)

Sets the current [GPURenderPipeline](#).

Called on: [GPURenderCommandsMixin](#) this.

Arguments:

Arguments for the [GPURenderCommandsMixin.setPipeline\(pipeline\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>pipeline</i>	GPURenderPipeline	✗	✗	The render pipeline to use for subsequent drawing commands.

Returns: [undefined](#)

[Content timeline](#) steps:

- Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

- [Validate the encoder state](#) of *this*. If it returns false, return.
- Let *pipelineTargetsLayout* be [derive render targets layout from pipeline](#)(*pipeline*.[\[\[descriptor\]\]](#)).
- If any of the following conditions are unsatisfied, [invalidate](#) *this* and return.
 - pipeline* is [valid to use with](#) *this*.
 - this*.[\[\[layout\]\]](#) [equals](#) *pipelineTargetsLayout*.
 - If *pipeline*.[\[\[writesDepth\]\]](#): *this*.[\[\[depthReadOnly\]\]](#) must be false.
 - If *pipeline*.[\[\[writesStencil\]\]](#): *this*.[\[\[stencilReadOnly\]\]](#) must be false.
- Set *this*.[\[\[pipeline\]\]](#) to be *pipeline*.

setIndexBuffer(buffer, indexFormat, offset, size)

Sets the current index buffer.

Called on: [GPURenderCommandsMixin](#) this.

Arguments:

Arguments for the [GPURenderCommandsMixin.setIndexBuffer\(buffer, indexFormat, offset, size\)](#) method.

Parameter	Type	Nullable	Optional	Description
-----------	------	----------	----------	-------------

Parameter	Type	Nullable	Optional	Description
<i>buffer</i>	GPUBuffer	✗	✗	Buffer containing index data to use for subsequent drawing commands.
<i>indexFormat</i>	GPUIndexFormat	✗	✗	Format of the index data contained in <i>buffer</i> .
<i>offset</i>	GPUSize64	✗	✓	Offset in bytes into <i>buffer</i> where the index data begins. Defaults to 0.
<i>size</i>	GPUSize64	✗	✓	Size in bytes of the index data in <i>buffer</i> . Defaults to the size of the buffer minus the offset.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of this.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of this. If it returns false, return.
2. If *size* is missing, set *size* to $\max(0, \text{buffer.size} - \text{offset})$.
3. If any of the following conditions are unsatisfied, [invalidate this](#) and return.
 - *buffer* is [valid to use with this](#).
 - *buffer.usage* contains [INDEX](#).
 - *offset* is a multiple of *indexFormat*'s byte size.
 - $\text{offset} + \text{size} \leq \text{buffer.size}$.
4. [Add buffer](#) to [\[\[usage_scope\]\]](#) with usage [input](#).
5. Set this.[\[\[index_buffer\]\]](#) to be *buffer*.
6. Set this.[\[\[index_format\]\]](#) to be *indexFormat*.
7. Set this.[\[\[index_buffer_offset\]\]](#) to be *offset*.
8. Set this.[\[\[index_buffer_size\]\]](#) to be *size*.

setVertexBuffer(slot, buffer, offset, size)

Sets the current vertex buffer for the given slot.

Called on: [GPURenderCommandsMixin](#) this.

Arguments:

Arguments for the [GPURenderCommandsMixin.setVertexBuffer\(slot, buffer, offset, size\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>slot</i>	GPUIndex32	✗	✗	The vertex buffer slot to set the vertex buffer for.
<i>buffer</i>	GPUBuffer?	✓	✗	Buffer containing vertex data to use for subsequent drawing commands.
<i>offset</i>	GPUSize64	✗	✓	Offset in bytes into <i>buffer</i> where the vertex data begins. Defaults to 0.
<i>size</i>	GPUSize64	✗	✓	Size in bytes of the vertex data in <i>buffer</i> . Defaults to the size of the buffer minus the offset.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of this.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of this. If it returns false, return.
2. Let *bufferSize* be 0 if *buffer* is `null`, or *buffer.size* if not.
3. If *size* is missing, set *size* to $\max(0, \text{bufferSize} - \text{offset})$.
4. If any of the following requirements are unmet, [invalidate this](#) and return.
 - *slot* must be $< \text{this}.\text{[[device]]}.\text{[[limits]]}.\text{maxVertexBuffers}$.
 - *offset* must be a multiple of 4.
 - $\text{offset} + \text{size}$ must be $\leq \text{bufferSize}$.
5. If *buffer* is `null`:
 1. [Remove this](#).[\[\[vertex_buffers\]\]](#)[*slot*].
 2. [Remove this](#).[\[\[vertex_buffer_sizes\]\]](#)[*slot*].

Otherwise:

1. If any of the following requirements are unmet, [invalidate this](#) and return.
 - *buffer* must be [valid to use with this](#).
 - *buffer.usage* must contain [VERTEX](#).
2. Add *buffer* to [\[\[usage_scope\]\]](#) with usage [input](#).
3. Set *this.[vertex_buffers][slot]* to be *buffer*.
4. Set *this.[vertex_buffer_sizes][slot]* to be *size*.

draw(vertexCount, instanceCount, firstVertex, firstInstance)

Draws primitives. See [§ 23.2 Rendering](#) for the detailed specification.

Called on: [GPURenderCommandsMixin](#) this.

Arguments:

Arguments for the [GPURenderCommandsMixin.draw\(vertexCount, instanceCount, firstVertex, firstInstance\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>vertexCount</i>	GPUSize32	✗	✗	The number of vertices to draw.
<i>instanceCount</i>	GPUSize32	✗	✓	The number of instances to draw.
<i>firstVertex</i>	GPUSize32	✗	✓	Offset into the vertex buffers, in vertices, to begin drawing from.
<i>firstInstance</i>	GPUSize32	✗	✓	First instance to draw.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this. [[device]]*.

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. All of the requirements in the following steps *must* be met. If any are unmet, [invalidate this](#) and return.
 1. It *must* be [valid to draw](#) with *this*.
 2. Let *buffers* be *this. [[pipeline]]. [[descriptor]]. vertex.buffers*.
 3. For each [GPUIndex32](#) *slot* from 0 to *buffers.size* (non-inclusive):
 1. If *buffers[slot]* is null, [continue](#).
 2. Let *bufferSize* be *this. [[vertex_buffer_sizes]][slot]*.
 3. Let *stride* be *buffers[slot].arrayStride*.
 4. Let *attributes* be *buffers[slot].attributes*
 5. Let *lastStride* be the maximum value of (*attribute.offset* + [byteSize\(attribute.format\)](#)) over each *attribute* in *attributes*, or 0 if *attributes* is [empty](#).
 6. Let *strideCount* be computed based on *buffers[slot].stepMode*:

["vertex"](#)

firstVertex + *vertexCount*

["instance"](#)

firstInstance + *instanceCount*

7. If *strideCount* ≠ 0:
 1. (*strideCount* − 1) × *stride* + *lastStride* *must* be ≤ *bufferSize*.
 3. Increment *this. [[drawCount]]* by 1.
 4. Let *bindingState* be a snapshot of *this*'s current state.
 5. [Enqueue a render command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) with *renderState* when executed.

[Queue timeline](#) steps:

1. Draw *instanceCount* instances, starting with instance *firstInstance*, of primitives consisting of *vertexCount* vertices, starting with vertex *firstVertex*, with the states from *bindingState* and *renderState*.

drawIndexed(indexCount, instanceCount, firstIndex, baseVertex, firstInstance)

Draws indexed primitives. See [§ 23.2 Rendering](#) for the detailed specification.

Called on: [GPURenderCommandsMixin](#) this.

Arguments:

Arguments for the [GPURenderCommandsMixin.drawIndexed\(indexCount, instanceCount, firstIndex, baseVertex, firstInstance\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>indexCount</i>	GPUSize32	✗	✗	The number of indices to draw.
<i>instanceCount</i>	GPUSize32	✗	✓	The number of instances to draw.
<i>firstIndex</i>	GPUSize32	✗	✓	Offset into the index buffer, in indices, begin drawing from.
<i>baseVertex</i>	GPUSignedOffset32	✗	✓	Added to each index value before indexing into the vertex buffers.
<i>firstInstance</i>	GPUSize32	✗	✓	First instance to draw.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. If any of the following conditions are unsatisfied, [invalidate](#) *this* and return.
 - It is [valid to draw indexed](#) with *this*.
 - $firstIndex + indexCount \leq this.[[index_buffer_size]] \div this.[[index_format]]$'s byte size;
 - Let *buffers* be *this*.[\[\[pipeline\]\].\[\[descriptor\]\].vertex.buffers](#).
 - For each [GPUIndex32](#) slot from 0 to *buffers.size* (non-inclusive):
 - If *buffers[slot]* is null, [continue](#).
 - Let *bufferSize* be *this*.[\[\[vertex_buffer_sizes\]\]\[slot\]](#).
 - Let *stride* be *buffers[slot]*.[arrayStride](#).
 - Let *lastStride* be $\max(attribute.offset + \text{byteSize}(attribute.format))$ for each *attribute* in *buffers[slot]*.[attributes](#).
 - Let *strideCount* be *firstInstance* + *instanceCount*.
 - If *buffers[slot]*.[stepMode](#) is ["instance"](#) and *strideCount* \neq 0:
 - Ensure $(strideCount - 1) \times stride + lastStride \leq bufferSize$.
3. Increment *this*.[\[\[drawCount\]\]](#) by 1.
4. Let *bindingState* be a snapshot of *this*'s current state.
5. [Enqueue a render command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) with *renderState* when executed.

[Queue timeline](#) steps:

1. Draw *instanceCount* instances, starting with instance *firstInstance*, of primitives consisting of *indexCount* indexed vertices, starting with index *firstIndex* from vertex *baseVertex*, with the states from *bindingState* and *renderState*.

Note: WebGPU applications should never use index data with indices out of bounds of any bound vertex buffer that has [GPUVertexStepMode "vertex"](#). WebGPU implementations have different ways of handling this, and therefore a range of behaviors is allowed. Either the whole draw call is discarded, or the access to those attributes out of bounds is described by WGSL's [invalid memory reference](#).

drawIndirect(indirectBuffer, indirectOffset)

Draws primitives using parameters read from a [GPUBuffer](#). See [§ 23.2 Rendering](#) for the detailed specification.

The *indirect draw parameters* encoded in the buffer must be a tightly packed block of **four 32-bit unsigned integer values (16 bytes total)**, given in the same order as the arguments for [draw\(\)](#). For example:

```
let drawIndirectParameters = new Uint32Array(4);
drawIndirectParameters[0] = vertexCount;
drawIndirectParameters[1] = instanceCount;
drawIndirectParameters[2] = firstVertex;
drawIndirectParameters[3] = firstInstance;
```

The value corresponding to *firstInstance* must be 0, unless the ["indirect-first-instance" feature](#) is enabled. If the ["indirect-first-instance" feature](#) is not enabled and *firstInstance* is not zero the [drawIndirect\(\)](#) call will be treated as a no-op.

Called on: [GPURenderCommandsMixin](#) this.

Arguments:

Arguments for the [GPURenderCommandsMixin.drawIndirect\(indirectBuffer, indirectOffset\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>indirectBuffer</i>	GPUBuffer	✗	✗	Buffer containing the indirect draw parameters .
<i>indirectOffset</i>	GPUSize64	✗	✗	Offset in bytes into <i>indirectBuffer</i> where the drawing data begins.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. If any of the following conditions are unsatisfied, [invalidate](#) *this* and return.
 - It is [valid to draw](#) with *this*.
 - *indirectBuffer* is [valid to use with](#) *this*.
 - *indirectBuffer*.[usage](#) contains [INDIRECT](#).
 - $\text{indirectOffset} + \text{sizeof}(\text{indirect draw parameters}) \leq \text{indirectBuffer}.\text{size}$.
 - *indirectOffset* is a multiple of 4.
3. [Add](#) *indirectBuffer* to [\[\[usage_scope\]\]](#) with usage [input](#).
4. Increment *this*.[\[\[drawCount\]\]](#) by 1.
5. Let *bindingState* be a snapshot of *this*'s current state.
6. [Enqueue a render command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) with *renderState* when executed.

[Queue timeline](#) steps:

1. Let *vertexCount* be an unsigned 32-bit integer read from *indirectBuffer* at *indirectOffset* bytes.
2. Let *instanceCount* be an unsigned 32-bit integer read from *indirectBuffer* at (*indirectOffset* + 4) bytes.
3. Let *firstVertex* be an unsigned 32-bit integer read from *indirectBuffer* at (*indirectOffset* + 8) bytes.
4. Let *firstInstance* be an unsigned 32-bit integer read from *indirectBuffer* at (*indirectOffset* + 12) bytes.
5. Draw *instanceCount* instances, starting with instance *firstInstance*, of primitives consisting of *vertexCount* vertices, starting with vertex *firstVertex*, with the states from *bindingState* and *renderState*.

drawIndexedIndirect(indirectBuffer, indirectOffset)

Draws indexed primitives using parameters read from a [GPUBuffer](#). See [§ 23.2 Rendering](#) for the detailed specification.

The *indirect drawIndexed parameters* encoded in the buffer must be a tightly packed block of **five 32-bit values (20 bytes total)**, given in the same order as the arguments for [drawIndexed\(\)](#). The value corresponding to **baseVertex** is a signed 32-bit integer, and all others are unsigned 32-bit integers. For example:

```
let drawIndexedIndirectParameters = new Uint32Array(5);
let drawIndexedIndirectParametersSigned = new Int32Array(drawIndexedIndirectParameters.buffer);
drawIndexedIndirectParameters[0] = indexCount;
drawIndexedIndirectParameters[1] = instanceCount;
drawIndexedIndirectParameters[2] = firstIndex;
// baseVertex is a signed value.
drawIndexedIndirectParametersSigned[3] = baseVertex;
drawIndexedIndirectParameters[4] = firstInstance;
```

The value corresponding to **firstInstance** must be 0, unless the ["indirect-first-instance" feature](#) is enabled. If the ["indirect-first-instance" feature](#) is not enabled and **firstInstance** is not zero the [drawIndexedIndirect\(\)](#) call will be treated as a no-op.

Called on: [GPURenderCommandsMixin](#) this.

Arguments:

Arguments for the [GPURenderCommandsMixin.drawIndexedIndirect\(indirectBuffer, indirectOffset\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>indirectBuffer</i>	GPUBuffer	✗	✗	Buffer containing the indirect drawIndexed parameters .
<i>indirectOffset</i>	GPUSize64	✗	✗	Offset in bytes into <i>indirectBuffer</i> where the drawing data begins.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. If any of the following conditions are unsatisfied, [invalidate](#) *this* and return.
 - It is [valid to draw indexed](#) with *this*.
 - *indirectBuffer* is [valid to use with](#) *this*.
 - *indirectBuffer*.[usage](#) contains [INDIRECT](#).
 - $\text{indirectOffset} + \text{sizeof}(\text{indirect drawIndexed parameters}) \leq \text{indirectBuffer.size}$.
 - *indirectOffset* is a multiple of 4.
3. [Add](#) *indirectBuffer* to [\[\[usage_scope\]\]](#) with usage [input](#).
4. Increment *this*.[\[\[drawCount\]\]](#) by 1.
5. Let *bindingState* be a snapshot of *this*'s current state.
6. [Enqueue a render command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) with *renderState* when executed.

[Queue timeline](#) steps:

1. Let *indexCount* be an unsigned 32-bit integer read from *indirectBuffer* at *indirectOffset* bytes.
2. Let *instanceCount* be an unsigned 32-bit integer read from *indirectBuffer* at (*indirectOffset* + 4) bytes.
3. Let *firstIndex* be an unsigned 32-bit integer read from *indirectBuffer* at (*indirectOffset* + 8) bytes.
4. Let *baseVertex* be a signed 32-bit integer read from *indirectBuffer* at (*indirectOffset* + 12) bytes.
5. Let *firstInstance* be an unsigned 32-bit integer read from *indirectBuffer* at (*indirectOffset* + 16) bytes.
6. Draw *instanceCount* instances, starting with instance *firstInstance*, of primitives consisting of *indexCount* indexed vertices, starting with index *firstIndex* from vertex *baseVertex*, with the states from *bindingState* and *renderState*.

To determine if it's *valid to draw* with [GPURenderCommandsMixin](#) encoder, run the following [device timeline](#) steps:

If any of the following conditions are unsatisfied, return **false**:

[Validate encoder bind groups](#)(*encoder*, *encoder*.[\[\[pipeline\]\]](#)) must be **true**.

Let *pipelineDescriptor* be *encoder*.[\[\[pipeline\]\].\[\[descriptor\]\]](#).

For each [GPUIndex32](#) slot 0 to *pipelineDescriptor*.[vertex.buffers.size](#):

If *pipelineDescriptor*.[vertex.buffers\[slot\]](#) is not **null**, *encoder*.[\[\[vertex_buffers\]\]](#) must [contain](#) *slot*.

Validate [maxBindGroupsPlusVertexBuffers](#):

Let *bindGroupSpaceUsed* be (the maximum key in *encoder*.[\[\[bind_groups\]\]](#)) + 1.

Let *vertexBufferSpaceUsed* be (the maximum key in *encoder*.[\[\[vertex_buffers\]\]](#)) + 1.

bindGroupSpaceUsed + *vertexBufferSpaceUsed* must be \leq *encoder*.[\[\[device\]\].\[\[limits\]\].maxBindGroupsPlusVertexBuffers](#).

Otherwise, return **true**.

To determine if it's *valid to draw indexed* with [GPURenderCommandsMixin](#) encoder, run the following [device timeline](#) steps:

If any of the following conditions are unsatisfied, return **false**:

It must be [valid to draw](#) with *encoder*.

encoder.[\[\[index_buffer\]\]](#) must not be **null**.

Let *topology* be *encoder*.[\[\[pipeline\]\].\[\[descriptor\]\].primitive.topology](#).

If *topology* is ["line-strip"](#) or ["triangle-strip"](#):

encoder.[\[\[index_format\]\]](#) must equal *encoder*.[\[\[pipeline\]\].\[\[descriptor\]\].primitive.stripIndexFormat](#).

Otherwise, return **true**.

17.2.2. Rasterization state

The [GPURenderPassEncoder](#) has several methods which affect how draw commands are rasterized to attachments used by this encoder.

setViewport(*x, y, width, height, minDepth, maxDepth*)

Sets the viewport used during the rasterization stage to linearly map from [normalized device coordinates](#) to [viewport coordinates](#).

Called on: [GPURenderPassEncoder](#) *this*.

Arguments:

Arguments for the [GPURenderPassEncoder.setViewport\(*x, y, width, height, minDepth, maxDepth*\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>x</i>	float	✗	✗	Minimum X value of the viewport in pixels.
<i>y</i>	float	✗	✗	Minimum Y value of the viewport in pixels.
<i>width</i>	float	✗	✗	Width of the viewport in pixels.
<i>height</i>	float	✗	✗	Height of the viewport in pixels.
<i>minDepth</i>	float	✗	✗	Minimum depth value of the viewport.
<i>maxDepth</i>	float	✗	✗	Maximum depth value of the viewport.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. Let *maxViewportRange* be *this*.[limits.maxTextureDimension2D](#) × 2.
3. If any of the following conditions are unsatisfied, [invalidate](#) *this* and return.
 - $x \geq -\text{maxViewportRange}$
 - $y \geq -\text{maxViewportRange}$
 - $0 \leq \text{width} \leq \text{this}.\text{limits.maxTextureDimension2D}$
 - $0 \leq \text{height} \leq \text{this}.\text{limits.maxTextureDimension2D}$
 - $x + \text{width} \leq \text{maxViewportRange} - 1$
 - $y + \text{height} \leq \text{maxViewportRange} - 1$
 - $0.0 \leq \text{minDepth} \leq 1.0$
 - $0.0 \leq \text{maxDepth} \leq 1.0$
 - $\text{minDepth} \leq \text{maxDepth}$
4. [Enqueue a render command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) with *renderState* when executed.

[Queue timeline](#) steps:

1. Round *x*, *y*, *width*, and *height* to some uniform precision, no less precise than integer rounding.
2. Set *renderState*.[\[\[viewport\]\]](#) to the extents *x*, *y*, *width*, *height*, *minDepth*, and *maxDepth*.

setScissorRect(*x, y, width, height*)

Sets the scissor rectangle used during the rasterization stage. After transformation into [viewport coordinates](#) any fragments which fall outside the scissor rectangle will be discarded.

Called on: [GPURenderPassEncoder](#) *this*.

Arguments:

Arguments for the [GPURenderPassEncoder.setScissorRect\(*x, y, width, height*\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>x</i>	GPUIntegerCoordinate	✗	✗	Minimum X value of the scissor rectangle in pixels.
<i>y</i>	GPUIntegerCoordinate	✗	✗	Minimum Y value of the scissor rectangle in pixels.
<i>width</i>	GPUIntegerCoordinate	✗	✗	Width of the scissor rectangle in pixels.
<i>height</i>	GPUIntegerCoordinate	✗	✗	Height of the scissor rectangle in pixels.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. If any of the following conditions are unsatisfied, [invalidate this](#) and return.
 - $x + \text{width} \leq \text{this}.\text{[[attachment_size]]}.\text{width}$.
 - $y + \text{height} \leq \text{this}.\text{[[attachment_size]]}.\text{height}$.
3. [Enqueue a render command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) with *renderState* when executed.

[Queue timeline](#) steps:

1. Set *renderState*.[\[\[scissorRect\]\]](#) to the extents *x*, *y*, *width*, and *height*.

setBlendConstant(color)

Sets the constant blend color and alpha values used with ["constant"](#) and ["one-minus-constant"](#) [GPUBlendFactor](#)s.

Called on: [GPURenderPassEncoder](#) *this*.

Arguments:

Arguments for the [GPURenderPassEncoder.setBlendConstant\(color\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>color</i>	GPUColor	✗	✗	The color to use when blending.

Returns: [undefined](#)

[Content timeline](#) steps:

1. [? validate GPUColor shape\(color\)](#).
2. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. [Enqueue a render command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) with *renderState* when executed.

[Queue timeline](#) steps:

1. Set *renderState*.[\[\[blendConstant\]\]](#) to *color*.

setStencilReference(reference)

Sets the [\[\[stencilReference\]\]](#) value used during stencil tests with the ["replace"](#) [GPUStencilOperation](#).

Called on: [GPURenderPassEncoder](#) *this*.

Arguments:

Arguments for the [GPURenderPassEncoder.setStencilReference\(reference\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>reference</i>	GPUStencilValue	✗	✗	The new stencil reference value.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. [Enqueue a render command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) with *renderState* when executed.

[Queue timeline](#) steps:

1. Set *renderState*.[\[\[stencilReference\]\]](#) to *reference*.

17.2.3. Queries

beginOcclusionQuery(queryIndex)

Called on: [GPURenderPassEncoder](#) *this*.

Arguments:

Arguments for the [GPURenderPassEncoder.beginOcclusionQuery\(queryIndex\)](#) method.

Parameter	Type	Nullable	Optional	Description
-----------	------	----------	----------	-------------

Parameter	Type	Nullable	Optional	Description
<i>queryIndex</i>	GPUSize32	✗	✗	The index of the query in the query set.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. If any of the following conditions are unsatisfied, [invalidate](#) *this* and return.
 - *this*.[\[\[occlusion_query_set\]\]](#) is not null.
 - *queryIndex* < *this*.[\[\[occlusion_query_set\]\].count](#).
 - The query at same *queryIndex* must not have been previously written to in this pass.
 - *this*.[\[\[occlusion_query_active\]\]](#) is false.
3. Set *this*.[\[\[occlusion_query_active\]\]](#) to true.
4. [Enqueue a render command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) with *renderState* when executed.

[Queue timeline](#) steps:

1. Set *renderState*.[\[\[occlusionQueryIndex\]\]](#) to *queryIndex*.

endOcclusionQuery()

Called on: [GPURenderPassEncoder](#) *this*.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.[\[\[device\]\]](#).

[Device timeline](#) steps:

1. [Validate the encoder state](#) of *this*. If it returns false, return.
2. If any of the following conditions are unsatisfied, [invalidate](#) *this* and return.
 - *this*.[\[\[occlusion_query_active\]\]](#) is true.
3. Set *this*.[\[\[occlusion_query_active\]\]](#) to false.
4. [Enqueue a render command](#) on *this* which issues the subsequent steps on the [Queue timeline](#) with *renderState* when executed.

[Queue timeline](#) steps:

1. Let *passingFragments* be non-zero if any fragment samples passed all per-fragment tests since the corresponding [beginOcclusionQuery\(\)](#) command was executed, and zero otherwise.

Note: If no draw calls occurred, *passingFragments* is zero.

2. Write *passingFragments* into *this*.[\[\[occlusion_query_set\]\]](#) at index *renderState*.[\[\[occlusionQueryIndex\]\]](#).

17.2.4. Bundles

executeBundles(bundles)

Executes the commands previously recorded into the given [GPURenderBundles](#) as part of this render pass.

When a [GPURenderBundle](#) is executed, it does not inherit the render pass's pipeline, bind groups, or vertex and index buffers. After a [GPURenderBundle](#) has executed, the render pass's pipeline, bind group, and vertex/index buffer state is cleared (to the initial, empty values).

Note: The state is cleared, not restored to the previous state. This occurs even if zero [GPURenderBundles](#) are executed.

Called on: [GPURenderPassEncoder](#) *this*.

Arguments:

Arguments for the [GPURenderPassEncoder.executeBundles\(bundles\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>bundles</i>	sequence < GPURenderBundle >	✗	✗	List of render bundles to execute.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of this.`[[device]]`.

[Device timeline](#) steps:

1. [Validate the encoder state](#) of this. If it returns false, return.

2. If any of the following conditions are unsatisfied, [invalidate](#) this and return.

- For each *bundle* in *bundles*:
- *bundle* must be [valid to use with](#) this.
- this.`[[layout]]` must equal *bundle*.`[[layout]]`.
- If this.`[[depthReadOnly]]` is true, *bundle*.`[[depthReadOnly]]` must be true.
- If this.`[[stencilReadOnly]]` is true, *bundle*.`[[stencilReadOnly]]` must be true.

3. For each *bundle* in *bundles*:

1. Increment this.`[[drawCount]]` by *bundle*.`[[drawCount]]`.

2. [Merge](#) *bundle*.`[[usage_scope]]` into this.`[[usage_scope]]`.

3. [Enqueue a render command](#) on this which issues the following steps on the [Queue timeline](#) with *renderState* when executed:

[Queue timeline](#) steps:

1. Execute each command in *bundle*.`[[command_list]]` with *renderState*.

Note: *renderState* cannot be changed by executing render bundles. Binding state was already captured at bundle encoding time, and so isn't used when executing bundles.

4. [Reset the render pass binding state](#) of this.

To [Reset the render pass binding state](#) of [GPURenderPassEncoder](#) *encoder* run the following [device timeline](#) steps:

[Clear](#) *encoder*.`[[bind_groups]]`.

Set *encoder*.`[[pipeline]]` to null.

Set *encoder*.`[[index_buffer]]` to null.

[Clear](#) *encoder*.`[[vertex_buffers]]`.

18. Bundles

A bundle is a partial, limited pass that is encoded once and can then be executed multiple times as part of future pass encoders without expiring after use like typical command buffers. This can reduce the overhead of encoding and submission of commands which are issued repeatedly without changing.

18.1. GPURenderBundle

`[[Exposed]]=(Window, Worker), SecureContext`

interface [GPURenderBundle](#) {

};

[GPURenderBundle](#) includes [GPUObjectBase](#);

`[[command_list]]`, of type `list<GPU command>`

A `list` of [GPU commands](#) to be submitted to the [GPURenderPassEncoder](#) when the [GPURenderBundle](#) is executed.

`[[usage_scope]]`, of type [usage_scope](#), initially empty

The [usage_scope](#) for this render bundle, stored for later merging into the [GPURenderPassEncoder](#)'s `[[usage_scope]]` in [executeBundles\(\)](#).

`[[layout]]`, of type [GPURenderPassLayout](#)

The layout of the render bundle.

`[[depthReadOnly]]`, of type [boolean](#)

If `true`, indicates that the depth component is not modified by executing this render bundle.

`[[stencilReadOnly]]`, of type [boolean](#)

If `true`, indicates that the stencil component is not modified by executing this render bundle.

`[[drawCount]]`, of type [GPUSize64](#)

The number of draw commands in this [GPURenderBundle](#).

18.1.1. Render Bundle Creation

dictionary

GPURenderBundleDescriptor

```
: GPUObjectDescriptorBase {  
};  
[Exposed=(Window, Worker), SecureContext]  
interface
```

GPURenderBundleEncoder

```
{  
  GPURenderBundle finish(optional GPURenderBundleDescriptor descriptor = {});  
};  
GPURenderBundleEncoder includes GPUObjectBase;  
GPURenderBundleEncoder includes GPUCommandsMixin;  
GPURenderBundleEncoder includes GPUDebugCommandsMixin;  
GPURenderBundleEncoder includes GPUBindingCommandsMixin;  
GPURenderBundleEncoder includes GPURenderCommandsMixin;
```

createRenderBundleEncoder(descriptor)

Creates a [GPURenderBundleEncoder](#).

Called on: [GPUDevice](#) *this*.

Arguments:

Arguments for the [GPUDevice.createRenderBundleEncoder\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPURenderBundleEncoderDescriptor	✗	✗	Description of the GPURenderBundleEncoder to create.

Returns: [GPURenderBundleEncoder](#)

[Content timeline](#) steps:

1. ? [Validate texture format required features](#) of each non-null element of *descriptor.colorFormats* with *this.[[device]]*.
2. If *descriptor.depthStencilFormat* is provided:
 1. ? [Validate texture format required features](#) of *descriptor.depthStencilFormat* with *this.[[device]]*.
3. Let *e* be ! [create a new WebGPU object](#)(*this*, [GPURenderBundleEncoder](#), *descriptor*).
4. Issue the *initialization steps* on the [Device timeline](#) of *this*.
5. Return *e*.

[Device timeline](#) initialization steps:

1. If any of the following conditions are unsatisfied [generate a validation error](#), [invalidate](#) *e* and return.
 - *this* must not be [lost](#).
 - *descriptor.colorFormats.size* must be ≤ *this.[[limits]].maxColorAttachments*.
 - For each non-null *colorFormat* in *descriptor.colorFormats*:
 - *colorFormat* must be a [color renderable format](#).
 - [Calculating color attachment bytes per sample](#)(*descriptor.colorFormats*) must be ≤ *this.[[limits]].maxColorAttachmentBytesPerSample*.
 - If *descriptor.depthStencilFormat* is provided:
 - *descriptor.depthStencilFormat* must be a [depth-or-stencil format](#).
 - There must exist at least one attachment, either:
 - A non-null value in *descriptor.colorFormats*, or
 - A *descriptor.depthStencilFormat*.
2. Set *e.[[layout]]* to a copy of *descriptor*'s included [GPURenderPassLayout](#) interface.
3. Set *e.[[depthReadOnly]]* to *descriptor.depthReadOnly*.
4. Set *e.[[stencilReadOnly]]* to *descriptor.stencilReadOnly*.
5. Set *e.[[state]]* to "open".
6. Set *e.[[drawCount]]* to 0.

18.1.2. Encoding

dictionary

GPURenderBundleEncoderDescriptor

```
: GPURenderPassLayout {
  boolean depthReadOnly = false;
  boolean stencilReadOnly = false;
};
```

depthReadOnly, of type [boolean](#), defaulting to `false`

If `true`, indicates that the render bundle does not modify the depth component of the [GPURenderPassDepthStencilAttachment](#) of any render pass the render bundle is executed in.

See [read-only depth-stencil](#).

stencilReadOnly, of type [boolean](#), defaulting to `false`

If `true`, indicates that the render bundle does not modify the stencil component of the [GPURenderPassDepthStencilAttachment](#) of any render pass the render bundle is executed in.

See [read-only depth-stencil](#).

18.1.3. Finalization

finish(descriptor)

Completes recording of the render bundle commands sequence.

Called on: [GPURenderBundleEncoder](#) this.

Arguments:

Arguments for the [GPURenderBundleEncoder.finish\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPURenderBundleDescriptor	✗	✓	

Returns: [GPURenderBundle](#)

[Content timeline](#) steps:

1. Let *renderBundle* be a new [GPURenderBundle](#).
2. Issue the *finish steps* on the [Device timeline](#) of *this.[[device]]*.
3. Return *renderBundle*.

[Device timeline](#) *finish steps*:

1. Let *validationSucceeded* be `true` if all of the following requirements are met, and `false` otherwise.
 - *this* must be [valid](#).
 - *this.[[usage_scope]]* must satisfy [usage scope validation](#).
 - *this.[[state]]* must be `"open"`.
 - *this.[[debug_group_stack]]* must [be empty](#).
2. Set *this.[[state]]* to `"ended"`.
3. If *validationSucceeded* is `false`, then:
 1. [Generate a validation error](#).
2. Return an [invalidated GPURenderBundle](#).
4. Set *renderBundle.[[command_list]]* to *this.[[commands]]*.
5. Set *renderBundle.[[usage_scope]]* to *this.[[usage_scope]]*.
6. Set *renderBundle.[[drawCount]]* to *this.[[drawCount]]*.

19. Queues

19.1. GPUQueueDescriptor

[GPUQueueDescriptor](#) describes a queue request.


```
dictionary GPUQueueDescriptor
: GPUObjectDescriptorBase {
};
```

19.2. GPUQueue

```
[Exposed=(Window, Worker), SecureContext]
interface GPUQueue {
  undefined submit(sequence<GPUCommandBuffer> commandBuffers);

  Promise<undefined> onSubmittedWorkDone();

  undefined writeBuffer(
    GPUBuffer buffer,
    GPUSize64 bufferOffset,
    AllowSharedBufferSource data,
    optional GPUSize64 dataOffset = 0,
    optional GPUSize64 size);

  undefined writeTexture(
    GPUTexelCopyTextureInfo destination,
    AllowSharedBufferSource data,
    GPUTexelCopyBufferLayout dataLayout,
    GPUExtent3D size);

  undefined copyExternalImageToTexture(
    GPUCopyExternalImageSourceInfo source,
    GPUCopyExternalImageDestInfo destination,
    GPUExtent3D copySize);
};
GPUQueue includes GPUObjectBase;
```

GPUQueue has the following methods:

writeBuffer(buffer, bufferOffset, data, dataOffset, size)

Issues a write operation of the provided data into a GPUBuffer.

Called on: GPUQueue this.

Arguments:

Arguments for the GPUQueue.writeBuffer(buffer, bufferOffset, data, dataOffset, size) method.

Parameter	Type	Nullable	Optional	Description
buffer	GPUBuffer	✗	✗	The buffer to write to.
bufferOffset	GPUSize64	✗	✗	Offset in bytes into buffer to begin writing at.
data	AllowSharedBufferSource	✗	✗	Data to write into buffer.
dataOffset	GPUSize64	✗	✓	Offset in into data to begin writing from. Given in elements if data is a TypedArray and bytes otherwise.
size	GPUSize64	✗	✓	Size of content to write from data to buffer. Given in elements if data is a TypedArray and bytes otherwise.

Returns: undefined

Content timeline steps:

- If data is an ArrayBuffer or DataView, let the element type be "byte". Otherwise, data is a TypedArray; let the element type be the type of the TypedArray.
- Let dataSize be the size of data, in elements.
- If size is missing, let contentsSize be dataSize – dataOffset. Otherwise, let contentsSize be size.
- If any of the following conditions are unsatisfied, throw an OperationError and return.
 - contentsSize ≥ 0.
 - dataOffset + contentsSize ≤ dataSize.
 - contentsSize, converted to bytes, is a multiple of 4 bytes.
- Let dataContents be a copy of the bytes held by the buffer source data.

6. Let *contents* be the *contentsSize* elements of *dataContents* starting at an offset of *dataOffset* elements.

7. Issue the subsequent steps on the [Device timeline](#) of *this*.

[Device timeline](#) steps:

1. If any of the following conditions are unsatisfied, [generate a validation error](#) and return.

- *buffer* is [valid to use with this](#).
- *buffer*.[\[\[internal state\]\]](#) is "[available](#)".
- *buffer*.[usage](#) includes [COPY_DST](#).
- *bufferOffset*, converted to bytes, is a multiple of 4 bytes.
- *bufferOffset* + *contentsSize*, converted to bytes, \leq *buffer*.[size](#) bytes.

2. Issue the subsequent steps on the [Queue timeline](#) of *this*.

[Queue timeline](#) steps:

1. Write *contents* into *buffer* starting at *bufferOffset*.

writeTexture(destination, data, dataLayout, size)

Issues a write operation of the provided data into a [GPUTexture](#).

Called on: [GPUQueue](#) *this*.

Arguments:

Arguments for the [GPUQueue.writeTexture\(destination, data, dataLayout, size\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>destination</i>	GPUTexelCopyTextureInfo	✗	✗	The texture subresource and origin to write to.
<i>data</i>	AllowSharedBufferSource	✗	✗	Data to write into <i>destination</i> .
<i>dataLayout</i>	GPUTexelCopyBufferLayout	✗	✗	Layout of the content in <i>data</i> .
<i>size</i>	GPUExtent3D	✗	✗	Extents of the content to write from <i>data</i> to <i>destination</i> .

Returns: [undefined](#)

[Content timeline](#) steps:

1. ? [validate GPUOrigin3D shape\(destination.origin\)](#).
2. ? [validate GPUExtent3D shape\(size\)](#).
3. Let *dataBytes* be [a copy of the bytes held by the buffer source data](#).

Note: This is described as copying all of *data* to the device timeline, but in practice *data* could be much larger than necessary. Implementations should optimize by copying only the necessary bytes.

4. Issue the subsequent steps on the [Device timeline](#) of *this*.

[Device timeline](#) steps:

1. Let *aligned* be **false**.
2. Let *dataLength* be *dataBytes*.[length](#).
3. If any of the following conditions are unsatisfied, [generate a validation error](#) and return.
 - *destination*.[texture](#).[\[\[destroyed\]\]](#) is **false**.
 - [validating texture buffer copy\(destination, dataLayout, dataLength, size, COPY_DST, aligned\)](#) returns **true**.

Note: unlike [GPUCommandEncoder.copyBufferToTexture\(\)](#), there is no alignment requirement on either *dataLayout*.[bytesPerRow](#) or *dataLayout*.[offset](#).

4. Issue the subsequent steps on the [Queue timeline](#) of *this*.

[Queue timeline](#) steps:

1. Let *blockWidth* be the [texel block width](#) of *destination*.[texture](#).
2. Let *blockHeight* be the [texel block height](#) of *destination*.[texture](#).
3. Let *dstOrigin* be *destination*.[origin](#);
4. Let *dstBlockOriginX* be (*dstOrigin*.[x](#) \div *blockWidth*).
5. Let *dstBlockOriginY* be (*dstOrigin*.[y](#) \div *blockHeight*).

6. Let *blockColumns* be (*copySize.width* ÷ *blockWidth*).
7. Let *blockRows* be (*copySize.height* ÷ *blockHeight*).
8. **Assert** that *dstBlockOriginX*, *dstBlockOriginY*, *blockColumns*, and *blockRows* are integers.
9. For each *z* in the range [0, *copySize.depthOrArrayLayers* − 1]:
 1. Let *dstSubregion* be [texture copy sub-region](#) (*z* + *dstOrigin.z*) of *destination*.
 2. For each *y* in the range [0, *blockRows* − 1]:
 1. For each *x* in the range [0, *blockColumns* − 1]:
 1. Let *blockOffset* be the [texel block byte offset](#) of *dataLayout* for (*x*, *y*, *z*) of *destination.texture*.
 2. Set [texel block](#) (*dstBlockOriginX* + *x*, *dstBlockOriginY* + *y*) of *dstSubregion* to be an [equivalent texel representation](#) to the [texel block](#) described by *dataBytes* at offset *blockOffset*.

copyExternalImageToTexture(source, destination, copySize)

Issues a copy operation of the contents of a platform image/canvas into the destination texture.

This operation performs [color encoding](#) into the destination encoding according to the parameters of [GPUCopyExternalImageDestInfo](#).

Copying into a -srgb texture results in the same texture bytes, not the same decoded values, as copying into the corresponding non-srgb format. Thus, after a copy operation, sampling the destination texture has different results depending on whether its format is -srgb, all else unchanged.

NOTE:

When copying from a "webgl"/"webgl2" context canvas, the [WebGL Drawing Buffer](#) may be not exist during certain points in the frame presentation cycle (after the image has been moved to the compositor for display). To avoid this, either:

- Issue [copyExternalImageToTexture\(\)](#) in the same [task](#) with WebGL rendering operation, to ensure the copy occurs before the WebGL canvas is presented.
- If not possible, set the `preserveDrawingBuffer` option in [WebGLContextAttributes](#) to `true`, so that the drawing buffer will still contain a copy of the frame contents after they've been presented. Note, this extra copy may have a performance cost.

Called on: [GPUQueue](#) *this*.

Arguments:

Arguments for the [GPUQueue.copyExternalImageToTexture\(source, destination, copySize\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>source</i>	GPUCopyExternalImageSourceInfo	✗	✗	source image and origin to copy to <i>destination</i> .
<i>destination</i>	GPUCopyExternalImageDestInfo	✗	✗	The texture subresource and origin to write to, and its encoding metadata.
<i>copySize</i>	GPUExtent3D	✗	✗	Extents of the content to write from <i>source</i> to <i>destination</i> .

Returns: [undefined](#)

[Content timeline](#) steps:

1. [? validate GPUOrigin2D shape](#)(*source.origin*).
2. [? validate GPUOrigin3D shape](#)(*destination.origin*).
3. [? validate GPUExtent3D shape](#)(*copySize*).
4. Let *sourceImage* be *source.source*
5. If *sourceImage* is not origin-clean, throw a [SecurityError](#) and return.
6. If any of the following requirements are unmet, throw an [OperationError](#) and return.
 - *source.origin.x* + *copySize.width* must be ≤ the width of *sourceImage*.
 - *source.origin.y* + *copySize.height* must be ≤ the height of *sourceImage*.
 - *copySize.depthOrArrayLayers* must be ≤ 1.
7. Let *usability* be [? check the usability of the image argument](#)(*source*).
8. Issue the subsequent steps on the [Device timeline](#) of *this*.

[Device timeline](#) steps:

1. Let *texture* be *destination.texture*.
2. If any of the following requirements are unmet, [generate a validation error](#) and return.
 - *usability* must be good.
 - *texture.[[destroyed]]* must be false.

- *texture* must be [valid to use with this](#).
- [validating GPU Texel Copy Texture Info](#)(*destination*, *copySize*) must return `true`.
- *texture.usage* must include both [RENDER_ATTACHMENT](#) and [COPY_DST](#).
- *texture.dimension* must be `"2d"`.
- *texture.sampleCount* must be 1.
- *texture.format* must be one of the following formats (which all support [RENDER_ATTACHMENT](#) usage):
 - ["r8unorm"](#)
 - ["r16float"](#)
 - ["r32float"](#)
 - ["rg8unorm"](#)
 - ["rg16float"](#)
 - ["rg32float"](#)
 - ["rgba8unorm"](#)
 - ["rgba8unorm-srgb"](#)
 - ["bgra8unorm"](#)
 - ["bgra8unorm-srgb"](#)
 - ["rgb10a2unorm"](#)
 - ["rgba16float"](#)
 - ["rgba32float"](#)

3. If *copySize.depthOrArrayLayers* is > 0, issue the subsequent steps on the [Queue timeline](#) of *this*.

[Queue timeline](#) steps:

1. [Assert](#) that the [texel block width](#) of *destination.texture* is 1, the [texel block height](#) of *destination.texture* is 1, and that *copySize.depthOrArrayLayers* is 1.
2. Let *srcOrigin* be *source.origin*.
3. Let *dstOrigin* be *destination.origin*.
4. Let *dstSubregion* be [texture copy sub-region](#) (*dstOrigin.z*) of *destination*.
5. For each *y* in the range [0, *copySize.height* – 1]:
 1. Let *srcY* be *y* if *source.flipY* is `false` and (*copySize.height* – 1 – *y*) otherwise.
 2. For each *x* in the range [0, *copySize.width* – 1]:
 1. Let *srcColor* be the [color-managed color value](#) of the pixel at (*srcOrigin.x* + *x*, *srcOrigin.y* + *srcY*) of *source.source*.
 2. Let *dstColor* be the numeric RGBA value resulting from applying any [color encoding](#) required by *destination.colorSpace* and *destination.premultipliedAlpha* to *srcColor*.
 3. If *texture.format* is an `-srgb` format:
 1. Set *dstColor* to the result of applying the sRGB non-linear-to-linear conversion to it.

Note: This cancels out the sRGB linear-to-non-linear conversion that occurs when writing an `-srgb` format in the next step, so that precision from an sRGB-like input image is not lost and the *linear* color values of the original image can be read from the texture (as is generally the purpose of using `-srgb` formats).
 4. Set [texel block](#) (*dstOrigin.x* + *x*, *dstOrigin.y* + *y*) of *dstSubregion* to an [equivalent texel representation](#) of *dstColor*.

submit(commandBuffers)

Schedules the execution of the command buffers by the GPU on this queue.

Submitted command buffers cannot be used again.

Called on: [GPUQueue](#) this.

Arguments:

Arguments for the [GPUQueue.submit\(commandBuffers\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>commandBuffers</i>	sequence<GPUCommandBuffer>	✗	✗	

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*:

[Device timeline](#) steps:

1. If any of the following requirements are unmet, [generate a validation error](#), [invalidate](#) each [GPUCommandBuffer](#) in *commandBuffers* and return.
 - Every [GPUCommandBuffer](#) in *commandBuffers* must be [valid to use with this](#).
 - Every [GPUCommandBuffer](#) in *commandBuffers* must be unique.
 - For each of the following types of resources used by any command in any element of *commandBuffers*:

[GPUBuffer](#) *b*

b.[\[\[internal_state\]\]](#) must be ["available"](#).

[GPUTexture](#) *t*

t.[\[\[destroyed\]\]](#) must be [false](#).

[GPUExternalTexture](#) *et*

et.[\[\[expired\]\]](#) must be [false](#).

[GPUQuerySet](#) *qs*

qs.[\[\[destroyed\]\]](#) must be [false](#).

Note: For occlusion queries, the [occlusionQuerySet](#) in [beginRenderPass\(\)](#) is not "used" unless it is also used by [beginOcclusionQuery\(\)](#).

2. For each *commandBuffer* in *commandBuffers*:

1. [Invalidate](#) *commandBuffer*.

3. Issue the subsequent steps on the [Queue timeline](#) of *this*:

[Queue timeline](#) steps:

1. For each *commandBuffer* in *commandBuffers*:

1. Execute each command in *commandBuffer*.[\[\[command_list\]\]](#).

onSubmittedWorkDone()

Returns a [Promise](#) that resolves once this queue finishes processing all the work submitted up to this moment.

Resolution of this [Promise](#) implies the completion of [mapAsync\(\)](#) calls made prior to that call, on [GPUBuffer](#)s last used exclusively on that queue.

Called on: [GPUQueue](#) *this*.

Returns: [Promise](#)<[undefined](#)>

[Content timeline](#) steps:

1. Let *contentTimeline* be the current [Content timeline](#).
2. Let *promise* be [a new promise](#).
3. Issue the [synchronization steps](#) on the [Device timeline](#) of *this*.
4. Return *promise*.

[Device timeline](#) [synchronization steps](#):

1. Let *event* occur upon the completion of all currently-enqueued operations.
2. [Listen for timeline event](#) *event* on *this*.[\[\[device\]\]](#), handled by the subsequent steps on *contentTimeline*.

20. Queries

20.1. GPUQuerySet

[[Exposed](#)=([Window](#), [Worker](#)), [SecureContext](#)]

```
interface GPUQuerySet {  
  undefined destroy();
```

```
  readonly attribute GPUQueryType type;
```

```
  readonly attribute GPUSize32Out count;
```

```
};
```

[GPUQuerySet](#) includes [GPUObjectBase](#);

[GPUQuerySet](#) has the following [immutable properties](#):

type, of type [GPUQueryType](#), readonly
The type of the queries managed by this [GPUQuerySet](#).

count, of type [GPUSize32Out](#), readonly
The number of queries managed by this [GPUQuerySet](#).

[GPUQuerySet](#) has the following [device timeline properties](#):

[[destroyed]], of type [boolean](#), initially `false`
If the query set is destroyed, it can no longer be used in any operation, and its underlying memory can be freed.

20.1.1.1. QuerySet Creation

A [GPUQuerySetDescriptor](#) specifies the options to use in creating a [GPUQuerySet](#).

dictionary

GPUQuerySetDescriptor

```
 : GPUObjectDescriptorBase {  
  required GPUQueryType type;  
  required GPUSize32 count;  
};
```

type, of type [GPUQueryType](#)
The type of queries managed by [GPUQuerySet](#).

count, of type [GPUSize32](#)
The number of queries managed by [GPUQuerySet](#).

createQuerySet(descriptor)

Creates a [GPUQuerySet](#).

Called on: [GPUDevice](#) this.

Arguments:

Arguments for the [GPUDevice.createQuerySet\(descriptor\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>descriptor</i>	GPUQuerySetDescriptor	✗	✗	Description of the GPUQuerySet to create.

Returns: [GPUQuerySet](#)

[Content timeline](#) steps:

- If *descriptor.type* is `"timestamp"`, but `"timestamp-query"` is not [enabled for this](#):
- Throw a [TypeError](#).
- Let *q* be [! create a new WebGPU object](#)(*this*, [GPUQuerySet](#), *descriptor*).
- Set *q.type* to *descriptor.type*.
- Set *q.count* to *descriptor.count*.
- Issue the *initialization steps* on the [Device timeline](#) of *this*.
- Return *q*.

[Device timeline](#) initialization steps:

- If any of the following requirements are unmet, [generate a validation error](#), [invalidate](#) *q* and return.
 - this* must not be [lost](#).
 - descriptor.count* must be ≤ 4096 .
- Create a device allocation for *q* where each entry in the query set is zero.

If the allocation fails without side-effects, [generate an out-of-memory error](#), [invalidate](#) *q*, and return.

Creating a [GPUQuerySet](#) which holds 32 occlusion query results.

```
const querySet = gpuDevice.createQuerySet({  
  type: 'occlusion',  
  count: 32  
});
```

20.1.2. Query Set Destruction

An application that no longer requires a [GPUQuerySet](#) can choose to lose access to it before garbage collection by calling [destroy\(\)](#).

[GPUQuerySet](#) has the following methods:

destroy()

Destroys the [GPUQuerySet](#).

Called on: [GPUQuerySet](#) *this*.

Returns: [undefined](#)

[Content timeline](#) steps:

1. Issue the subsequent steps on the [device timeline](#).

[Device timeline](#) steps:

1. Set *this*.[\[\[destroyed\]\]](#) to `true`.

20.2. QueryType

enum

[GPUQueryType](#)

{

["occlusion"](#)

,

["timestamp"](#)

,

};

20.3. Occlusion Query

Occlusion query is only available on render passes, to query the number of fragment samples that pass all the per-fragment tests for a set of drawing commands, including scissor, sample mask, alpha to coverage, stencil, and depth tests. Any non-zero result value for the query indicates that at least one sample passed the tests and reached the output merging stage of the render pipeline, 0 indicates that no samples passed the tests.

When beginning a render pass, [GPURenderPassDescriptor.occlusionQuerySet](#) must be set to be able to use occlusion queries during the pass. An occlusion query is begun and ended by calling [beginOcclusionQuery\(\)](#) and [endOcclusionQuery\(\)](#) in pairs that cannot be nested, and resolved into a [GPUBuffer](#) as a [64-bit unsigned integer](#) by [GPUCommandEncoder.resolveQuerySet\(\)](#).

20.4. Timestamp Query

Timestamp queries allow applications to write timestamps to a [GPUQuerySet](#), using:

[GPUComputePassDescriptor.timestampWrites](#)

[GPURenderPassDescriptor.timestampWrites](#)

and then resolve timestamp values (in nanoseconds as a [64-bit unsigned integer](#)) into a [GPUBuffer](#), using [GPUCommandEncoder.resolveQuerySet\(\)](#).

Timestamp values are [implementation-defined](#). Applications must handle arbitrary timestamp results, and should not be written in such a way that unexpected timestamps cause an application failure.

Note: The physical device may reset the timestamp counter occasionally, which can result in unexpected values such as negative deltas from one timestamp to the next. These instances should be rare, and these data points can safely be discarded.

Timestamp queries are implemented using high-resolution timers (see [§ 2.1.7.2 Device/queue-timeline timing](#)). To mitigate security and privacy concerns, their precision must be reduced:

To get the *current queue timestamp*, run the following [queue timeline](#) steps:

Let *fineTimestamp* be the current timestamp value of the current [queue timeline](#), in nanoseconds, relative to an [implementation-defined](#) point in the past.

Return the result of calling [coarsen time](#) on *fineTimestamp* with `crossOriginIsolatedCapability` set to `false`.

Note: Cross-origin isolation never applies to the [device timeline](#) or [queue timeline](#), so `crossOriginIsolatedCapability` is never set to `true`.

Validate timestampWrites(device, timestampWrites)

Arguments:

[GPUDevice](#) *device*

([GPUComputePassTimestampWrites](#) or [GPURenderPassTimestampWrites](#)) *timestampWrites*

[Device timeline](#) steps:

Return `true` if the following requirements are met, and `false` if not:

`"timestamp-query"` must be [enabled for](#) *device*.

timestampWrites.`querySet` must be [valid to use with](#) *device*.

timestampWrites.`querySet.type` must be `"timestamp"`.

Of the write index members in *timestampWrites* (`beginningOfPassWriteIndex`, `endOfPassWriteIndex`):

At least one must be [provided](#).

Of those which are [provided](#):

No two may be equal.

Each must be $< \textit{timestampWrites.querySet.count}$.

21. Canvas Rendering

21.1. [HTMLCanvasElement.getContext\(\)](#)

A [GPUCanvasContext](#) object is [created](#) via the [getContext\(\)](#) method of an [HTMLCanvasElement](#) instance by passing the string literal `'webgpu'` as its `contextType` argument.

Get a [GPUCanvasContext](#) from an offscreen [HTMLCanvasElement](#):

```
const canvas = document.createElement('canvas');
const context = canvas.getContext('webgpu');
```

Unlike WebGL or 2D context creation, the second argument of [HTMLCanvasElement.getContext\(\)](#) or [OffscreenCanvas.getContext\(\)](#), the context creation attribute dictionary `options`, is ignored. Instead, use [GPUCanvasContext.configure\(\)](#), which allows changing the canvas configuration without replacing the canvas.

To create a 'webgpu' context on a canvas ([HTMLCanvasElement](#) or [OffscreenCanvas](#)) canvas, run the following [content timeline](#) steps:

Let *context* be a new [GPUCanvasContext](#).

Set *context*.`canvas` to *canvas*.

[Replace the drawing buffer](#) of *context*.

Return *context*.

Note: User agents should consider issuing developer-visible warnings when an ignored `options` argument is provided when calling `getContext()` to get a WebGPU canvas context.

21.2. GPUCanvasContext

[[Exposed](#)=(Window, Worker), [SecureContext](#)]

interface

[GPUCanvasContext](#)

```
{
  readonly attribute (HTMLCanvasElement or OffscreenCanvas) canvas;
```

```
  undefined configure(GPUCanvasConfiguration configuration);
  undefined unconfigure();
```

```
  GPUCanvasConfiguration? getConfiguration();
  GPUTexture getCurrentTexture();
```

```
};
```

[GPUCanvasContext](#) has the following [content timeline properties](#):

canvas, of type ([HTMLCanvasElement](#) or [OffscreenCanvas](#)), readonly
The canvas this context was created from.

`[[configuration]]`, of type [GPUCanvasConfiguration](#)?, initially `null`

The options this context is currently configured with.

`null` if the context has not been configured or has been [unconfigured](#).

`[[textureDescriptor]]`, of type [GPUTextureDescriptor?](#), initially `null`

The currently configured texture descriptor, derived from the `[[configuration]]` and canvas.

`null` if the context has not been configured or has been [unconfigured](#).

`[[drawingBuffer]]`, an image, initially a transparent black image with the same size as the canvas

The drawing buffer is the working-copy image data of the canvas. It is exposed as writable by `[[currentTexture]]` (returned by `getCurrentTexture()`).

The drawing buffer is used to [get a copy of the image contents of a context](#), which occurs when the canvas is displayed or otherwise read. It may be transparent, even if `[[configuration]].alphaMode` is `"opaque"`. The `alphaMode` only affects the result of the ["get a copy of the image contents of a context"](#) algorithm.

The drawing buffer outlives the `[[currentTexture]]` and contains the previously-rendered contents even after the canvas has been presented. It is only cleared in [Replace the drawing buffer](#).

Any time the drawing buffer is read, implementations must ensure that all previously submitted work (e.g. queue submissions) have completed writing to it via `[[currentTexture]]`.

`[[currentTexture]]`, of type [GPUTexture?](#), initially `null`

The [GPUTexture](#) to draw into for the current frame. It exposes a writable view onto the underlying `[[drawingBuffer]]`. `getCurrentTexture()` populates this slot if `null`, then returns it.

In the steady-state of a visible canvas, any changes to the drawing buffer made through the `currentTexture` get presented when [updating the rendering of a WebGPU canvas](#). At or before that point, the texture is also destroyed and `[[currentTexture]]` is set to `null`, signalling that a new one is to be created by the next call to `getCurrentTexture()`.

[Destroying](#) the `currentTexture` has no effect on the drawing buffer contents; it only terminates write-access to the drawing buffer early. During the same frame, `getCurrentTexture()` continues returning the same destroyed texture.

[Expire the current texture](#) sets the `currentTexture` to `null`. It is called by `configure()`, resizing the canvas, presentation, `transferToImageBitmap()`, and others.

`[[lastPresentedImage]]`, of type `(readonly image)?`, initially `null`

The image most recently presented for this canvas in ["updating the rendering of a WebGPU canvas"](#). If the device is lost or destroyed, this image **may** be used as a fallback in ["get a copy of the image contents of a context"](#) in order to prevent the canvas from going blank.

Note: This property only needs to exist in implementations which implement the fallback, which is optional.

[GPUCanvasContext](#) has the following methods:

`configure(configuration)`

Configures the context for this canvas. This clears the drawing buffer to transparent black (in [Replace the drawing buffer](#)).

See `getConfiguration()` for information on [feature detection](#).

Called on: [GPUCanvasContext](#) *this*.

Arguments:

Arguments for the [GPUCanvasContext.configure\(configuration\)](#) method.

Parameter	Type	Nullable	Optional	Description
<code>configuration</code>	GPUCanvasConfiguration	✗	✗	Desired configuration for the context.

Returns: undefined

[Content timeline](#) steps:

1. Let *device* be `configuration.device`.
2. ? [Validate texture format required features](#) of `configuration.format` with `device.[[device]]`.
3. ? [Validate texture format required features](#) of each element of `configuration.viewFormats` with `device.[[device]]`.
4. If [Supported context formats](#) does not [contain](#) `configuration.format`, throw a [TypeError](#).
5. Let *descriptor* be the [GPUTextureDescriptor for the canvas and configuration](#)(`this.canvas`, `configuration`).
6. Set `this.[[configuration]]` to `configuration`.

Note: This exposes only the members defined in an implementation's definition of [GPUCanvasConfiguration](#). See the specifications of those members for notes about [feature detection](#).

7. Set `this.[[textureDescriptor]]` to `descriptor`.
8. [Replace the drawing buffer](#) of *this*.

9. Issue the subsequent steps on the [Device timeline](#) of device.

[Device timeline](#) steps:

1. If any of the following requirements are unmet, [generate a validation error](#) and return.
 - [validating GPUTextureDescriptor](#)(device, descriptor) must return true.

Note: This early validation remains valid until the next [configure\(\)](#) call, **except** for validation of the [size](#), which changes when the canvas is resized.

unconfigure()

Removes the context configuration. Destroys any textures produced while configured.

Called on: [GPUCanvasContext](#) *this*.

Returns: undefined

[Content timeline](#) steps:

1. Set *this*.[\[\[configuration\]\]](#) to null.
2. Set *this*.[\[\[textureDescriptor\]\]](#) to null.
3. [Replace the drawing buffer](#) of *this*.

getConfiguration()

Returns the context configuration, or null if the context is not configured.

Note: This method exists primarily for [feature detection](#) of members (and sub-members) of [GPUCanvasConfiguration](#); see those members for details. For supported members, it returns the originally-supplied values.

Called on: [GPUCanvasContext](#) *this*.

Returns: [GPUCanvasConfiguration](#) or null

[Content timeline](#) steps:

1. Let *configuration* be a copy of *this*.[\[\[configuration\]\]](#).
2. Return *configuration*.

getCurrentTexture()

Get the [GPUTexture](#) that will be composited to the document by the [GPUCanvasContext](#) next.

NOTE:

An application **should** call [getCurrentTexture\(\)](#) in the same task that renders to the canvas texture. Otherwise, the texture could get destroyed by these steps before the application is finished rendering to it.

The expiry task (defined below) is optional to implement. Even if implemented, task source priority is not normatively defined, so may happen as early as the next task, or as late as after all other task sources are empty (see [automatic expiry task source](#)). Expiry is only guaranteed when a visible canvas is displayed ([updating the rendering of a WebGPU canvas](#)) and in other callers of "Expire the current texture".

Called on: [GPUCanvasContext](#) *this*.

Returns: [GPUTexture](#)

[Content timeline](#) steps:

1. If *this*.[\[\[configuration\]\]](#) is null, throw an [InvalidStateError](#) and return.
2. [Assert](#) *this*.[\[\[textureDescriptor\]\]](#) is not null.
3. Let *device* be *this*.[\[\[configuration\]\].device](#).
4. If *this*.[\[\[currentTexture\]\]](#) is null:
 1. [Replace the drawing buffer](#) of *this*.
2. Set *this*.[\[\[currentTexture\]\]](#) to the result of calling *device*.[createTexture\(\)](#) with *this*.[\[\[textureDescriptor\]\]](#), except with the [GPUTexture](#)'s underlying storage pointing to *this*.[\[\[drawingBuffer\]\]](#).

Note: If the texture can't be created (e.g. due to validation failure or out-of-memory), this generates an error and returns an [invalidated GPUTexture](#). Some validation here is redundant with that done in [configure\(\)](#). Implementations **must not** skip this redundant validation.

5. **Optionally**, [queue an automatic expiry task](#) with device *device* and the following steps:

1. [Expire the current texture](#) of *this*.

Note: If this already happened when [updating the rendering of a WebGPU canvas](#), it has no effect.

6. Return *this*.[\[\[currentTexture\]\]](#).

Note: The same [GPUTexture](#) object will be returned by every call to [getCurrentTexture\(\)](#) until "[Expire the current texture](#)" runs, even if that [GPUTexture](#) is destroyed, failed validation, or failed to allocate.

To get a copy of the image contents of a context:

Arguments:

context: the [GPUCanvasContext](#)

Returns: image contents

[Content timeline](#) steps:

Let *snapshot* be a transparent black image of the same size as *context.canvas*.

Let *configuration* be *context.[]configuration[]*.

If *configuration* is `null`:

Return *snapshot*.

Note: The configuration will be `null` if the context has not been configured or has been [unconfigured](#). This is identical to the behavior when the canvas has no context.

Ensure that all submitted work items (e.g. queue submissions) have completed writing to the image (via *context.[]currentTexture[]*).

If *configuration.device* is found to be [valid](#):

Set *snapshot* to a copy of the *context.[]drawingBuffer[]*.

Otherwise, if *context.[]lastPresentedImage[]* is not `null`:

Optionally, set *snapshot* to a copy of *context.[]lastPresentedImage[]*.

Note: This is optional because the *[]lastPresentedImage[]* may no longer exist, depending on what caused device loss. Implementations may choose to skip it even if do they still have access to that image.

Let *alphaMode* be *configuration.alphaMode*.

If *alphaMode* is "[opaque](#)":

Clear the alpha channel of *snapshot* to 1.0.

Note: If the *[]currentTexture[]*, if any, has been destroyed (for example in "[Expire the current texture](#)"), the alpha channel is unobservable, and implementations may clear the alpha channel in-place.

Tag *snapshot* as being opaque.

Otherwise:

Tag *snapshot* with *alphaMode*.

Tag *snapshot* with the [colorSpace](#) and [toneMapping](#) of *configuration*.

Return *snapshot*.

To *Replace the drawing buffer* of a [GPUCanvasContext](#) *context*, run the following [content timeline](#) steps:

[Expire the current texture](#) of *context*.

Let *configuration* be *context.[]configuration[]*.

Set *context.[]drawingBuffer[]* to a transparent black image of the same size as *context.canvas*.

If *configuration* is null, the drawing buffer is tagged with the color space "[srgb](#)". In this case, the drawing buffer will remain blank until the context is configured.

If not, the drawing buffer has the specified *configuration.format* and is tagged with the specified *configuration.colorSpace* and *configuration.toneMapping*.

Note: *configuration.alphaMode* is ignored until "[get a copy of the image contents of a context](#)".

NOTE:

A newly replaced drawing buffer image behaves as if it is cleared to transparent black, but, like after "[discard](#)", an implementation can clear it lazily only if it becomes necessary.

Note: This will often be a no-op, if the drawing buffer is already cleared and has the correct configuration.

To *Expire the current texture* of a [GPUCanvasContext](#) *context*, run the following [content timeline](#) steps:

If *context.[]currentTexture[]* is not `null`:

Call *context.[]currentTexture[]*.[destroy\(\)](#) (without destroying *context.[]drawingBuffer[]*) to terminate write access to the image.

Set *context.[]currentTexture[]* to `null`.

21.3. HTML Specification Hooks

The following algorithms "hook" into algorithms in the HTML specification, and must run at the specified points.

When the "bitmap" is read from an [HTMLCanvasElement](#) or [OffscreenCanvas](#) with a [GPUCanvasContext](#) *context*, run the following [content timeline](#) steps:

Return [a copy of the image contents](#) of *context*.

NOTE:

This occurs in many places, including:

When an [HTMLCanvasElement](#) has its rendering updated.

Including when the canvas is the [placeholder canvas element](#) of an [OffscreenCanvas](#).

When [transferToImageBitmap\(\)](#) creates an [ImageBitmap](#) from the bitmap. (See also [transferToImageBitmap from WebGPU](#).)

When WebGPU canvas contents are read using other Web APIs, like [drawImage\(\)](#), [texImage2D\(\)](#), [texSubImage2D\(\)](#), [toDataURL\(\)](#), [toBlob\(\)](#), and so on.

If [alphaMode](#) is ["opaque"](#), this incurs a clear of the alpha channel. Implementations may skip this step when they are able to read or display images in a way that ignores the alpha channel.

If an application needs a canvas only for interop (not presentation), avoid ["opaque"](#) if it is not needed.

transferToImageBitmap from WebGPU:

When [transferToImageBitmap\(\)](#) is called on a canvas with [GPUCanvasContext](#) *context*, after creating an [ImageBitmap](#) from the canvas's bitmap, run the following [content timeline](#) steps:

[Replace the drawing buffer](#) of *context*.

Note: This makes [transferToImageBitmap\(\)](#) equivalent to "moving" (and possibly alpha-clearing) the image contents into the ImageBitmap, without a copy.

The [update the canvas size](#) algorithm.

21.4. GPUCanvasConfiguration

The *supported context formats* are the [set](#) of [GPUTextureFormats](#): «["bgra8unorm"](#), ["rgba8unorm"](#), ["rgba16float"](#)».. These formats must be supported when specified as a [GPUCanvasConfiguration.format](#) regardless of the given [GPUCanvasConfiguration.device](#).

Note: Canvas configuration cannot use [srgb](#) formats like ["bgra8unorm-srgb"](#). Instead, use the non-[srgb](#) equivalent (["bgra8unorm"](#)), specify the [srgb](#) format in the [viewFormats](#), and use [createView\(\)](#) to create a view with an [srgb](#) format.

```
enum GPUCanvasAlphaMode {
```

```
    "opaque",  
    "premultiplied",
```

```
};
```

```
enum GPUCanvasToneMappingMode {
```

```
    "standard",  
    "extended",
```

```
};
```

dictionary

```
GPUCanvasToneMapping
```

```
{
```

```
    GPUCanvasToneMappingMode
```

```
mode
```

```
    = "standard";
```

```
};
```

dictionary

```
GPUCanvasConfiguration
```

```
{
```

```
    required GPUDevice device;
```

```
    required GPUTextureFormat format;
```

```
    GPUTextureUsageFlags usage = 0x10; // GPUTextureUsage.RENDER_ATTACHMENT
```

```
    sequence<GPUTextureFormat> viewFormats = [];
```

```
    PredefinedColorSpace colorSpace = "srgb";
```

```
    GPUCanvasToneMapping toneMapping = {};
```

```
GPUCanvasAlphaMode alphaMode = "opaque";  
};
```

[GPUCanvasConfiguration](#) has the following members:

device, of type [GPUDevice](#)

The [GPUDevice](#) that textures returned by [getCurrentTexture\(\)](#) will be compatible with.

format, of type [GPUTextureFormat](#)

The format that textures returned by [getCurrentTexture\(\)](#) will have. Must be one of the [Supported context formats](#).

usage, of type [GPUTextureUsageFlags](#), defaulting to 0x10

The usage that textures returned by [getCurrentTexture\(\)](#) will have. [RENDER_ATTACHMENT](#) is the default, but is not automatically included if the usage is explicitly set. Be sure to include [RENDER_ATTACHMENT](#) when setting a custom usage if you wish to use textures returned by [getCurrentTexture\(\)](#) as color targets for a render pass.

viewFormats, of type sequence<[GPUTextureFormat](#)>, defaulting to []

The formats that views created from textures returned by [getCurrentTexture\(\)](#) may use.

colorSpace, of type [PredefinedColorSpace](#), defaulting to "srgb"

The color space that values written into textures returned by [getCurrentTexture\(\)](#) should be displayed with.

toneMapping, of type [GPUCanvasToneMapping](#), defaulting to { }

The tone mapping determines how the content of textures returned by [getCurrentTexture\(\)](#) are to be displayed.

NOTE:

This is a required feature, but user agents might not yet implement it, effectively supporting only the default [GPUCanvasToneMapping](#). In such implementations, this member **should not** exist in its implementation of [GPUCanvasConfiguration](#), to make [feature detection](#) possible using [getConfiguration\(\)](#).

This is especially important in implementations which otherwise have HDR capabilities (where a [dynamic-range](#) of [high](#) would be exposed).

If an implementation exposes this member and a [high](#) dynamic range, it **should** render the canvas as an HDR element, not clamp values to the SDR range of the HDR display.

alphaMode, of type [GPUCanvasAlphaMode](#), defaulting to "opaque"

Determines the effect that alpha values will have on the content of textures returned by [getCurrentTexture\(\)](#) when read, displayed, or used as an image source.

Configure a [GPUCanvasContext](#) to be used with a specific [GPUDevice](#), using the preferred format for this context:

```
const canvas = document.createElement('canvas');  
const context = canvas.getContext('webgpu');  
  
context.configure({  
  device: gpuDevice,  
  format: navigator.gpu.getPreferredCanvasFormat(),  
});
```

21.4.1. Canvas Color Space

During presentation, the color values in the canvas are converted to the color space of the screen.

The [toneMapping](#) determines the handling of values outside of the [0, 1] interval in the color space of the screen.

21.4.2. Canvas Context sizing

All canvas configuration is set in [configure\(\)](#) except for the resolution of the canvas, which is set by the canvas's [width](#) and [height](#).

Note: Like WebGL and 2d canvas, resizing a WebGPU canvas loses the current contents of the drawing buffer. In WebGPU, it does so by [replacing the drawing buffer](#).

When an [HTMLCanvasElement](#) or [OffscreenCanvas](#) *canvas* with a [GPUCanvasContext](#) *context* has its [width](#) or [height](#) attributes set, *update the canvas size* by running the following [content timeline](#) steps:

[Replace the drawing buffer](#) of *context*.

Let *configuration* be *context*.[\[\[configuration\]\]](#).

If *configuration* is not null:

Set *context*.[\[\[textureDescriptor\]\]](#) to the [GPUTextureDescriptor for the canvas and configuration](#)(*canvas*, *configuration*).

Note: This may result in a [GPUTextureDescriptor](#) which exceeds the [maxTextureDimension2D](#) of the device. In this case, validation will fail inside [getCurrentTexture\(\)](#).

Note: This algorithm is run any time the *canvas width* or *height* attributes are set, even if their value is not changed.

21.5. GPUCanvasToneMappingMode

This enum specifies how color values are displayed to the screen.

"standard"

Color values within the standard dynamic range of the screen are unchanged, and all other color values are projected to the standard dynamic range of the screen.

Note: This projection is often accomplished by clamping color values in the color space of the screen to the $[0, 1]$ interval.

For example, suppose that the value $(1.035, -0.175, -0.140)$ is written to an 'srgb' canvas.

If this is presented to an sRGB screen, then this will be converted to sRGB (which is a no-op, because the canvas is sRGB), then projected into the display's space. Using component-wise clamping, this results in the sRGB value $(1.0, 0.0, 0.0)$.

If this is presented to a Display P3 screen, then this will be converted to the value $(0.948, 0.106, 0.01)$ in the Display P3 color space, and no clamping will be needed.

"extended"

Color values in the extended dynamic range of the screen are unchanged, and all other color values are projected to the extended dynamic range of the screen.

Note: This projection is often accomplished by clamping color values in the color space of the screen to the interval of values that the screen is capable of displaying, which may include values greater than 1.

For example, suppose that the value $(2.5, -0.15, -0.15)$ is written to an 'srgb' canvas.

If this is presented to an sRGB screen that is capable of displaying values in the $[0, 4]$ interval in sRGB space, then this will be converted to sRGB (which is a no-op, because the canvas is sRGB), then projected into the display's space. If using component-wise clamping, this results in the sRGB value $(2.5, 0.0, 0.0)$.

If this is presented to a Display P3 screen that is capable of displaying values in the $[0, 2]$ interval in Display P3 space, then this will be converted to the value $(2.3, 0.545, 0.386)$ in the Display P3 color space, then projected into the display's space. If using component-wise clamping, this results in the Display P3 value $(2.0, 0.545, 0.386)$.

21.6. GPUCanvasAlphaMode

This enum selects how the contents of the canvas will be interpreted when read, when [displayed to the screen or used as an image source](#) (in `drawImage`, `toDataURL`, etc.)

Below, `src` is a value in the canvas texture, and `dst` is an image that the canvas is being composited into (e.g. an HTML page rendering, or a 2D canvas).

"opaque"

Read RGB as opaque and ignore alpha values. If the content is not already opaque, the alpha channel is cleared to 1.0 in "[get a copy of the image contents of a context](#)".

"premultiplied"

Read RGBA as premultiplied: color values are premultiplied by their alpha value. 100% red at 50% alpha is $[0.5, 0, 0, 0.5]$.

If the canvas texture contains [out-of-gamut premultiplied RGBA values](#) at the time the canvas contents are read, the behavior depends on whether the canvas is: [used as an image source](#)

Values are preserved, as described in [color space conversion](#).

displayed to the screen

Compositing results are undefined.

Note: This is true even if color space conversion would produce in-gamut values before compositing, because the intermediate format for compositing is not specified.

22. Errors & Debugging

During the normal course of operation of WebGPU, errors are raised via [dispatch error](#).

After a device is [lost](#), errors are no longer surfaced, where possible. After this point, implementations do not need to run validation or error tracking:

The validity of objects on the device becomes unobservable.

[popErrorScope\(\)](#) and [uncapturederror](#) stop reporting errors. (No errors are generated by the device loss itself. Instead, the [GPUDevice.lost](#) promise resolves to indicate the device is lost.)

All operations which send a message back to the [content timeline](#) will skip their usual steps. Most will appear to succeed, except for [mapAsync\(\)](#), which produces an error because it is impossible to provide the correct mapped data after the device has been lost.

This makes it unobservable whether other types of operations (that don't send messages back) actually execute or not.

22.1. Fatal Errors

enum

```
GPUDeviceLostReason
```

```
{
```

```
"unknown"
```

```
,
```

```
"destroyed"
```

```
,
```

```
};
```

[Exposed=(Window, Worker), SecureContext]

interface

```
GPUDeviceLostInfo
```

```
{
```

```
  readonly attribute GPUDeviceLostReason
```

```
  reason
```

```
;
```

```
  readonly attribute DOMString
```

```
  message
```

```
;
```

```
};
```

partial interface GPUDevice {

```
  readonly attribute Promise<GPUDeviceLostInfo> lost;
```

```
};
```

[GPUDevice](#) has the following additional attributes:

lost, of type Promise<GPUDeviceLostInfo>, readonly

A [slot-backed attribute](#) holding a promise which is created with the device, remains pending for the lifetime of the device, then resolves when the device is lost.

Upon initialization, it is set to [a new promise](#).

22.2. GPUError

[Exposed=(Window, Worker), SecureContext]

interface GPUError {

```
  readonly attribute DOMString message;
```

```
};
```

[GPUError](#) is the base interface for all errors surfaced from [popErrorScope\(\)](#) and the [uncapturederror](#) event.

Errors must only be generated for operations that explicitly state the conditions one may be generated under in their respective algorithms, and the subtype of error that is generated.

No errors are generated from a device which is lost. See [§ 22 Errors & Debugging](#).

Note: [GPUError](#) may gain new subtypes in future versions of this spec. Applications should handle this possibility, using only the error's [message](#) when possible, and specializing using `instanceof`. Use `error.constructor.name` when it's necessary to serialize an error (e.g. into JSON, for a debug report).

[GPUError](#) has the following [immutable properties](#):

message, of type [DOMString](#), readonly

A human-readable, [localizable text](#) message providing information about the error that occurred.

Note: This message is generally intended for application developers to debug their applications and capture information for debug reports, not to be surfaced to end-users.

Note: User agents should not include potentially machine-parsable details in this message, such as free system memory on ["out-of-memory"](#) or other details about the conditions under which memory was exhausted.

Note: The [message](#) should follow the [best practices for language and direction information](#). This includes making use of any future standards which may emerge regarding the reporting of string language and direction metadata.

Editorial note: At the time of this writing, no language/direction recommendation is available that provides compatibility and consistency with legacy APIs, but when

there is, adopt it formally.

[Exposed=(Window, Worker), SecureContext]

interface

GPUValidationError

: GPUError {

constructor

(DOMString

message

);

};

[GPUValidationError](#) is a subtype of [GPUError](#) which indicates that an operation did not satisfy all validation requirements. Validation errors are always indicative of an application error, and is expected to fail the same way across all devices assuming the same [\[\[features\]\]](#) and [\[\[limits\]\]](#) are in use.

To generate a validation error for [GPUDevice](#) device, run the following steps:

[Device timeline](#) steps:

Let *error* be a new [GPUValidationError](#) with an appropriate error message.

[Dispatch error](#) *error* to *device*.

[Exposed=(Window, Worker), SecureContext]

interface

GPUOutOfMemoryError

: GPUError {

constructor

(DOMString

message

);

};

[GPUOutOfMemoryError](#) is a subtype of [GPUError](#) which indicates that there was not enough free memory to complete the requested operation. The operation may succeed if attempted again with a lower memory requirement (like using smaller texture dimensions), or if memory used by other resources is released first.

To generate an out-of-memory error for [GPUDevice](#) device, run the following steps:

[Device timeline](#) steps:

Let *error* be a new [GPUOutOfMemoryError](#) with an appropriate error message.

[Dispatch error](#) *error* to *device*.

[Exposed=(Window, Worker), SecureContext]

interface

GPUInternalError

: GPUError {

constructor

(DOMString

message

);

};

[GPUInternalError](#) is a subtype of [GPUError](#) which indicates than an operation failed for a system or implementation-specific reason even when all validation requirements have been satisfied. For example, the operation may exceed the capabilities of the implementation in a way not easily captured by the [supported limits](#). The same operation may succeed on other devices or under difference circumstances.

To generate an internal error for [GPUDevice](#) device, run the following steps:

[Device timeline](#) steps:

Let *error* be a new [GPUInternalError](#) with an appropriate error message.

[Dispatch error](#) *error* to *device*.

22.3. Error Scopes

A *GPU error scope* captures [GPUError](#)s that were generated while the [GPU error scope](#) was current. Error scopes are used to isolate errors that occur within a set of WebGPU calls, typically for debugging purposes or to make an operation more fault tolerant.

[GPU error scope](#) has the following [device timeline properties](#):

[[errors]], of type [list](#)<[GPUError](#)>, initially []

The [GPUError](#)s, if any, observed while the [GPU error scope](#) was current.

[[filter]], of type [GPUErrorFilter](#)

Determines what type of [GPUError](#) this [GPU error scope](#) observes.

enum

[GPUErrorFilter](#)

```
{  
  "validation",  
  "out-of-memory",  
  "internal",  
};
```

```
partial interface GPUDevice {  
  undefined pushErrorScope(GPUErrorFilter filter);  
  Promise<GPUError?> popErrorScope();  
};
```

[GPUErrorFilter](#) defines the type of errors that should be caught when calling [pushErrorScope\(\)](#):

"*validation*"

Indicates that the error scope will catch a [GPUValidationError](#).

"*out-of-memory*"

Indicates that the error scope will catch a [GPUOutOfMemoryError](#).

"*internal*"

Indicates that the error scope will catch a [GPUInternalError](#).

[GPUDevice](#) has the following [device timeline properties](#):

[[errorScopeStack]], of type [stack](#)<[GPU error scope](#)>

A [stack](#) of [GPU error scopes](#) that have been pushed to the [GPUDevice](#).

The *current error scope* for a [GPUError](#) error and [GPUDevice](#) device is determined by issuing the following steps to the [device timeline](#) of *device*:

[Device timeline](#) steps:

If *error* is an instance of:

[GPUValidationError](#)

Let *type* be "validation".

[GPUOutOfMemoryError](#)

Let *type* be "out-of-memory".

[GPUInternalError](#)

Let *type* be "internal".

Let *scope* be the last [item](#) of *device*.*[[errorScopeStack]]*.

While *scope* is not *undefined*:

If *scope*.*[[filter]]* is *type*, return *scope*.

Set *scope* to the previous [item](#) of *device*.*[[errorScopeStack]]*.

Return *undefined*.

To *dispatch an error* [GPUError](#) *error* on [GPUDevice](#) *device*, run the following [device timeline](#) steps:

[Device timeline](#) steps:

Note: No errors are generated from a device which is lost. If this algorithm is called while *device* is [lost](#), it will not be observable to the application. See [§ 22 Errors & Debugging](#).

Let *scope* be the [current error scope](#) for *error* and *device*.

If *scope* is not `undefined`:

[Append](#) *error* to *scope*.[\[\[errors\]\]](#).

Return.

Otherwise, issue the following steps to the [content timeline](#):

[Content timeline](#) steps:

If the user agent chooses, [queue a global task for GPUDevice](#) *device* with the following steps:

Fire a [GPUUncapturedErrorEvent](#) named "[uncapturederror](#)" on *device*, with an [error](#) of *error*.

Note: After dispatching the event, user agents **should** surface uncaptured errors to developers, for example as warnings in the browser’s developer console, unless the event’s [defaultPrevented](#) is true. In other words, calling [preventDefault\(\)](#) on the event should silence the console warning.

Note: The user agent may choose to throttle or limit the number of [GPUUncapturedErrorEvent](#)s that a [GPUDevice](#) can raise to prevent an excessive amount of error handling or logging from impacting performance.

pushErrorScope(*filter*)

Pushes a new [GPU error scope](#) onto the [\[\[errorScopeStack\]\]](#) for *this*.

Called on: [GPUDevice](#) *this*.

Arguments:

Arguments for the [GPUDevice.pushErrorScope\(filter\)](#) method.

Parameter	Type	Nullable	Optional	Description
<i>filter</i>	GPUErrorFilter	×	×	Which class of errors this error scope observes.

Returns: `undefined`

[Content timeline](#) steps:

1. Issue the subsequent steps on the [Device timeline](#) of *this*.

[Device timeline](#) steps:

1. Let *scope* be a new [GPU error scope](#).
2. Set *scope*.[\[\[filter\]\]](#) to *filter*.
3. [Push](#) *scope* onto *this*.[\[\[errorScopeStack\]\]](#).

popErrorScope()

Pops a [GPU error scope](#) off the [\[\[errorScopeStack\]\]](#) for *this* and resolves to **any** [GPUError](#) observed by the error scope, or `null` if none.

There is no guarantee of the ordering of promise resolution.

Called on: [GPUDevice](#) *this*.

Returns: [Promise](#)<[GPUError](#)?>

[Content timeline](#) steps:

1. Let *contentTimeline* be the current [Content timeline](#).
2. Let *promise* be [a new promise](#).
3. Issue the *check* steps on the [Device timeline](#) of *this*.
4. Return *promise*.

[Device timeline](#) *check* steps:

1. If *this* is [lost](#):
 1. Issue the following steps on *contentTimeline*:
 2. Return.

Note: No errors are generated from a device which is lost. See [§ 22 Errors & Debugging](#).

2. If any of the following requirements are unmet:

- `this.[[errorScopeStack]].size` must be > 0 .

Then issue the following steps on *contentTimeline* and return:

3. Let *scope* be the result of [popping](#) an [item](#) off of `this.[[errorScopeStack]]`.
4. Let *error* be **any** one of the items in `scope.[[errors]]`, or `null` if there are none.

For any two errors E1 and E2 in the list, if E2 was caused by E1, E2 **should not** be the one selected.

Note: For example, if E1 comes from `t = createTexture()`, and E2 comes from `t.createView()` because `t` was [invalid](#), E1 should be preferred since it will be easier for a developer to understand what went wrong. Since both of these are [GPUValidationErrors](#), the only difference will be in the [message](#) field, which is meant only to be read by humans anyway.

5. At an **unspecified point now or in the future**, issue the subsequent steps on *contentTimeline*.

Note: By allowing [popErrorScope\(\)](#) calls to resolve in any order, with any of the errors observed by the scope, this spec allows validation to complete out of order, as long as any state observations are made at the appropriate point in adherence to this spec. For example, this allows implementations to perform shader compilation, which depends only on non-stateful inputs, to be completed on a background thread in parallel with other device-timeline work, and report any resulting errors later.

Using error scopes to capture validation errors from a [GPUDevice](#) operation that may fail:

```
gpuDevice.pushErrorScope('validation');
```

```
let sampler = gpuDevice.createSampler({
  maxAnisotropy: 0, // Invalid, maxAnisotropy must be at least 1.
});
```

```
gpuDevice.popErrorScope().then((error) => {
  if (error) {
    // There was an error creating the sampler, so discard it.
    sampler = null;
    console.error(`An error occurred while creating sampler: ${error.message}`);
  }
});
```

NOTE:

Error scopes can encompass as many commands as needed. The number of commands an error scope covers will generally be correlated to what sort of action the application intends to take in response to an error occurring.

For example: An error scope that only contains the creation of a single resource, such as a texture or buffer, can be used to detect failures such as out of memory conditions, in which case the application may try freeing some resources and trying the allocation again.

Error scopes do not identify which command failed, however. So, for instance, wrapping all the commands executed while loading a model in a single error scope will not offer enough granularity to determine if the issue was due to memory constraints. As a result freeing resources would usually not be a productive response to a failure of that scope. A more appropriate response would be to allow the application to fall back to a different model or produce a warning that the model could not be loaded. If responding to memory constraints is desired, the operations allocating memory can always be wrapped in a smaller nested error scope.

22.4. Telemetry

When a [GPUError](#) is generated that is not observed by any [GPU error scope](#), the user agent **may** [fire an event](#) named *uncapturederror* at a [GPUDevice](#) using [GPUUncapturedErrorEvent](#).

Note: [uncapturederror](#) events are intended to be used for telemetry and reporting unexpected errors. They won't necessarily be dispatched for all uncaptured errors (for example, there may be a limit on the number of errors surfaced), so they should not be used for handling known error cases that may occur during normal operation of an application. Prefer using [pushErrorScope\(\)](#) and [popErrorScope\(\)](#) in those cases.

[Exposed=(Window, Worker), [SecureContext](#)]

interface
<code>GPUUncapturedErrorEvent</code>
<code>: Event</code> {
<i>constructor</i>
(
DOMString
<i>type</i>
,

[GPUUncapturedErrorEventInit](#)

gpuUncapturedErrorEventInitDict

```
);
[SameObject] readonly attribute GPUError error;
};
```

dictionary

GPUUncapturedErrorEventInit

```
: EventInit {
    required GPUError
```

error

```
;
};
```

[GPUUncapturedErrorEvent](#) has the following attributes:

error, of type [GPUError](#), readonly

A [slot-backed attribute](#) holding an object representing the error that was uncaptured. This has the same type as errors returned by [popErrorScope\(\)](#).

partial interface [GPUDevice](#) {

attribute [EventHandler onuncapturederror](#);

```
};
```

[GPUDevice](#) has the following [content timeline properties](#):

onuncapturederror, of type [EventHandler](#)

An [event handler IDL attribute](#) for the [uncapturederror](#) event type.

Listening for uncaptured errors from a [GPUDevice](#):

```
gpuDevice.addEventListener('uncapturederror', (event) => {
    // Re-surface the error, because adding an event listener may silence console logs.
    console.error('A WebGPU error was not captured:', event.error);

    myEngineDebugReport.uncapturedErrors.push({
        type: event.error.constructor.name,
        message: event.error.message,
    });
});
```

23. Detailed Operations

This section describes the details of various GPU operations.

23.1. Computing

Computing operations provide direct access to GPU's programmable hardware. Compute shaders do not have shader stage inputs or outputs; their results are side effects from writing data into storage bindings bound either as [GPUBufferBindingLayout](#) with [GPUBufferBindingType "storage"](#) or as [GPUStorageTextureBindingLayout](#). These operations are encoded within [GPUComputePassEncoder](#) as:

[dispatchWorkgroups\(\)](#)

[dispatchWorkgroupsIndirect\(\)](#)

The main compute algorithm:

compute(descriptor, dispatchCall)

Arguments:

descriptor: Description of the current [GPUComputePipeline](#).

dispatchCall: The dispatch call parameters. May come from function arguments or an [INDIRECT](#) buffer.

Let *computeInvocations* be an [empty list](#).

Let *computeStage* be *descriptor.compute*.

Let *workgroupSize* be the computed workgroup size for *computeStage.entryPoint* after applying *computeStage.constants* to *computeStage.module*.

```

For workgroupX in range [0, dispatchCall.workgroupCountX):
For workgroupY in range [0, dispatchCall.workgroupCountY):
For workgroupZ in range [0, dispatchCall.workgroupCountZ):
For localX in range [0, workgroupSize.x):
For localY in range [0, workgroupSize.y):
For localZ in range [0, workgroupSize.z):
Let invocation be { computeStage, workgroupX, workgroupY, workgroupZ, localX, localY, localZ }

```

[Append](#) *invocation* to *computeInvocations*.

For every *invocation* in *computeInvocations*, in any order the [device](#) chooses, including in parallel:

Set the shader [builtins](#):

```

Set the num\_workgroups builtin, if any, to (
dispatchCall.workgroupCountX,
dispatchCall.workgroupCountY,
dispatchCall.workgroupCountZ
)

```

```

Set the workgroup\_id builtin, if any, to (
invocation.workgroupX,
invocation.workgroupY,
invocation.workgroupZ
)

```

```

Set the local\_invocation\_id builtin, if any, to (
invocation.localX,
invocation.localY,
invocation.localZ
)

```

```

Set the global\_invocation\_id builtin, if any, to (
invocation.workgroupX * workgroupSize.x + invocation.localX,
invocation.workgroupY * workgroupSize.y + invocation.localY,
invocation.workgroupZ * workgroupSize.z + invocation.localZ
).

```

```

Set the local\_invocation\_index builtin, if any, to invocation.localX + (invocation.localY * workgroupSize.x) + (invocation.localZ * workgroupSize.x * workgroupSize.y)

```

Invoke the compute shader entry point described by *invocation.computeStage*.

Note: Shader invocations have no guaranteed order, and will generally run in parallel according to device capabilities. Developers should not assume that any given invocation or workgroup will complete before any other one is started. Some devices may appear to execute in a consistent order, but this behavior should not be relied on as it will not perform identically across all devices. Shaders that require synchronization across invocations must use [Synchronization Built-in Functions](#) to coordinate execution.

The [device](#) may become [lost](#) if [shader execution does not end](#) in a reasonable amount of time, as determined by the user agent.

23.2. Rendering

Rendering is done by a set of GPU operations that are executed within [GPURenderPassEncoder](#), and result in modifications of the texture data, viewed by the render pass attachments. These operations are encoded with:

```

draw\(\)
drawIndexed\(\),
drawIndirect\(\)
drawIndexedIndirect\(\).

```

Note: rendering is the traditional use of GPUs, and is supported by multiple fixed-function blocks in hardware.

The main rendering algorithm:

```
render(pipeline, drawCall, state)
```

Arguments:

pipeline: The current [GPURenderPipeline](#).

drawCall: The draw call parameters. May come from function arguments or an [INDIRECT](#) buffer.

state: [RenderState](#) of the [GPURenderCommandsMixin](#) where the draw call is issued.

Let *descriptor* be *pipeline*.[\[\[descriptor\]\]](#).

Resolve indices. See [§ 23.2.1 Index Resolution](#).

Let *vertexList* be the result of [resolve indices](#)(*drawCall*, *state*).

Process vertices. See [§ 23.2.2 Vertex Processing](#).

Execute [process vertices](#)(*vertexList*, *drawCall*, *descriptor*.[vertex](#), *state*).

Assemble primitives. See [§ 23.2.3 Primitive Assembly](#).

Execute [assemble primitives](#)(*vertexList*, *drawCall*, *descriptor*.[primitive](#)).

Clip primitives. See [§ 23.2.4 Primitive Clipping](#).

Let *primitiveList* be the result of this stage.

Rasterize. See [§ 23.2.5 Rasterization](#).

Let *rasterizationList* be the result of [rasterize](#)(*primitiveList*, *state*).

Process fragments. See [§ 23.2.6 Fragment Processing](#).

Gather a list of *fragments*, resulting from executing [process fragment](#)(*rasterPoint*, *descriptor*, *state*) for each *rasterPoint* in *rasterizationList*.

Write pixels. See [§ 23.2.7 Output Merging](#).

For each non-null *fragment* of *fragments*:

Execute [process depth stencil](#)(*fragment*, *pipeline*, *state*).

Execute [process color attachments](#)(*fragment*, *pipeline*, *state*).

23.2.1. Index Resolution

At the first stage of rendering, the pipeline builds a list of vertices to process for each instance.

resolve indices(*drawCall*, *state*)

Arguments:

drawCall: The draw call parameters. May come from function arguments or an [INDIRECT](#) buffer.

state: The snapshot of the [GPURenderCommandsMixin](#) state at the time of the draw call.

Returns: list of integer indices.

Let *vertexIndexList* be an empty list of indices.

If *drawCall* is an indexed draw call:

Initialize the *vertexIndexList* with *drawCall.indexCount* integers.

For *i* in range 0 .. *drawCall.indexCount* (non-inclusive):

Let *relativeVertexIndex* be [fetch index](#)(*i* + *drawCall.firstIndex*, *state*.[\[\[index_buffer\]\]](#)).

If *relativeVertexIndex* has the special value "out of bounds", return the empty list.

Note: Implementations may choose to display a warning when this occurs, especially when it is easy to detect (like in non-indirect indexed draw calls).

Append *drawCall.baseVertex* + *relativeVertexIndex* to the *vertexIndexList*.

Otherwise:

Initialize the *vertexIndexList* with *drawCall.vertexCount* integers.

Set each *vertexIndexList* item *i* to the value *drawCall.firstVertex* + *i*.

Return *vertexIndexList*.

Note: in the case of indirect draw calls, the *indexCount*, *vertexCount*, and other properties of *drawCall* are read from the indirect buffer instead of the draw command itself.

fetch index(*i*, *buffer*, *offset*, *format*)

Arguments:

i: Index of a vertex index to fetch.

state: The snapshot of the [GPURenderCommandsMixin](#) state at the time of the draw call.

Returns: unsigned integer or "out of bounds"

Let *indexSize* be defined by the *state*.[\[\[index_format\]\]](#):

["uint16"](#)

2

["uint32"](#)

4

If *state*.[\[\[index_buffer_offset\]\]](#) + |*i* + 1| × *indexSize* > *state*.[\[\[index_buffer_size\]\]](#), return the special value "out of bounds".

Interpret the data in *state*.[\[\[index_buffer\]\]](#), starting at offset *state*.[\[\[index_buffer_offset\]\]](#) + *i* × *indexSize*, of size *indexSize* bytes, as an unsigned integer and return it.

23.2.2. Vertex Processing

Vertex processing stage is a programmable stage of the render [pipeline](#) that processes the vertex attribute data, and produces clip space positions for [§ 23.2.4 Primitive Clipping](#), as well as other data for the [§ 23.2.6 Fragment Processing](#).

process vertices(vertexIndexList, drawCall, desc, state)

Arguments:

vertexIndexList: List of vertex indices to process (mutable, passed by reference).

drawCall: The draw call parameters. May come from function arguments or an [INDIRECT](#) buffer.

desc: The descriptor of type [GPUVertexState](#).

state: The snapshot of the [GPURenderCommandsMixin](#) state at the time of the draw call.

Each vertex *vertexIndex* in the *vertexIndexList*, in each instance of index *rawInstanceIndex*, is processed independently. The *rawInstanceIndex* is in range from 0 to *drawCall*.instanceCount - 1, inclusive. This processing happens in parallel, and any side effects, such as writes into [GPUBufferBindingType "storage"](#) bindings, may happen in any order.

Let *instanceIndex* be *rawInstanceIndex* + *drawCall*.firstInstance.

For each non-null *vertexBufferLayout* in the list of *desc*.[buffers](#):

Let *i* be the index of the buffer layout in this list.

Let *vertexBuffer*, *vertexBufferOffset*, and *vertexBufferBindingSize* be the buffer, offset, and size at slot *i* of *state*.[\[\[vertex_buffers\]\]](#).

Let *vertexElementIndex* be dependent on *vertexBufferLayout*.[stepMode](#):

["vertex"](#)

vertexIndex

["instance"](#)

instanceIndex

Let *drawCallOutOfBounds* be **false**.

For each *attributeDesc* in *vertexBufferLayout*.[attributes](#):

Let *attributeOffset* be *vertexBufferOffset* + *vertexElementIndex* * *vertexBufferLayout*.[arrayStride](#) + *attributeDesc*.[offset](#).

If *attributeOffset* + [byteSize](#)(*attributeDesc*.[format](#)) > *vertexBufferOffset* + *vertexBufferBindingSize*:

Set *drawCallOutOfBounds* to **true**.

Optionally ([implementation-defined](#)), [empty](#) *vertexIndexList* and return, cancelling the draw call.

Note: This allows implementations to detect out-of-bounds values in the index buffer before issuing a draw call, instead of using [invalid memory reference](#) behavior.

For each *attributeDesc* in *vertexBufferLayout*.[attributes](#):

If *drawCallOutOfBounds* is **true**:

Load the attribute *data* according to WGSL's [invalid memory reference](#) behavior, from *vertexBuffer*.

Note: [Invalid memory reference](#) allows several behaviors, including actually loading the "correct" result for an attribute that is in-bounds, even when the draw-call-wide *drawCallOutOfBounds* is **true**.

Otherwise:

Let *attributeOffset* be *vertexBufferOffset* + *vertexElementIndex* * *vertexBufferLayout.arrayStride* + *attributeDesc.offset*.

Load the attribute *data* of format *attributeDesc.format* from *vertexBuffer* starting at offset *attributeOffset*. The components are loaded in the order x, y, z, w from buffer memory.

Convert the *data* into a shader-visible format, according to [channel formats](#) rules.

An attribute of type ["snorm8x2"](#) and byte values of [0x70, 0xD0] will be converted to `vec2<f32>(0.88, -0.38)` in WGSL.

Adjust the *data* size to the shader type:

if both are scalar, or both are vectors of the same dimensionality, no adjustment is needed.

if *data* is vector but the shader type is scalar, then only the first component is extracted.

if both are vectors, and *data* has a higher dimension, the extra components are dropped.

An attribute of type ["float32x3"](#) and value `vec3<f32>(1.0, 2.0, 3.0)` will be exposed to the shader as `vec2<f32>(1.0, 2.0)` if a 2-component vector is expected.

if the shader type is a vector of higher dimensionality, or the *data* is a scalar, then the missing components are filled from `vec4<*>(0, 0, 0, 1)` value.

An attribute of type ["sint32"](#) and value 5 will be exposed to the shader as `vec4<i32>(5, 0, 0, 1)` if a 4-component vector is expected.

Bind the *data* to vertex shader input location *attributeDesc.shaderLocation*.

For each [GPUBindGroup](#) group at *index* in *state.[[bind_groups]]*:

For each resource [GPUBindingResource](#) in the bind group:

Let *entry* be the corresponding [GPUBindGroupLayoutEntry](#) for this resource.

If *entry.visibility* includes [VERTEX](#):

Bind the resource to the shader under group *index* and binding [GPUBindGroupLayoutEntry.binding](#).

Set the shader [builtins](#):

Set the `vertex_index` builtin, if any, to *vertexIndex*.

Set the `instance_index` builtin, if any, to *instanceIndex*.

Invoke the vertex shader entry point described by *desc*.

Note: The target platform caches the results of vertex shader invocations. There is no guarantee that any *vertexIndex* that repeats more than once will result in multiple invocations. Similarly, there is no guarantee that a single *vertexIndex* will only be processed once.

The [device](#) may become [lost](#) if [shader execution does not end](#) in a reasonable amount of time, as determined by the user agent.

23.2.3. Primitive Assembly

Primitives are assembled by a fixed-function stage of GPUs.

assemble primitives(*vertexIndexList*, *drawCall*, *desc*)

Arguments:

vertexIndexList: List of vertex indices to process.

drawCall: The draw call parameters. May come from function arguments or an [INDIRECT](#) buffer.

desc: The descriptor of type [GPUPrimitiveState](#).

For each instance, the primitives get assembled from the vertices that have been processed by the shaders, based on the *vertexIndexList*.

First, if the primitive topology is a strip, (which means that *desc.stripIndexFormat* is not undefined) and the *drawCall* is indexed, the *vertexIndexList* is split into sub-lists using the maximum value of *desc.stripIndexFormat* as a separator.

Example: a *vertexIndexList* with values [1, 2, 65535, 4, 5, 6] of type ["uint16"](#) will be split in sub-lists [1, 2] and [4, 5, 6].

For each of the sub-lists *vl*, primitive generation is done according to the *desc.topology*:

["line-list"](#)

Line primitives are composed from (*vl.0*, *vl.1*), then (*vl.2*, *vl.3*), then (*vl.4* to *vl.5*), etc. Each subsequent primitive takes 2 vertices.

["line-strip"](#)

Line primitives are composed from (*vl.0*, *vl.1*), then (*vl.1*, *vl.2*), then (*vl.2*, *vl.3*), etc. Each subsequent primitive takes 1 vertex.

["triangle-list"](#)

Triangle primitives are composed from (*vl.0*, *vl.1*, *vl.2*), then (*vl.3*, *vl.4*, *vl.5*), then (*vl.6*, *vl.7*, *vl.8*), etc. Each subsequent primitive takes 3 vertices.

["triangle-strip"](#)

Triangle primitives are composed from (v1.0, v1.1, v1.2), then (v1.2, v1.1, v1.3), then (v1.2, v1.3, v1.4), then (v1.4, v1.3, v1.5), etc. Each subsequent primitive takes 1 vertices.

Any incomplete primitives are dropped.

23.2.4. Primitive Clipping

Vertex shaders have to produce a built-in [position](#) (of type `vec4<f32>`), which denotes the *clip position* of a vertex in [clip space coordinates](#).

Primitives are clipped to the *clip volume*, which, for any [clip position](#) p inside a primitive, is defined by the following inequalities:

$$-p.w \leq p.x \leq p.w$$

$$-p.w \leq p.y \leq p.w$$

$$0 \leq p.z \leq p.w \text{ (depth clipping)}$$

When the "[clip-distances](#)" feature is enabled, this [clip volume](#) can be further restricted by user-defined half-spaces by declaring [clip_distances](#) in the output of vertex stage. Each value in the [clip_distances](#) array will be linearly interpolated across the primitive, and the portion of the primitive with interpolated distances less than 0 will be clipped.

If [descriptor.primitive.unclippedDepth](#) is `true`, [depth clipping](#) is not applied: the [clip volume](#) is not bounded in the z dimension.

A primitive passes through this stage unchanged if every one of its edges lie entirely inside the [clip volume](#). If the edges of a primitives intersect the boundary of the [clip volume](#), the intersecting edges are reconnected by new edges that lie along the boundary of the [clip volume](#). For triangular primitives ([descriptor.primitive.topology](#) is "[triangle-list](#)" or "[triangle-strip](#)"), this reconnection may result in introduction of new vertices into the polygon, internally.

If a primitive intersects an edge of the [clip volume](#)'s boundary, the clipped polygon must include a point on this boundary edge.

If the vertex shader outputs other floating-point values (scalars and vectors), qualified with "perspective" interpolation, they also get clipped. The output values associated with a vertex that lies within the clip volume are unaffected by clipping. If a primitive is clipped, however, the output values assigned to vertices produced by clipping are clipped.

Considering an edge between vertices a and b that got clipped, resulting in the vertex c , let's define t to be the ratio between the edge vertices: $c.p = t \times a.p + (1 - t) \times b.p$, where $x.p$ is the output [clip position](#) of a vertex x .

For each vertex output value "v" with a corresponding fragment input, $a.v$ and $b.v$ would be the outputs for a and b vertices respectively. The clipped shader output $c.v$ is produced based on the interpolation qualifier:

[flat](#)

Flat interpolation is unaffected, and is based on the *provoking vertex*, which is determined by the [interpolation sampling](#) mode declared in the shader. The output value is the same for the whole primitive, and matches the vertex output of the [provoking vertex](#).

[linear](#)

The interpolation ratio gets adjusted against the perspective coordinates of the [clip positions](#), so that the result of interpolation is linear in screen space.

[perspective](#)

The value is linearly interpolated in clip space, producing perspective-correct values.

The result of primitive clipping is a new set of primitives, which are contained within the [clip volume](#).

23.2.5. Rasterization

Rasterization is the hardware processing stage that maps the generated primitives to the 2-dimensional rendering area of the *framebuffer* - the set of render attachments in the current [GPURenderPassEncoder](#). This rendering area is split into an even grid of pixels.

The [framebuffer](#) coordinates start from the top-left corner of the render targets. Each unit corresponds exactly to one pixel. See [§ 3.3 Coordinate Systems](#) for more information.

Rasterization determines the set of pixels affected by a primitive. In case of multi-sampling, each pixel is further split into [descriptor.multisample.count](#) samples. The *standard sample patterns* are as follows, with positions in framebuffer coordinates relative to the top-left corner of the pixel, such that the pixel ranges from (0, 0) to (1, 1):

multisample.count	Sample positions
1	Sample 0: (0.5, 0.5)
4	Sample 0: (0.375, 0.125) Sample 1: (0.875, 0.375) Sample 2: (0.125, 0.625) Sample 3: (0.625, 0.875)

Implementations must use the [standard sample pattern](#) for the given [multisample.count](#) when performing rasterization.

Let's define a *FragmentDestination* to contain:

position

the 2D pixel position using [framebuffer coordinates](#)

sampleIndex

an integer in case [§ 23.2.10 Per-Sample Shading](#) is active, or `null` otherwise

We'll also use a notion of [normalized device coordinates](#), or NDC. In this coordinate system, the viewport bounds range in X and Y from -1 to 1, and in Z from 0 to 1.

Rasterization produces a list of *RasterizationPoints*, each containing the following data:

destination

refers to [FragmentDestination](#)

coverageMask

refers to multisample coverage mask (see [§ 23.2.11 Sample Masking](#))

frontFacing

is true if it's a point on the front face of a primitive

perspectiveDivisor

refers to interpolated $1.0 \div W$ across the primitive

depth

refers to the depth in [viewport coordinates](#), i.e. between the `[[viewport]].minDepth` and `maxDepth`.

primitiveVertices

refers to the list of vertex outputs forming the primitive

barycentricCoordinates

refers to [§ 23.2.5.3 Barycentric coordinates](#)

rasterize(primitiveList, state)

Arguments:

primitiveList: List of primitives to rasterize.

state: The active [RenderState](#).

Returns: list of [RasterizationPoint](#).

Each primitive in *primitiveList* is processed independently. However, the order of primitives affects later stages, such as depth/stencil operations and pixel writes.

First, the clipped vertices are transformed into [NDC](#) - normalized device coordinates. Given the output position *p*, the [NDC](#) position and perspective divisor are:

$$\text{ndc}(p) = \text{vector}(p.x \div p.w, p.y \div p.w, p.z \div p.w)$$

$$\text{divisor}(p) = 1.0 \div p.w$$

Let *vp* be *state*.`[[viewport]]`. Map the [NDC](#) position *n* into [viewport coordinates](#):

Compute [framebuffer](#) coordinates from the render target offset and size:

$$\text{framebufferCoords}(n) = \text{vector}(vp.x + 0.5 \times (n.x + 1) \times vp.width, vp.y + 0.5 \times (-n.y + 1) \times vp.height)$$

Compute depth by linearly mapping [0,1] to the viewport depth range:

$$\text{depth}(n) = vp.minDepth + n.z \times (vp.maxDepth - vp.minDepth)$$

Let *rasterizationPoints* be the list of points, each having its attributes (`divisor(p)`, `framebufferCoords(n)`, `depth(n)`, etc.) interpolated according to its position on the primitive, using the same interpolation as [§ 23.2.4 Primitive Clipping](#). If the attribute is user-defined (not a [built-in output value](#)) then the [interpolation type](#) specified by the [@interpolate](#) WGSL attribute is used.

Proceed with a specific rasterization algorithm, depending on [primitive.topology](#):

["point-list"](#)

The point, if not filtered by [§ 23.2.4 Primitive Clipping](#), goes into [§ 23.2.5.1 Point Rasterization](#).

["line-list"](#) or ["line-strip"](#)

The line cut by [§ 23.2.4 Primitive Clipping](#) goes into [§ 23.2.5.2 Line Rasterization](#).

["triangle-list"](#) or ["triangle-strip"](#)

The polygon produced in [§ 23.2.4 Primitive Clipping](#) goes into [§ 23.2.5.4 Polygon Rasterization](#).

Remove all the points *rp* from *rasterizationPoints* that have *rp*.`destination.position` outside of *state*.`[[scissorRect]]`.

Return *rasterizationPoints*.

23.2.5.1. Point Rasterization

A single [FragmentDestination](#) is selected within the pixel containing the [framebuffer](#) coordinates of the point.

The coverage mask depends on multi-sampling mode:

sample-frequency

$\text{coverageMask} = 1 \ll \text{sampleIndex}$

pixel-frequency multi-sampling

$\text{coverageMask} = 1 \ll \text{descriptor.multisampleCount} - 1$

no multi-sampling

$\text{coverageMask} = 1$

23.2.5.2. Line Rasterization

The exact algorithm used for line rasterization is not defined, and may differ between implementations. For example, the line may be drawn using [§ 23.2.5.4 Polygon Rasterization](#) of a 1px-width rectangle around the line segment, or using Bresenham's line algorithm to select the [FragmentDestinations](#).

Note: See [Basic Line Segment Rasterization](#) and [Bresenham Line Segment Rasterization](#) in the [Vulkan 1.3](#) spec for more details of how line these line rasterization algorithms may be implemented.

23.2.5.3. Barycentric coordinates

Barycentric coordinates is a list of n numbers b_i , defined for a point p inside a convex polygon with n vertices v_i in [framebuffer](#) space. Each b_i is in range 0 to 1, inclusive, and represents the proximity to vertex v_i . Their sum is always constant:

$$\sum (b_i) = 1$$

These coordinates uniquely specify any point p within the polygon (or on its boundary) as:

$$p = \sum (b_i \times p_i)$$

For a polygon with 3 vertices - a triangle, barycentric coordinates of any point p can be computed as follows:

$$A_{\text{polygon}} = A(v_1, v_2, v_3) \quad b_1 = A(p, v_2, v_3) \div A_{\text{polygon}} \quad b_2 = A(v_1, p, v_3) \div A_{\text{polygon}} \quad b_3 = A(v_1, v_2, p) \div A_{\text{polygon}}$$

Where $A(\text{list of points})$ is the area of the polygon with the given set of vertices.

For polygons with more than 3 vertices, the exact algorithm is implementation-dependent. One of the possible implementations is to triangulate the polygon and compute the barycentrics of a point based on the triangle it falls into.

23.2.5.4. Polygon Rasterization

A polygon is *front-facing* if it's oriented towards the projection. Otherwise, the polygon is *back-facing*.

`rasterize polygon()`

Arguments:

Returns: list of [RasterizationPoint](#).

Let `rasterizationPoints` be an empty list.

Let $v(i)$ be the [framebuffer](#) coordinates for the clipped vertex number i (starting with 1) in a rasterized polygon of n vertices.

Note: this section uses the term "polygon" instead of a "triangle", since [§ 23.2.4 Primitive Clipping](#) stage may have introduced additional vertices. This is non-observable by the application.

Determine if the polygon is front-facing, which depends on the sign of the *area* occupied by the polygon in [framebuffer](#) coordinates:

$$\text{area} = 0.5 \times ((v_1.x \times v_n.y - v_n.x \times v_1.y) + \sum (v_{i+1}.x \times v_i.y - v_i.x \times v_{i+1}.y))$$

The sign of *area* is interpreted based on the [primitive.frontFace](#):

["ccw"](#)

$\text{area} > 0$ is considered [front-facing](#), otherwise [back-facing](#)

["cw"](#)

$\text{area} < 0$ is considered [front-facing](#), otherwise [back-facing](#)

Cull based on [primitive.cullMode](#):

["none"](#)

All polygons pass this test.

["front"](#)

The [front-facing](#) polygons are discarded, and do not process in later stages of the render pipeline.

"back"

The [back-facing](#) polygons are discarded.

Determine a set of [fragments](#) inside the polygon in [framebuffer](#) space - these are locations scheduled for the per-fragment operations. This operation is known as "point sampling". The logic is based on *descriptor.multisample*:

disabled

[Fragments](#) are associated with pixel centers. That is, all the points with coordinates C , where $\text{fract}(C) = \text{vector2}(0.5, 0.5)$ in the [framebuffer](#) space, enclosed into the polygon, are included. If a pixel center is on the edge of the polygon, whether or not it's included is not defined.

Note: this becomes a subject of precision for the rasterizer.

enabled

Each pixel is associated with *descriptor.multisample.count* locations, which are [implementation-defined](#). The locations are ordered, and the list is the same for each pixel of the [framebuffer](#). Each location corresponds to one fragment in the multisampled [framebuffer](#).

The rasterizer builds a mask of locations being hit inside each pixel and provides it as "sample-mask" built-in to the fragment shader.

For each produced fragment of type [FragmentDestination](#):

Let rp be a new [RasterizationPoint](#) object

Compute the list b as [§ 23.2.5.3 Barycentric coordinates](#) of that fragment. Set $rp.barycentricCoordinates$ to b .

Let d_i be the depth value of v_i .

Set $rp.depth$ to $\sum (b_i \times d_i)$

Append rp to *rasterizationPoints*.

Return *rasterizationPoints*.

23.2.6. Fragment Processing

The fragment processing stage is a programmable stage of the render [pipeline](#) that computes the fragment data (often a color) to be written into render targets.

This stage produces a *Fragment* for each [RasterizationPoint](#):

destination refers to [FragmentDestination](#).

frontFacing is true if it's a fragment on the front face of a primitive.

coverageMask refers to multisample coverage mask (see [§ 23.2.11 Sample Masking](#)).

depth refers to the depth in [viewport coordinates](#), i.e. between the [\[\[viewport\]\]](#) *minDepth* and *maxDepth*.

colors refers to the list of color values, one for each target in [colorAttachments](#).

depthPassed is `true` if the fragment passed the [depthCompare](#) operation.

stencilPassed is `true` if the fragment passed the stencil [compare](#) operation.

process fragment(rp , *descriptor*, *state*)

Arguments:

rp: The [RasterizationPoint](#), produced by [§ 23.2.5 Rasterization](#).

descriptor: The descriptor of type [GPURenderPipelineDescriptor](#).

state: The active [RenderState](#).

Returns: [Fragment](#) or `null`.

Let *fragmentDesc* be *descriptor.fragment*.

Let *depthStencilDesc* be *descriptor.depthStencil*.

Let *fragment* be a new [Fragment](#) object.

Set *fragment.destination* to *rp.destination*.

Set *fragment.frontFacing* to *rp.frontFacing*.

Set *fragment.coverageMask* to *rp.coverageMask*.

Set *fragment.depth* to *rp.depth*.

If *frag_depth builtin* is not produced by the shader:

Set *fragment.depthPassed* to the result of `compare fragment(fragment.destination, fragment.depth, "depth", state.[[depthStencilAttachment]], depthStencilDesc?.depthCompare)`.

Set *stencilState* to *depthStencilDesc?.stencilFront* if *rp.frontFacing* is `true` and *depthStencilDesc?.stencilBack* otherwise.

Set *fragment.stencilPassed* to the result of `compare fragment(fragment.destination, state.[[stencilReference]], "stencil", state.[[depthStencilAttachment]], stencilState?.compare)`.

If *fragmentDesc* is not `null`:

If *fragment.depthPassed* is `false`, the `frag_depth builtin` is not produced by the shader entry point, and the shader entry point does not write to any `storage` bindings, the following steps may be skipped.

Set the shader input `builtins`. For each non-composite argument of the entry point, annotated as a `builtin`, set its value based on the annotation:

`position`

`vec4<f32>(rp.destination.position, rp.depth, rp.perspectiveDivisor)`

`front_facing`

`rp.frontFacing`

`sample_index`

`rp.destination.sampleIndex`

`sample_mask`

`rp.coverageMask`

For each user-specified `shader stage input` of the fragment stage:

Let *value* be the interpolated fragment input, based on *rp.barycentricCoordinates*, *rp.primitiveVertices*, and the `interpolation` qualifier on the input.

Set the corresponding fragment shader `location` input to *value*.

Invoke the fragment shader entry point described by *fragmentDesc*.

The `device` may become `lost` if `shader execution does not end` in a reasonable amount of time, as determined by the user agent.

If the fragment issued `discard`, return `null`.

Set *fragment.colors* to the user-specified `shader stage output` values from the shader.

Take the shader output `builtins`:

If `frag_depth builtin` is produced by the shader as *value*:

Let *vp* be *state.[[viewport]]*.

Set *fragment.depth* to `clamp(value, vp.minDepth, vp.maxDepth)`.

Set *fragment.depthPassed* to the result of `compare fragment(fragment.destination, fragment.depth, "depth", state.[[depthStencilAttachment]], depthStencilDesc?.depthCompare)`.

If `sample_mask builtin` is produced by the shader as *value*:

Set *fragment.coverageMask* to *fragment.coverageMask* \wedge *value*.

Otherwise we are in § 23.2.8 No Color Output mode, and *fragment.colors* is empty.

Return *fragment*.

`compare fragment(destination, value, aspect, attachment, compareFunc)`

Arguments:

destination: The `FragmentDestination`.

value: The value to be compared.

aspect: The `aspect` of *attachment* to sample values from.

attachment: The attachment to be compared against.

compareFunc: The `GPUCompareFunction` to use, or `undefined`.

Returns: `true` if the comparison passes, or `false` otherwise

If *attachment* is `undefined` or does not have *aspect*, return `true`.

If *compareFunc* is `undefined` or `"always"`, return `true`.

Let *attachmentValue* be the value of *aspect* of *attachment* at *destination*.

Return `true` if comparing *value* with *attachmentValue* using *compareFunc* succeeds, and `false` otherwise.

Processing of fragments happens in parallel, while any side effects, such as writes into [GPUBufferBindingType "storage"](#) bindings, may happen in any order.

23.2.7. Output Merging

Output merging is a fixed-function stage of the render [pipeline](#) that outputs the fragment color, depth and stencil data to be written into the render pass attachments.

process depth stencil(fragment, pipeline, state)

Arguments:

fragment: The [Fragment](#), produced by § 23.2.6 [Fragment Processing](#).

pipeline: The current [GPURenderPipeline](#).

state: The active [RenderState](#).

Let *depthStencilDesc* be *pipeline*.[\[\[descriptor\]\].depthStencil](#).

If *pipeline*.[\[\[writesDepth\]\]](#) is `true` and *fragment*.[depthPassed](#) is `true`:

Set the value of the depth aspect of *state*.[\[\[depthStencilAttachment\]\]](#) at *fragment*.[destination](#) to *fragment*.[depth](#).

If *pipeline*.[\[\[writesStencil\]\]](#) is `true`:

Set *stencilState* to *depthStencilDesc*.[stencilFront](#) if *fragment*.[frontFacing](#) is `true` and *depthStencilDesc*.[stencilBack](#) otherwise.

If *fragment*.[stencilPassed](#) is `false`:

Let *stencilOp* be *stencilState*.[failOp](#).

Otherwise, if *fragment*.[depthPassed](#) is `false`:

Let *stencilOp* be *stencilState*.[depthFailOp](#).

Otherwise:

Let *stencilOp* be *stencilState*.[passOp](#).

Update the value of the stencil aspect of *state*.[\[\[depthStencilAttachment\]\]](#) at *fragment*.[destination](#) by performing the operation described by *stencilOp*.

The depth input to this stage, if any, is clamped to the current [\[\[viewport\]\]](#) depth range (regardless of whether the fragment shader stage writes the `frag_depth` builtin).

process color attachments(fragment, pipeline, state)

Arguments:

fragment: The [Fragment](#), produced by § 23.2.6 [Fragment Processing](#).

pipeline: The current [GPURenderPipeline](#).

state: The active [RenderState](#).

If *fragment*.[depthPassed](#) is `false` or *fragment*.[stencilPassed](#) is `false`, return.

Let *targets* be *pipeline*.[\[\[descriptor\]\].fragment.targets](#).

For each *attachment* of *state*.[\[\[colorAttachments\]\]](#):

Let *color* be the value from *fragment*.[colors](#) that corresponds with *attachment*.

Let *targetDesc* be the *targets* entry that corresponds with *attachment*.

If *targetDesc*.[blend](#) is `provided`:

Let *colorBlend* be *targetDesc*.[blend.color](#).

Let *alphaBlend* be *targetDesc*.[blend.alpha](#).

Set the RGB components of *color* to the value computed by performing the operation described by *colorBlend*.[operation](#) with the values described by *colorBlend*.[srcFactor](#) and *colorBlend*.[dstFactor](#).

Set the alpha component of *color* to the value computed by performing the operation described by *alphaBlend*.[operation](#) with the values described by *alphaBlend*.[srcFactor](#) and *alphaBlend*.[dstFactor](#).

Set the value of *attachment* at *fragment*.[destination](#) to *color*.

23.2.8. No Color Output

In no-color-output mode, [pipeline](#) does not produce any color attachment outputs.

The [pipeline](#) still performs rasterization and produces depth values based on the vertex position output. The depth testing and stencil operations can still be used.

23.2.9. Alpha to Coverage

In alpha-to-coverage mode, an additional *alpha-to-coverage mask* of MSAA samples is generated based on the *alpha* component of the fragment shader output value at `@location(0)`.

The algorithm of producing the extra mask is platform-dependent and can vary for different pixels. It guarantees that:

if $\alpha \leq 0.0$, the result is 0x0

if $\alpha \geq 1.0$, the result is 0xFFFFFFFF

intermediate *alpha* values should result in a proportionate number of bits set to 1 in the mask. Not all platforms guarantee that the number of bits set to 1 in the mask monotonically increases as alpha increases for a given pixel.

23.2.10. Per-Sample Shading

When rendering into multisampled render attachments, fragment shaders can be run once per-pixel or once per-sample. Fragment shaders **must** run once per-sample if either the `sample_index` [builtin](#) or `sample` [interpolation sampling](#) is used and contributes to the shader output. Otherwise fragment shaders **may** run once per-pixel with the result broadcast out to each of the samples included in the [final sample mask](#).

When using per-sample shading, the color output for sample *N* is produced by the fragment shader execution with `sample_index == N` for the current pixel.

23.2.11. Sample Masking

The *final sample mask* for a pixel is computed as: `rasterization_mask & mask & shader-output_mask`.

Only the lower `count` bits of the mask are considered.

If the least-significant bit at position *N* of the [final sample mask](#) has value of "0", the sample color outputs (corresponding to sample *N*) to all attachments of the fragment shader are discarded. Also, no depth test or stencil operations are executed on the relevant samples of the depth-stencil attachment.

The *rasterization mask* is produced by the rasterization stage, based on the shape of the rasterized polygon. The samples included in the shape get the relevant bits 1 in the mask.

The *shader-output mask* takes the output value of "sample_mask" [builtin](#) in the fragment shader. If the builtin is not output from the fragment shader, and [alphaToCoverageEnabled](#) is enabled, the `shader-output_mask` becomes the [alpha-to-coverage mask](#). Otherwise, it defaults to 0xFFFFFFFF.

24. Type Definitions

```
typedef [EnforceRange] unsigned long
```

```
GPUBufferDynamicOffset
```

```
;
```

```
typedef [EnforceRange] unsigned long
```

```
GPUStencilValue
```

```
;
```

```
typedef [EnforceRange] unsigned long
```

```
GPUSampleMask
```

```
;
```

```
typedef [EnforceRange] long
```

```
GPUDepthBias
```

```
;
```

```
typedef [EnforceRange] unsigned long long
```

```
GPUSize64
```

```
;
```

```
typedef [EnforceRange] unsigned long
```

```
GPUIntegerCoordinate
```

```
;
```

```
typedef [EnforceRange] unsigned long
```

```
GPUIndex32
```

```
;
```

```
typedef [EnforceRange] unsigned long
```

```
GPUSize32
```

```
;
```

```
typedef [EnforceRange] long
```

```
GPUSignedOffset32
```

```
;
```

```
typedef unsigned long long
```

```
GPUSize640ut
```

```
;
```

```
typedef unsigned long
```

```
GPUIntegerCoordinateOut
```

```
;
```

```
typedef unsigned long
```

```
GPUSize320ut
```

```
;
```

```
typedef unsigned long
```

```
GPUFlagsConstant
```

```
;
```

24.1. Colors & Vectors

```
dictionary
```

```
GPUColorDict
```

```
{
```

```
    required double r;
```

```
    required double g;
```

```
    required double b;
```

```
    required double a;
```

```
};
```

```
typedef (sequence<double> or GPUColorDict)
```

```
GPUColor
```

```
;
```

Note: `double` is large enough to precisely hold 32-bit signed/unsigned integers and single-precision floats.

r, of type `double`

The red channel value.

g, of type `double`

The green channel value.

b, of type `double`

The blue channel value.

a, of type `double`

The alpha channel value.

For a given `GPUColor` value *color*, depending on its type, the syntax:

color.r refers to either `GPUColorDict.r` or the first item of the sequence (`asserting` there is such an item).

color.g refers to either `GPUColorDict.g` or the second item of the sequence (`asserting` there is such an item).

color.b refers to either `GPUColorDict.b` or the third item of the sequence (`asserting` there is such an item).

color.a refers to either `GPUColorDict.a` or the fourth item of the sequence (`asserting` there is such an item).

`validate GPUColor shape(color)`

Arguments:

color: The `GPUColor` to validate.

Returns: `undefined`

[Content timeline](#) steps:

Throw a `TypeError` if *color* is a sequence and *color.size* \neq 4.


```
dictionary
GPUOrigin2DDict
{
  GPUIntegerCoordinate
x
= 0;
  GPUIntegerCoordinate
y
= 0;
};
typedef (sequence<GPUIntegerCoordinate> or GPUOrigin2DDict)
GPUOrigin2D
;
```

For a given GPUOrigin2D value *origin*, depending on its type, the syntax:

origin.x refers to either GPUOrigin2DDict.x or the first item of the sequence (0 if not present).

origin.y refers to either GPUOrigin2DDict.y or the second item of the sequence (0 if not present).

validate GPUOrigin2D shape(*origin*)

Arguments:

origin: The GPUOrigin2D to validate.

Returns: undefined

Content timeline steps:

Throw a TypeError if *origin* is a sequence and *origin.size* > 2.

```
dictionary
GPUOrigin3DDict
{
  GPUIntegerCoordinate
x
= 0;
  GPUIntegerCoordinate
y
= 0;
  GPUIntegerCoordinate
z
= 0;
};
typedef (sequence<GPUIntegerCoordinate> or GPUOrigin3DDict)
GPUOrigin3D
;
```

For a given GPUOrigin3D value *origin*, depending on its type, the syntax:

origin.x refers to either GPUOrigin3DDict.x or the first item of the sequence (0 if not present).

origin.y refers to either GPUOrigin3DDict.y or the second item of the sequence (0 if not present).

origin.z refers to either GPUOrigin3DDict.z or the third item of the sequence (0 if not present).

validate GPUOrigin3D shape(*origin*)

Arguments:

origin: The GPUOrigin3D to validate.

Returns: undefined

Content timeline steps:

Throw a TypeError if *origin* is a sequence and *origin.size* > 3.

dictionary

GPUExtent3DDict

```
{
  required GPUIntegerCoordinate width;
  GPUIntegerCoordinate height = 1;
  GPUIntegerCoordinate depthOrArrayLayers = 1;
};
typedef (sequence<GPUIntegerCoordinate> or GPUExtent3DDict)
```

GPUExtent3D

;

width, of type [GPUIntegerCoordinate](#)

The width of the extent.

height, of type [GPUIntegerCoordinate](#), defaulting to 1

The height of the extent.

depthOrArrayLayers, of type [GPUIntegerCoordinate](#), defaulting to 1

The depth of the extent or the number of array layers it contains. If used with a [GPUTexture](#) with a [GPUTextureDimension](#) of ["3d"](#) defines the depth of the texture. If used with a [GPUTexture](#) with a [GPUTextureDimension](#) of ["2d"](#) defines the number of array layers in the texture.

For a given [GPUExtent3D](#) value *extent*, depending on its type, the syntax:

extent.width refers to either [GPUExtent3DDict.width](#) or the first item of the sequence ([asserting](#) there is such an item).

extent.height refers to either [GPUExtent3DDict.height](#) or the second item of the sequence (1 if not present).

extent.depthOrArrayLayers refers to either [GPUExtent3DDict.depthOrArrayLayers](#) or the third item of the sequence (1 if not present).

validate GPUExtent3D shape(*extent*)

Arguments:

extent: The [GPUExtent3D](#) to validate.

Returns: [undefined](#)

[Content timeline](#) steps:

Throw a [TypeError](#) if:

extent is a sequence, and

extent.size < 1 or *extent.size* > 3.

25. Feature Index

25.1. "core-features-and-limits"

Allows all Core WebGPU features and limits to be used.

Note: This is currently available on all adapters and enabled automatically on all devices even if not requested.

25.2. "depth-clip-control"

Allows [depth clipping](#) to be disabled.

This feature adds the following [optional API surfaces](#):

New [GPUPrimitiveState](#) dictionary members:

[unclippedDepth](#)

25.3. "depth32float-stencil8"

Allows for explicit creation of textures of format ["depth32float-stencil8"](#).

This feature adds the following [optional API surfaces](#):

New [GPUTextureFormat](#) enum values:

["depth32float-stencil8"](#)

25.4. "texture-compression-bc"

Allows for explicit creation of textures of [BC compressed formats](#) which include the "S3TC", "RGTC", and "BPTC" formats. Only supports 2D textures.

Note: Adapters which support ["texture-compression-bc"](#) do not always support ["texture-compression-bc-sliced-3d"](#). To use ["texture-compression-bc-sliced-3d"](#), ["texture-compression-bc"](#) must be enabled explicitly as this feature does not enable the BC formats.

This feature adds the following [optional API surfaces](#):

New [GPUTextureFormat](#) enum values:

["bc1-rgba-unorm"](#)

["bc1-rgba-unorm-srgb"](#)

["bc2-rgba-unorm"](#)

["bc2-rgba-unorm-srgb"](#)

["bc3-rgba-unorm"](#)

["bc3-rgba-unorm-srgb"](#)

["bc4-r-unorm"](#)

["bc4-r-snorm"](#)

["bc5-rg-unorm"](#)

["bc5-rg-snorm"](#)

["bc6h-rgb-ufloat"](#)

["bc6h-rgb-float"](#)

["bc7-rgba-unorm"](#)

["bc7-rgba-unorm-srgb"](#)

25.5. "texture-compression-bc-sliced-3d"

Allows the [3d](#) dimension for textures with [BC compressed formats](#).

Note: Adapters which support ["texture-compression-bc"](#) do not always support ["texture-compression-bc-sliced-3d"](#). To use ["texture-compression-bc-sliced-3d"](#), ["texture-compression-bc"](#) must be enabled explicitly as this feature does not enable the BC formats.

This feature adds no [optional API surfaces](#).

25.6. "texture-compression-etc2"

Allows for explicit creation of textures of [ETC2 compressed formats](#). Only supports 2D textures.

This feature adds the following [optional API surfaces](#):

New [GPUTextureFormat](#) enum values:

["etc2-rgb8unorm"](#)

["etc2-rgb8unorm-srgb"](#)

["etc2-rgb8a1unorm"](#)

["etc2-rgb8a1unorm-srgb"](#)

["etc2-rgba8unorm"](#)

["etc2-rgba8unorm-srgb"](#)

["eac-r11unorm"](#)

["eac-r11snorm"](#)

["eac-rg11unorm"](#)

["eac-rg11snorm"](#)

25.7. "texture-compression-astc"

Allows for explicit creation of textures of [ASTC compressed formats](#). Only supports 2D textures.

This feature adds the following [optional API surfaces](#):

New [GPUTextureFormat](#) enum values:

["astc-4x4-unorm"](#)

["astc-4x4-unorm-srgb"](#)

["astc-5x4-unorm"](#)

["astc-5x4-unorm-srgb"](#)

["astc-5x5-unorm"](#)

["astc-5x5-unorm-srgb"](#)

["astc-6x5-unorm"](#)

["astc-6x5-unorm-srgb"](#)

["astc-6x6-unorm"](#)

["astc-6x6-unorm-srgb"](#)

["astc-8x5-unorm"](#)

["astc-8x5-unorm-srgb"](#)

["astc-8x6-unorm"](#)

["astc-8x6-unorm-srgb"](#)

["astc-8x8-unorm"](#)

["astc-8x8-unorm-srgb"](#)

["astc-10x5-unorm"](#)

["astc-10x5-unorm-srgb"](#)

["astc-10x6-unorm"](#)

["astc-10x6-unorm-srgb"](#)

["astc-10x8-unorm"](#)

["astc-10x8-unorm-srgb"](#)

["astc-10x10-unorm"](#)

["astc-10x10-unorm-srgb"](#)

["astc-12x10-unorm"](#)

["astc-12x10-unorm-srgb"](#)

["astc-12x12-unorm"](#)

["astc-12x12-unorm-srgb"](#)

25.8. "texture-compression-astc-sliced-3d"

Allows the [3d](#) dimension for textures with [ASTC compressed formats](#).

Note: Adapters which support ["texture-compression-astc"](#) do not always support ["texture-compression-astc-sliced-3d"](#). To use ["texture-compression-astc-sliced-3d"](#), ["texture-compression-astc"](#) must be enabled explicitly as this feature does not enable the ASTC formats.

This feature adds no [optional API surfaces](#).

25.9. "timestamp-query"

Adds the ability to query timestamps from GPU command buffers. See [§ 20.4 Timestamp Query](#).

This feature adds the following [optional API surfaces](#):

New [GPUQueryType](#) values:

["timestamp"](#)

New [GPUComputePassDescriptor](#) members:

[timestampWrites](#)

New [GPURenderPassDescriptor](#) members:

[timestampWrites](#)

25.10. "indirect-first-instance"

Allows the use of non-zero `firstInstance` values in [indirect draw parameters](#) and [indirect drawIndexed parameters](#).

This feature adds no [optional API surfaces](#).

25.11. "shader-f16"

Allows the use of the half-precision floating-point type [f16](#) in WGSL.

This feature adds the following [optional API surfaces](#):

New WGSL extensions:

[f16](#)

25.12. "rg11b10float-renderable"

Allows the [RENDER_ATTACHMENT](#) usage on textures with format ["rg11b10float"](#), and also allows textures of that format to be blended, multisampled, and resolved.

This feature adds no [optional API surfaces](#).

Enabling ["texture-formats-tier1"](#) at device creation will also enable ["rg11b10float-renderable"](#).

25.13. "bgra8unorm-storage"

Allows the [STORAGE_BINDING](#) usage on textures with format ["bgra8unorm"](#).

This feature adds no [optional API surfaces](#).

25.14. "float32-filterable"

Makes textures with formats ["r32float"](#), ["rg32float"](#), and ["rgba32float"](#) [filterable](#).

25.15. "float32-blendable"

Makes textures with formats ["r32float"](#), ["rg32float"](#), and ["rgba32float"](#) [blendable](#).

25.16. "clip-distances"

Allows the use of [clip_distances](#) in WGSL.

This feature adds the following [optional API surfaces](#):

New WGSL extensions:

[clip_distances](#)

25.17. "dual-source-blending"

Allows the use of [blend_src](#) in WGSL and simultaneously using both pixel shader outputs ([@blend_src\(0\)](#) and [@blend_src\(1\)](#)) as inputs to a blending operation with the single color attachment at [location 0](#).

This feature adds the following [optional API surfaces](#):

Allows the use of the below [GPUBlendFactors](#):

["src1"](#)

["one-minus-src1"](#)

["src1-alpha"](#)

["one-minus-src1-alpha"](#)

New WGSL extensions:

[dual_source_blending](#)

25.18. "subgroups"

Allows the use of the subgroup and quad operations in WGSL.

This feature adds no [optional API surfaces](#), but the following entries of [GPUAdapterInfo](#) expose real values whenever the feature is available on the adapter:

[subgroupMinSize](#)

[subgroupMaxSize](#)

New WGSL extensions:

[subgroups](#)

25.19. "texture-formats-tier1"

Supports the below new [GPUTextureFormats](#) with the [RENDER_ATTACHMENT](#), [blendable](#), multisampling capabilities and the [STORAGE_BINDING](#) capability with the ["read-only"](#) and ["write-only"](#) [GPUStorageTextureAccesses](#):

["r16unorm"](#)

["r16snorm"](#)

["rg16unorm"](#)

["rg16snorm"](#)

["rgba16unorm"](#)

["rgba16snorm"](#)

Allows the [RENDER_ATTACHMENT](#), [blendable](#), multisampling and resolve capabilities on below [GPUTextureFormats](#):

["r8snorm"](#)

["rg8snorm"](#)

["rgba8snorm"](#)

Allows the ["read-only"](#) or ["write-only"](#) [GPUStorageTextureAccess](#) on below [GPUTextureFormats](#):

["r8unorm"](#)

["r8snorm"](#)

["r8uint"](#)

["r8sint"](#)

["rg8unorm"](#)

["rg8snorm"](#)

["rg8uint"](#)

["rg8sint"](#)

["r16uint"](#)

["r16sint"](#)

["r16float"](#)

["rg16uint"](#)

["rg16sint"](#)

["rg16float"](#)

["rgb10a2uint"](#)

["rgb10a2unorm"](#)

["rg11b10ufloat"](#)

Enabling ["texture-formats-tier2"](#) at device creation will also enable ["texture-formats-tier1"](#).

Enabling ["texture-formats-tier1"](#) at device creation will also enable ["rg11b10ufloat-renderable"](#).

25.20. "texture-formats-tier2"

Allows the ["read-write"](#) [GPUStorageTextureAccess](#) on below [GPUTextureFormats](#):

["r8unorm"](#)

["r8uint"](#)

["r8sint"](#)

["rgba8unorm"](#)

["rgba8uint"](#)

["rgba8sint"](#)

["r16uint"](#)

["r16sint"](#)

["r16float"](#)

["rgba16uint"](#)

["rgba16sint"](#)
["rgba16float"](#)
["rgba32uint"](#)
["rgba32sint"](#)
["rgba32float"](#)

Enabling ["texture-formats-tier2"](#) at device creation will also enable ["texture-formats-tier1"](#).

25.21. "primitive-index"

Allows the use of [primitive_index](#) in WGSL.

This feature adds the following [optional API surfaces](#):

New WGSL extensions:

[primitive_index](#)

26. Appendices

26.1. Texture Format Capabilities

26.1.1. Plain color formats

All [supported](#) plain color formats support usages [COPY_SRC](#), [COPY_DST](#), and [TEXTURE_BINDING](#), and dimension ["3d"](#).

The [RENDER_ATTACHMENT](#) and [STORAGE_BINDING](#) columns specify support for [GPUTextureUsage.RENDER_ATTACHMENT](#) and [GPUTextureUsage.STORAGE_BINDING](#) usage respectively.

The *render target pixel byte cost* and *render target component alignment* are used to validate the [maxColorAttachmentBytesPerSample](#) limit.

Note: The [texel block memory cost](#) of each of these formats is the same as its [texel block copy footprint](#).

Format	Required Feature	GPUTextureSampleType	RENDER_ATTACHMENT	blendable	multisampling	resolve	STORAGE_BINDING			Texel block copy footprint (Bytes)	Render target pixel byte cost (Bytes)
							"write-only"	"read-only"	"read-write"		
8 bits per component (1-byte render target component alignment)											
r8unorm		"float", "unfilterable-float"	✓	✓	✓	✓	If "texture-formats-tier1" is enabled	If "texture-formats-tier2" is enabled	1		
r8snorm		"float", "unfilterable-float"	If "texture-formats-tier1" is enabled						1	–	
r8uint		"uint"	✓		✓		If "texture-formats-tier1" is enabled	If "texture-formats-tier2" is enabled	1		
r8sint		"sint"	✓		✓		If "texture-formats-tier1" is enabled	If "texture-formats-tier2" is enabled	1		
rg8unorm		"float", "unfilterable-float"	✓	✓	✓	✓	If "texture-formats-tier1" is enabled		2		
rg8snorm		"float", "unfilterable-float"	If "texture-formats-tier1" is enabled						2	–	
rg8uint		"uint"	✓		✓		If "texture-formats-tier1" is enabled		2		
rg8sint		"sint"	✓		✓		If "texture-formats-tier1" is enabled		2		
rgba8unorm		"float", "unfilterable-float"	✓	✓	✓	✓	✓	✓	If "texture-formats-tier2" is enabled	4	8
rgba8unorm-srgb		"float", "unfilterable-float"	✓	✓	✓	✓				4	8

Format	Required Feature	GPUTextureSampleType	RENDER_ATTACHMENT	blendable	multisampling	resolve	STORAGE_BINDING			Texel block copy footprint (Bytes)	Render target pixel byte cost (Bytes)
							"write-only"	"read-only"	"read-write"		
rgba8snorm		"float", "unfilterable-float"	If "texture-formats-tier1" is enabled			✓	✓		4	–	
rgba8uint		"uint"	✓		✓		✓	✓	If "texture-formats-tier2" is enabled	4	
rgba8sint		"sint"	✓		✓		✓	✓	If "texture-formats-tier2" is enabled	4	
bgra8unorm		"float", "unfilterable-float"	✓	✓	✓	✓	If "bgra8unorm-storage" is enabled			4	8
bgra8unorm-srgb		"float", "unfilterable-float"	✓	✓	✓	✓				4	8
16 bits per component (2-byte render target component alignment)											
r16unorm	"texture-formats-tier1"	"unfilterable-float"	✓	✓	✓		✓	✓		2	
r16snorm	"texture-formats-tier1"	"unfilterable-float"	✓	✓	✓		✓	✓		2	
r16uint		"uint"	✓		✓		If "texture-formats-tier1" is enabled		If "texture-formats-tier2" is enabled	2	
r16sint		"sint"	✓		✓		If "texture-formats-tier1" is enabled		If "texture-formats-tier2" is enabled	2	
r16float		"float", "unfilterable-float"	✓	✓	✓	✓	If "texture-formats-tier1" is enabled		If "texture-formats-tier2" is enabled	2	
rg16unorm	"texture-formats-tier1"	"unfilterable-float"	✓	✓	✓		✓	✓		4	
rg16snorm	"texture-formats-tier1"	"unfilterable-float"	✓	✓	✓		✓	✓		4	
rg16uint		"uint"	✓		✓		If "texture-formats-tier1" is enabled			4	
rg16sint		"sint"	✓		✓		If "texture-formats-tier1" is enabled			4	
rg16float		"float", "unfilterable-float"	✓	✓	✓	✓	If "texture-formats-tier1" is enabled			4	
rgba16unorm	"texture-formats-tier1"	"unfilterable-float"	✓	✓	✓		✓	✓		8	
rgba16snorm	"texture-formats-tier1"	"unfilterable-float"	✓	✓	✓		✓	✓		8	
rgba16uint		"uint"	✓		✓		✓	✓	If "texture-formats-tier2" is enabled	8	
rgba16sint		"sint"	✓		✓		✓	✓	If "texture-formats-tier2" is enabled	8	
rgba16float		"float", "unfilterable-float"	✓	✓	✓	✓	✓	✓	If "texture-formats-tier2" is enabled	8	

Format	Required Feature	GPUTextureSampleType	RENDER_ATTACHMENT	blendable	multisampling	resolve	STORAGE_BINDING			Texel block copy footprint (Bytes)	Render target pixel byte cost (Bytes)
							"write-only"	"read-only"	"read-write"		
									tier2 is enabled		
32 bits per component (4-byte render target component alignment)											
r32uint		"uint"	✓				✓	✓	✓	4	
r32sint		"sint"	✓				✓	✓	✓	4	
r32float		"float" if " float32-filterable " is enabled "unfilterable-float"	✓	If " float32-blendable " is enabled	✓		✓	✓	✓	4	
rg32uint		"uint"	✓				✓	✓		8	
rg32sint		"sint"	✓				✓	✓		8	
rg32float		"float" if " float32-filterable " is enabled "unfilterable-float"	✓	If " float32-blendable " is enabled			✓	✓		8	
rgba32uint		"uint"	✓				✓	✓	If " texture-formats-tier2 " is enabled	16	
rgba32sint		"sint"	✓				✓	✓	If " texture-formats-tier2 " is enabled	16	
rgba32float		"float" if " float32-filterable " is enabled "unfilterable-float"	✓	If " float32-blendable " is enabled			✓	✓	If " texture-formats-tier2 " is enabled	16	
mixed component width, 32 bits per texel (4-byte render target component alignment)											
rgb10a2uint		"uint"	✓		✓		If " texture-formats-tier1 " is enabled			4	8
rgb10a2unorm		"float", "unfilterable-float"	✓	✓	✓	✓	If " texture-formats-tier1 " is enabled			4	8
rg11b10ufloat		"float", "unfilterable-float"	If " rg11b10ufloat-renderable " is enabled				If " texture-formats-tier1 " is enabled			4	8

26.1.2. Depth-stencil formats

A *depth-or-stencil format* is any format with depth and/or stencil aspects. A *combined depth-stencil format* is a [depth-or-stencil format](#) that has both depth and stencil aspects.

All [depth-or-stencil formats](#) support the [COPY_SRC](#), [COPY_DST](#), [TEXTURE_BINDING](#), and [RENDER_ATTACHMENT](#) usages. All of these formats support multisampling. However, certain copy operations also restrict the source and destination formats, and none of these formats support textures with "[3d](#)" dimension.

Depth textures cannot be used with "[filtering](#)" samplers, but can always be used with "[comparison](#)" samplers even if they use filtering.

Format	NOTE: Texel block memory cost (Bytes)	Aspect	GPUTextureSampleType	Valid texel copy source	Valid texel copy destination	Texel block copy footprint (Bytes)	Aspect-specific format
stencil8	1 – 4	stencil	"uint"	✓		1	stencil8
depth16unorm	2	depth	"depth" , "unfilterable-float"	✓		2	depth16unorm
depth24plus	4	depth	"depth" , "unfilterable-float"	✗		–	depth24plus
depth24plus-stencil8	4 – 8	depth	"depth" , "unfilterable-float"	✗		–	depth24plus
		stencil	"uint"	✓		1	stencil8
depth32float	4	depth	"depth" , "unfilterable-float"	✓	✗	4	depth32float

Format	NOTE: Texel block memory cost (Bytes)	Aspect	GPUTextureSampleType	Valid texel copy source	Valid texel copy destination	Texel block copy footprint (Bytes)	Aspect-specific format
depth32float-stencil8	5 – 8	depth	"depth", "unfilterable-float"	✓	✗	4	depth32float
		stencil	"uint"	✓		1	stencil8

24-bit *depth* refers to a 24-bit unsigned normalized depth format with a range from 0.0 to 1.0, which would be spelled "depth24unorm" if exposed.

26.1.2.1. Reading and Sampling Depth/Stencil Textures

It is [possible](#) to bind a depth-aspect [GPUTextureView](#) to either a `texture_depth_*` binding or a binding with other non-depth 2d/cube texture types.

A stencil-aspect [GPUTextureView](#) must be bound to a normal texture binding type. The [sampleType](#) in the [GPUBindGroupLayout](#) must be ["uint"](#).

Reading or sampling the depth or stencil aspect of a texture behaves as if the texture contains the values (V, X, X, X), where V is the actual depth or stencil value, and each X is an [implementation-defined](#) unspecified value.

For depth-aspect bindings, the unspecified values are not visible through bindings with `texture_depth_*` types.

If a depth texture is bound to `tex` with type `texture_2d<f32>`:

`textureSample(tex, ...)` will return `vec4<f32>(D, X, X, X)`.

`textureGather(0, tex, ...)` will return `vec4<f32>(D1, D2, D3, D4)`.

`textureGather(2, tex, ...)` will return `vec4<f32>(X1, X2, X3, X4)` (a completely unspecified value).

Note: Short of adding a new more constrained stencil sampler type (like depth), it’s infeasible for implementations to efficiently paper over the driver differences for depth/stencil reads. As this was not a portability pain point for WebGL, it’s not expected to be problematic in WebGPU. In practice, expect either (V, V, V, V) or (V, 0, 0, 1) (where V is the depth or stencil value), depending on hardware.

26.1.2.2. Copying Depth/Stencil Textures

The depth aspects of depth32float formats (["depth32float"](#) and ["depth32float-stencil8"](#) have a limited range. As a result, copies into such textures are only valid from other textures of the same format.

The depth aspects of depth24plus formats (["depth24plus"](#) and ["depth24plus-stencil8"](#)) have opaque representations (implemented as either [24-bit depth](#) or ["depth32float"](#)). As a result, depth-aspect [texel copies](#) are not allowed with these formats.

NOTE:

It is possible to imitate these disallowed copies:

All of these formats can be written in a render pass using a fragment shader that outputs depth values via the `frag_depth` output.

Textures with "depth24plus" formats can be read as shader textures, and written to a texture (as a render pass attachment) or buffer (via a storage buffer binding in a compute shader).

26.1.3. Packed formats

All packed texture formats support [COPY_SRC](#), [COPY_DST](#), and [TEXTURE_BINDING](#) usages. All of these formats are [filterable](#). None of these formats are [renderable](#) or support multisampling.

A *compressed format* is any format with a block size greater than 1×1.

Note: The [texel block memory cost](#) of each of these formats is the same as its [texel block copy footprint](#).

Format	Texel block copy footprint (Bytes)	GPUTextureSampleType	Texel block width/height	"3d"	Feature
rgb9e5ufloat	4	"float", "unfilterable-float"	1 × 1	✓	
bcl-rgba-unorm	8	"float", "unfilterable-float"	4 × 4	If "texture-compression-bc-sliced-3d" is enabled	texture-compression-bc
bc1-rgba-unorm-srgb					
bc2-rgba-unorm	16				
bc2-rgba-unorm-srgb					
bc3-rgba-unorm	16				
bc3-rgba-unorm-srgb					
bc4-r-unorm	8				
bc4-r-snorm					

Format	Texel block copy footprint (Bytes)	GPUTextureSampleType	Texel block width/height	"3d"	Feature
bc5-rg-unorm	16				
bc5-rg-snorm					
bc6h-rgb-ufloat	16				
bc6h-rgb-float					
bc7-rgba-unorm	16				
bc7-rgba-unorm-srgb					
etc2-rgb8unorm	8	"float", "unfilterable-float"	4 × 4		texture-compression-etc2
etc2-rgb8unorm-srgb					
etc2-rgb8a1unorm	8				
etc2-rgb8a1unorm-srgb					
etc2-rgba8unorm	16				
etc2-rgba8unorm-srgb					
eac-r11unorm	8				
eac-r11snorm					
eac-rg11unorm	16				
eac-rg11snorm					
astc-4x4-unorm	16	"float", "unfilterable-float"	4 × 4	If "texture-compression-astc-sliced-3d" is enabled	texture-compression-astc
astc-4x4-unorm-srgb					
astc-5x4-unorm	5 × 4				
astc-5x4-unorm-srgb					
astc-5x5-unorm	5 × 5				
astc-5x5-unorm-srgb					
astc-6x5-unorm	6 × 5				
astc-6x5-unorm-srgb					
astc-6x6-unorm	6 × 6				
astc-6x6-unorm-srgb					
astc-8x5-unorm	8 × 5				
astc-8x5-unorm-srgb					
astc-8x6-unorm	8 × 6				
astc-8x6-unorm-srgb					
astc-8x8-unorm	8 × 8				
astc-8x8-unorm-srgb					
astc-10x5-unorm	10 × 5				
astc-10x5-unorm-srgb					
astc-10x6-unorm	10 × 6				
astc-10x6-unorm-srgb					
astc-10x8-unorm	10 × 8				
astc-10x8-unorm-srgb					
astc-10x10-unorm	10 × 10				
astc-10x10-unorm-srgb					
astc-12x10-unorm	12 × 10				
astc-12x10-unorm-srgb					
astc-12x12-unorm	12 × 12				
astc-12x12-unorm-srgb					